



MySQL Upgrade

Best Practices

Peter Zaitsev, CEO

Percona Inc

April 18, 2012

MySQL Upgrade: Agenda

- Look at Different types of Upgrade
- Discuss when to upgrade and not
- Issues you can discover during Upgrade
- Look through the process you can use for upgrade

What is your Experience ?

- How many have done MySQL Upgrades ?
- How many of you have done 1 major version transition ?
 - Two ? More ?
- Did you get any “surprises” during upgrade ?
- Did you ever have to roll back ?

Upgrade as Version Changes

- Issues are similar for upgrade and over version changes
- Moving between MySQL, Percona Server, MariaDB, proprietary solution ?
- Downgrading ?

Reasons to Upgrade

- Bugs in old versions
- Security problems in old version
- New Version benefits
 - Features, scalability, performance, new bugs
- Hard to find people familiar with old versions
 - Anyone remembers MySQL 3.23 restrictions ?
- Support Becomes limited

Case to stay on old MySQL version

- Application is on isolated network
- Application is not being actively developed
- Application is not growing actively
- No changes in Platform
 - Hardware or Operating system

Which Upgrades are Risky ?

- Major upgrades, such as 5.1 to 5.5 are riskier
- MySQL to Percona Server or MariaDB
- Jumping over many minor versions at once
 - 5.1.20 to 5.1.61
- Upgrades from PreGA versions
- Skipping one or more major releases
 - 4.1 to 5.5 is tested less than 5.1 to 5.5

Typical reasons for MySQL Upgrade project

- Running into bug or performance issue
- Concern with bug which can potentially affect system
 - Checking change logs for relevant bugs
- Security Concerns
- Generally keeping up with fresh versions

Issues to consider with Upgrade

- Data
- Queries
 - Reads and Writes
- Performance and Scalability
- Replication
- Resource Usage
- Advanced Features

MySQL Upgrade Data Issues

- On Disk format changes
- Changes to MySQL Types
- Sorting order Changes
- Data Presentation Issues
 - TIMESTAMP text format changed in
- Changes to limits
- Reserved Words
- Statistics

Issues with Queries

- Syntax might have changed
- Query has different result
 - Changes in query “meaning”
 - Non deterministic queries
- Query Produces warnings or Error

Performance

- Query Execution Time
- Query Plan Changes
- Performance at Concurrency
 - Scalability
 - Locking

Replication

- Can workload be still replicated
 - And are there any warnings in error logs ?
 - Any “data drift” causing inconsistent slave
- Is replication performance adequate ?

Resource Usage

- Memory is main resource to care
- Check memory usage for new version
 - Can be memory leaks or simply more memory needed

Advanced Features

- Complex Features = More things to break
- Pay special attention to
 - Stored Procedures/Functions
 - Plugins
 - Triggers
 - Events
 - Views

Upgrading Multiple Environments

- Development, Staging, Production
- Do we upgrade Development first or last ?
 - Both have their problems
- Creating special Upgraded Development environment is best

Drop In or Replication Based Upgrades

- Drop-In for small systems
 - Which can handle downtime or risks
 - LVM snapshot great for quick rollback
 - <http://bit.ly/yQwm2v>
- Drop in for “shards”
 - When in place upgrade is well tested
- Replication based upgrades are best default

Upgrading in the Cloud

- Great Possibilities !
- Inexpensive way to temporary increase your infrastructure size
 - Can build a “clone” on updated version and test thoroughly

Many Changes at Once


- Are you upgrading Hardware, OS, Application, Changing Configuration at the same time ?
- Good
 - Usually can invest more time in testing
- Bad
 - If something breaks you would not be sure what to blame.
 - Many more moving parts. Riskier.

Reckless or Paranoid

- For some upgrade is as simple as “yum update”
- Others take months to upgrade
- A lot depends on application complexity and your risk tolerance
 - We’ll look at reasonable conservative approach in this talk.

Upgrading Replication

- Upgrade Slaves first
 - Upgrade Inactive Master In Master-Master Pair
- Setup extra Slaves or “Replication Tree if possible”



Upgrading Sharded Environment

- You can take shortcuts
- No full testing for each shard might be needed
- Upgrade one shard, create “script” and monitor
- Update couple of more shards with full testing
- Update over shards with reduced testing

Doing an Upgrade

MySQL Upgrade Process in Details

Read the Release Notes

- Understanding changes helps you to know what to expect
- Major Releases Come with great Summary documentation
 - <http://dev.mysql.com/doc/refman/5.5/en/mysql-nutshell.htm>

Perform Preliminary testing

- Setup QA environment with new release
- See if it works
 - Helps to understand amount of effort needed for upgrade

Adjust MySQL Configuration

- Focus on limited changes, you can fine tune later.
 - remove depreciated options
 - Adjust defaults you might relay on
 - **storage_engine=MYISAM**
 - Set options which is hard to change later
 - **innodb_file_format=Barracuda**
 - New compatibility requirements
 - READ-COMMITTED does not work with Statement replication in MySQL 5.5

Moving The Data

- Mysqldump and import back
 - The most clean way to move data, but slowest
 - Helpful for upgrading very old databases or when you want to move to new features
- Mysql_upgrade
 - Will check tables for compatibility and “fix” incompatible tables
 - Run this in any case as it upgrade system tables

Check the data is the same

- In both cases data may start to “look” differently
- Check data is the same in both versions
 - Compare restored backup to upgraded restored backup
 - CHECKSUM TABLE may report false positives
 - Pt-table-checksum
 - Mk-table-checksum can compare standalone servers.

Replication Old->New

- You will run this some way if you do not schedule downtime
 - Let replication run for 24 hours or so
 - Check it is in sync with pt-table-checksum
 - Check error logs on the Slave for errors or warnings

Replication New->Old

- Needs to be tested too if this is roll back strategy
 - Old->New->Old chain good for testing
 - Does not 100% represent production use case
 - Run New->Old in QA to be sure
- Unsupported but works in most cases
 - MySQL 5.0-> MySQL 5.5 works in most cases.

Replication New->New

- Is not expected to give problems if previous cases work
- Might wish to check it for extra safety
- Often validated post factum after upgrade
 - As normal replication consistency check practice

Validate Replication Performance

- You want to ensure replication is not slower in new version
- Measure catch up speed in All Versions
- See replication thread utilization in Percona Server or MariaDB

Getting Queries for Test

- Get representative live traffic from production
 - Full query log/ slow query log with `long_query_time=0`
 - Tcpdump (pt-query-digest)
- Get several samples for each query
 - `pt-query-digest --sample 5 --print --no-report queries.log > samples.log`

Checking Queries

- Single connection test
 - Run pt-upgrade on sample of queries
- High concurrency test
 - Run pt-log-player with full traffic
 - Use 2 sets; one for warmup over for real run
- Looking at pt-query-digest report on full query log is good to check in both cases
 - Outliers; number of examined rows etc.

Do you have load testing setup ?

- Full application load testing is great final step
- Few companies have it setup well

Consider Getting Help

- Upgrades is something you do not get a lot of training working on single applications
- Common upgrade issues are different all the time 4.0->4.1 and 5.1->5.5 are very different
- We at Percona can help !

Thank You !

- pz@percona.com
- <http://www.percona.com>
- @percona
 - Use #perconalive
- <http://www.facebook.com/Percona>
- Slides will be published on the conference site