



[Introduction to] Writing non-blocking code ... in Node.js and Perl

Jay Janssen
Percona, Inc.

Problems solved

- Standard programs assume sequential execution of statements
- Long running statements block the program while waiting for a response
- Non-blocking programming allows your process to handle other tasks while waiting for long operations to complete

Good use cases

- Servers that need to handle lots of concurrency, but that often block on network or IO operations
 - example: node.js
- A process that coordinates a lot of tasks, in particular timed tasks
 - example: execute an arbitrary set of tasks, each with their own schedule
- Any process that interacts with an external system
 - example: a process that interacts heavily with Disk, Database, webservices, etc.

Overview

- High level Languages: Perl AnyEvent and Node.js, not C/C++ libevent level programming
- Agenda:
 - Intro to non-blocking language concepts
 - Example code in Node.js and Perl
 - Useful libraries
 - Tradeoffs



Intro to non-blocking language concepts

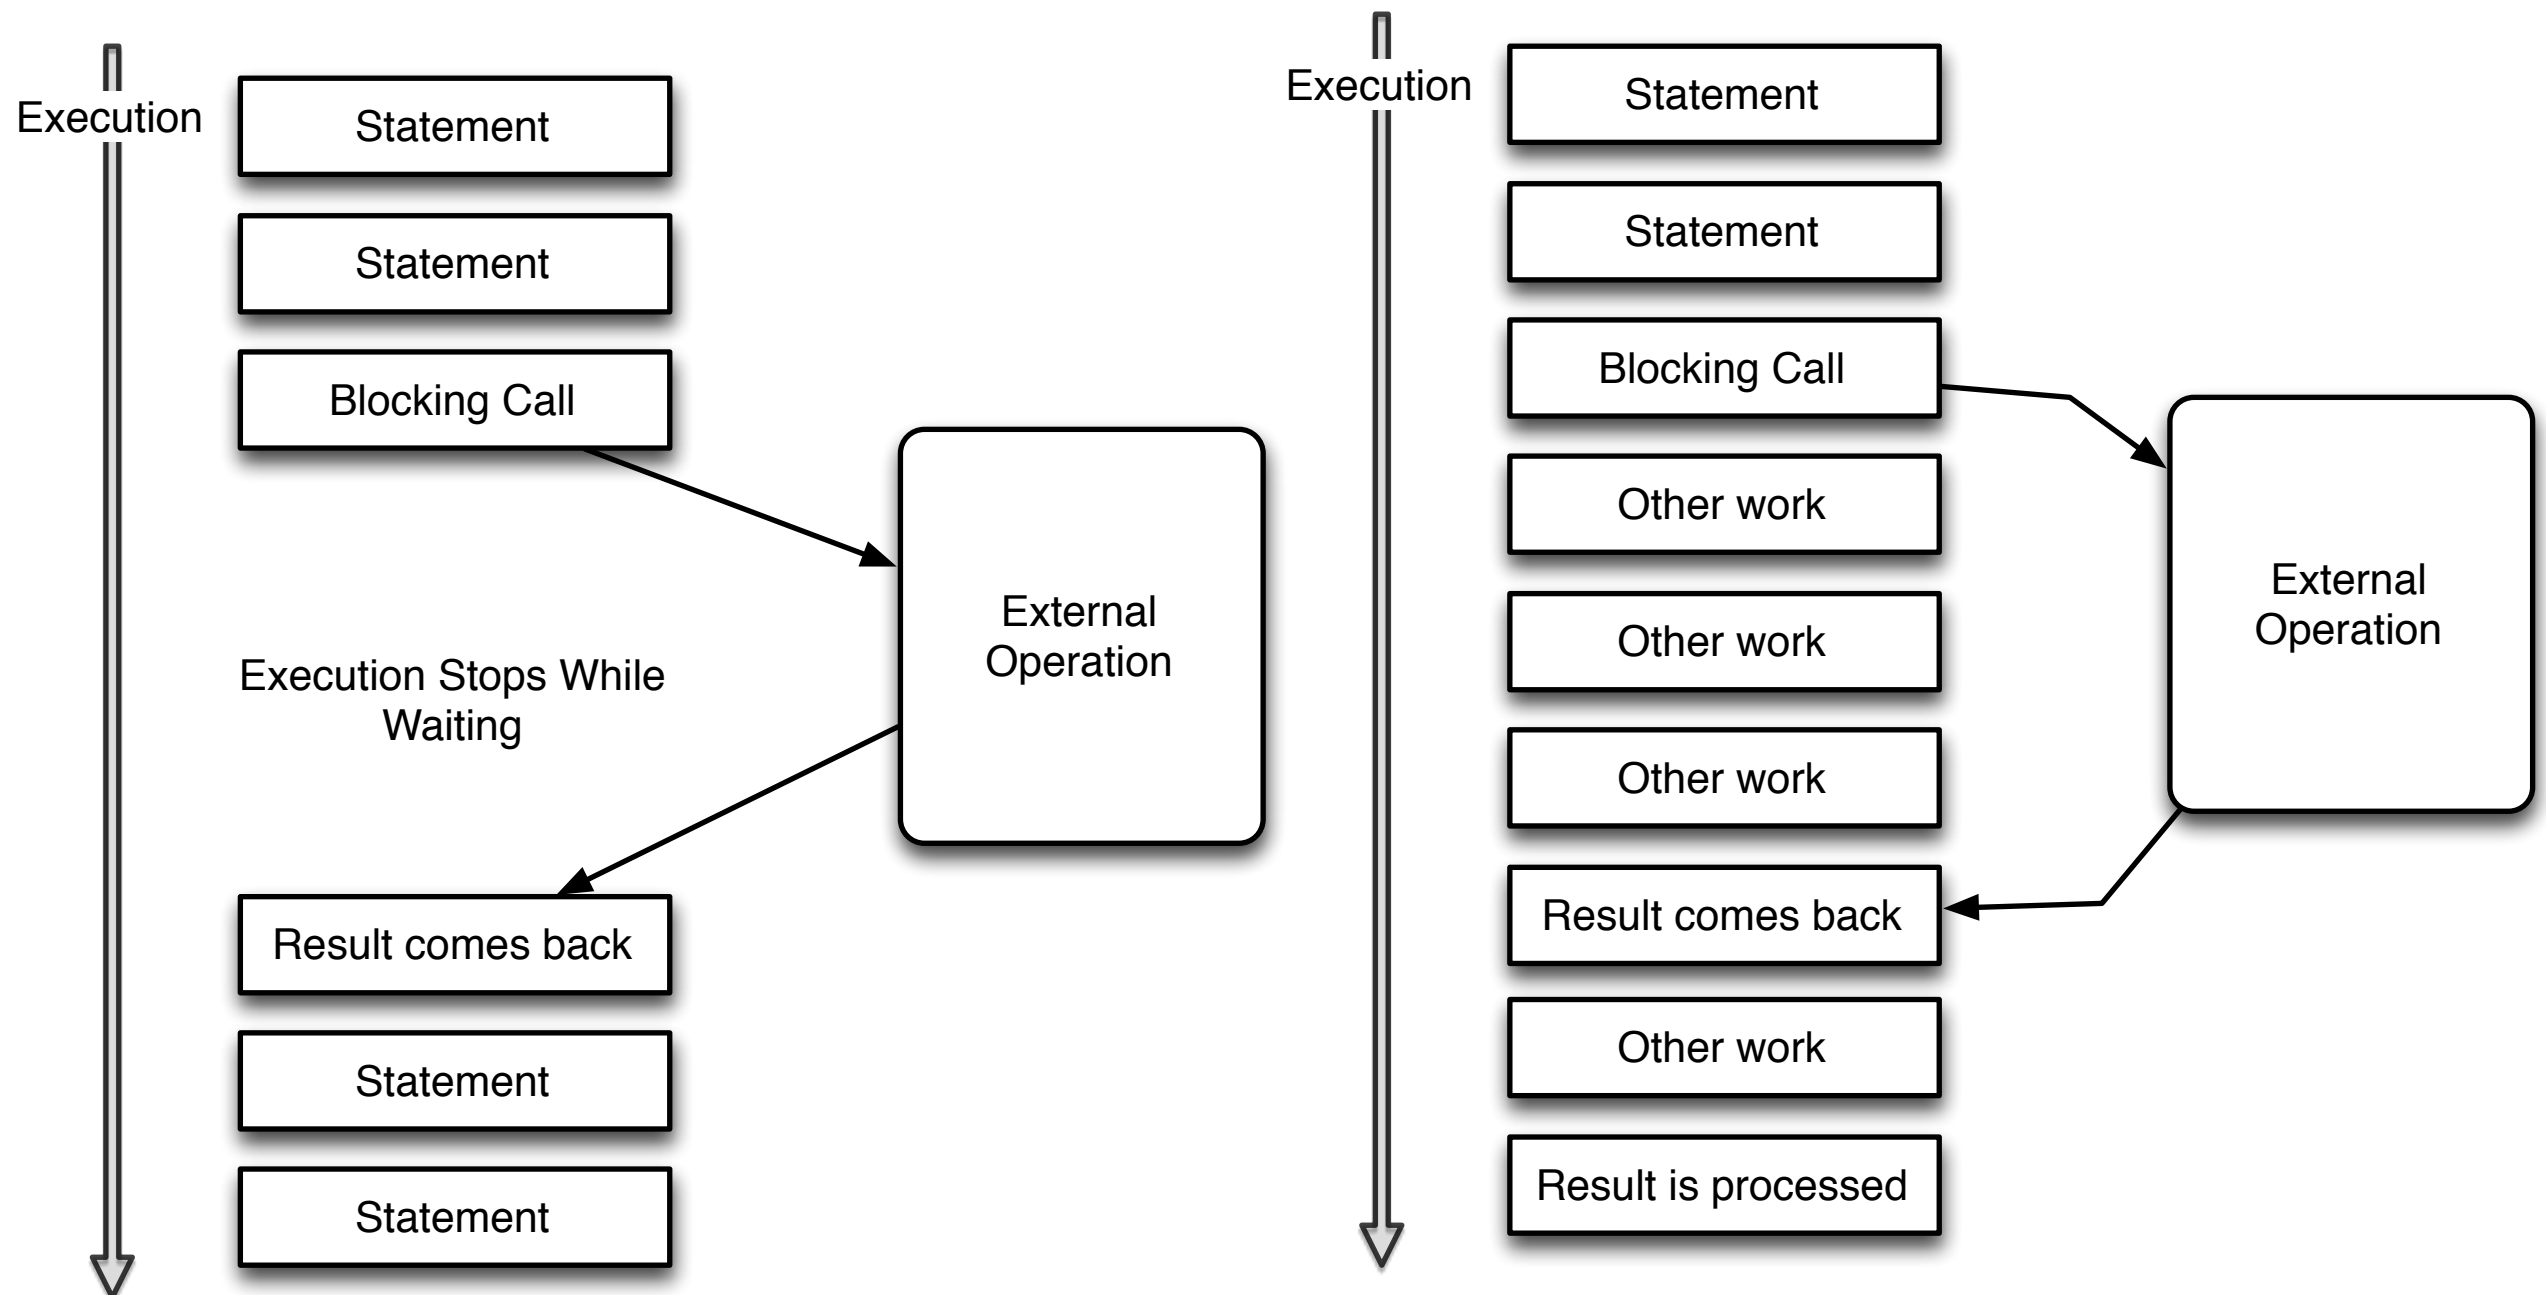
Definitions

- **procedural single-process:** a process that runs a single execution path
- **procedural threaded:** process running multiple execution paths simultaneously across CPU cores (ithreads)
- **non-blocking process:** process runs multiple execution paths that are executed serially (but interwoven) on a single CPU core

More Definitions

- **blocking call**: a function that waits for an answer before continuing execution.
- **non-blocking call**: a function that starts some operation (e.g., send a message to a server) and then returns immediately continuing execution. The response from the call is handled via a callback passed to the function.

Blocking vs Non-blocking flow



What does it look like?

```
1.  #!/usr/bin/perl -w

3.  use strict;

5.  use AnyEvent;
6.  use AnyEvent::HTTP;

8.  my $cv = AnyEvent->condvar;

10. http_get "http://forecast.weather.gov/MapClick.php?  

lat=42.12100&lon=-77.94750&FcstType=dwml", sub {
11.     my( $data, $headers ) = @_;
12.
13.     print "Got HTTP response: " . $headers->{Status} . "\n";
14.
15.     $cv->send;
16. };

18. print "Request sent\n";

20. $cv->wait;

22. print "Data received\n";
```

How does it really execute?

- True execution is serial
- Non-blocking calls allow execution to continue while the call completes
- All pending tasks are tracked and resumed when the response callback is ready to be processed
- Code paths cannot be interrupted until they enter the event loop , so precise scheduling is not possible

Advantages and Disadvantages

- Eliminates busy loops with 'sleep', write event handlers instead
- Brings scheduling inside of user-space
- Less/No locking needed compared with a threaded process
- Can fully utilize a CPU core and no more
- Less available libraries (i.e., Perl)



Node.js and Perl Examples

Timers -- AnyEvent

```
1.  #!/usr/bin/perl -w

3.  use strict;

5.  use AnyEvent;

7.  my $cv = AnyEvent->condvar;

9.  my $t = AnyEvent->timer( after => 5, cb => sub {
10.      print "5 seconds passed\n";
11.      $cv->send;
12.  });

14. my $t2 = AnyEvent->timer( after => 1, interval => 1, cb =>
    sub {
15.      print "1 second passed\n";
16.  });

18. $cv->wait;
```

Timers - Node.js

```
1. var intervalID = setInterval( function() {  
2.     console.log( "1 second passed" );  
3. }, 1000 );  
  
5. setTimeout( function() {  
6.     console.log( "5 seconds passed" );  
7.     clearInterval( intervalID );  
8. }, 5000 );
```

HTTP curl calls -- AnyEvent::HTTP

```
1.  #!/usr/bin/perl -w
3.  use strict;
5.  use AnyEvent;
6.  use AnyEvent::HTTP;
8.  my $cv = AnyEvent->condvar;
10. http_get "http://forecast.weather.gov/MapClick.php?  
    lat=42.12100&lon=-77.94750&FcstType=dwml", sub {
11.     my( $data, $headers ) = @_;
12.
13.     print "Got HTTP response: " . $headers->{Status} . "\n";
14.     print "Data is " . length( $data ) . " bytes long\n";
15.
16.     $cv->send;
17. };
19. print "Request sent\n";
21. $cv->wait;
23. print "Data received\n";
```

HTTP calls -- Node

```
1.  var http = require('http');

3.  http.get(
4.      {
5.          host: 'forecast.weather.gov',
6.          path: '/MapClick.php?lat=42.12100&lon=-77.94750&FcstType=dwml'
7.      },
8.      function( res ) {
9.          var data;
10.
11.         console.log( "Got HTTP response " + res.statusCode );
12.
13.         res.on( 'data', function( chunk ) {
14.             data += chunk;
15.         });
16.
17.         res.on( 'end', function() {
18.             console.log( "Data is " + data.length + " bytes long" );
19.             console.log('Data received');
20.         })
21.     }
22. );

24. console.log('Request sent');
```




AnyEvent::HTTPD vs Node.js



AIO -- Perl

Serve large files - Node.js

```
1. var http = require('http');
2. var fs = require('fs');

4. var i = 1;
5. http.createServer(function (request, response) {
6.     console.log('starting #' + i++);
7.     var stream = fs.createReadStream('data.bin',
    { bufferSize: 64 * 1024 });
8.     stream.pipe(response);
9. }).listen(8000);

11. console.log('Server running at http://
    127.0.0.1:8000/');
```

Database -- AnyEvent::DBI

```
1.  #!/usr/bin/perl -w

3.  use strict;

5.  use AnyEvent;
6.  use AnyEvent::DBI;

8.  my $cv = AnyEvent->condvar;

10. my $dbh = new AnyEvent::DBI "DBI:mysql:dbname=test", 'root', '';

12. $dbh->exec ("select * from test where id=?", 10, sub {
13.     my ($dbh, $rows, $rv) = @_;

15.     $_[0] or die "failure: $_[1]";
16.
17.     my $arr = $rows->[0];
18.     print "(";
19.     print join( ', ', @$arr );
20.     print ")\n";

22.     $cv->send;
23. });

25. $cv->wait;
```


Database - Node

```
1. var mysql = require( 'mysql' );

3. var client = mysql.createClient( {
4.     user: 'root',
5.     database: 'test'
6. });

8. client.query(
9.     'select * from test where id = ?', [ 10 ],
10.    function ( err, results, fields ) {
11.        if( err ) {
12.            throw err;
13.        }
14.        console.log( results );
15.        client.end();
16.    }
17. );
```

Interesting AnyEvent modules

- Memcached
- Curl::Multi
- AIO
- Redis
- Handle
- Filesys::Notify
- Gearman
- XMLRPC
- DNS
- MongoDB
- Pg
- IRC
- BDB
- Run
- JSONRPC
- Twitter
- Worker
- CouchDB
- Riak
- RabbitMQ
- Kanye



The tradeoff

Issues with Non-blocking Libraries

- Less common and unsupported by central development teams in general
- Concurrency tends to be limited -- limitations of the client protocol
- Transactional support tends to not “just work”.

Debugging Non-blocking Code

- Hard-to-reproduce race condition crashes
- Traditional “stack trace” debugging much harder to do
- Some non-blocking frameworks offer debugging tools (e.g., Perl’s “Coro”)
- Context variables may be modified/undefined/etc. between the initial call and the callback.
- This is what tripped me up the most.

Example Crash

```
1. my $temporary_object = Obj->new();

3. $object->do_query_nb( "SELECT * FROM
    my_table", callback => sub {
4.     my( $result ) = @_;
5.     $temporary_object->print( $result );
6.     # ^^ We get a crash on an undefined
7.     # object here, why?
8. } );

10.undef $temporary_object;
```

Code Organization

- Code gets nested deeply very quickly
- A set of non-blocking tasks that each depend on the previous to finish can be tedious to code.
- Some helpers exist to work-around this:
 - Node.js 'async' library
 - Perl - Coro and rouse_cb/rouse_wait
- Otherwise, think carefully about your Object structure and build helper functions to help

Coro rouse_cb example

```
1.  #!/usr/bin/perl -w
3.  use strict;
5.  use Coro;
7.  use AnyEvent;
8.  use AnyEvent::DBI;
10. my $cv = AnyEvent->condvar;
12. my $dbh = new AnyEvent::DBI "DBI:mysql:dbname=test", 'root', '';
14. # non-blocking call to the DB, defer handling response until later.
15. $dbh->exec( "select * from test where id=?", 10, Coro::rouse_cb );
17. # do something else here
19. # now block this execution path until the query results come back.
20. my ($dbh2, $rows, $rv) = Coro::rouse_wait;
22. $_[0] or die "failure: $_[0]";
24. my $arr = $rows->[0];
25. print "(";
26. print join( ', ', @$arr );
27. print ")\n";
```


Node.js 'async' library example

```
1.  var async = require( 'async' );
2.  //...

4.  // Start a "transaction"
5.  async.waterfall([
6.      // statement 1
7.      function( callback ) {
8.          client.query( 'select * from test where id = ?', [ 10 ], callback );
9.      },
10.     // // statement 2, gets all arguments passed to last to last callback
11.     function( last_results, last_fields, callback ) {
12.         console.log( last_results );
13.         client.query( 'select * from test where id = ?', [ 5 ], callback );
14.     }
15. ],

17. // This function gets called if any statement produces an error, or all statements
    complete successfully
18. function( err, result ) {
19.     if( err ) {
20.         throw err;
21.     } else {
22.         console.log( "Done successfully" );
23.         console.log( result );
24.         client.end();
25.     }
26. });
```

Conclusion

- Non-blocking programming made coding fun for me again
- See all my code examples in a working VM:
- <https://github.com/jayjanssen/non-blocking-code-examples>