



OLTP and OLAP

Have Your Cake and Eat it Too!

Lawrence Schwartz

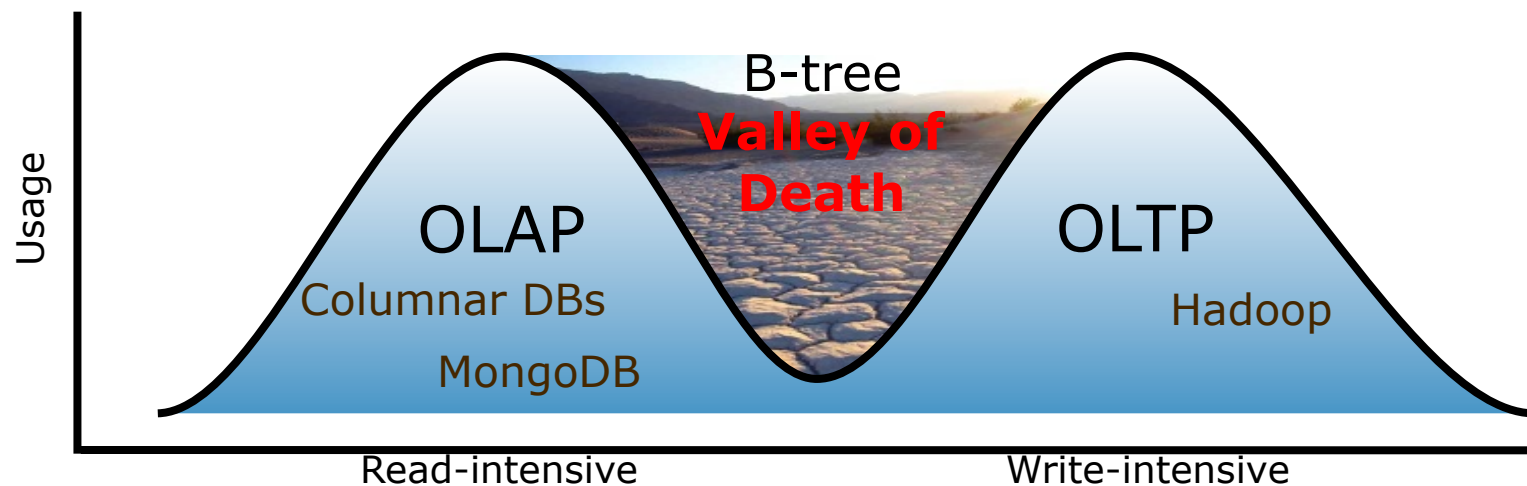


Avoiding the B-tree Valley of Death



“The real end-game for Big Data is to have transactional and analytic data on the same database.”


However, Today's database categories - OLAP, OLTP are compromises to get around dated, rigid indexing technology



B-trees Perform Poorly on Disk

MySQL + InnoDB

Based on **B-Tree** Indexes




- **Random Workloads Lead To Frequent Drive Head Movement**
- **MySQL Can't Take Advantage of Full Drive I/O Potential**



MySQL + TokuDB


Based on **Fractal Tree** Indexes




- **Continuous Rebalancing and Aggregation at Nodes**
- **Eliminates Fragmentation, Maximizes Disk I/O**

B-trees Can't Handle Big Data


MySQL + InnoDB



- **Performance Drops to a Trickle above 50-100M Rows**
- **Delayed Checkpoints Diminish Performance**



MySQL + TokuDB



- **20x Insertion Rates at One Billion Rows**
- **Stand Up More Indexes for Better Query Performance**

B-trees Aren't Agile

MySQL + InnoDB

InnoDB
Index Creation

00:31:34

InnoDB
Column Addition

17:44:41

- Big Database Schema Changes Take Hours To Complete
- Changes Have to Wait for Scheduled Downtime



MySQL + TokuDB

TokuDB
Hot Indexing

00:00:02

TokuDB
Hot Column Addition

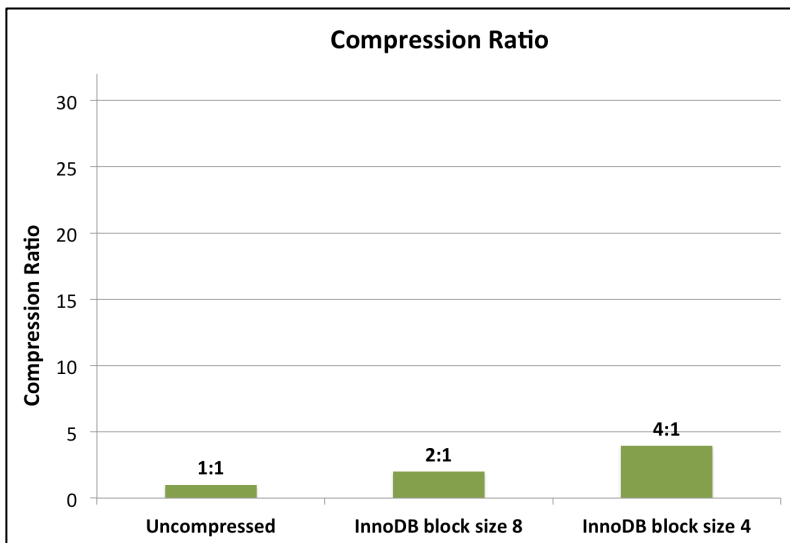
00:00:03

- Changes Complete in Seconds
- Allows Business' Data Models to Immediately Adapt

Tokutek

B-trees Don't Compress Well

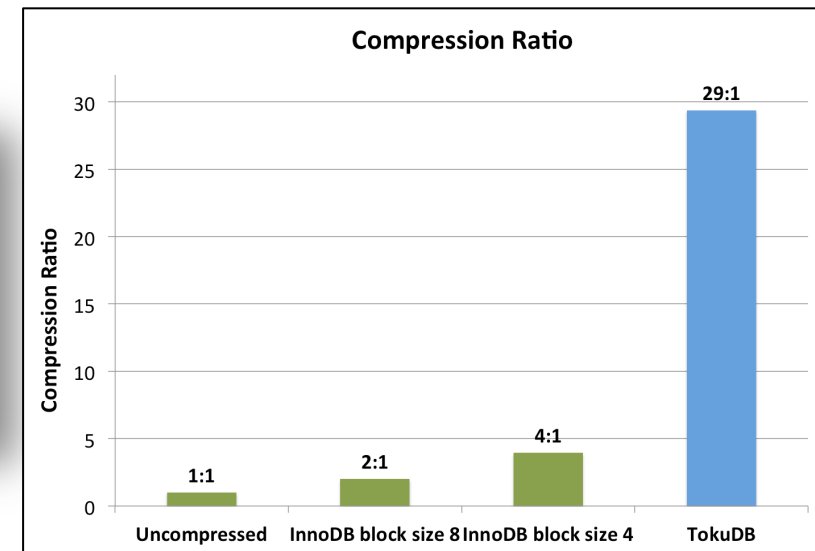
MySQL + InnoDB



- **Compression ~2x to ~2.5x Typical**
- **Compression Impacts Performance**



MySQL + TokuDB

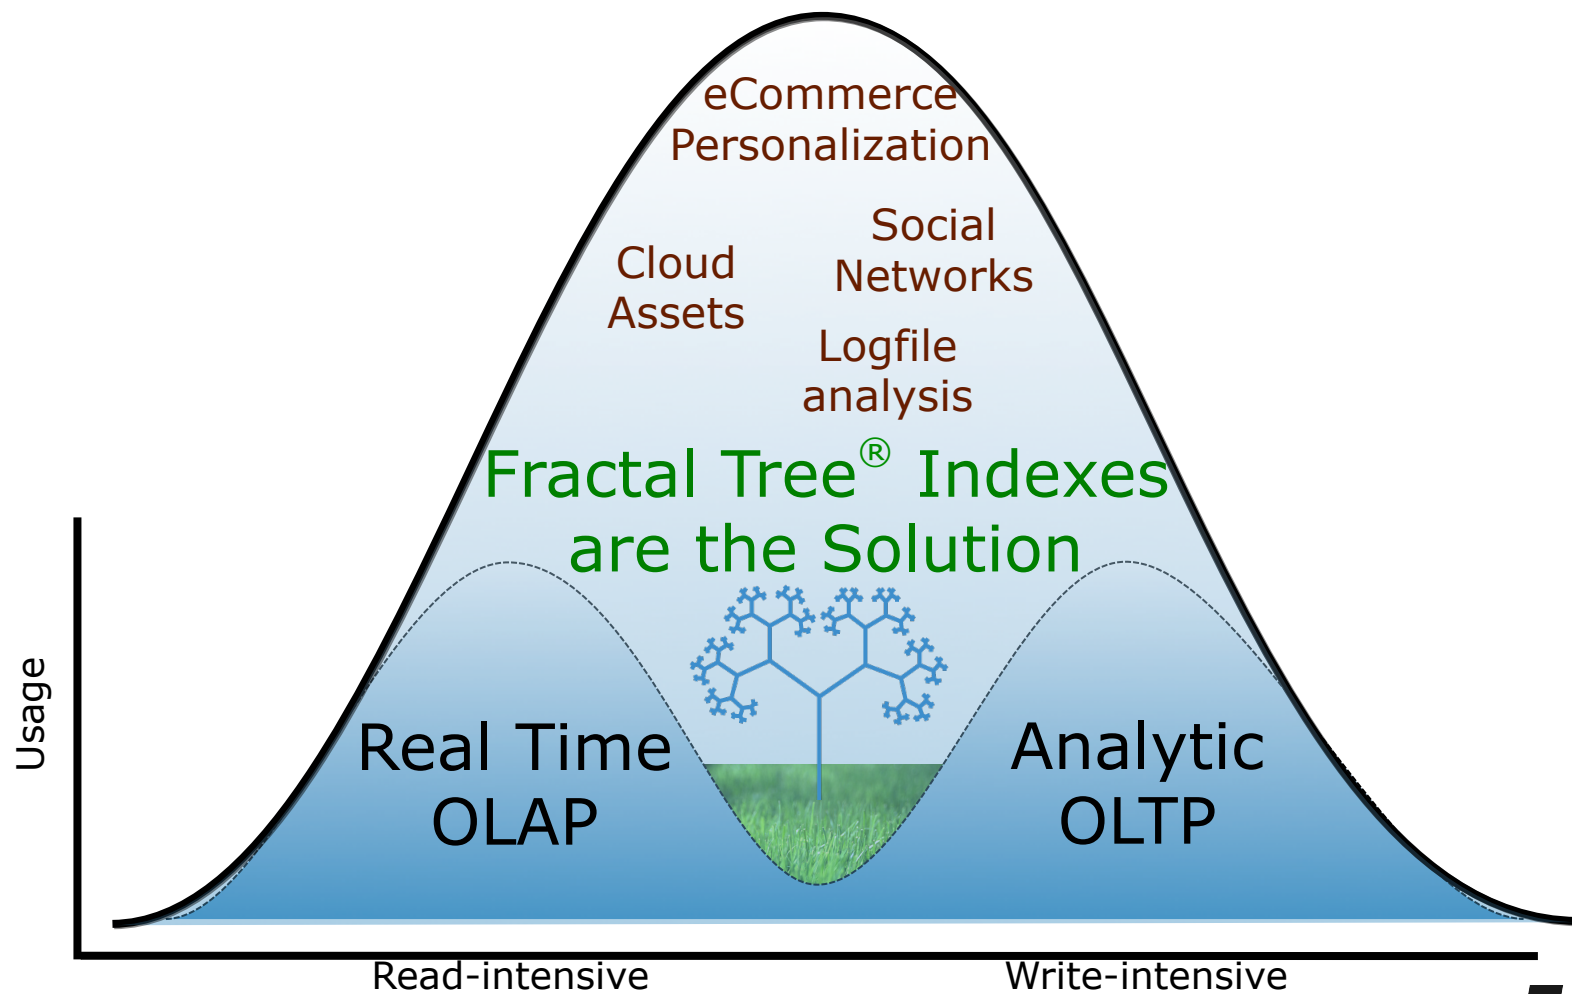


- **High Compression**
- **No Impact to Performance**

TokuDB Eliminates Death Valley



"While TokuDB® is effectively an operational database technology, it does blur the lines between operations and analytics"



Tokutek®