



## ***A Global In-memory Data System for MySQL***

***Daniel Austin, PayPal Technical Staff***

Percona Live! MySQL Conference  
Santa Clara, April 12th, 2012 v1.3

**Intro: Globalizing NDB**

Proposed Architecture

What We Learned

Q&A

# ***Mission YesSQL***

***"Develop a globally distributed DB For user-related data."***

- Must Not Fail (99.999%)
- Must Not Lose Data. Period.
- Must Support Transactions
- Must Support (some) SQL
- Must WriteRead 32-bit integer globally in 1000ms
- MinMax Data Volume: 10-100 TB (working)
- Must Scale Linearly with Costs

# ***THE FUNDAMENTAL PROBLEM IN DISTRIBUTED DATA SYSTEMS***

“How Do We Manage Reliable  
Distribution of Data Across Geographical  
Distances?”




## ***One Approach: The NoSQL Solution***

- NoSQL Systems provide a solution that relaxes many of the common constraints of typical RDBMS systems
  - Slow - RDBMS has not scaled with CPUs
  - Often require complex data management (SOX, SOR)
  - Costly to build and maintain, slow to change and adapt
  - Intolerant of CAP models (more on this later)
- Non-relational models, usually key-value
- May be batched or streaming
- Not necessarily distributed geographically

# ***Big Data Myth #1: Big Data = NoSQL***

- 'Big Data' Refers to a Common Set of Problems
  - Large Volumes
  - High Rates of Change
    - Of Data
    - Of Data Models
    - Of Data Presentation and Output
  - Often Require 'Fast Data' as well as 'Big'
    - Near-real Time Analytics
    - Mapping Complex Structures



Takeaway: Big Data is the problem, NoSQL is one (proposed) solution

## NoSQL Hype Cycle: Where Are We Now?

There are currently more than **120+** NoSQL databases listed at [nosql-database.org](http://nosql-database.org)!



As the pace of new technology solutions has slowed, some clear winners have emerged.

### ***3 Essential Kinds of NoSQL Systems***

1. Columnar K-V Systems
  - Hadoop, Hbase, Cassandra, PNUTs
2. Document-Based
  - MongoDB, TerraCotta
3. Graph-Based
  - FlockDB, Voldemort

Takeaway: These were originally designed as solutions to specific problems because no commercial solution would work.



Intro: Globalizing Big Data

**Proposed Architecture**

What We Learned

Q&A

# ***OUR APPROACH: THE YESSQL SOLUTION***

Use MySQL/NDB

Use AWS

Key Tradeoffs:

Cost vs. Performance

HA vs. CAP

Complexity vs. RDBMS features



# WHY NDB?

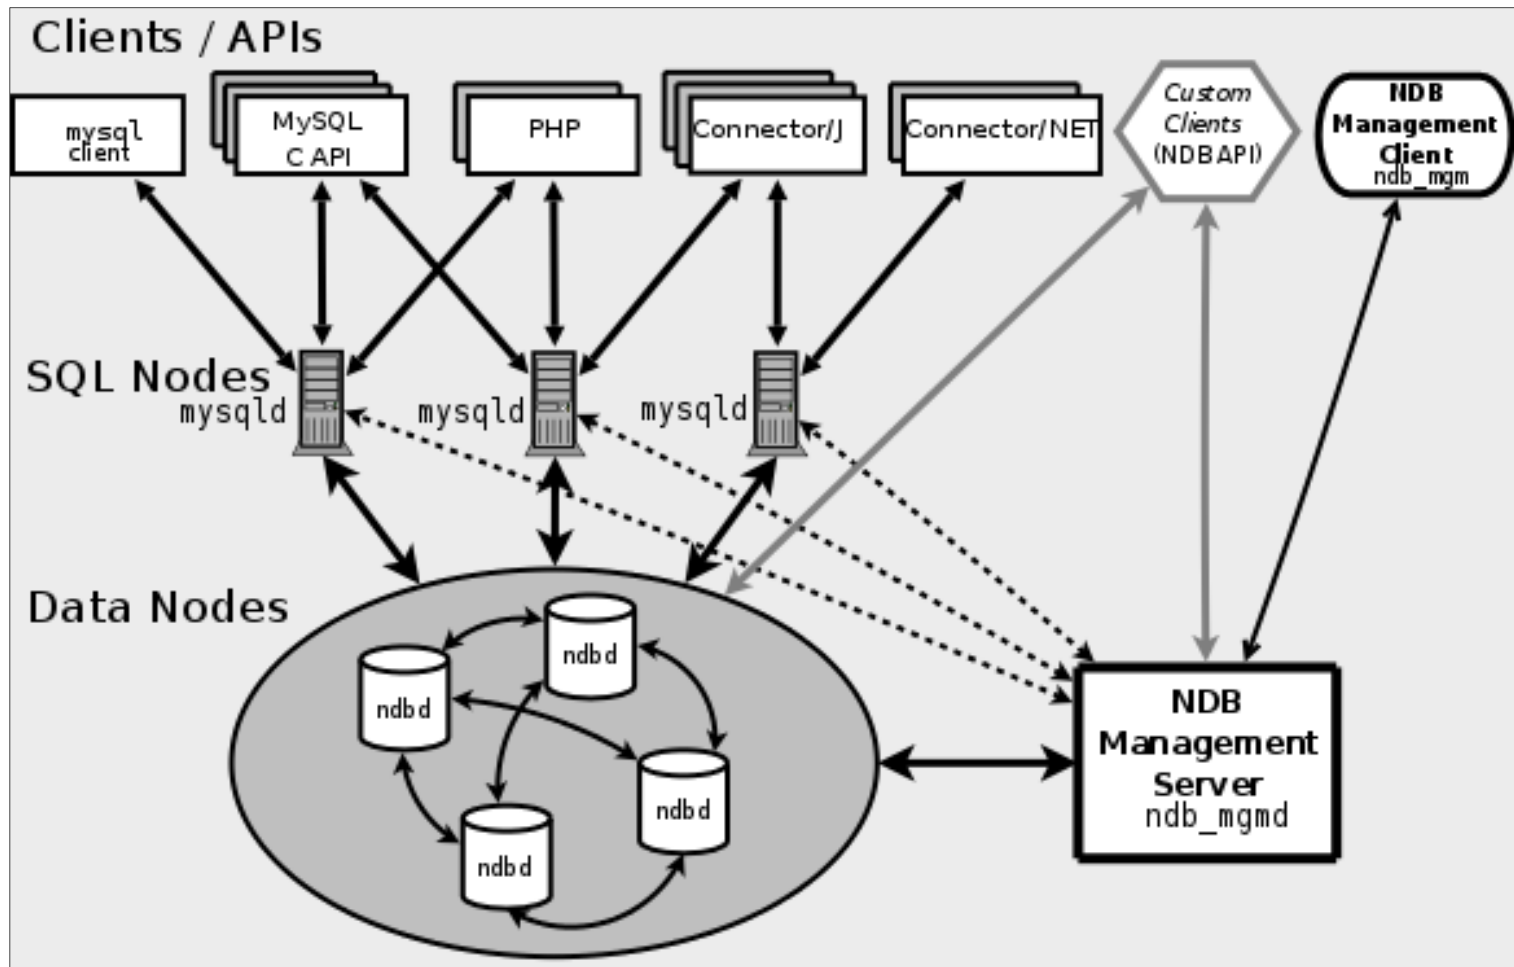
## Pro

- True HA by design
  - Fast recovery
- Supports (some) X-actions
- Relational Model
- In-memory architecture = high performance
- Disk storage for non-indexed data (since 5.1)
- APIs, APIs, APIs

## Con



- Some semantic limitations on fields
- Size constraints (2 TB?)
  - Hardware limits also
- Higher cost/byte
- Requires reasonable data partitioning
- Higher complexity

# How NDB Works in One Slide



Graphics courtesy dev.mysql.com

# CIRCULAR REPLICATION/FAILOVER



Graphics courtesy O'Reilly OnLamp.com


# SYSTEM AVAILABILITY DEFINED

- Availability of the entire system:

$$A_{sys} = 1 - \prod_{i=1}^n (1 - \prod_{j=1}^m r_i)_j$$

- Number of Parallel Components Needed to Achieve Availability  $A_{min}$ :

$$N_{min} = \lceil \ln(1 - A_{min}) / \ln(1 - r) \rceil$$



## ***Big Data Myth #2: The CAP Theorem Doesn't Say What You Think It Does***

- Consistency, Availability, (Network) Partition
  - Pick any two? Not really.
- The Real Story: These are not Independent Variables
- $AP = CP$  (Um, what? But...  $A \neq C$  )
- Variations:
  - PACELC (adds latency tolerance)
  - VPAC (adds variability)


Takeaway: the real story here is about the tradeoffs made by designers of different systems, and the main tradeoff is between consistency and availability, usually in favor of the latter.

# ***What about "High Performance"?***

- Maximum lightspeed distance on Earth's Surface: ~**67** ms
- Target: data available worldwide in  $< 1000$  ms

**Sound Easy?**

**Think Again!**





# ***AWS Meets NDB***

- Why AWS?
  - Cheap and easy infrastructure-in-a-box  
(Or so we thought! Ha!)
- Services Used:
  - EC2 (Centos 5.3, small instances for mgm & query nodes, XL for data)
  - Elastic IPs/ELB
  - EBS Volumes
  - S3
  - Cloudwatch

# ARCHITECTURAL TILES

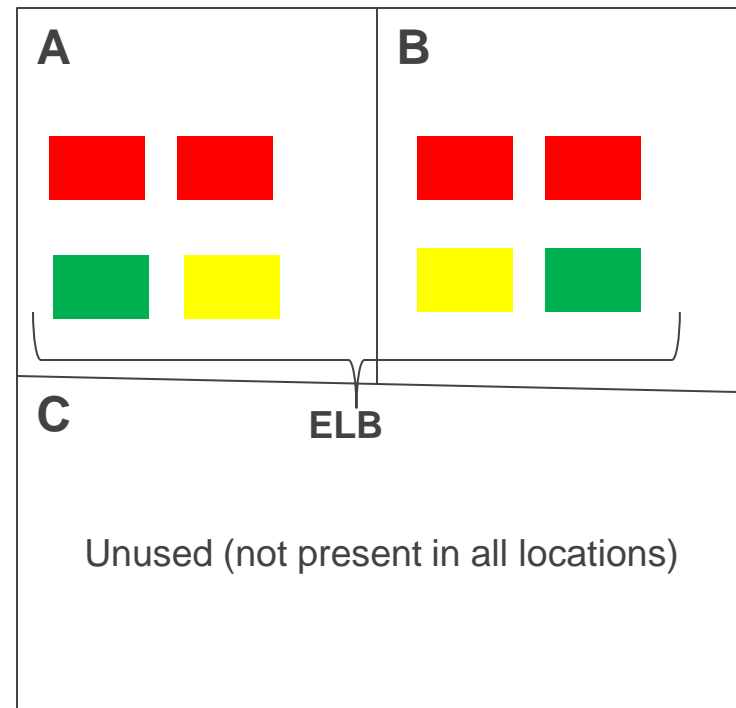
## Tiling Rules

- Never separate NDB & SQL
- Ndb:2-SQL:1-MGM:1
- Scale by adding more tiles
- Failover 1<sup>st</sup> to nearest AZ
- Then to nearest DC
- At least 1 replica/AZ
- Don't share nodes
- Mgmt nodes are redundant

## Limitations

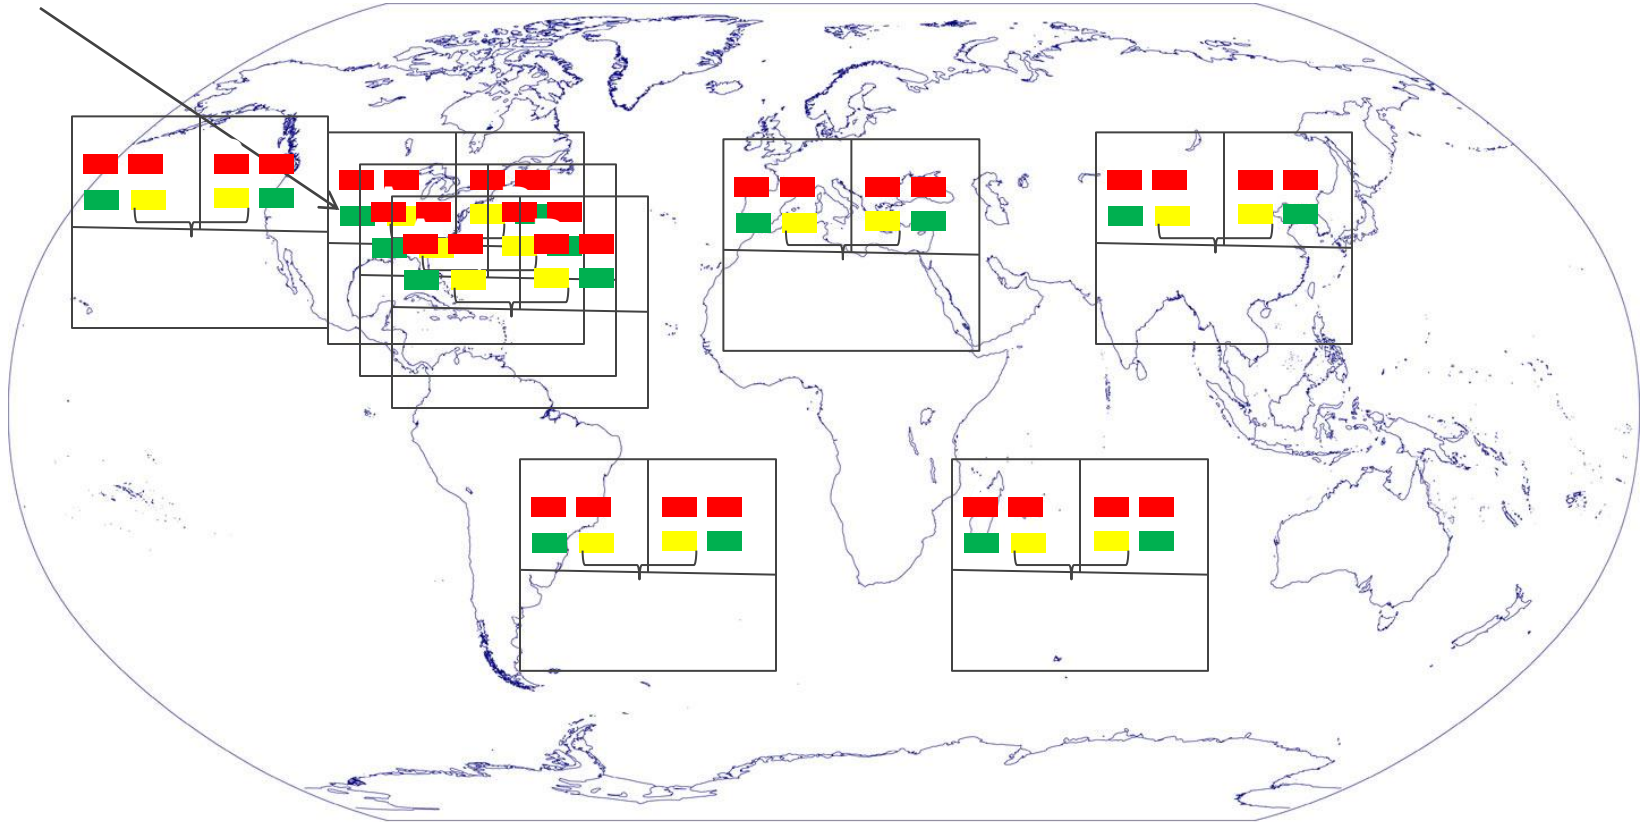
- AWS is network-bound @ 250 MBPS – ouch!
- Need specific ACL across AZ boundaries
- AZs not uniform!
- No GSLB
- Dynamic IPs
- ELB sticky sessions !reliable

## AWS Availability Zones



# Architecture Stack

Scale by Tiling



**7 AWS Data Centers:**  
US-E, US-W, US-W2, TK, EU, AS, SA-B

# ***Other Technologies Considered***

- Paxos
  - Elegant-but-complex consensus-based messaging protocol
  - Used in Google Megastore, Bing metadata
- Java Query Caching
  - Queries as serialized objects
  - Not yet working
- Multiple Ring Architectures
  - Even more complicated = no way

Intro: Globalizing Big Data

Proposed Architecture

**What We Learned**

Q&A

# SYSTEM READ/WRITE PERFORMANCE (!)


## What we tested:

- 32 & 256 byte char fields
- Reads, writes, query speed vs. volume
- Data replication speeds

## Results:

- Global replication < 350 ms
- 256 byte writeread < 1000ms worldwide

In-region replication tests



# ***The Commit Ordering Problem***

- Why does commit ordering matter?
- Write operators are non-commutative

$$[W(d,t1),W(d,t2)] \neq 0 \text{ unless } t1=t2$$

- Can lead to inconsistency
- Can lead to timestamp corruption
- Forcing sequential writes defeats Amdahl's rule
- Can show up in GSLB scenarios

# ***Dark Side of AWS***

- Deploying NDB at scale on AWS is hard
    - Dynamic IPs (use hostfile)
    - DNS issues
    - Security groups (ec2-authorize)
    - Inconsistent EC2 deployments
      - Are availability zones independent network segments?
    - No GSLB (!) (rent or buy)
- Be prepared to struggle a bit!



# ***Future Directions for YesSQL***

- Alternate solution using Pacemaker, Heartbeat
  - From Yves Trudeau @ Percona
  - Uses InnoDB, not NDB
- Implement Memcached plugin
  - To test NoSQL functionality, APIs
- Add simple connection-based persistence to preserve connections during failover
- Can we increase the number of data nodes?
- As the number of datacenters goes up, sync times go down. What's the relation?
- So far only limited testing of hard configurations
  - Production systems will likely need more RAM for MGMs



# ***Hard Lessons, Shared***

- Be Careful...
  - With “Eventual Consistency”-related concepts
  - ACID, CAP are not really as well-defined as we’d like considering how often we invoke them
- NDB is a good solution
  - Real HA, real performance, real SQL
  - Notable limitations around fields, datatypes
  - Successfully competes with NoSQL systems for most use cases – better in many cases
- AWS makes the hard things possible
  - But not so much the easy things easier

**Takeaway: You can achieve high performance and availability without giving up relational models and read consistency!**



***"In the long run, we are all ~~dead~~  
eventually consistent."***

***Maynard Keynes on NoSQL Databases***

Twitter: @daniel\_b\_austin  
Email: daaustin@paypal.com

*With apologies and thanks to the real DB experts, Andrew Goodman, Yves  
Trudeau, Clement Frazer, Daniel Abadi, and everyone else!*