




# Improving MySQL/MariaDB query performance through optimizer tuning

Sergey Petrunya  
Timour Katchaounov  
{psergey,timour}@montyprogram.com

# What is an optimizer and why do I need to tune it

# Why does one need optimizer tuning

- Database performance is affected by many factors



- One of them is the **query optimizer** (and query executor, together **query processor**)

# What is a query optimizer ?

- Converts SQL into execution instructions
- Like GPS gives driving directions from A to B



id	select_type	table	type	possible_keys	key	key_len
1	PRIMARY	part	ALL	PRIMARY	NULL	NULL
1	PRIMARY	partsupp	ref	PRIMARY,i_ps_partkey,i_ps_suppkey	i_ps_partkey	4
1	PRIMARY	supplier	eq_ref	PRIMARY,i_s_nationkey	PRIMARY	4
1	PRIMARY	nation	eq_ref	PRIMARY,i_n_regionkey	PRIMARY	4
1	PRIMARY	region	ALL	PRIMARY	NULL	NULL
2	DEPENDENT SUBQUERY	partsupp	ref	PRIMARY,i_ps_partkey,i_ps_suppkey	i_ps_partkey	4
2	DEPENDENT SUBQUERY	supplier	eq_ref	PRIMARY,i_s_nationkey	PRIMARY	4
2	DEPENDENT SUBQUERY	nation	eq_ref	PRIMARY,i_n_regionkey	PRIMARY	4
2	DEPENDENT SUBQUERY	region	eq_ref	PRIMARY	PRIMARY	4

# Can the optimizer work poorly?

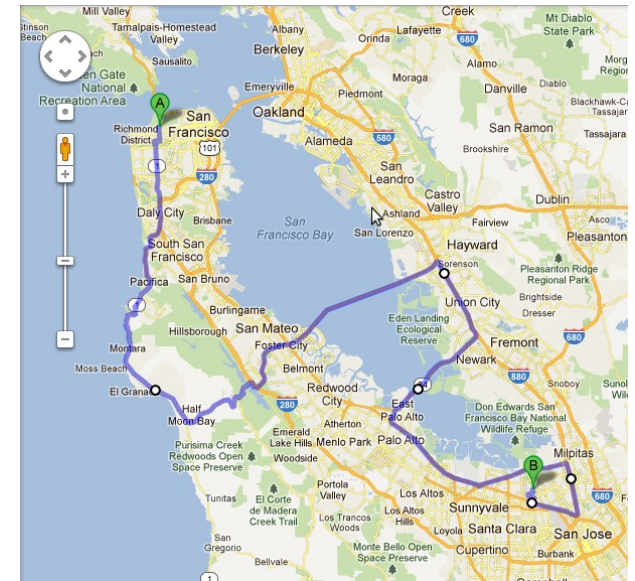
Yes, just like your GPS

It can produce a sub-optimal query plan (compare to GPS's directions)

Unneeded CPU/IO usage in the database is the same extra mileage

How can one tell if they have a problem in query optimizer?

- Like with the GPS:
  - Knowledge of possible routes
  - Common sense



This is what this tutorial is about



# Start from a database performance problem

- Database response time is too high
  - or the server denies connections altogether
    - because running queries have occupied all @@max\_connections threads.
- Monitoring (cacti, etc) shows high CPU/disk utilization
- Is this an optimizer problem?
  - Maybe.

# Signs that it's caused by the optimizer

- Only certain kinds of queries are affected
  - Unlike concurrency/disk/binlog problems that cause slowdown for all queries
- These queries can be caught with
  - Slow query log
    - pt-query-digest
  - SHOW PROCESSLIST
  - SHOW PROFILE

# Optimizer problems in Slow query log

- Percona server/MariaDB:  
--log\_slow\_verbosity=query\_plan

# Thread\_id: 1 Schema: dbt3sf10 QC\_hit: No

# Query\_time: 2.452373 Lock\_time: 0.000113 Rows\_sent: 0 Rows\_examined: 1500000

# Full\_scan: Yes Full\_join: No Tmp\_table: No Tmp\_table\_on\_disk: No

# Filesort: No Filesort\_on\_disk: No Merge\_passes: 0

SET timestamp=1333385770;

select \* from customer where c\_acctbal < -1000;

- Look for big Query\_time (>> Lock\_time)
- Look for Rows\_examined >> Rows\_sent
  - Although there are “legitimate” cases for this
- Query\_plan members (violet), if present, also show  
\*possible\* problem areas



# SHOW PROCESSLIST;

```
MariaDB [dbt3sf10]> show processlist;
```

Id	User	Host	db	Command	Time	State	Info
1	root	localhost	dbt3sf10	Query	176	Sending data	select
2	root	localhost	dbt3sf10	Query	0	NULL	show pr
3	root	localhost	dbt3sf10	Query	19	freeing items	insert
4	root	localhost	dbt3sf10	Sleep	77		NULL
5	root	localhost	dbt3sf10	Query	4	Updating	update
6	root	localhost	dbt3sf10	Query	5	statistics	select

- States to look for:
  - optimizing
  - executing
  - Statistics
  - Sending data

# SHOW PROFILE

- Uses the same set of states
- Shows where exactly the query has spent time

```
mysql> show profile for query 1;
```

+-----+-----+	
Status	Duration
+-----+-----+	
starting	0.000097
Opening tables	0.000038
System lock	0.000014
Table lock	0.000019
init	0.000027
optimizing	0.000029
statistics	0.000036
preparing	0.000023
executing	0.000014
Sending data	18.501940
end	0.000015
query end	0.000006
freeing items	0.000029
logging slow query	0.000004
logging slow query	0.000050
cleaning up	0.000006
+-----+-----+	

## Phases that can occupy lots of time:

Storage engine: transaction start, etc

Optimization (search for query plan)

Optimization and sometimes InnoDB locking

Execution (of the query plan)

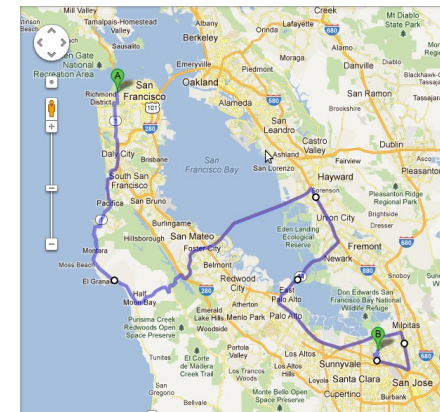
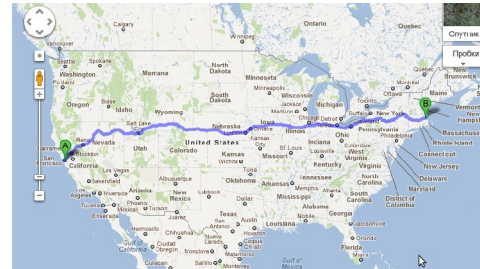
Storage engine: transaction commit,  
binary logging

# Global observations summary

- A query [pattern] that
  - Spends a lot of time in optimization/execution phase
  - Examines many more records than you think it should
  - Uses Full\_scan, Full\_join, etc
  - .. or all of the above
- Is a candidate for further investigation

# Reasons why a query is “heavy”

- Some queries are inherently hard to compute
- Some get a plan which looks good but turns out bad due to circumstances that were not accounted for
- And some are just bad query plan choices



# Know thy optimizer

- How to tell which case you're dealing with?

- Need to know “the distance”

- The lower bound of amount of effort one needs to expend to arrive at the destination

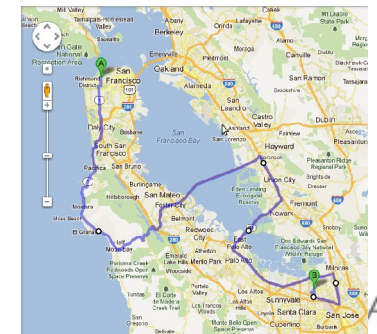
- Easy in geography, not so easy in SQL



- Need to know “the roads”

- Available table access strategies

- How they are composed into query plans





# Optimizer troubleshooting

From simple to complex queries



# Let's take a simple, slow query

```
select * from orders
where o_orderDate='1992-06-06' and o_clerk='Clerk#000009506';
```

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	orders	ALL	NULL	NULL	NULL	NULL	15084733	Using where

# Thread\_id: 2 Schema: dbt3sf10 QC\_hit: No

# Query\_time: 41.129178 Lock\_time: 0.000174 Rows\_sent: 2 Rows\_examined: 15000000

# Full\_scan: Yes Full\_join: No Tmp\_table: No Tmp\_table\_on\_disk: No


# Filesort: No Filesort\_on\_disk: No Merge\_passes: 0

use dbt3sf10;

SET timestamp=1333432984;

select \* from orders where o\_orderDate='1992-06-06' and o\_clerk='Clerk#000009506';

- Full scan examined lots of rows
- But only a few (two) of them “have contributed to the result set”
  - It was unnecessary to read the rest
- Query plan is poor, it seems.



# Status variables

```
MariaDB [dbt3sf10]> flush status;  
Query OK, 0 rows affected (0.00 sec)
```

```
MariaDB [dbt3sf10]> select * from orders where o_orderDate='1992-06-06' and o_clerk='Clerk#000009506';  
...
```

```
MariaDB [dbt3sf10]> show status like 'Handler%';
```

Variable_name	Value
Handler_commit	1
Handler_delete	0
Handler_discover	0
Handler_icp_attempts	0
Handler_icp_match	0
Handler_mrr_init	0
Handler_mrr_key_refills	0
Handler_mrr_rowid_refills	0
Handler_prepare	0
Handler_read_first	0
Handler_read_key	0
Handler_read_next	0
Handler_read_prev	0
Handler_read_rnd	0
Handler_read_rnd_deleted	0
Handler_read_rnd_next	15000001
Handler_rollback	0
Handler_savepoint	0
Handler_savepoint_rollback	0
Handler_tmp_update	0
Handler_tmp_write	0
Handler_update	0
Handler_write	0

Indeed, 15M sequential reads

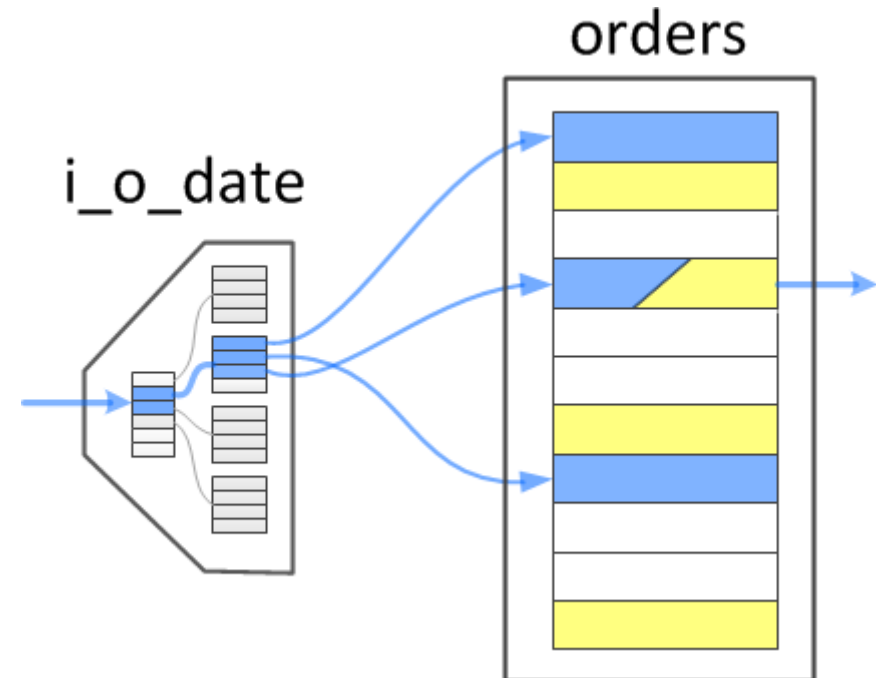
```
23 rows in set (0.01 sec)
```

# With index, we'll read fewer rows

```
alter table orders add key i_o_orderdate (o_orderdate);  
select * from orders  
where o_orderDate='1992-06-06' and o_clerk='Clerk#000009506';
```

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	orders	ref	i_o_orderdate	i_o_orderdate	4	const	6303	Using where

- The index allows to only read records with i\_o\_orderdate='Clerk..'
  - (At the cost of having to read index entries also)
- Still not perfect: we're reading records that have the wrong o\_orderDate
- Let's check how it executes:



# Status variables

```
MariaDB [dbt3sf10]> flush status;
```

```
MariaDB [dbt3sf10]> select * from orders where o_orderDate='1992-06-06' and o_clerk='Clerk#000009506';  
... (0.39 sec)
```

```
MariaDB [dbt3sf10]> show status like 'Handler%';
```

Variable_name	Value
Handler_commit	1
Handler_delete	0
Handler_discover	0
Handler_icp_attempts	0
Handler_icp_match	0
Handler_mrr_init	0
Handler_mrr_key_refills	0
Handler_mrr_rowid_refills	0
Handler_prepare	0
Handler_read_first	0
Handler_read_key	1
Handler_read_next	6122
Handler_read_prev	0
Handler_read_rnd	0
Handler_read_rnd_deleted	0
Handler_read_rnd_next	0
Handler_rollback	0
Handler_savepoint	0
Handler_savepoint_rollback	0
Handler_tmp_update	0
Handler_tmp_write	0
Handler_update	0
Handler_write	0

1 index lookup  
6K forward index reads


```
23 rows in set (0.00 sec)
```

# Can have an even better index

```
alter table orders add key i_o_date_clerk (o_orderdate, o_clerk);  
select * from orders  
where o_orderDate='1992-06-06' and o_clerk='Clerk#000009506';
```

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	orders	ref	i_o_orderdate,i_o_date_clerk	i_o_date_clerk	20	const,const	1	Using index condition

- Added a wide index covering both columns
- Now, only records that match the whole WHERE will be read
  - Perfect!



# Status variables

```
MariaDB [dbt3sf10]> select * from orders where o_orderDate='1992-06-06' and o_clerk='Clerk#000009506';  
... (0.00 sec)
```

```
MariaDB [dbt3sf10]> show status like 'Handler%';
```

Variable_name	Value
Handler_commit	1
Handler_delete	0
Handler_discover	0
Handler_icp_attempts	0
Handler_icp_match	0
Handler_mrr_init	0
Handler_mrr_key_refills	0
Handler_mrr_rowid_refills	0
Handler_prepare	0
Handler_read_first	0
Handler_read_key	1
Handler_read_next	0
Handler_read_prev	0
Handler_read_rnd	0
Handler_read_rnd_deleted	0
Handler_read_rnd_next	0
Handler_rollback	0
Handler_savepoint	0
Handler_savepoint_rollback	0
Handler_tmp_update	0
Handler_tmp_write	0
Handler_update	0
Handler_write	0

```
23 rows in set (0.00 sec)
```

1 index lookup, and that's it


- Let's try to generalize..



# Index fundamentals

# ALTER TABLE tbl ADD INDEX(column1, column2, column3)

## On the physical level, it's a B-TREE



On a logical level: an ordered directory

[illegible]

# Querying data through index: 'ref'

alter table orders add index(o\_orderDATE, o\_shippriority, o\_custkey)

- **ref** access can use any prefix:

o\_orderDATE='2011-11-02'

o\_orderDATE='2011-11-02' AND  
o\_shippriority=2

o\_orderDATE='2011-11-08' AND  
o\_shippriority=3 AND o\_custkey= 4

o_orderDATE	o_shipprio	o_custkey
2011-11-01	1	10
2011-11-02	1	10
2011-11-02	1	10
2011-11-02	1	11
2011-11-02	2	10
2011-11-03	2	10
2011-11-03	2	11
2011-11-04	3	9
2011-11-04	4	10
2011-11-05	3	11
2011-11-06	1	10
2011-11-08	2	4
2011-11-09	1	10

- Not allowed:

o\_shipprioirity=2

o\_shipprioirity=2

# Ref(const) access summary

- Uses equalities to make index lookups
- #rows estimate is usually precise
- ANALYZE will not help

EXPLAIN: type=ref, ref=const, key\_len

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	orders	ref	i_o_orderdate	i_o_orderdate	4	const	6303	Using where


- Status variables:
  - 1 x Handler\_read\_key
  - N x Handler\_read\_next
- Estimates and control.. will talk later

# Range access


# Range access

- Can use equality and non-equality comparisons
- Think of index as a directed axis.
- and “ranges” on it where the WHERE condition is TRUE:

... WHERE o\_orderDATE > '2010-01-01'




... WHERE o\_orderDATE > '2010-01-01' AND o\_orderDATE < '2010-01-07'



# Range access (2)

- Arbitrarily deep AND/OR formulas are supported
- The optimizer will split out the “useful” parts of WHERE
- And produce a list of disjoint ranges to be scanned:



- “Useful” conditions compare key with constant:

`tbl.key > const`

`tbl.key = const`

`tbl.key BETWEEN c1 AND c2`

`tbl.key LIKE 'foo%'`


`tbl.key IS NULL`

`tbl.key IN (const1, const2 ...)`



# Range access for multi-part keys

- Much more complex
- Ranges over (keypart1, keypart2, ...) tuple space
  - Think of it as axis with sub-axis at each point:

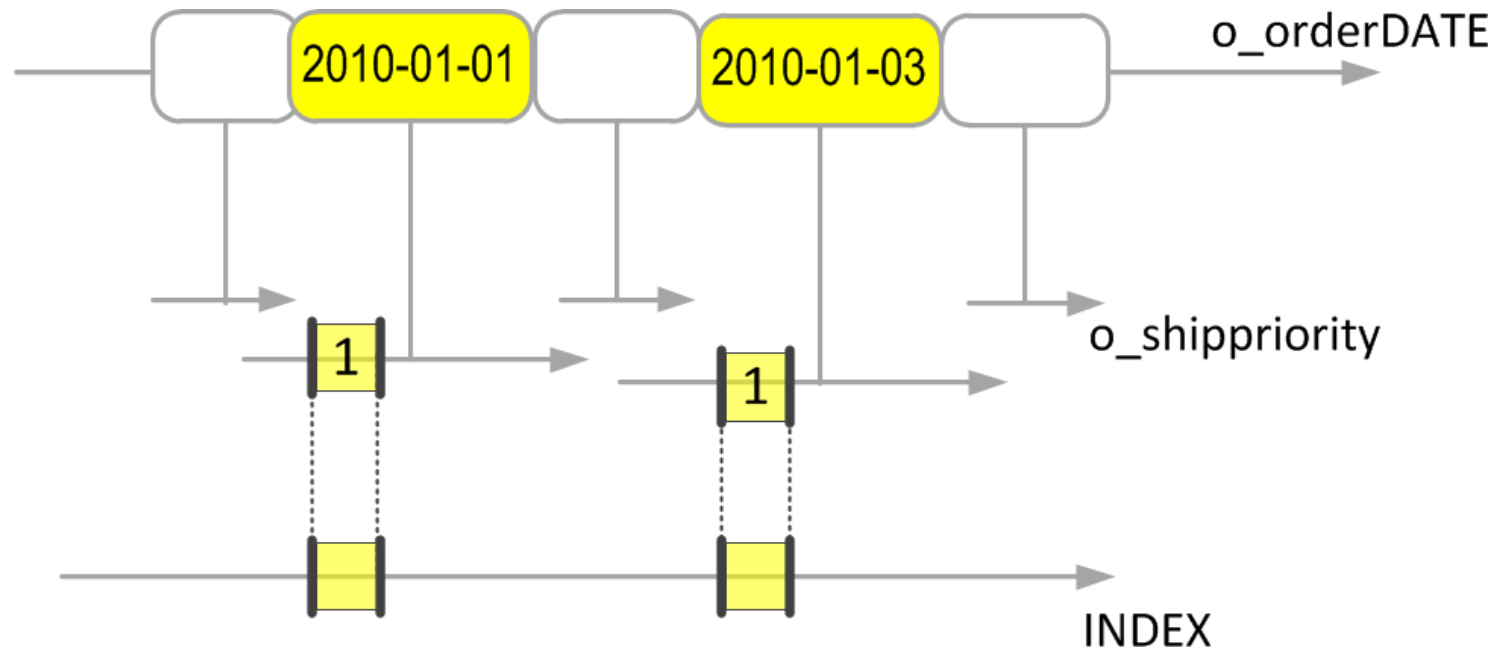


- Range building rules:
  - Number of ranges is a function of WHERE, not of data in the table
  - Do not use knowledge about domain density (e.g. no knowledge like “there is no integers between 1 and 2”)

# Multi-key-part range #1

- Equality on 1<sup>st</sup> key part? Can put condition on 2<sup>nd</sup> key part into sub-axis

```
alter table orders add key i_o_date_clerk (o_orderdate, o_clerk);  
select * from orders  
where o_orderDATE IN ('2010-01-01', '2010-01-03') AND o_shippriority=1
```



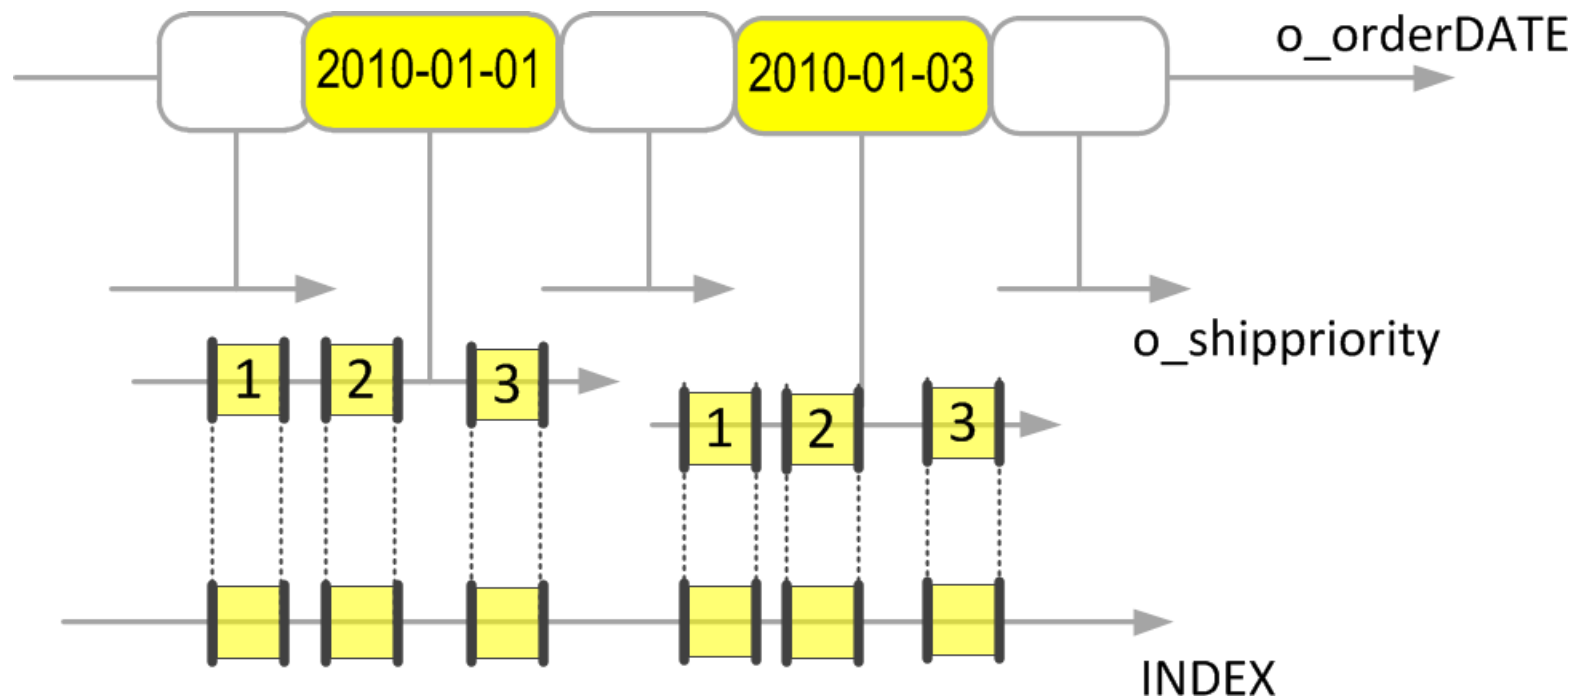
# Multi-key-part range #2

- This can cause big number of ranges. That's ok

select \* from orders

where

o\_orderDATE IN ('2010-01-01', '2010-01-03') AND  
(o\_shippriority =1 OR o\_shippriority=2 OR o\_shippriority=3)

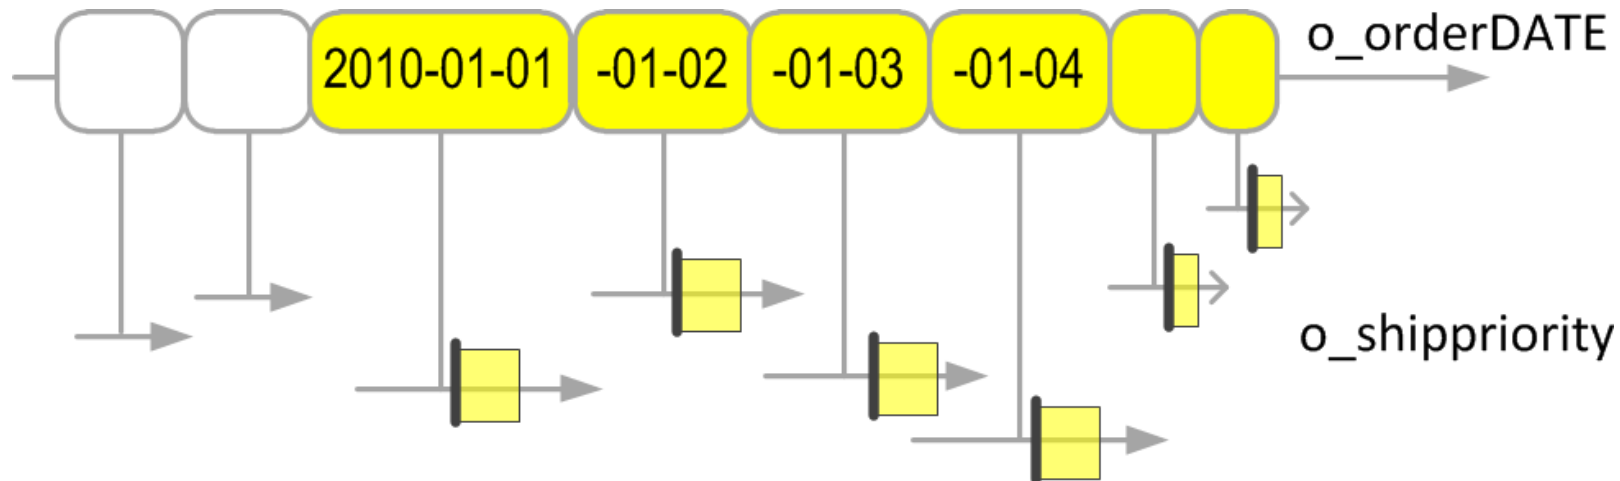


# Multi-key-part range #3

- Non-equality on 1<sup>st</sup> keypart:

select \* from orders

where o\_orderDATE >= '2010-01-01' AND o\_shippriority>3



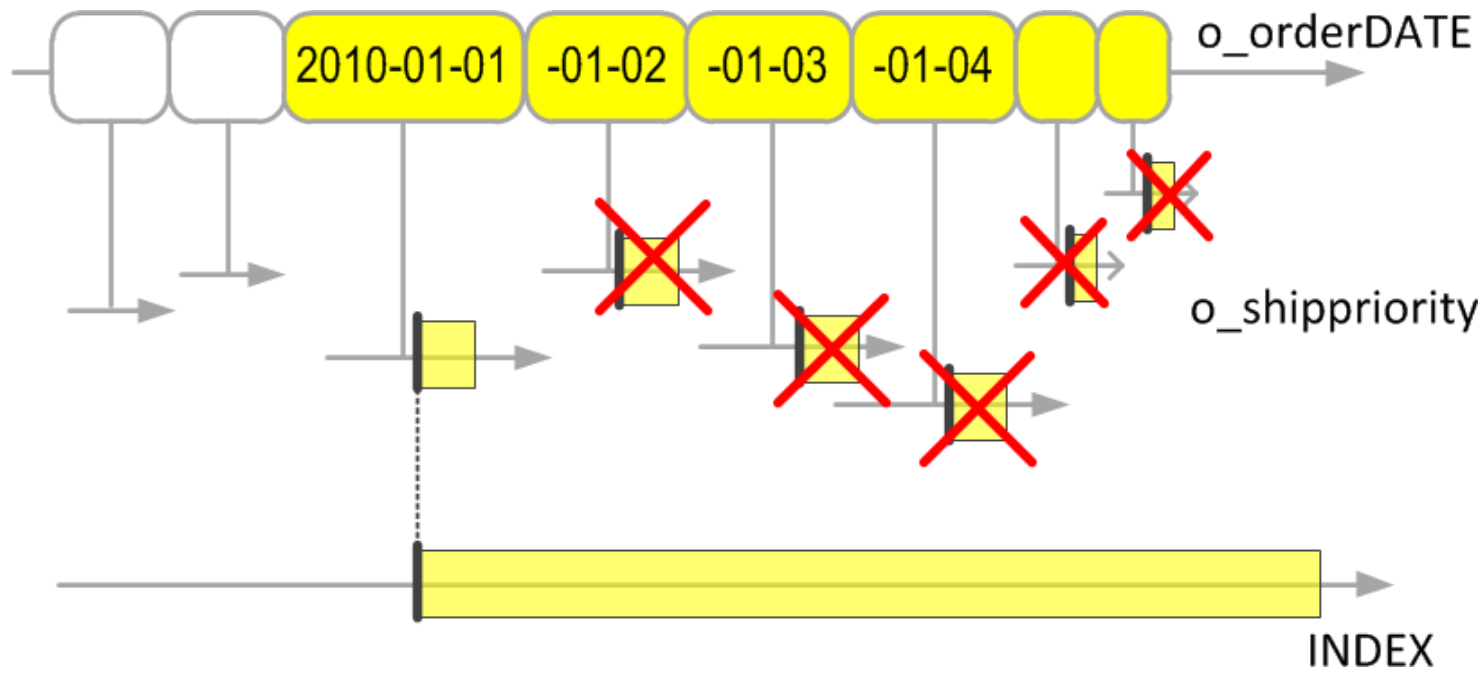
# Multi-key-part range #3

- Non-equality on 1<sup>st</sup> keypart:


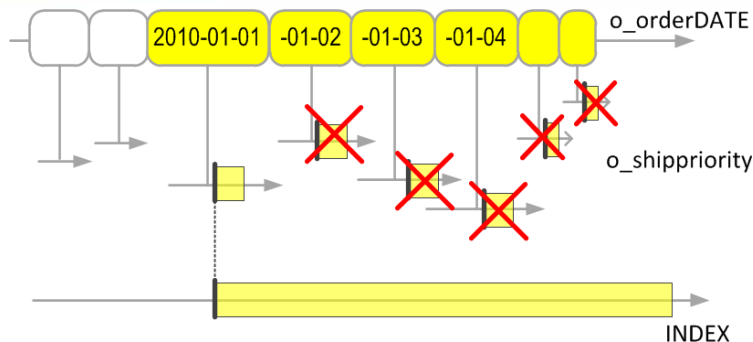
select \* from orders

where o\_orderDATE >= '2010-01-01' AND o\_shippriority>3

- Can attach the second only to the edge “cell”:



# Multi-key-part range summary



- keyparts for which queries have equalities should go first
- multiple keyparts with non-equality ranges will not reduce the number of rows read

e.g. `t.keypart1 BETWEEN ... AND t.keypart2 BETWEEN ...`  
won't take advantage of condition on keypart2

- When keypartX is followed by other keyparts, then  
`keypartX IN (1,... N)` will work better than `keypartX BETWEEN 1 AND N`



# Can I look at the ranges it has constructed?

- Normally, no. But there is a last-resort method:
  - Get a really small subset of your data onto separate server
  - Get a debug build there. Start it with `-debug`
  - Run your query with `FORCE INDEX`:  
`explain select * from orders force index(i_o_datepriokey)`  
`where (o_orderDATE IN ('2010-02-01','2010-02-03') AND o_shippriority < 3) OR`  
`(o_orderDATE >= '2010-03-10' AND o_shippriority < 4)`
- Go into `mysqld.trace`
  - Grep for “`query:`” until you find the query
  - Grep for “`>print_quick`”

# Ranges in mysqld.trace

## mysqld.trace

```
T@10 : | | query: explain select * from orders force index(i_o_datepriokey) where o_orderDATE
IN ('2010-02-01','2010-02-03') AND o_shippriority < 3 OR (o_orderDATE > '2010-03-10' AND o
_shippriority < 4)
T@10 : | | >PROFILING::set_query_source
T@10 : | | <PROFILING::set_query_source
T@10 : | | >mysql_parse
...
T@10 : | | | | | | | | | | >print_quick
quick range select, key i_o_datepriokey, length: 9
  2010-02-01/NULL < X < 2010-02-01/3
  2010-02-03/NULL < X < 2010-02-03/3
  2010-03-10 < X
other_keys: 0x0:
```

- EXPLAIN only shows key\_len which is max # keyparts used

id	select_type	table	type	possible_keys	key	key_len	ref	rows	E
1	SIMPLE	orders	range	i_o_datepriokey	i_o_datepriokey	9	NULL	3	U

# Range access estimates

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	orders	range	i_o_datepriokey	i_o_datepriokey	9	NULL	3	Using where

- Are obtained by doing dives into the index
  - Dive at the start of the range
  - Dive at the end
  - Estimate the difference
- Usually are precise (not more than 2x wrong)
- Are not affected/do not need ANALYZE TABLE
- Cost you disk IO!
  - Indexes that are never worth using will still be probed!
  - Check your buffer pool + EXPLAINs to see if this happens!
    - Quick fix: IGNORE INDEX (unusable\_index)

# Range access: Status variables

Variable_name	Value
Handler_commit	1
Handler_delete	0
Handler_discover	0
Handler_icp_attempts	9
Handler_icp_match	9
Handler_mrr_init	0
Handler_mrr_key_refills	0
Handler_mrr_rowid_refills	0
Handler_prepare	0
Handler_read_first	0
Handler_read_key	9
Handler_read_next	9
Handler_read_prev	0
Handler_read_rnd	0
Handler_read_rnd_deleted	0
Handler_read_rnd_next	0
Handler_rollback	0
Handler_savepoint	0
Handler_savepoint_rollback	0
Handler_tmp_update	0
Handler_tmp_write	0
Handler_update	0
Handler_write	0

23 rows in set (0.05 sec)


## Increments

- index lookup counter
- walk-forward-in-index counter

# Index-only scans

# Index-only scans

- Besides lookups, indexes have another use
- If the index has all columns you need, it is possible to avoid accessing table rows altogether
- This is called “index only” scan
- Can be applied to **ref**, and **range** accesses.





# Index-only read example

```
explain select sum(o_totalprice)
from orders
where o_orderdate between '1995-01-01' and '1995-01-07';
```

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	orders	range	i_o_orderdate	i_o_orderdate	4	NULL	4358	Using index condition

```
alter table orders add key o_date_price (o_orderDATE, o_totalprice);
# re-run the EXPLAIN:
```

```
alter table orders add key o_date_price (o_orderDATE, o_totalprice);
explain select sum(o_totalprice) from orders where o_orderdate between '1995-01-01' and '1995-02-01';
```

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	orders	range	...	o_date_price	4	NULL	39424	Using where; Using index

- scale=1, in-mem: from 0.2 to 0.1 sec

# Index-only reads and status variables

- Surprisingly, no difference

```
select sum(o_totalprice)
from orders use index(o_date_price)
where o_orderdate between '1995-01-01' and
                        '1995-02-01';
```

Variable_name	Value
Handler_commit	1
Handler_delete	0
Handler_discover	0
Handler_icp_attempts	0
Handler_icp_match	0
Handler_mrr_init	0
Handler_mrr_key_refills	0
Handler_mrr_rowid_refills	0
Handler_prepare	0
Handler_read_first	0
Handler_read_key	1
Handler_read_next	20137
Handler_read_prev	0
Handler_read_rnd	0
Handler_read_rnd_deleted	0
Handler_read_rnd_next	0
Handler_rollback	0
Handler_savepoint	0
Handler_savepoint_rollback	0
Handler_tmp_update	0
Handler_tmp_write	0
Handler_update	0
Handler_write	0

```
select sum(o_totalprice)
from orders use index(i_o_orderdate)
where o_orderdate between '1995-01-01' and
                        '1995-02-01';
```

Variable_name	Value
Handler_commit	1
Handler_delete	0
Handler_discover	0
Handler_icp_attempts	20137
Handler_icp_match	20137
Handler_mrr_init	0
Handler_mrr_key_refills	0
Handler_mrr_rowid_refills	0
Handler_prepare	0
Handler_read_first	0
Handler_read_key	1
Handler_read_next	20137
Handler_read_prev	0
Handler_read_rnd	0
Handler_read_rnd_deleted	0
Handler_read_rnd_next	0
Handler_rollback	0
Handler_savepoint	0
Handler_savepoint_rollback	0
Handler_tmp_update	0
Handler_tmp_write	0
Handler_update	0
Handler_write	0

# Index-only reads and status variables

- Let's look at InnoDB “lower-level” counters

```
show status like 'Innodb%';
```

Variable_name	Value	
...		
Innodb_buffer_pool_read_requests	413168	+68929
Innodb_buffer_pool_reads	10034	
Innodb_buffer_pool_wait_free	0	
...		
Innodb_rows_inserted	0	
Innodb_rows_read	181237	+20137
Innodb_rows_updated	0	

```
show status like 'Innodb%';
```

Variable_name	Value	
...		
Innodb_buffer_pool_read_requests	422426	+2566
Innodb_buffer_pool_reads	10034	
Innodb_buffer_pool_wait_free	0	
...		
Innodb_rows_inserted	0	
Innodb_rows_read	241651	+20138
Innodb_rows_updated	0	

- These are global non FLUSHable counters, so we manually calculate increment and show it on the right
- \_rows\_read are the same, too
- \_buffer\_pool\_reads are different
- A bit sense-less counting
  - Have plans to fix in MariaDB.

# Index-only reads summary

- Can be used with any index-based access methods
- Allows to skip accessing table records
  - Some savings for CPU-bound loads
  - Huge savings for IO-bound loads
  - Extra columns in the index a cost, though
    - Index is bigger
    - [Potentially] more frequently updated
- EXPLAIN shows “Using index”
- Counting in for Handler\_xxx and Innodb\_row\_reads status variables
  - Watch lower-level counters.

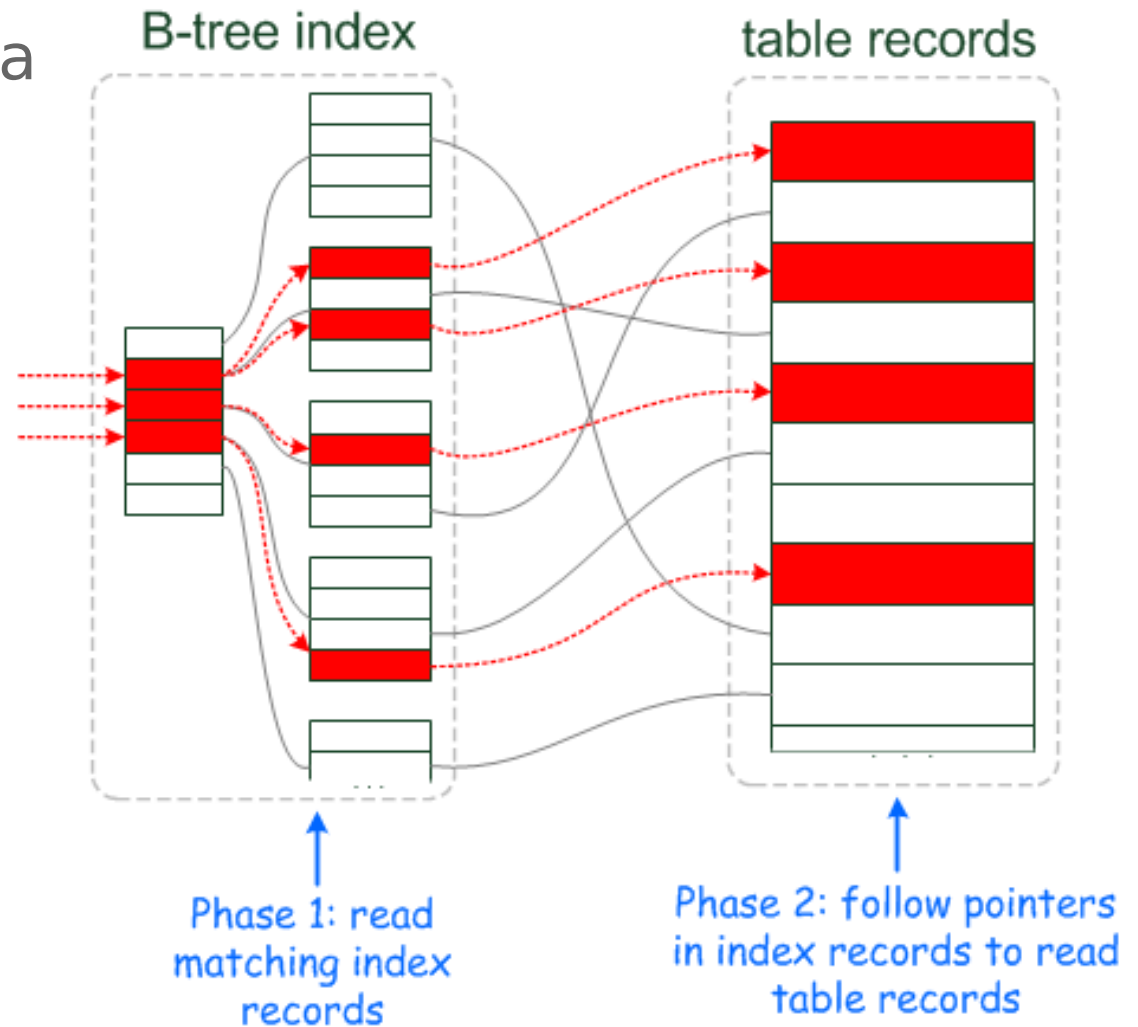
# Index Condition Pushdown

MySQL 5.6+, MariaDB 5.3+

# Index Condition Pushdown idea

A non-index only read is a two-step process:

1. Read index
2. Read record
3. Check the WHERE condition.






# Index Condition Pushdown idea

## Index Condition Pushdown

1. Read index
2. Check condition on index columns
3. Read record
4. Check the WHERE condition



# Index condition pushdown example

- Start without ICP

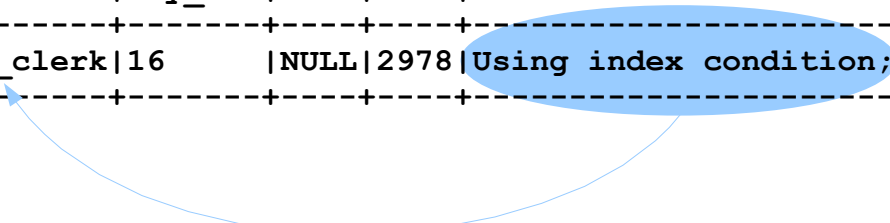
```
alter table orders add key i_o_clerk(o_clerk);
select o_clerk, o_orderkey, o_shippriority
from orders
where o_clerk IN ('Clerk#000000001', 'Clerk#00000201', 'Clerk#000003001') and
      dayofweek(o_orderDATE)=1;
```

# MariaDB 5.2.x/ MySQL 5.5.x

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	orders	range	i_o_clerk	i_o_clerk	16	NULL	2978	Using where

# MariaDB 5.3+, MySQL 5.6+

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	orders	range	i_o_clerk	i_o_clerk	16	NULL	2978	Using index condition; Using where



# Counters

```
alter table orders add key i_o_clerk(o_clerk);
select o_clerk, o_orderkey, o_shippriority
from orders
where o_clerk IN ('Clerk#000000001', 'Clerk#00000201', 'Clerk#000003001') and
      dayofweek(o_orderDATE)=1;
```

# MariaDB 5.3+, MySQL 5.6+

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	orders	range	i_o_clerk	i_o_clerk	16	NULL	2978	Using index condition; Using where

Variable_name	Value
Handler_commit	1
Handler_delete	0
Handler_discover	0
Handler_icp_attempts	2979
Handler_icp_match	2979
Handler_mrr_init	0
Handler_mrr_key_refills	0
Handler_mrr_rowid_refills	0
Handler_prepare	0
Handler_read_first	0
Handler_read_key	3
Handler_read_next	2979
Handler_read_prev	0
Handler_read_rnd	0

...

MariaDB [(none)]> show status like 'Innodb%';

Variable_name	Value
Innodb_buffer_pool_read_requests	32103 + 24416
Innodb_rows_read	3406 + 2979

# Index condition pushdown example

- Now, add an index that ICP could use:

```
alter table orders add key i_o_clerk_date(o_clerk, o_orderDATE);
```

```
select o_clerk, o_orderkey, o_shippriority  
from orders
```

```
where o_clerk IN ('Clerk#000000001', 'Clerk#00000201', 'Clerk#000003001') and  
dayofweek(o_orderDATE)=1;
```

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	orders	range	i_o_clerk_date, i_o_clerk	i_o_clerk_date	16	NULL	2978	Using index condition

# Counters

```
alter table orders add key i_o_clerk_date(o_clerk, o_orderDATE);
select o_clerk, o_orderkey, o_shippriority
from orders
where o_clerk IN ('Clerk#000000001', 'Clerk#00000201', 'Clerk#000003001') and
      dayofweek(o_orderDATE)=1;
```

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	orders	range	i_o_clerk_date, i_o_clerk	i_o_clerk_date	16	NULL	2978	Using index condition

Variable_name	Value	
Handler_commit	2	
Handler_delete	0	
Handler_discover	0	
Handler_icp_attempts	2979	
Handler_icp_match	427	= 1/7th
Handler_mrr_init	0	
Handler_mrr_key_refills	0	
Handler_mrr_rowid_refills	0	
Handler_prepare	0	
Handler_read_first	0	
Handler_read_key	3	
Handler_read_next	427	= 1/7th
Handler_read_prev	0	
Handler_read_rnd	0	

MariaDB [(none)]> show status like 'Innodb%';

Variable_name	Value	
Innodb_buffer_pool_read_requests	7687	+2585
Innodb_rows_read	427	+427

# Index Condition Pushdown optimization ... or lack thereof

```
alter table orders add key i_o_clerk_date(o_clerk, o_orderDATE);  
select o_clerk, o_orderkey, o_shippriority  
from orders  
where o_clerk IN ('Clerk#000000001', 'Clerk#00000201', 'Clerk#000003001') and  
dayofweek(o_orderDATE)=1;
```

INDEX (o\_clerk)

- Range selectivity on 1<sup>st</sup> keypart is (or, should be) the same for both
- Index entries are smaller

INDEX (o\_clerk, o\_orderDATE);

- ICP is applicable
  - But is condition dayofweek(...)=1 selective?

Current optimizer actions:

- #1. Choose an index to use
- #2. If possible, use IndexConditionPushdown with it
  - => Possibility of ICP doesn't affect index choice.



# Index Condition Pushdown optimization ... or lack thereof

What to do, then?


- USE/IGNORE INDEX hints
- Create only indexes such that ICP will be used
  - Include appropriate extra columns at the end
  - In our example:
    - Create INDEX (o\_clerk, o\_orderDATE)
    - Don't create INDEX (o\_clerk)

# Index Merge

# Index Merge Union

- ORed conditions on different columns cannot be resolved with any single index
- Example: airline on-time performance data, 15M flights:

```
select * from ontime where (Origin='SEA' or Dest='SEA');
```




# Index Merge union


```
select * from ontime where (Origin='SEA' or Dest='SEA');
```

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	ontime	index_merge	Origin, Dest	Origin, Dest	6, 6	NULL	92850	Using union(Origin, Dest); Using where

- Optimization is challenging:



V.S.




...no way to tell.

Current assumption:  
pessimistic (top)

# Index merge: statistics counters


select count(\*) from ontime where (Origin='SEA' or Dest='SEA'); -- 46412 rows

Variable_name	Value
...	
Handler_read_first	0
Handler_read_key	2
Handler_read_next	46412
Handler_read_prev	0
Handler_read_rnd	46412
Handler_read_rnd_deleted	0
Handler_read_rnd_next	0
...	
Innodb_rows_read	464120 +92824



select count(\*) from ontime where Origin='SEA' -- 23207 rows

Variable_name	Value
...	
Handler_read_first	0
Handler_read_key	1
Handler_read_next	23207
Handler_read_prev	0
Handler_read_rnd	0
Handler_read_rnd_deleted	0
Handler_read_rnd_next	0
...	
Innodb_rows_read	556943 +23207



1.5X reads look like 3x!

# Index merge properties

[MySQL, Percona, MariaDB 5.2x]

index\_merge plans are removed from consideration if a range access is possible:

```
MySQL [ontime]> explain select * from ontime where (Origin='SEA' or Dest='SEA') and securitydelay=0;
```

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	ontime	ref	Origin, Dest, SecurityDelay	SecurityDelay	5	const	791546	Using where

```
MySQL [ontime]> explain select * from ontime where (Origin='SEA' or Dest='SEA') and depdelay < 12*60;
```

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	ontime	ALL	Origin, DepDelay, Dest	NULL	NULL	NULL	1583093	Using where

Solutions:

- Upgrade to MariaDB 5.3+ :-)
- IGNORE INDEX(*\$index\_used\_for\_range*)
  - there is no “FORCE INDEX MERGE” hint



# Index Merge/Union summary

- Used for ORs spanning multiple indexes
  - Each part of OR must have an index it could use
- Counter increments not quite comparable with single-index access methods
- Can be turned off globally with `set optimizer_switch='index_merge=off'`
- Can be turned off locally with `IGNORE INDEX(some_merged_index)`
- Can be blocked by potential ref/range accesses
  - Fix1: Upgrade to MariaDB 5.3
  - Fix2: `IGNORE INDEX(other_range_indexes)`  
`USE INDEX (index_merge_indexes)`

# Index Merge/Intersection

- Like Index Merge/union, but for ANDs
- Aimed at cases when there is no composite index

select avg(arrdelay) from ontime where depdel15=1 and OriginState='CA';


select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
SIMPLE	ontime	index_merge	OriginState,DepDel15	OriginState,DepDel15	3,5	NULL	76952	Using intersect (OriginState,DepDel15); Us

- MySQL, MariaDB 5.2: only support equality conditions
- MariaDB 5.3: any range condition supported
  - Must be manually enabled:  
set optimizer\_switch='index\_merge\_sort\_intersection=on'
  - EXPLAIN: Using sort-intersect

# Index Merge/intersect optimization

- Optimization challenges are similar to UNION

select avg(arrdelay) from ontime where depdel15=1 and OriginState='CA';



- There is no way to force index\_merge/intersect
- To disable
  - set optimizer\_switch='index\_merge\_intersection=off'
  - IGNORE INDEX(one\_of\_used\_indexes)
- Tip: showeling through lots of data?  
set optimizer\_switch='index\_merge\_sort\_intersection=on'  
set sort\_buffer\_size=...

# JOINS

# Consider a join

```
select count(*)
from
  customer, orders
where
  c_custkey=o_custkey and c_acctbal < -500 and o_orderpriority='1-URGENT';
```

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	customer	range	...	c_acctbal	9	NULL	6802	Using where; Using index
1	SIMPLE	orders	ref	i_o_custkey	i_o_custkey	5	customer.c_custkey	7	Using where

# Nested loops join

```
select count(*)
from
  customer, orders
where
  c_custkey=o_custkey and c_acctbal < -500 and o_orderpriority='1-URGENT';
```

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	customer	range	...	c_acctbal	9	NULL	6802	Using where; Using index
1	SIMPLE	orders	ref	i_o_custkey	i_o_custkey	5	customer.c_custkey	7	Using where

for each record R1 in customer

{

for each matching record R2 in orders

{

pass (R1, R2) into join output

}

}

- EXPLAIN shows the loops, from outer to inner

# Nested loops join

```
select count(*)
from
  customer, orders
where
  c_custkey=o_custkey and c_acctbal < -500 and o_orderpriority='1-URGENT';
```

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	customer	range	...	c_acctbal	9	NULL	6802	Using where; Using index
1	SIMPLE	orders	ref	i_o_custkey	i_o_custkey	5	customer.c_custkey	7	Using where

```
for each record R1 in customer
{
  if (where[customer] is satisfied)
  {
    for each matching record R2 in orders
    {
      if (where[orders] is satisfied)
        pass (R1, R2) into join output
    }
  }
}
```

- “Using where” means checking a part of WHERE
- Which part?



# Nested loops join

```
select count(*)
from
  customer, orders
where
  c_custkey=o_custkey and c_acctbal < -500 and o_orderpriority='1-URGENT';
```

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	customer	range	...	c_acctbal	9	NULL	6802	Using where; Using index
1	SIMPLE	orders	ref	i_o_custkey	i_o_custkey	5	customer.c_custkey	7	Using where

```
for each record R1 in customer
{
  if (where[customer] is satisfied)
  {
    for each matching record R2 in orders
    {
      if (where[orders] is satisfied)
        pass (R1, R2) into join output
    }
  }
}
```

- The part that refers to the tables for which we know  $R\{n\}$ , the “current record”.

# Nested loops join

```
select count(*)
from
  customer, orders
where
  c_custkey=o_custkey and c_acctbal < -500 and o_orderpriority='1-URGENT';
```

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	customer	range	...	c_acctbal	9	NULL	6802	Using where; Using index
1	SIMPLE	orders	ref	i_o_custkey	i_o_custkey	5	customer.c_custkey	7	Using where

```
for each record R1 in customer // 6802 loops
{
  if (where[customer] is satisfied)
  {
    for each matching record R2 in orders // 7 loops
    {
      if (where[orders] is satisfied)
        pass (R1, R2) into join output
    }
  }
}
```

- “rows” shows how many rows are read from each table
- For `orders`, 7 rows will be read 6802 times

# Nested loops join

```
select count(*)
from
  customer, orders
where
  c_custkey=o_custkey and c_acctbal < -500 and o_orderpriority='1-URGENT';
```

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	customer	range	...	c_acctbal	9	NULL	6802	Using where; Using index
1	SIMPLE	orders	ref	i_o_custkey	i_o_custkey	5	customer.c_custkey	7	Using where

```
for each record R1 in customer // 6802 loops
{
  if (where[customer] is satisfied)
  {
    for each matching record R2 in orders // 7 loops
    {
      if (where[orders] is satisfied)
        pass (R1, R2) into join output
    }
  }
}
```

- Wait, what about this 'if'?
- It may reject some of the 6802 values of R1

# Nested loops join

explain extended

```
select * from customer, orders
where
  c_custkey=o_custkey and c_acctbal < 1000 and o_orderpriority='1-URGENT';
```

id	select_type	table	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	customer	ALL	...	NULL	NULL	NULL	132987	39.44	Using where
1	SIMPLE	orders	ref	i_o_custkey	i_o_custkey	5	customer.c_custkey	7	100.00	Using where

```
for each record R1 in customer // 6802 loops
{
  if (where[customer] is satisfied)
  {
    for each matching record R2 in orders // 7 loops
    {
      if (where[orders] is satisfied)
        pass (R1, R2) into join output
    }
  }
}
```

- “filtered” column shows the percentage of records that are expected to pass the “if”.
- 100% is the most frequent (optimizer puts 100% when there are no reliable estimates for filtering)

# Ref access

```
select count(*)
from
  customer, orders
where
  c_custkey=o_custkey and c_acctbal < -500 and o_orderpriority='1-URGENT';
```

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	customer	range	...	c_acctbal	9	NULL	6802	Using where; Using index
1	SIMPLE	orders	ref	i_o_custkey	i_o_custkey	5	customer.c_custkey	7	Using where

- `orders` is accessed with **ref** access: index lookups on index **i\_o\_custkey** using value of customer.c\_custkey
- This is similar to key=const ref access we've discussed earlier
  - But there is no "const", the value of customer.c\_custkey is different for each lookup.
- Each lookup is expected to produce 7 matching rows.
  - This information is from **Index Statistics**. We'll cover it later.

How good a join can/should be

# Let's collect some data

```
select count(*)
from
  customer, orders
where
  c_custkey=o_custkey and c_acctbal < -500 and o_orderpriority='1-URGENT';
```

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	customer	range	...	c_acctbal	9	NULL	6802	Using where; Using index
1	SIMPLE	orders	ref	i_o_custkey	i_o_custkey	5	customer.c_custkey	7	Using where

```
[explain] select count(*) from customer; -- 150K
```

```
[explain] select count(*) from customer where c_acctbal < -500 -- 6802
```

```
[explain] select count(*) from orders -- 1.5M
```

```
[explain] select count(*) from orders where o_orderpriority='1-URGENT' - 300K
```

```
One customer (c_custkey) - 7 orders
```

- And use it to analyze the join:




# Analyzing the join

```
select count(*) from customer, orders
```

```
where
```

```
c_custkey=o_custkey and c_acctbal < 1000 and o_orderpriority='1-URGENT';
```




There are three ways to compute the join:

- Yellow → Green
- Blue → Green
- Yellow, Blue

# Join plan analysis

id	select_type	table	type	possible_keys	key	key_len	ref	rows	Extra
1	SIMPLE	customer	range	...	c_acctbal	9	NULL	6802	Using where; Using in
1	SIMPLE	orders	ref	i_o_custkey	i_o_custkey	5	customer.c_custkey	7	Using where



- We're using the Yellow->Green read
- The yellow part is read efficiently
  - Range access is used, it scans 6K rows
- The green part read efficiently
  - ref access, 6K lookups, each producing 7 rows on average
- This is probably better than the two other variants

# Tracking query plan execution

- EXPLAINS are what the optimizer \*expects\* to happen
- What actually happens?
- MySQL has no EXPLAIN ANALYZE
- Possible solutions
  - Status variables
  - “userstat”
  - performance\_schema?

# Check status variables

```
MariaDB [dbt3sf1]> show status like 'Handler%';
```

Variable_name	Value
Handler_commit	1
Handler_delete	0
Handler_discover	0
Handler_icp_attempts	0
Handler_icp_match	0
Handler_mrr_init	0
Handler_mrr_key_refills	0
Handler_mrr_rowid_refills	0
Handler_prepare	0
Handler_read_first	0
Handler_read_key	6805
Handler_read_next	75951
Handler_read_prev	0
Handler_read_rnd	0
Handler_read_rnd_deleted	0
Handler_read_rnd_next	0
Handler_rollback	0
Handler_savepoint	0

...

- Accesses to all tables are summed together!
- It is still possible to do analysis:
  - Run a “sub-join” with the 1<sup>st</sup> table; note the counters
  - Run a “sub-join” with the 1<sup>st</sup> and 2<sup>nd</sup> tables from the join order; note the counters; subtract counters from #1
  - ...
- This is slow and painful.

# More powerful: userstatv2

- Percona Server and MariaDB:

```
MariaDB [dbt3sf1]> show table_statistics;
```

Table_schema	Table_name	Rows_read	Rows_changed	Rows_changed_x_indexes
dbt3sf1	customer	6805	0	0
dbt3sf1	orders	69147	0	0

2 rows in set (0.00 sec)

```
MariaDB [dbt3sf1]> show index_statistics;
```

Table_schema	Table_name	Index_name	Rows_read
dbt3sf1	customer	c_acctbal	6805
dbt3sf1	orders	i_o_custkey	69147

2 rows in set (0.00 sec)


- Much better but self-joins are still painful

# Recap

- Ok we now know
  - Single table access methods
    - ref(const)
    - Range
    - index\_merge
  - Nested loop join algorithm
  - EXPLAIN
  - How to check if execution follows EXPLAIN

# Typical reasons of poor join performance


- Poor access strategy for the 1<sup>st</sup> table
  - Highly-selective non-indexed conditions? Create an index!
- Joins on non-equality conditions
  - e.g. “same day” joins on DATETIME columns
  - Tip: Convert to equality join at all costs
- No index to do ref access for join
  - Cross-product joins inherently slow
- Correlated data (...)
- Wrong index statistics (...)





# Non-uniform distributions

- Some data is highly non-uniform
- Performance depends on hitting the popular values
- No one-for-all solution



# Index and Table Statistics

# Database statistics and plan stability (InnoDB/XtraDB)

- Sample [8] random index leaf pages
- Table statistics (stored)
  - **rows** - estimated number of rows in a table
  - **clust\_size** - cluster index (table/primary key) size in number of pages
  - **other\_size** - other index (non primary key) size in number of pages
- Index statistics (stored)
  - **fields** - #fields in the index
  - **rows\_per\_key** - rows per 1 key value, per prefix fields ([1 column value], [2 columns value], [3 columns value], ...)
  - **index\_total\_pages** - number of index pages
  - **index\_leaf\_pages** - number of leaf pages

# Table statistics updates

- Statistics updated when:
  - ANALYZE TABLE tbl\_name [, tbl\_name] ...
  - SHOW TABLE STATUS, SHOW INDEX
  - Access to INFORMATION\_SCHEMA.[TABLES| STATISTICS]
  - A table is opened for the first time (that is, after server restart)
  - A table has changed a lot (1/16th of the table updated/deleted/inserted)
  - When InnoDB Monitor is turned ON
  - Others (?)

# Displaying statistics

- MySQL 5.5, MariaDB 5.3, and older (GA versions)
  - Issue SQL statements to count rows/keys
  - Indirectly, look at EXPLAIN for simple queries
- MariaDB 5.5, Percona Server 5.5 (using XtraDB) (GA versions)
  - `information_schema.[innodb_index_stats, innodb_table_stats]`
  - Read-only, always visible
- MySQL 5.6 (development)
  - `mysql.[innodb_index_stats, innodb_table_stats]`
  - User updatetable
  - Only available if `innodb_analyze_is_persistent=ON`
- MariaDB \$NEXT (development)
  - MyISAM persistent updateable tables
  - + current MySQL and Percona mechanisms

# Plan [in]stability

- Statistics may vary a lot (orders)

MariaDB [dbt3]> select \* from information\_schema.innodb\_index\_stats;

table_name	index_name	rows_per_key	rows_per_key	error (actual)
partsupp	PRIMARY	3, 1	4, 1	25%
partsupp	i_ps_partkey	3, 0	4, 1	25% (4)
partsupp	i_ps_suppkey	64, 0	91, 1	30% (80)
orders	i_o_orderdate	9597, 1	1660956, 0	99% (6234)
orders	i_o_custkey	15, 1	15, 0	0% (15)
lineitem	i_l_receiptdate	7425, 1, 1	6665850, 1, 1	99.9% (23477)

MariaDB [dbt3]> select \* from information\_schema.innodb\_table\_stats;

table_name	rows	rows	
partsupp	6524766	9101065	28% (8000000)
orders	15039855	14948612	0.6% (15000000)
lineitem	60062904	59992655	0.1% (59986052)

# Controlling statistics (GA versions)

- MySQL 5.5, MariaDB 5.3, and older
  - Manual tuning: optimizer hints, system variables
- MySQL 5.5, MariaDB 5.5 with InnoDB plugin
  - `innodb_stats_on_metadata = OFF` – update only on restart, `ANALYZE`
  - `innodb_stats_sample_pages = 8` is default - increase precision
  - No way to “freeze” db statistics in all cases
- MariaDB 5.5, Percona Server 5.5 (using XtraDB)
  - Can “freeze” the current InnoDB statistics
  - `innodb_use_sys_stats_table=ON` – use `I_S.INNODB_SYS_STATS`
  - `innodb_stats_auto_update=OFF` – recalculate except for “first open” and “`ANALYZE TABLE`”
  - No manual control over statistics, only refresh by random sampling



# Plan [in]stability

- Same query, different statistics

=>

- Different access methods and/or
- Different JOIN orders

=> different query plans

- Query performance may change a lot when statistics changes
- BEWARE WHEN BENCHMARKING THE OPTIMIZER

# Controlling statistics (MySQL dev)

MySQL 5.6 (development version, public code)

- Persistent and user-updatetable InnoDB statistics
  - `innodb_analyze_is_persistent = ON`,
  - updated only on `ANALYZE TABLE`
- Control the precision of sampling [default 8]
  - `innodb_stats_persistent_sample_pages`,
  - `innodb_stats_persistent_sample_pages`
- No new statistics compared to older versions

# Controlling statistics (MariaDB dev)

MariaDB \$NEXT (development version, public code)

- Current MySQL and Percona InnoDB statistics +
- Engine-independent, persistent, user-updateable statistics
- Precise
- Additional statistics per column (even when there is no index):
  - min\_value, max\_value: minimum/maximum value per column
  - nulls\_ratio: fraction of null values in a column
  - avg\_length: average size of values in a column
  - avg\_frequency: average number of rows with the same value

=> better query plans

Code: `bzr branch lp:~maria-captains/maria/maria-5.5-mwl248`

# Resources on InnoDB statistics

From the InnoDB/XtraDB developers:

[https://mysqlperformanceblog.com/doc/percona-server/5.5/diagnostics/innodb\\_stats.html](https://mysqlperformanceblog.com/doc/percona-server/5.5/diagnostics/innodb_stats.html)

<http://dev.mysql.com/doc/refman/5.6/en/innodb-other-changes-statistics-estimation.html>

<http://dev.mysql.com/doc/refman/5.6/en/innodb-parameters.html>

<http://dev.mysql.com/doc/refman/5.6/en/innodb-performance.html#innodb-persistent-stats>

<http://blogs.innodb.com/wp/2011/04/innodb-persistent-statistics-at-last/>

<http://dev.mysql.com/doc/refman/5.6/en/analyze-table.html>

# Optimizer differences between MySQL branches

# Optimizer in different MySQL branches

- Base version: MySQL 5.5
  - Optimizer is not different from MySQL 5.1
- Percona Server X.Y
  - Closely follows MySQL X.Y
  - Some improvements in diagnostics (eg. userstats)
- MariaDB 5.3
  - Lots of new features
- MariaDB 5.5
  - For the most part, is a merge of Maria-5.3 with MySQL-5.5
  - But has a couple of small improvements
- MySQL 5.6
  - Lots of new features, intersection with MariaDB 5.3

# Optimizer feature comparison (1)

Feature	MariaDB 5.3/5.5	MySQL 5.5	MySQL (5.6 dev)
<b>Disk access optimizations</b>			
Index Condition Pushdown (ICP)	YES	-	YES
Disk-sweep Multi-range read (DS-MRR)	YES	-	YES
DS-MRR with Key-ordered retrieval	YES	-	-
Index_merge / Sort_intersection	YES	-	-
Cost-based choice of range vs. index_merge	YES	-	-
ORDER BY .. LIMIT <small_limit>	-	-	YES
Use extended (hidden) primary keys for innodb/xtradb	YES (5.5)	-	-
<b>Join optimizations</b>			
Batched key access (BKA)	YES	-	YES
Block hash join	YES	-	-
User-set memory limits on all join buffers	YES	-	-
Apply early outer table ON conditions	YES	-	-
Null-rejecting conditions tested early for NULLs	YES	-	-



# Optimizer feature comparison (2)

Feature	MariaDB 5.5	MySQL 5.5	MySQL (5.6 dev)
<b>Subquery optimizations</b>			
In-to-exists	YES	YES	YES
Semi-join	YES	-	-YES
Materialization	YES	-	YES
NULL-aware Materialization	YES	-	-
Cost choice of materialization vs in-to-exists	YES	-	-
Subquery cache	YES	-	-
Fast explain with subqueries	YES		-
<b>Optimizations for derived tables / views</b>			
Delayed materialization of derived tables / materialized views	YES	-	YES
Instant EXPLAIN for derived tables	YES	-	YES
Derived Table with Keys optimization	YES	-	YES
Fields of merge-able views and derived tables used in equality optimizations	YES	-	-

# Optimizer feature comparison (3)

Feature	MariaDB 5.3/5.5	MySQL 5.5	MySQL (5.6 dev)
<b>Execution control</b>			
LIMIT ROWS EXAMINED rows_limit	YES (5.5)	-	-
<b>Optimizer control (optimizer switch)</b>			
Systematic control of all optimizer strategies	YES	-	partial
<b>EXPLAIN improvements</b>			
Explain for DELETE, INSERT, REPLACE, and UPDATE	-	-	YES
EXPLAIN in JSON format	-	-	YES
More detailed and consistent EXPLAIN for subqueries	YES	-	-

# Optimizer switch (2)

## MariaDB 5.5

### Index merge:

```
index_merge=on,  
index_merge_union=on,  
index_merge_sort_union=on,  
index_merge_intersection=on,  
index_merge_sort_intersection=off
```

### Condition pushdown:

```
engine_condition_pushdown=off,  
index_condition_pushdown=on,  
  
mrr=off,  
mrr_cost_based=off,  
mrr_sort_keys=off
```

## MySQL 5.6

### Index merge:

```
index_merge=on,  
index_merge_union=on,  
index_merge_sort_union=on,  
index_merge_intersection=on
```

### Condition pushdown:

```
engine_condition_pushdown=on,  
index_condition_pushdown=on,  
  
mrr=on,  
mrr_cost_based=on
```

# Optimizer switch (2)

## MariaDB 5.5

### **Subqueries:**

materialization=on,  
semijoin=on,  
loosescan=on,  
firstmatch=on,  
**in\_to\_exists=on,**  
**partial\_match\_rowid\_merge=on,**  
**partial\_match\_table\_scan=on,**  
**subquery\_cache=on,**

### **Derived tables:**

**derived\_merge=on,**  
**derived\_with\_keys=on**

## MySQL 5.6

### **Subqueries:**

materialization=on,  
semijoin=on,  
loosescan=on,  
firstmatch=on

# Optimizer switch (3)

## MariaDB 5.5

### Joins

`join_cache_bka=on,`  
`join_cache_incremental=on,`  
`join_cache_hashed=on,`  
`outer_join_with_cache=on,`  
`semijoin_with_cache=on,`  
`optimize_join_buffer_size=off`  
+ other system variables

### Other features

`table_elimination=on,`  
`extended_keys=off`

## MySQL 5.6

### Joins:

`batched_key_access=off`  
`block_nested_loop=on,`

Thank You!

Q & A