



continuent

MySQL Replication 101

Giuseppe Maxia
Continuent, Inc

about me - Giuseppe Maxia

- a.k.a. The Data Charmer
- QA Director at Continuent, Inc
- Long time hacking with MySQL features
- Formerly, community manager, db consultant, designer, coder.
- A passion for QA and open source
- Blogger
- <http://datacharmer.blogspot.com>



AGENDA

Why replication

- The web economy
- Scaling out

How to set replication

- From single server to replication
- Adding a slave

Using Replication

- Binary log formats
- What gets replicated and how
- Replication awareness

Managing replication

- Monitoring
- Log management
- Replacing a slave
- Replacing a master

What Replication is for

- master replacement
- backup
- better reads

Gotchas, tips, and tricks

More info

- What to read
- More replication sessions

AGENDA

Why replication

- **The web economy**

- Scaling out

- From single server to replication

- Adding a slave

- Binary log formats

- What gets replicated and how

- Replication awareness

- Monitoring

- Log management

- Replacing a slave

- Replacing a master

- master replacement

- backup

- better reads

Managing replication

What Replication is for

Gotchas, tips, and tricks

More info

- What to read

- More replication sessions



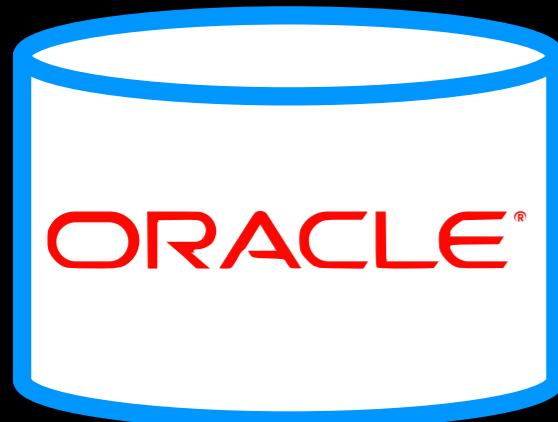
The world today is
dominated by the
web economy



Databases are
the backbone
of the web
economy

What database for the web?

What database for the web?

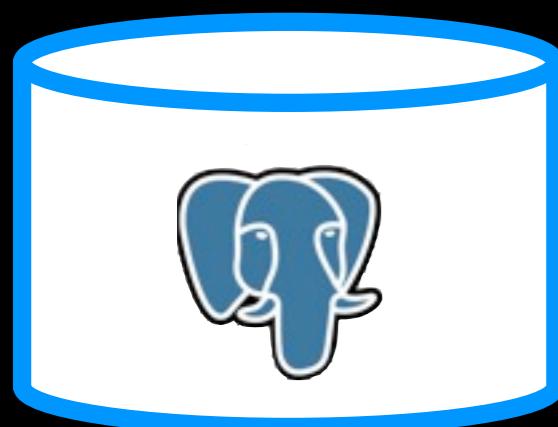


The most **powerful** database

What database for the web?



The most **powerful** database

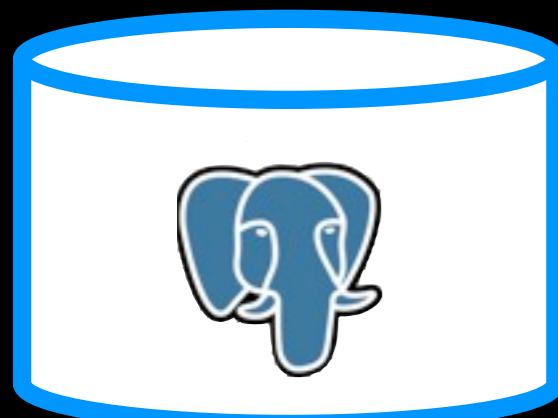


The most **advanced** open source database

What database for the web?



The most **powerful** database



The most **advanced** open source database

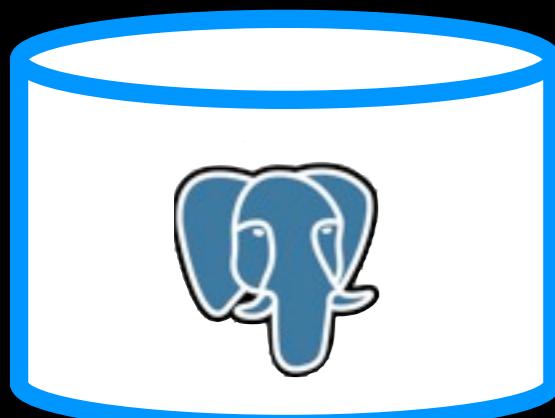


The most **deployed** open source database

What database for the web?



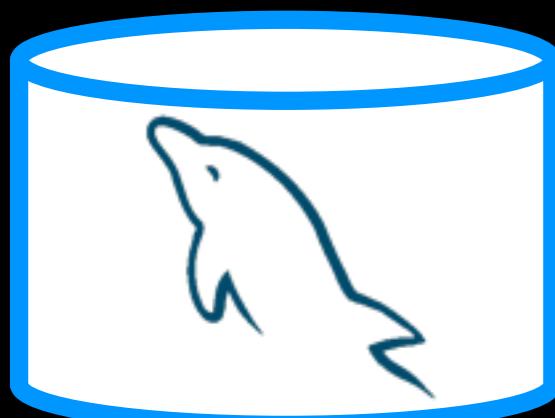
The most **powerful** database



The most **advanced** open source database



The most **deployed** open source database



The most **popular** open source database

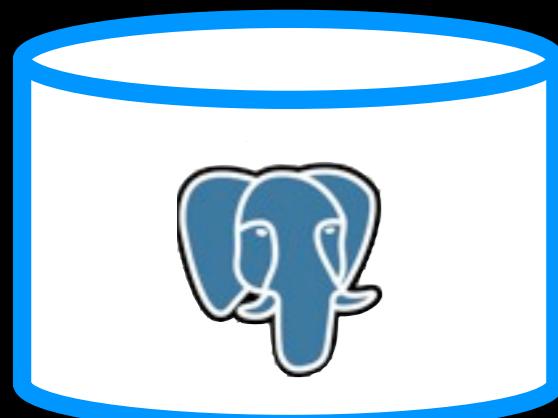


Actually, **MySQL**
databases are the
backbone of the web
economy

What database for the web?



No built-in replication



No built-in replication



No built-in replication



Built-in replication





More precisely, **MySQL REPLICATION** is the backbone of the web economy

AGENDA

Why replication

- The web economy
- **Scaling out**

How to set replication

- From single server to replication
- Adding a slave

Using Replication

- Binary log formats
- What gets replicated and how
- Replication awareness

Managing replication

- Monitoring
- Log management
- Replacing a slave
- Replacing a master

What Replication is for

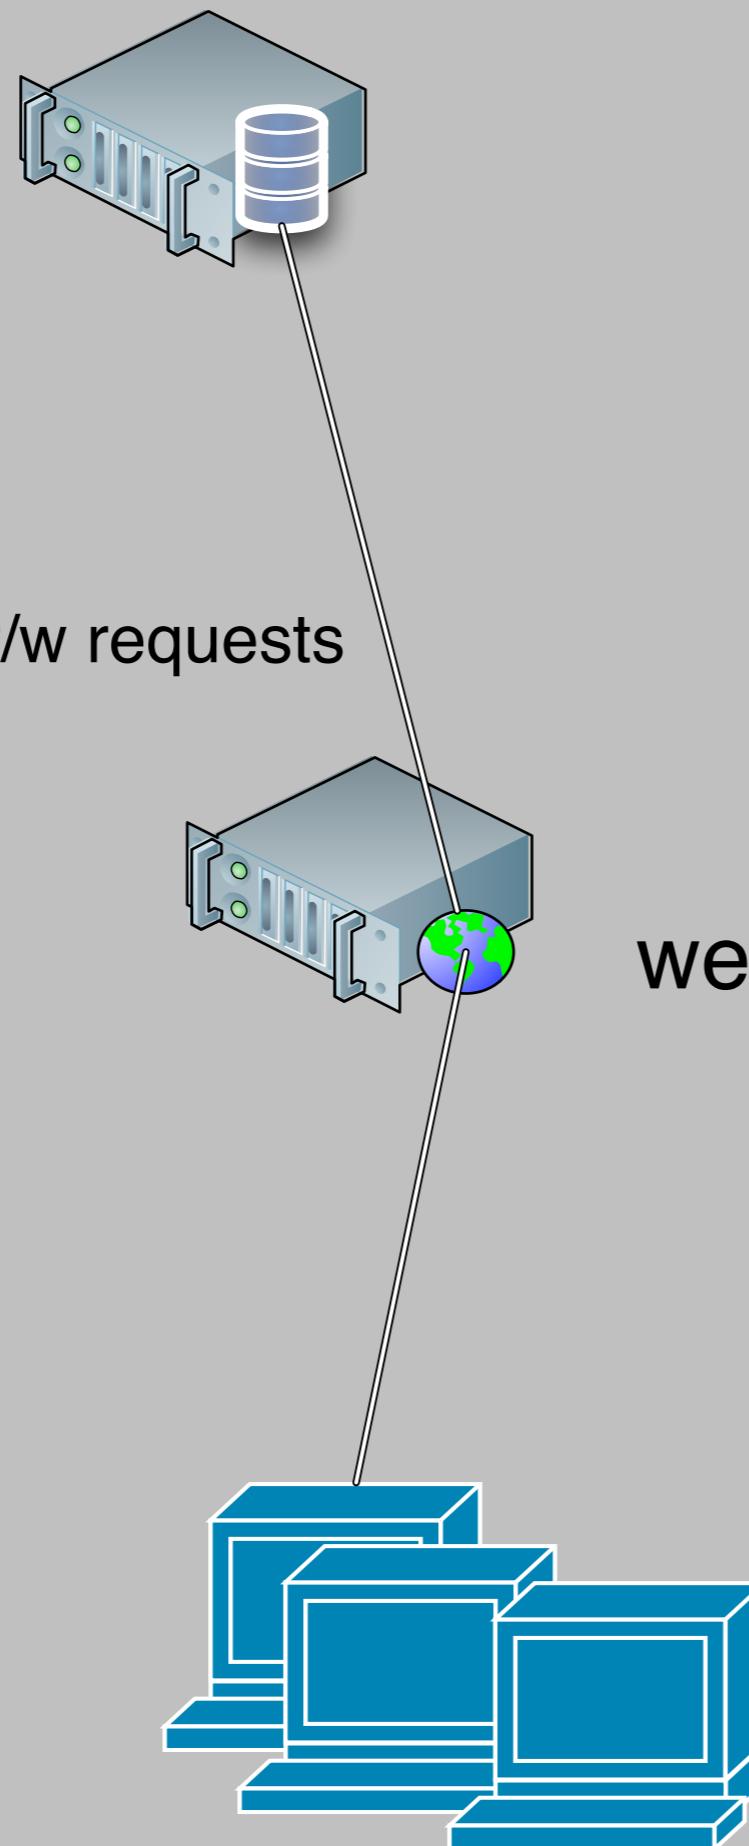
- master replacement
- backup
- better reads

Gotchas, tips, and tricks

- What to read
- More replication sessions

More info

database server

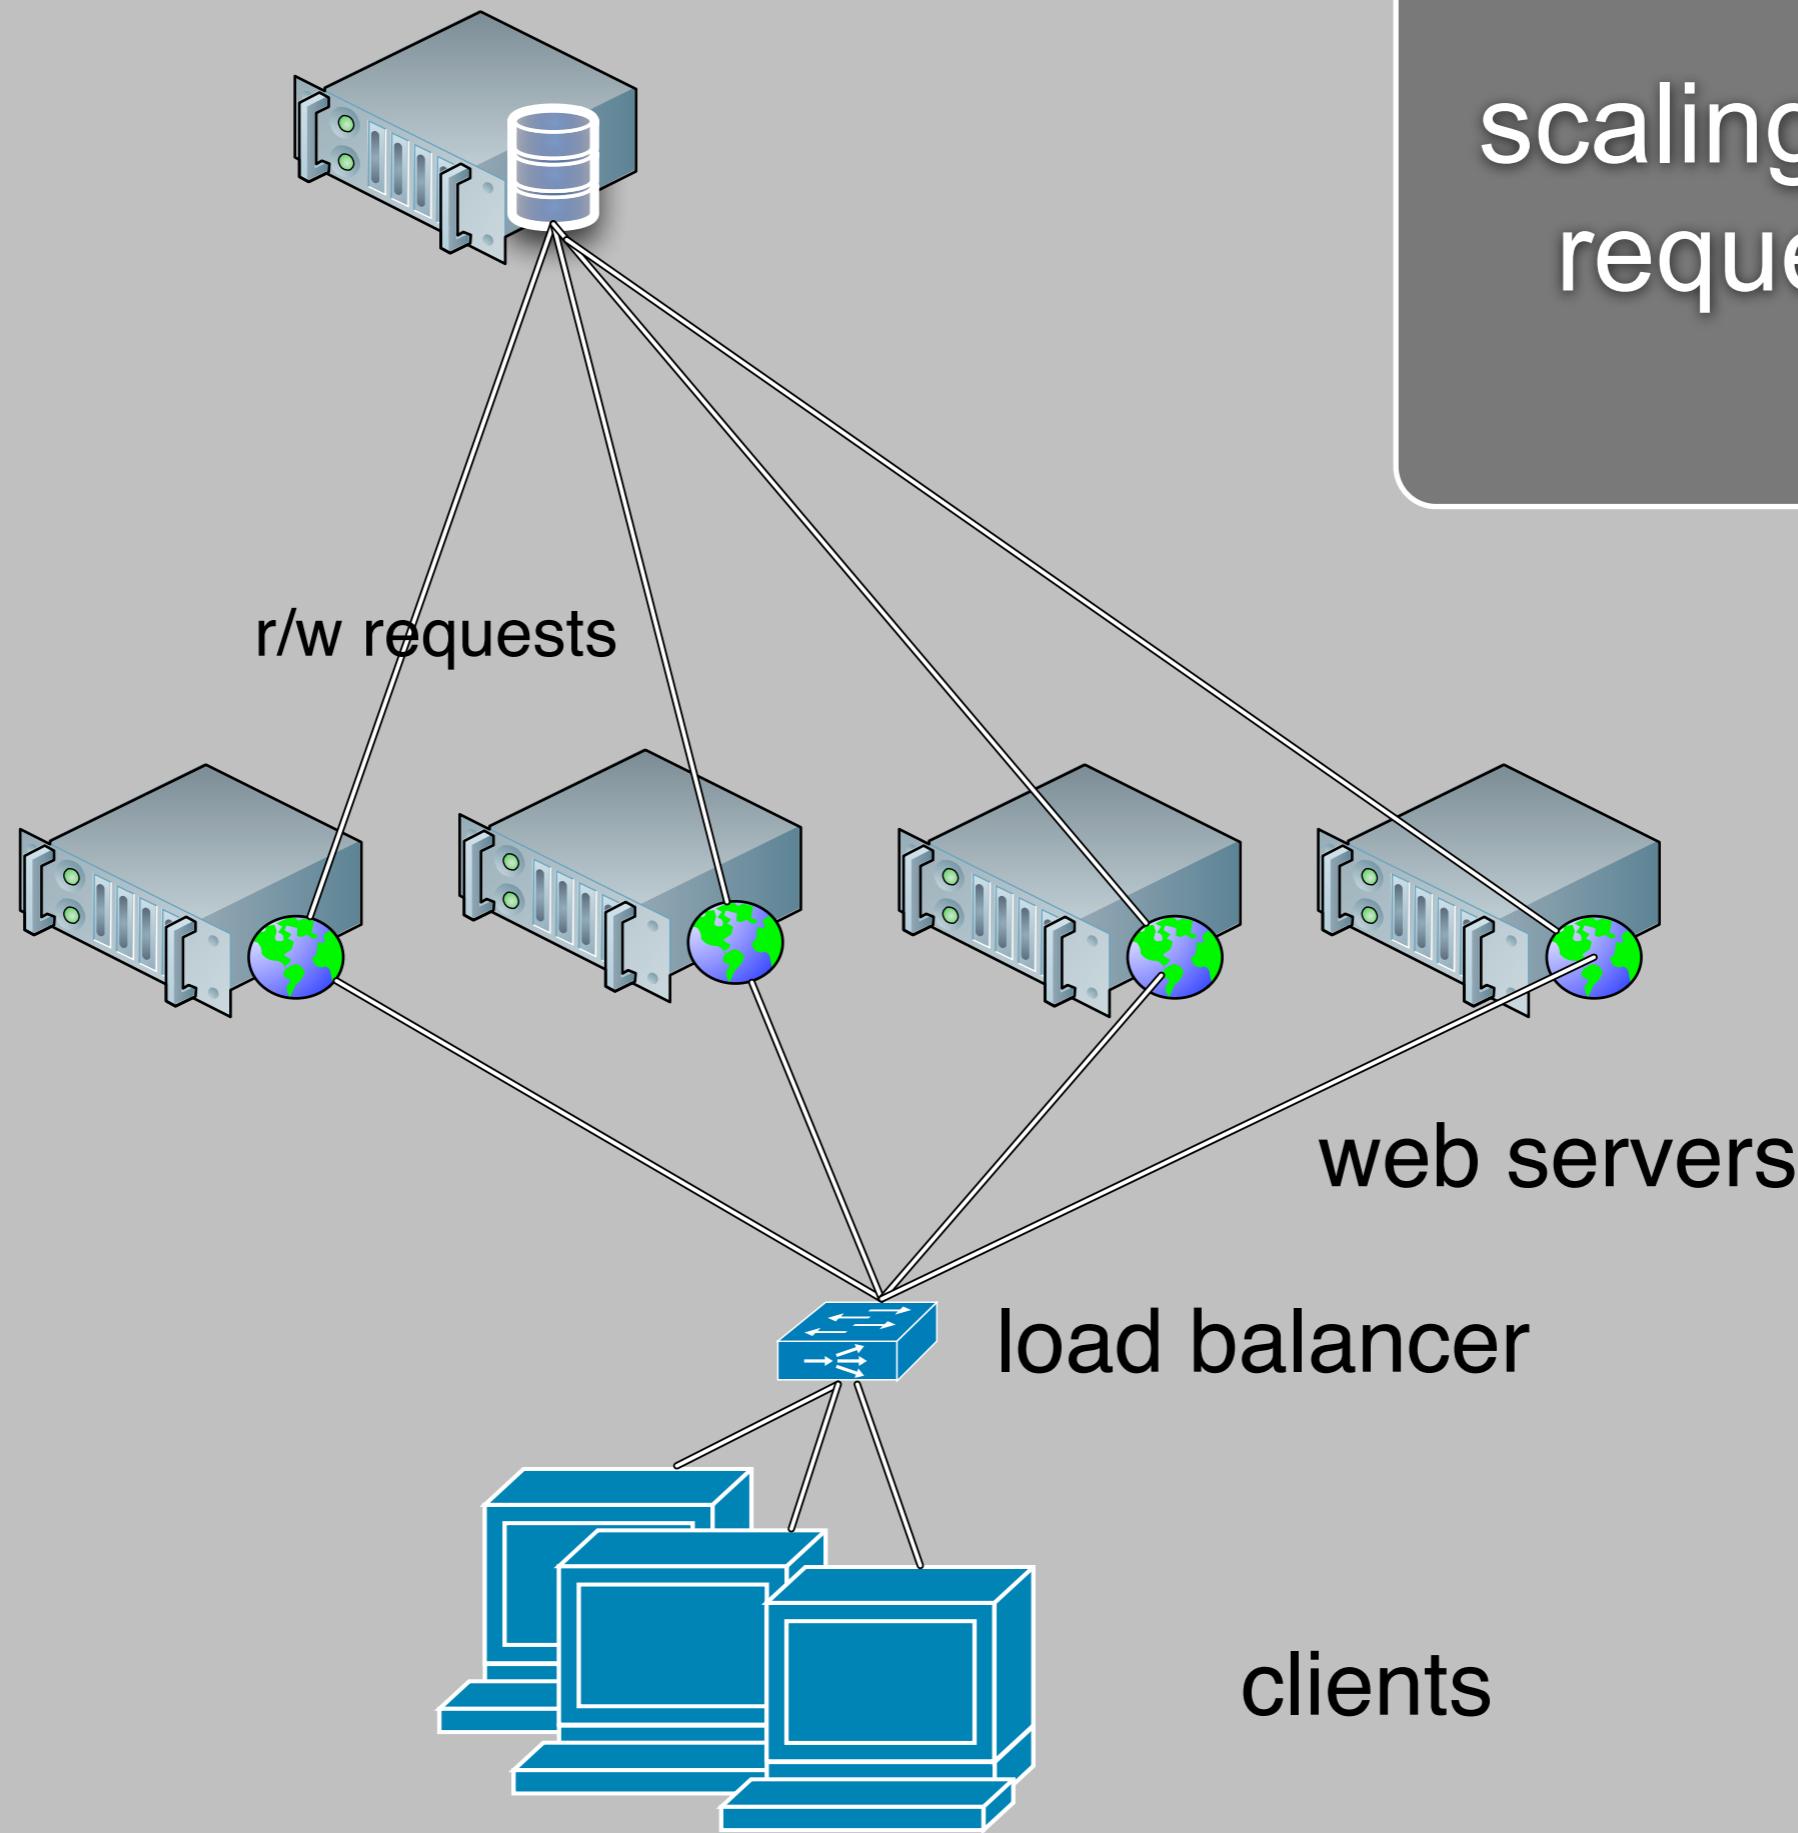


a simple web
application
scheme

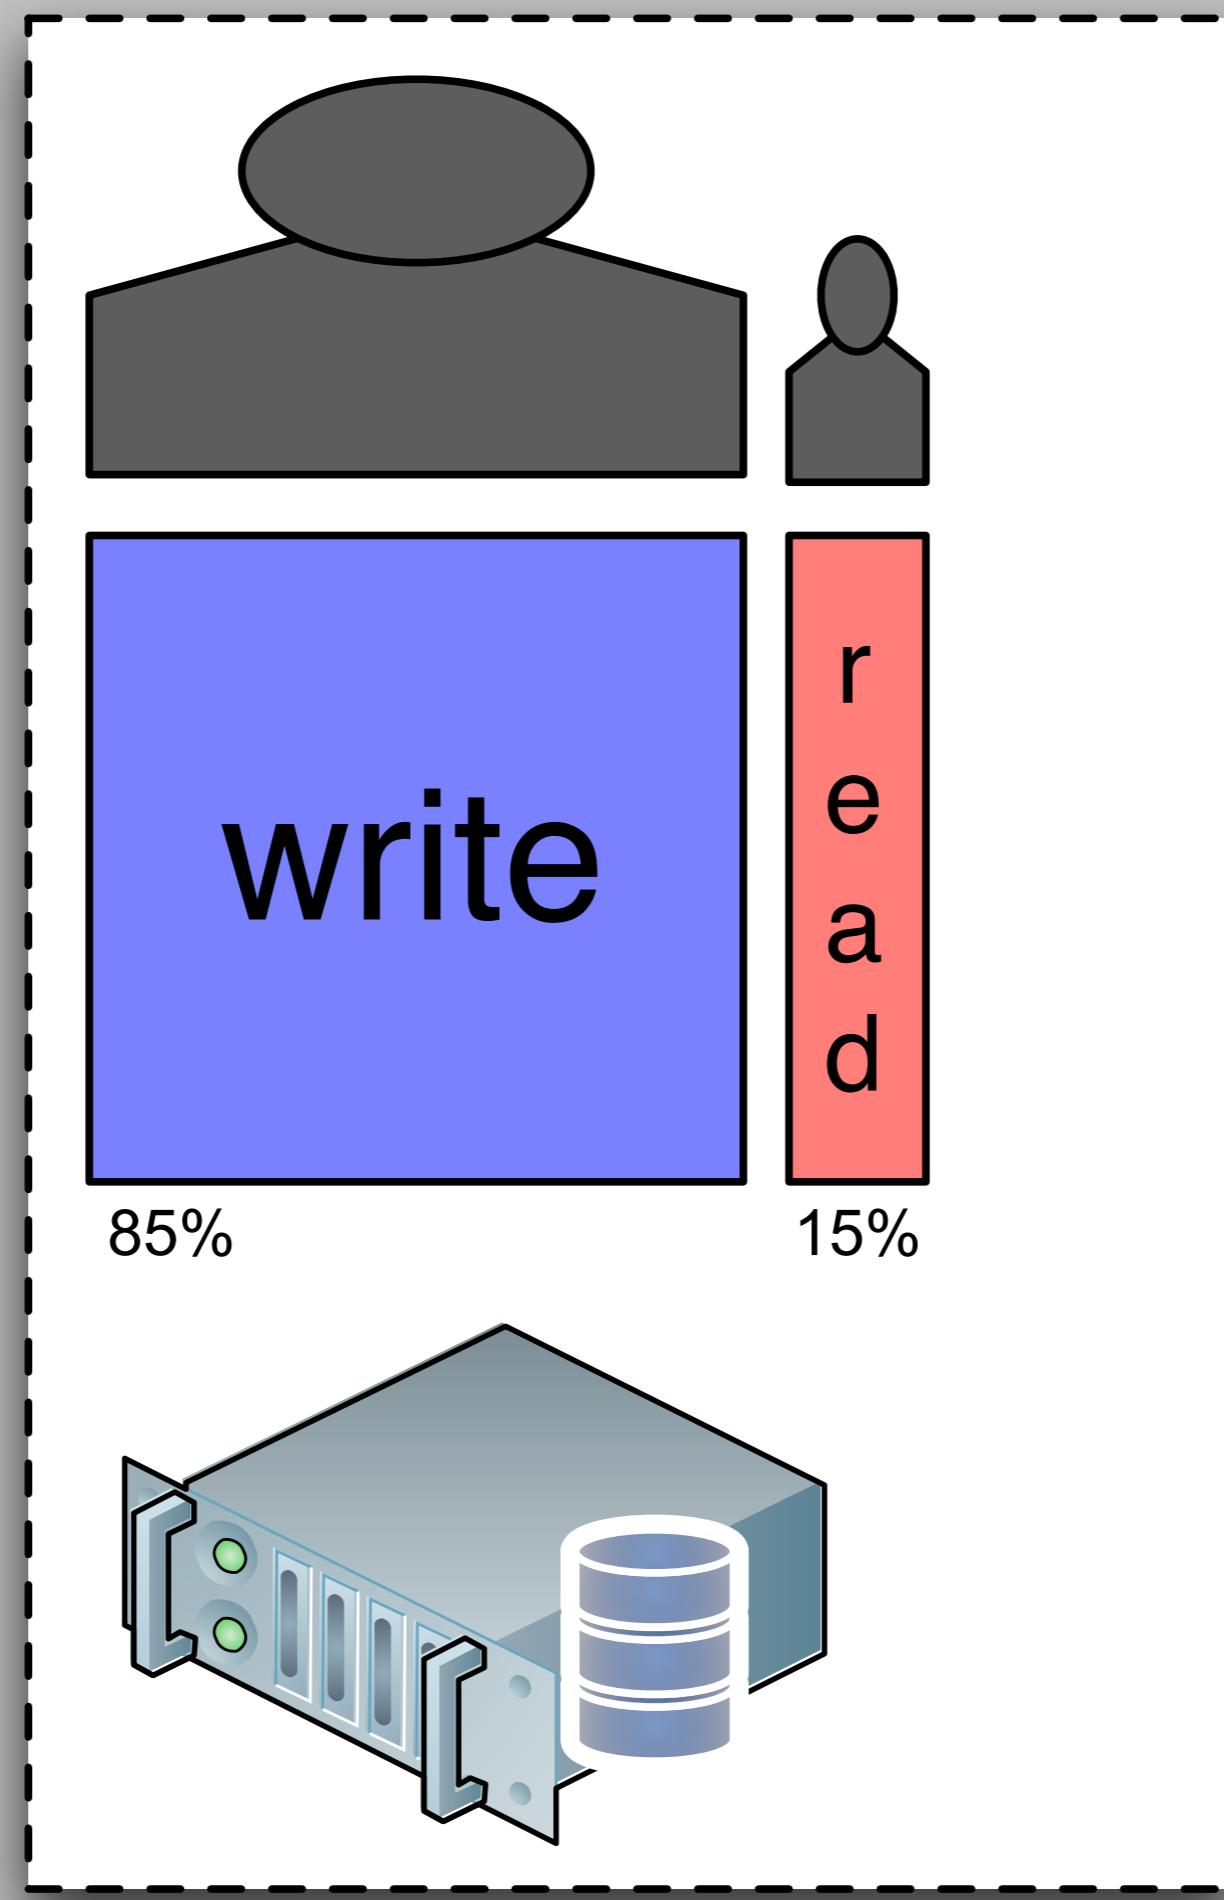
web server

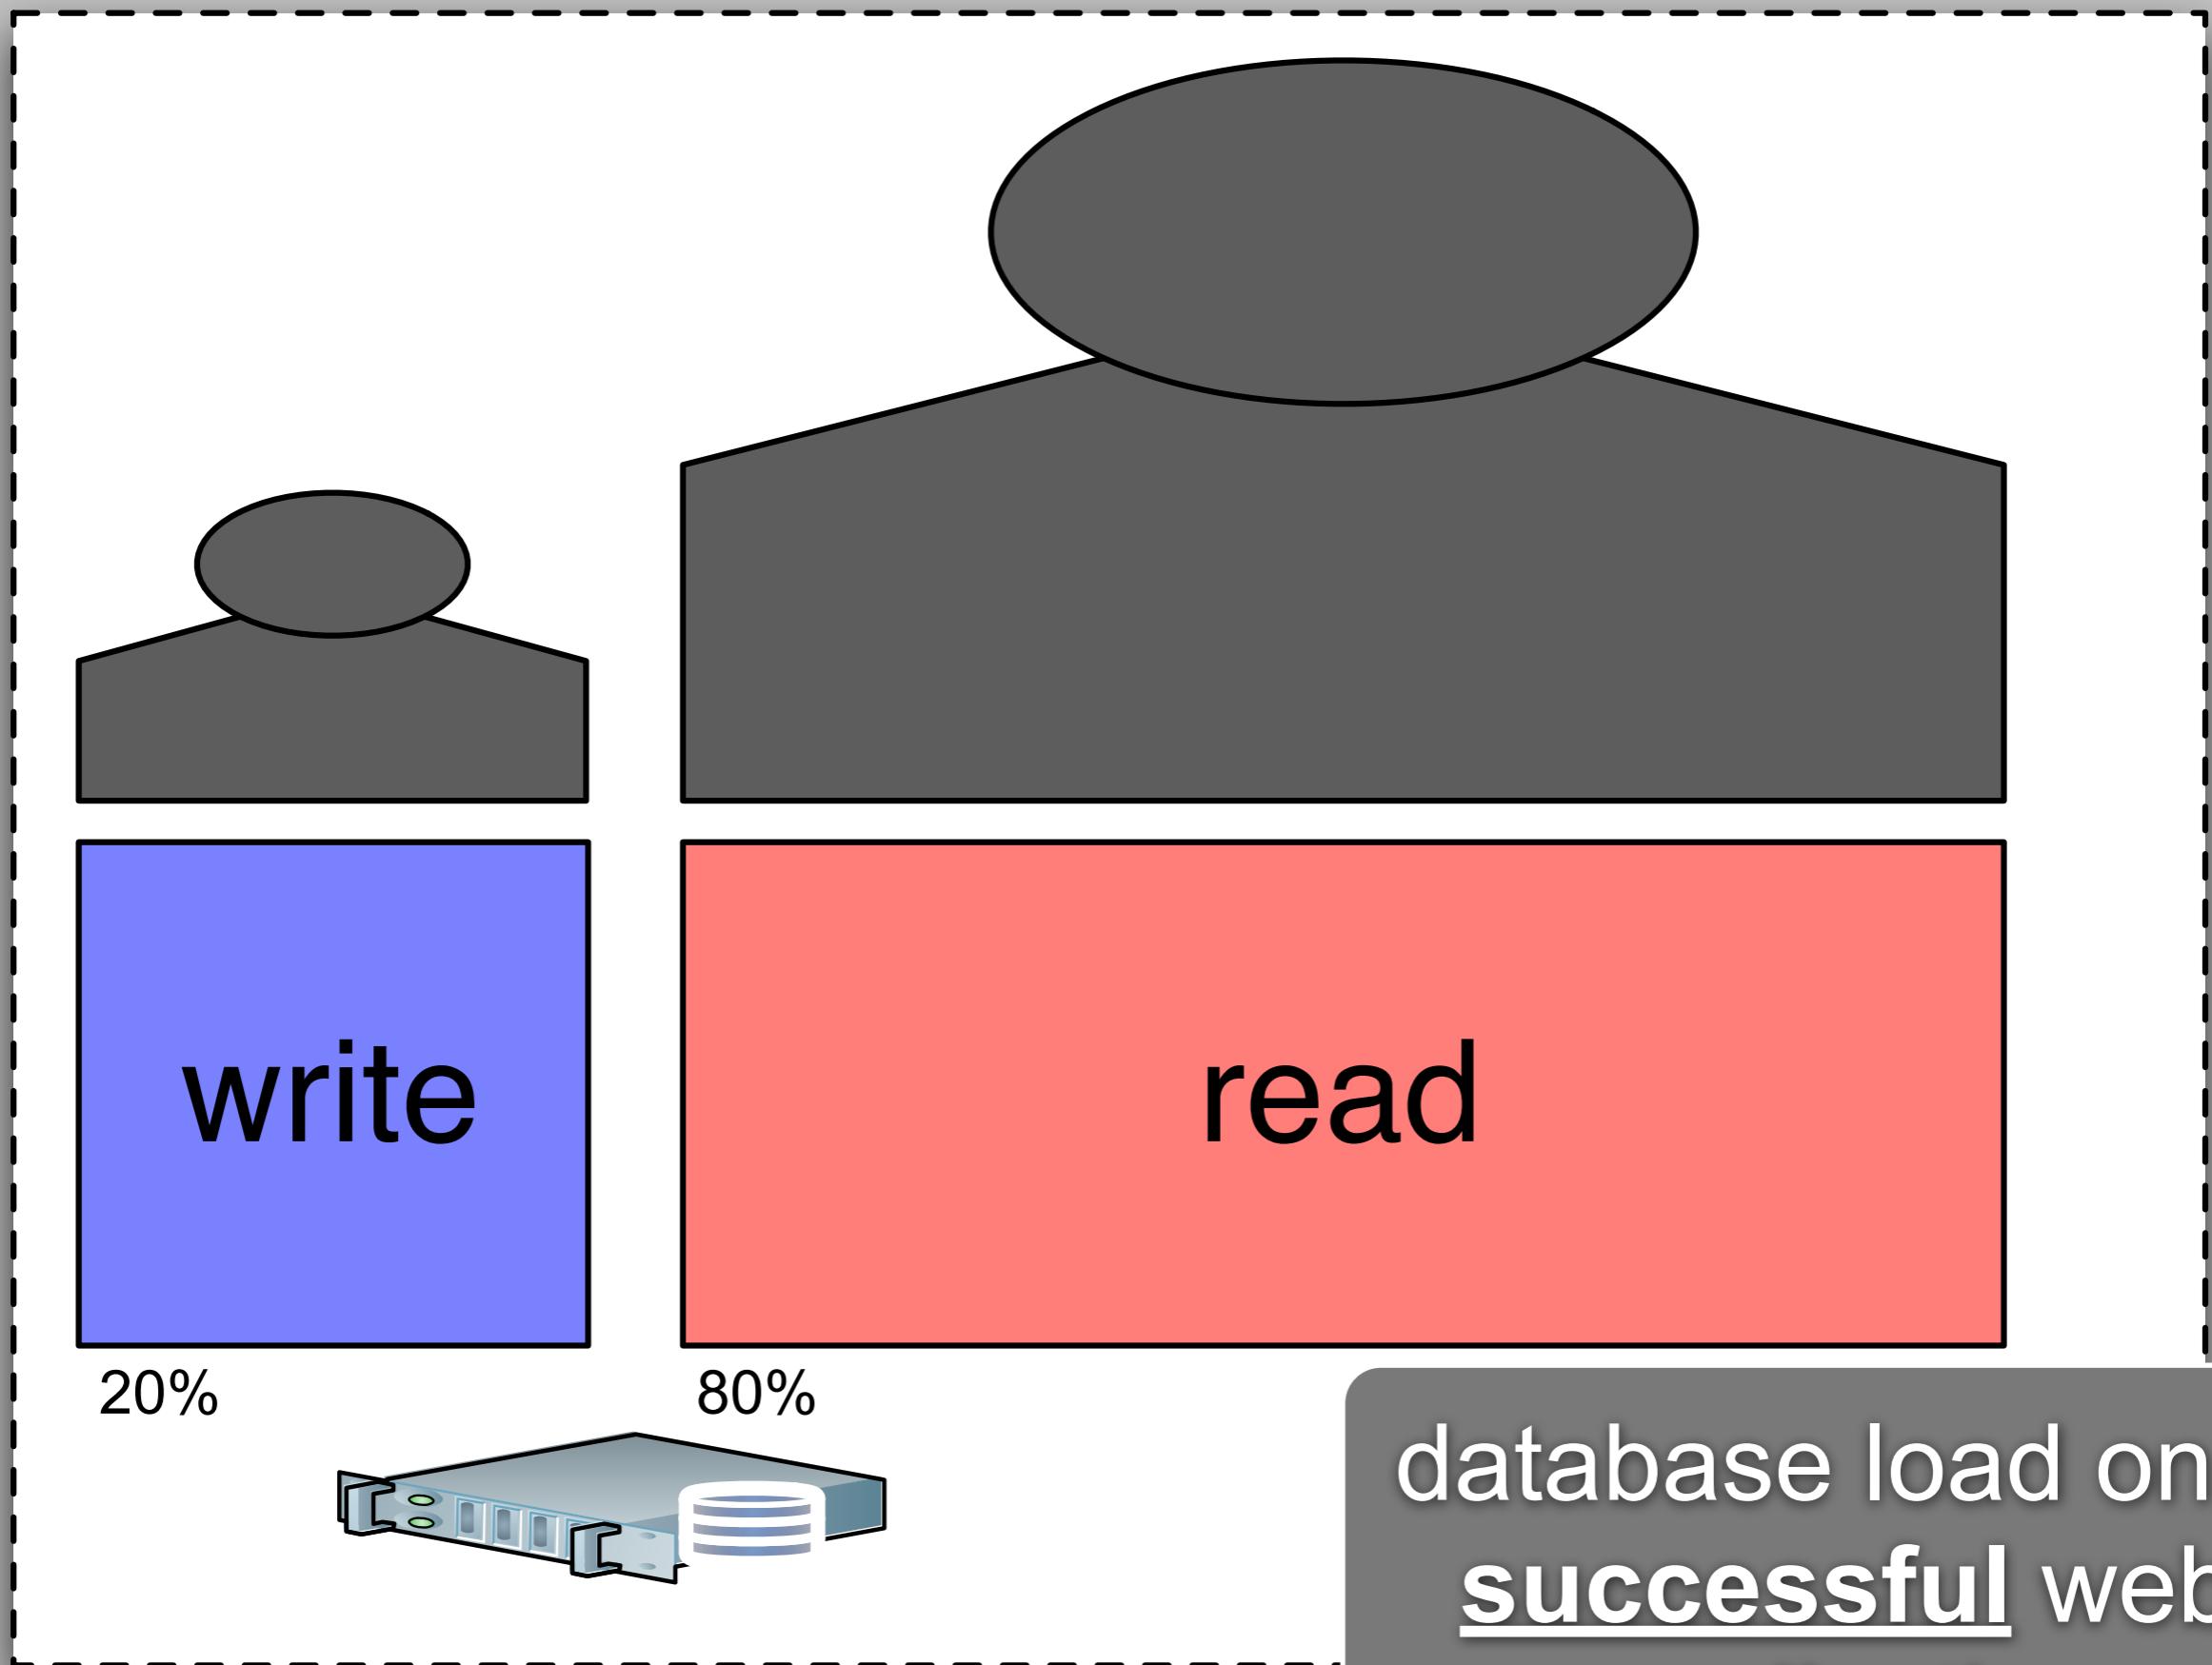
clients

database server

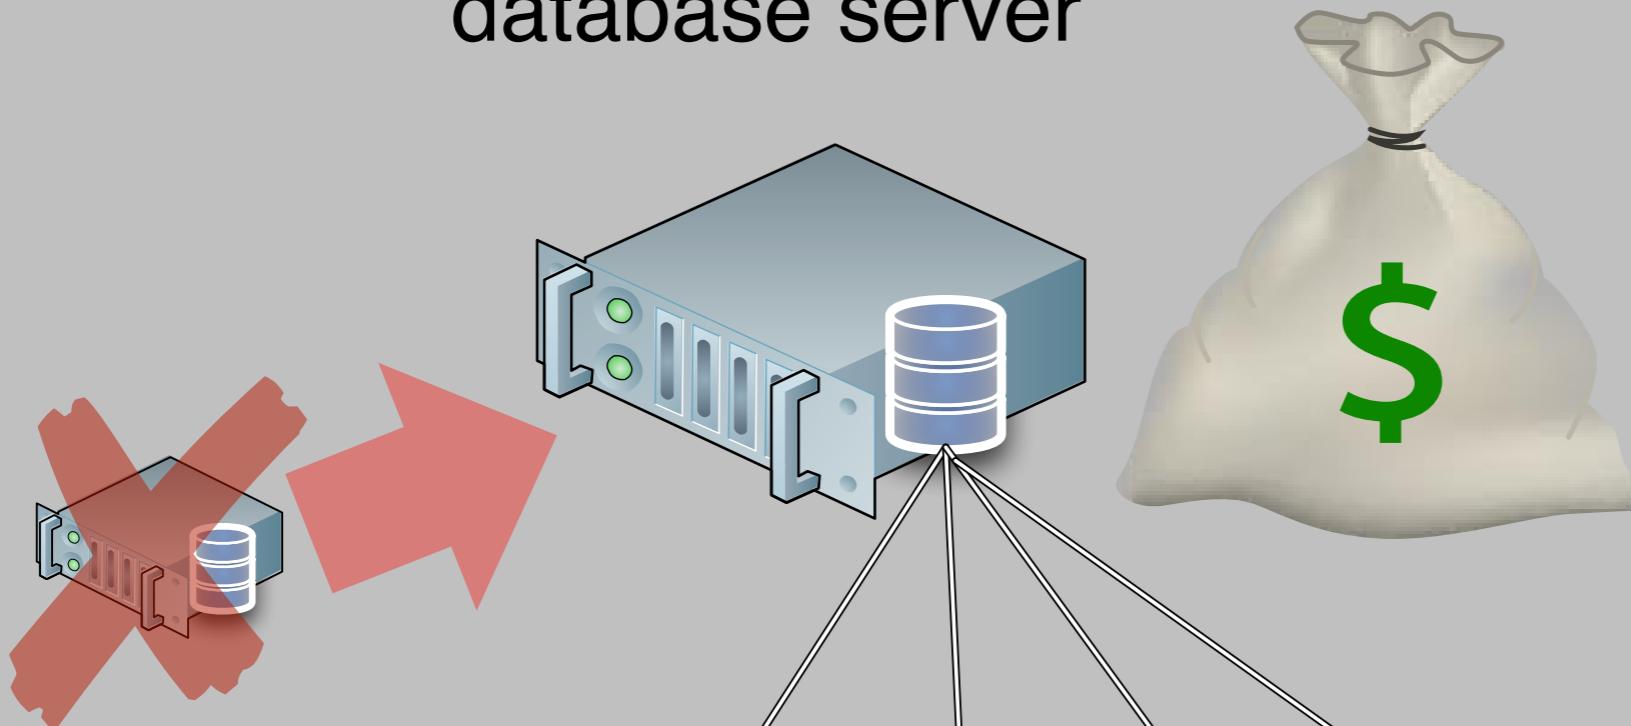


database load
on a simple
web
application



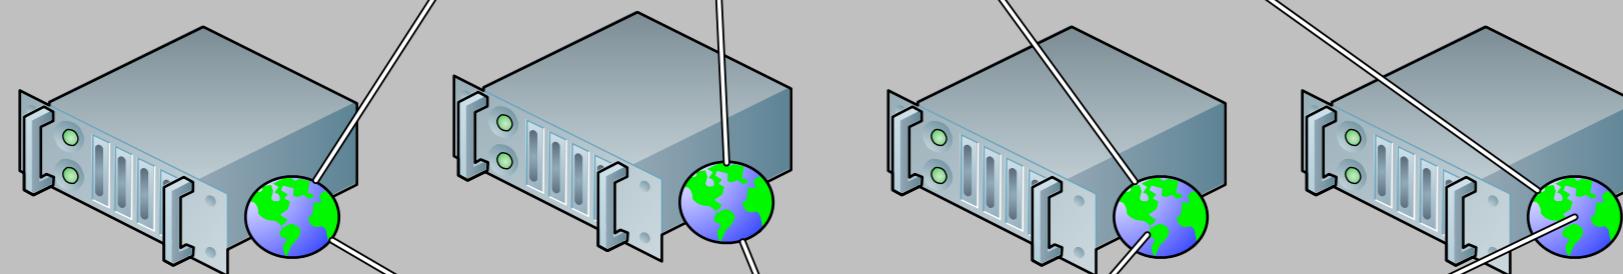


database server



scaling up
means buying
a bigger
database
server

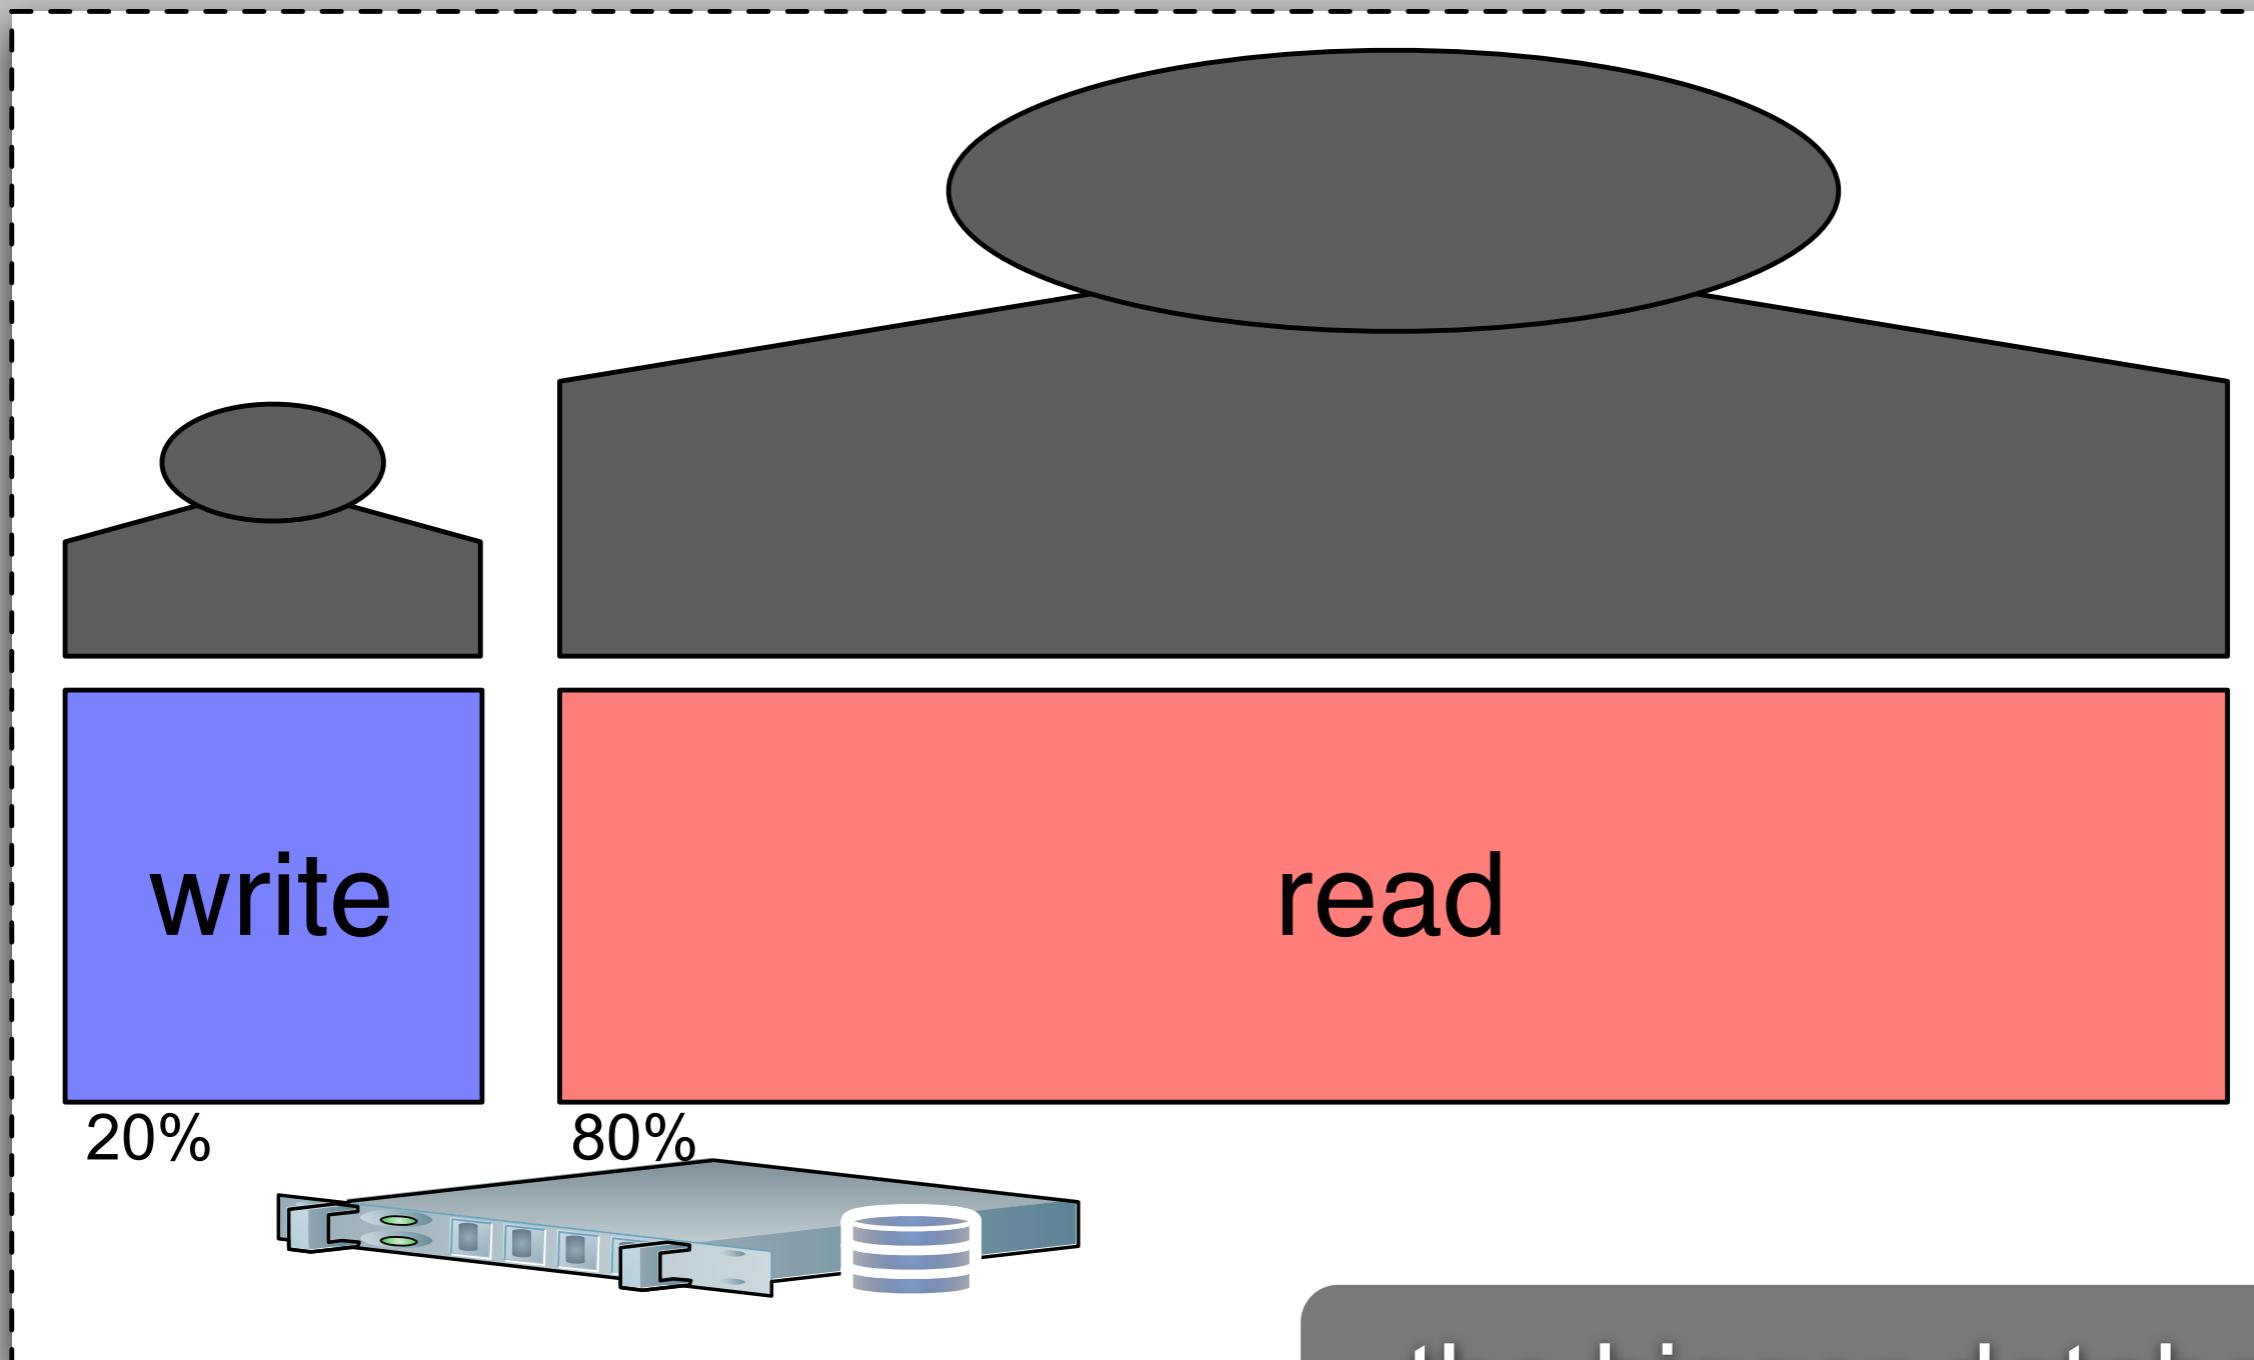
r/w requests



web servers

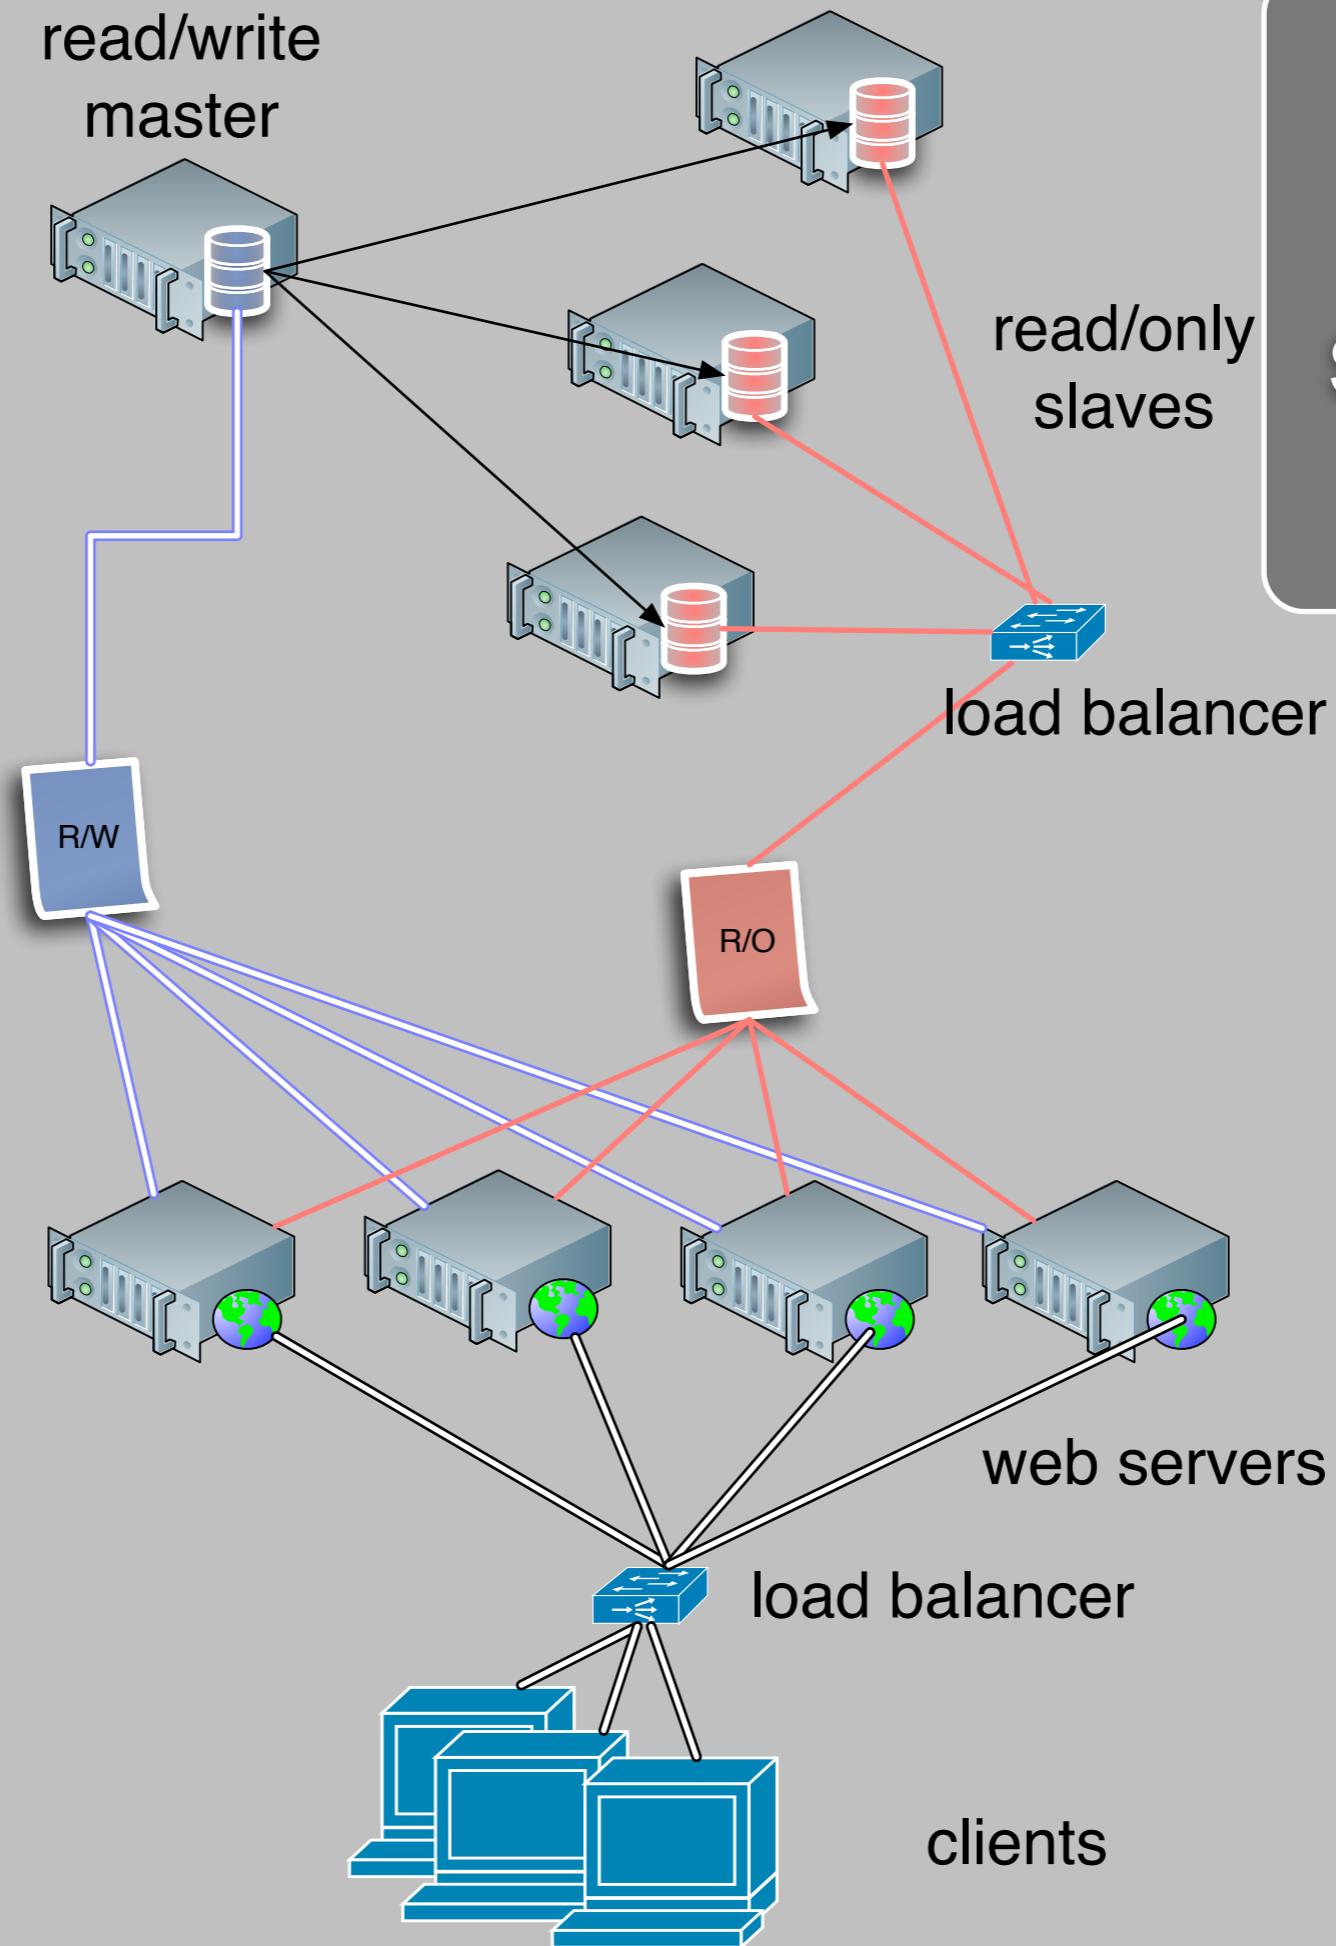
load balancer

clients

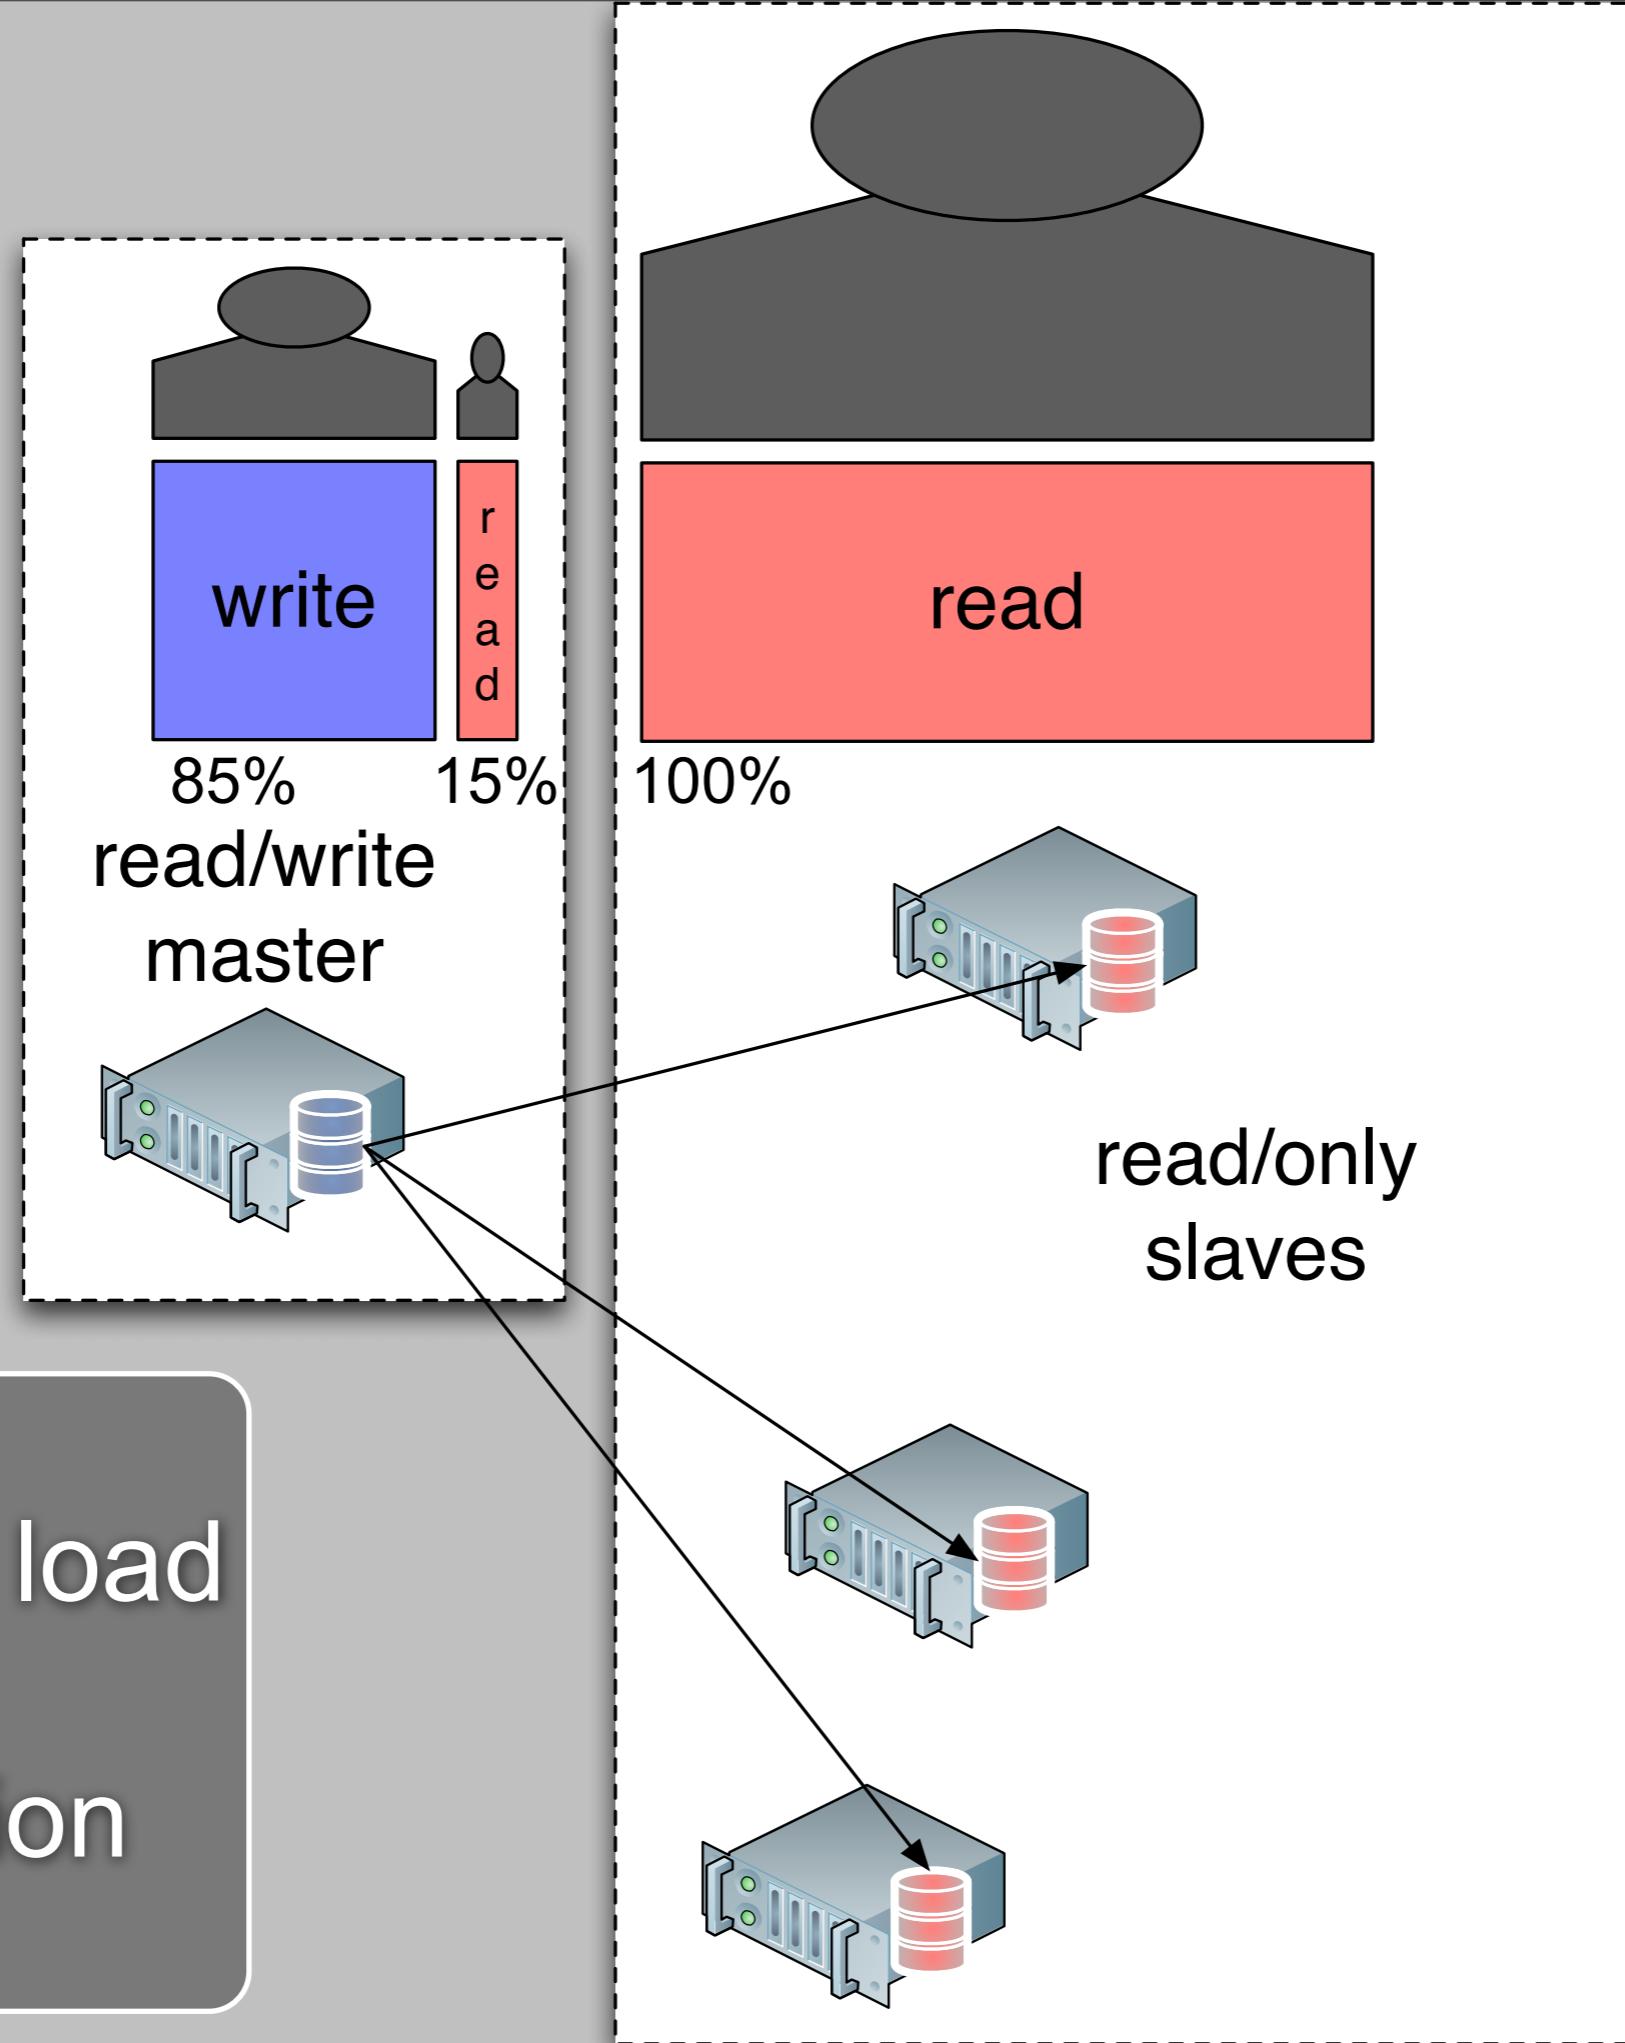


the bigger database
server will eventually
have the same problem

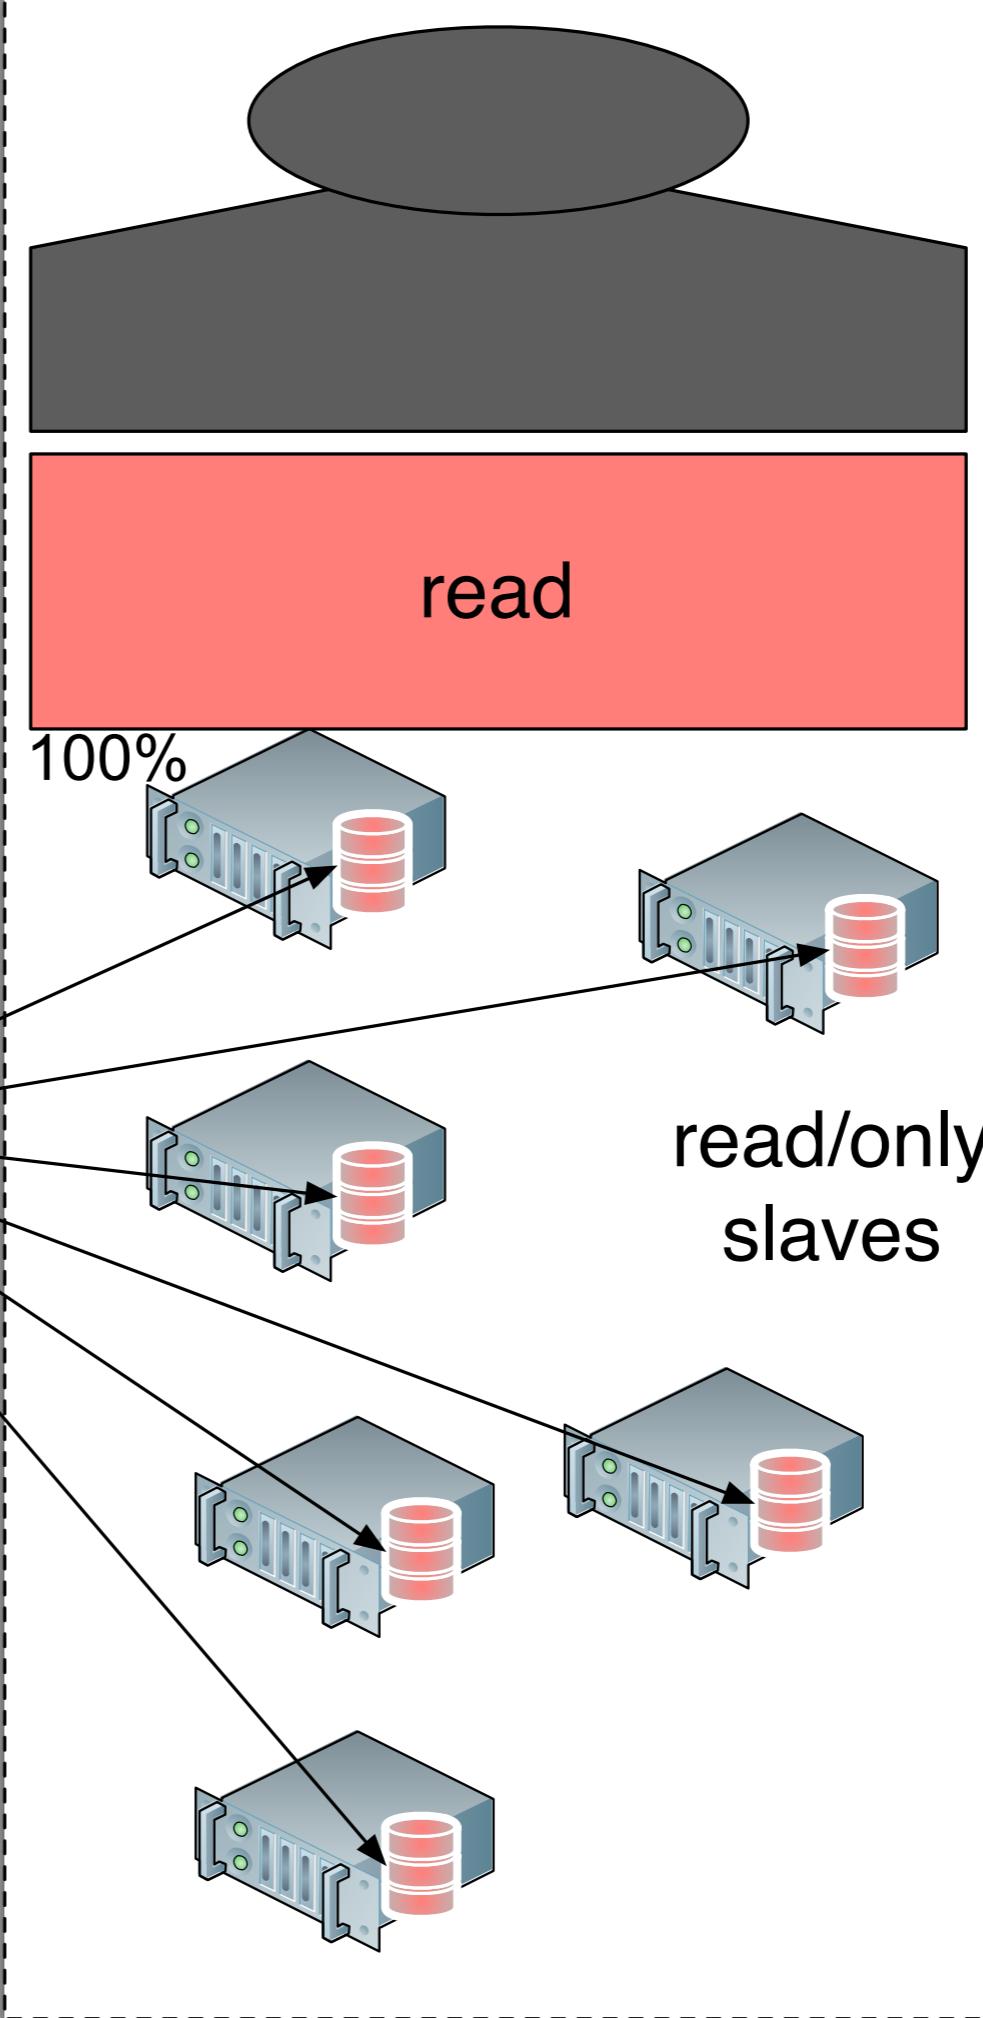
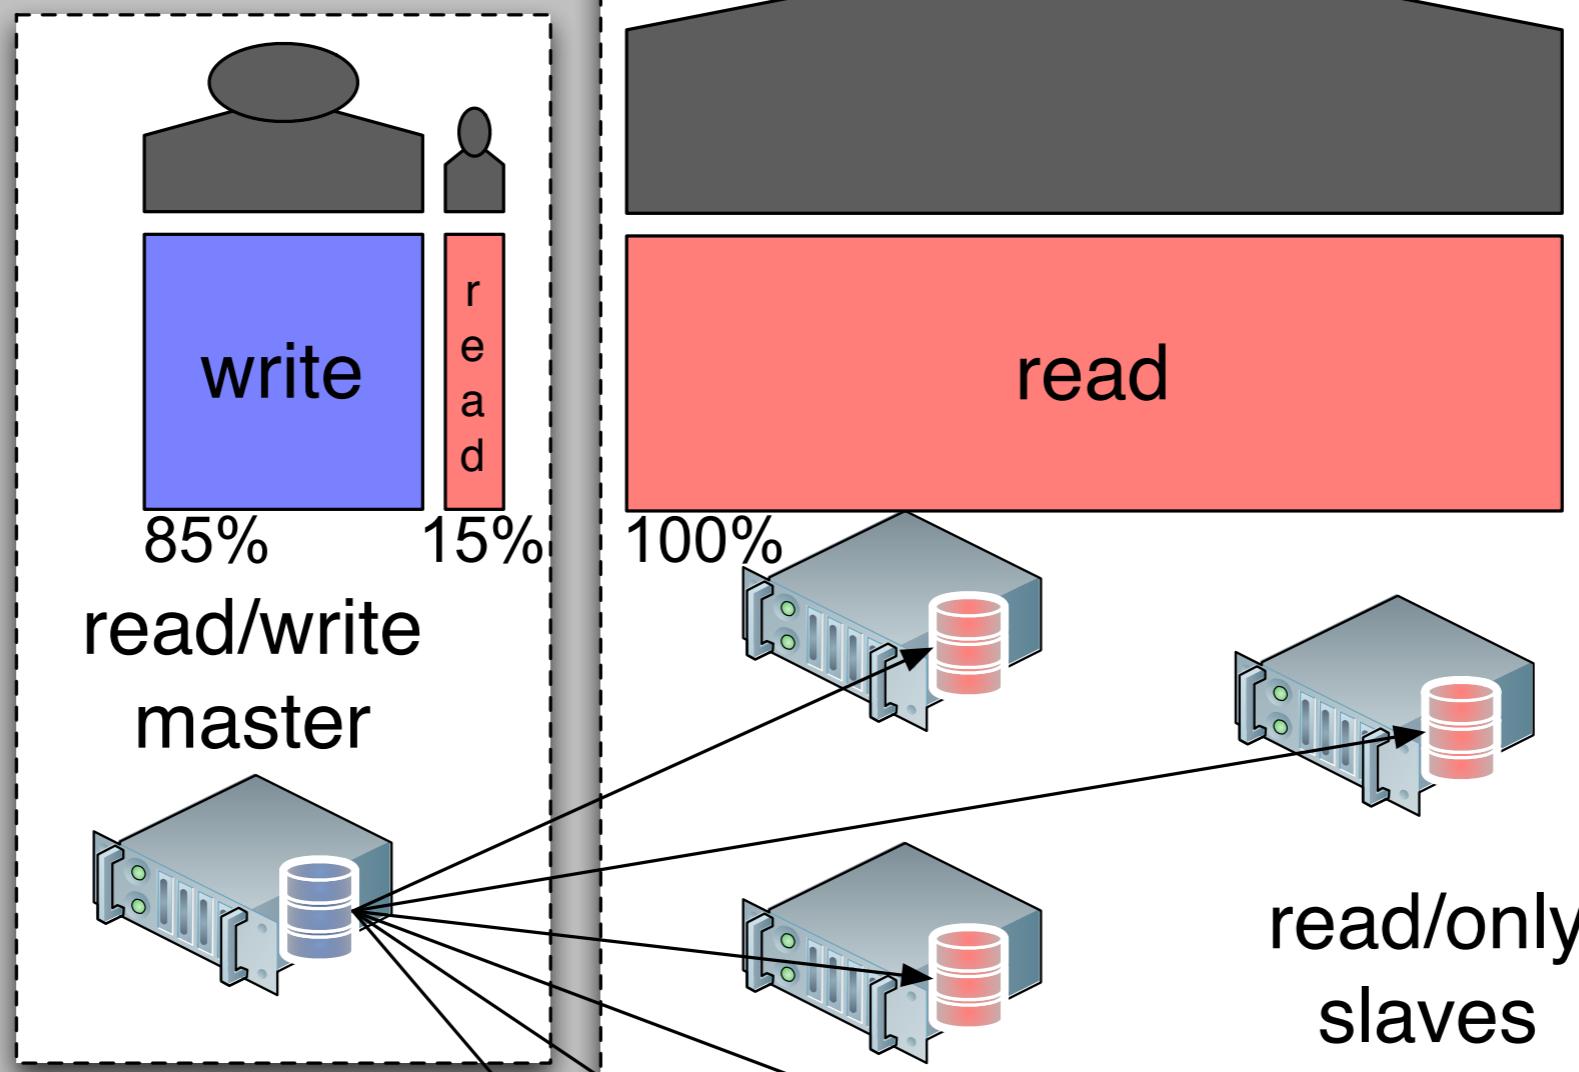
a web application scheme with replication



database load with replication



scaling
database load
with
replication



Replication assessment

	without replication	with replication
database handling	easy	harder
performance	high	lower (binary logs)
Point in Time recovery	none	easy
failover	none	possible
write scaling	none	minimal
backup	with downtime	without downtime
read scaling	none	easy

AGENDA

Why replication

- The web economy
- Scaling out

How to set replication

- From single server to replication
- Adding a slave

Using Replication

- Binary log formats
- What gets replicated and how
- Replication awareness

Managing replication

- Monitoring
- Log management
- Replacing a slave
- Replacing a master

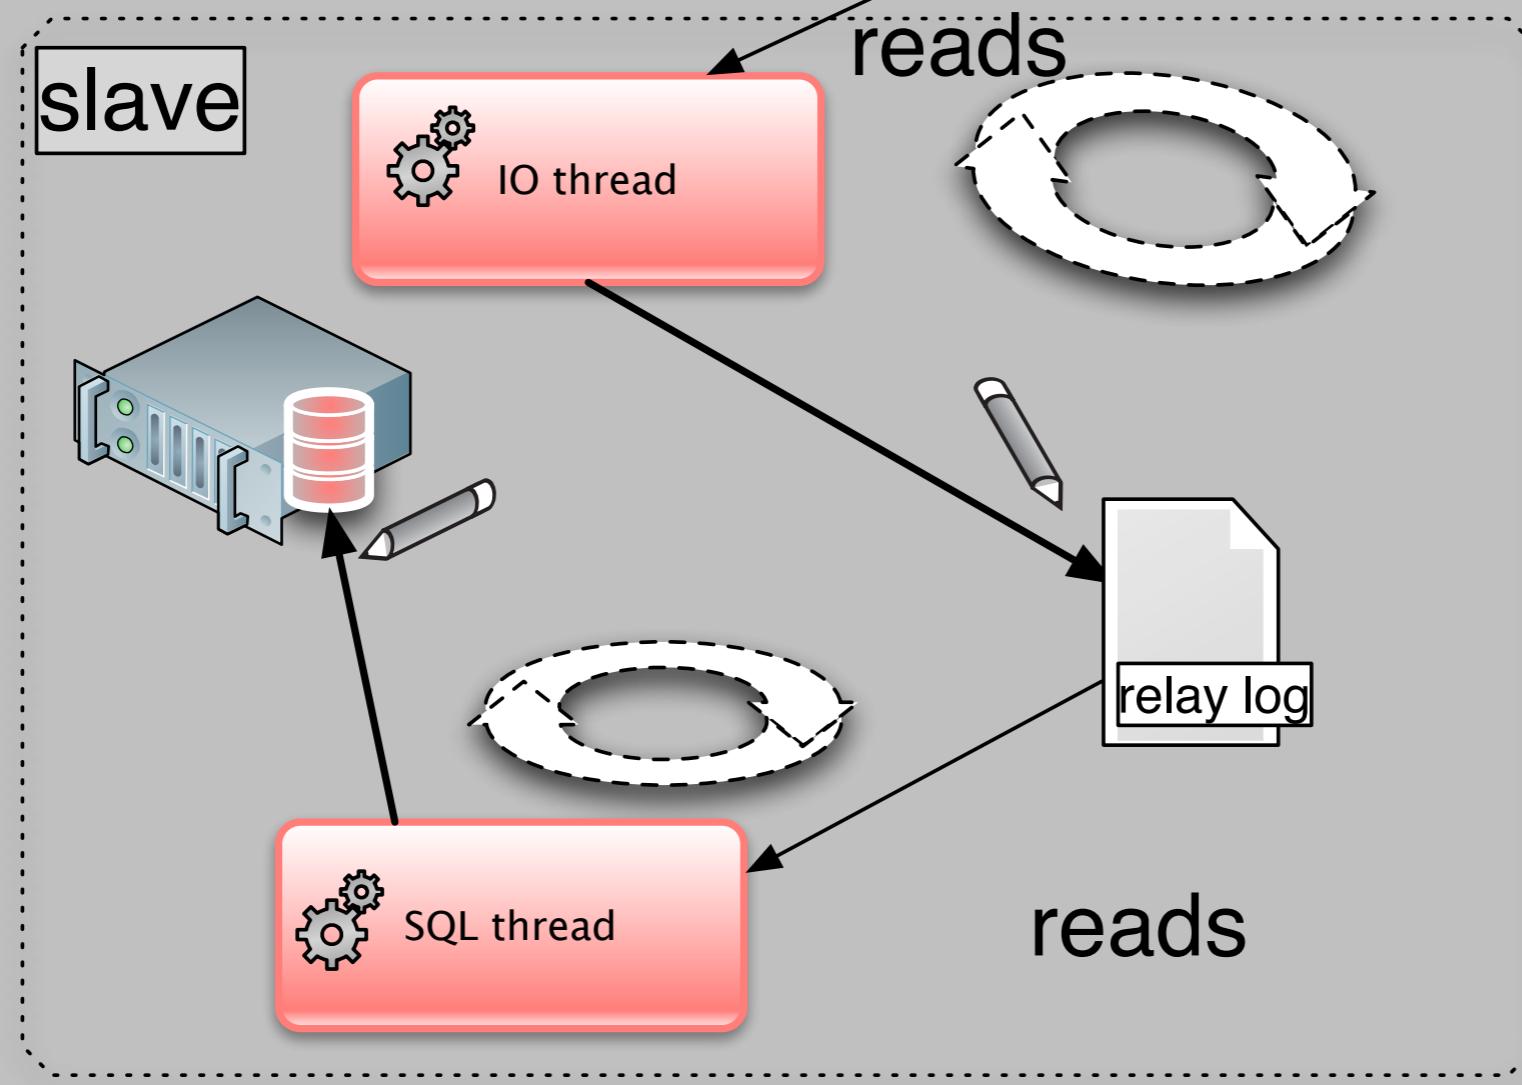
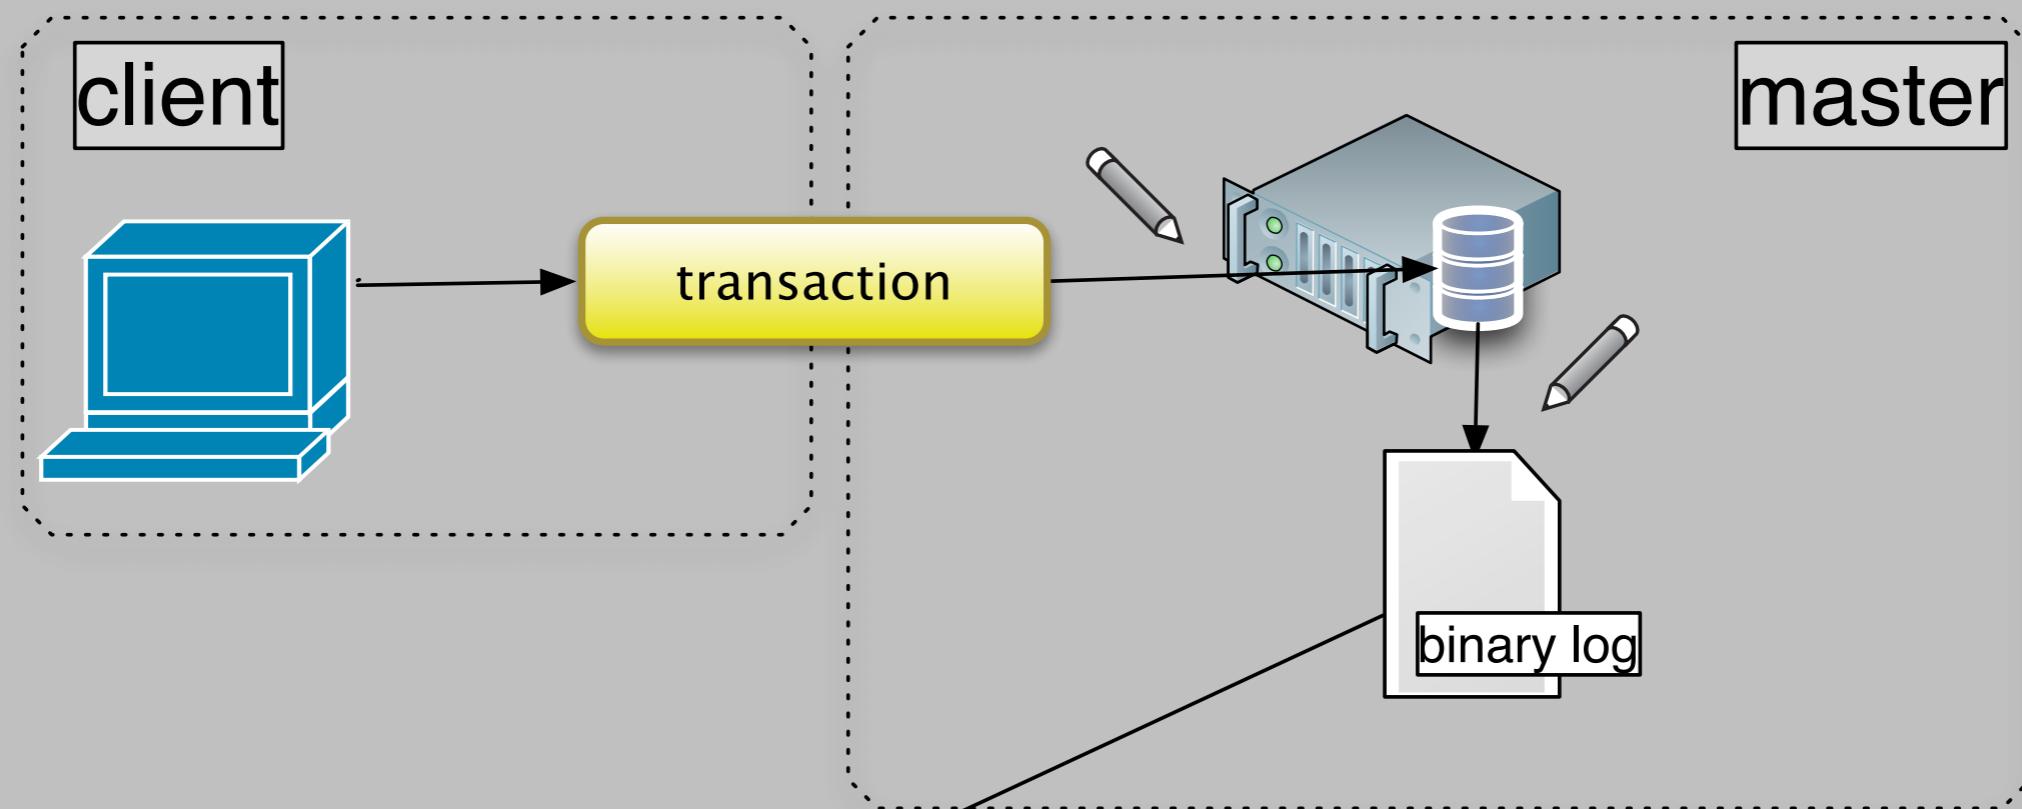
What Replication is for

- master replacement
- backup
- better reads

Gotchas, tips, and tricks

More info

- What to read
- More replication sessions



replication
concepts

AGENDA

Why replication

- The web economy
- Scaling out

How to set replication

- **From single server to replication**
- Adding a slave

Using Replication

- Binary log formats
- What gets replicated and how
- Replication awareness

Managing replication

- Monitoring
- Log management
- Replacing a slave
- Replacing a master

What Replication is for

- master replacement
- backup
- better reads

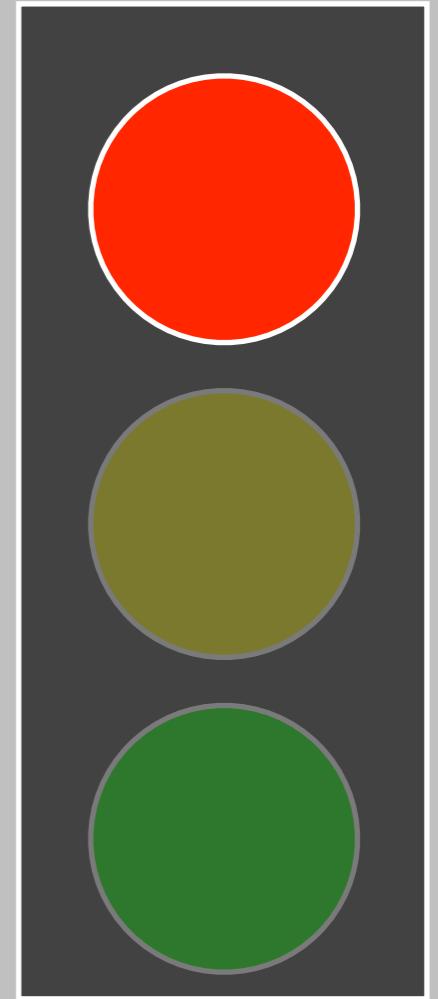
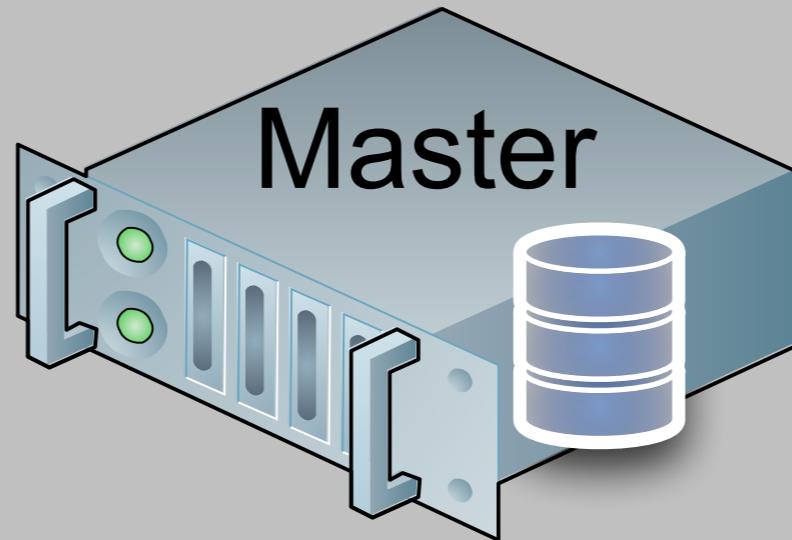
Gotchas, tips, and tricks

- What to read
- More replication sessions

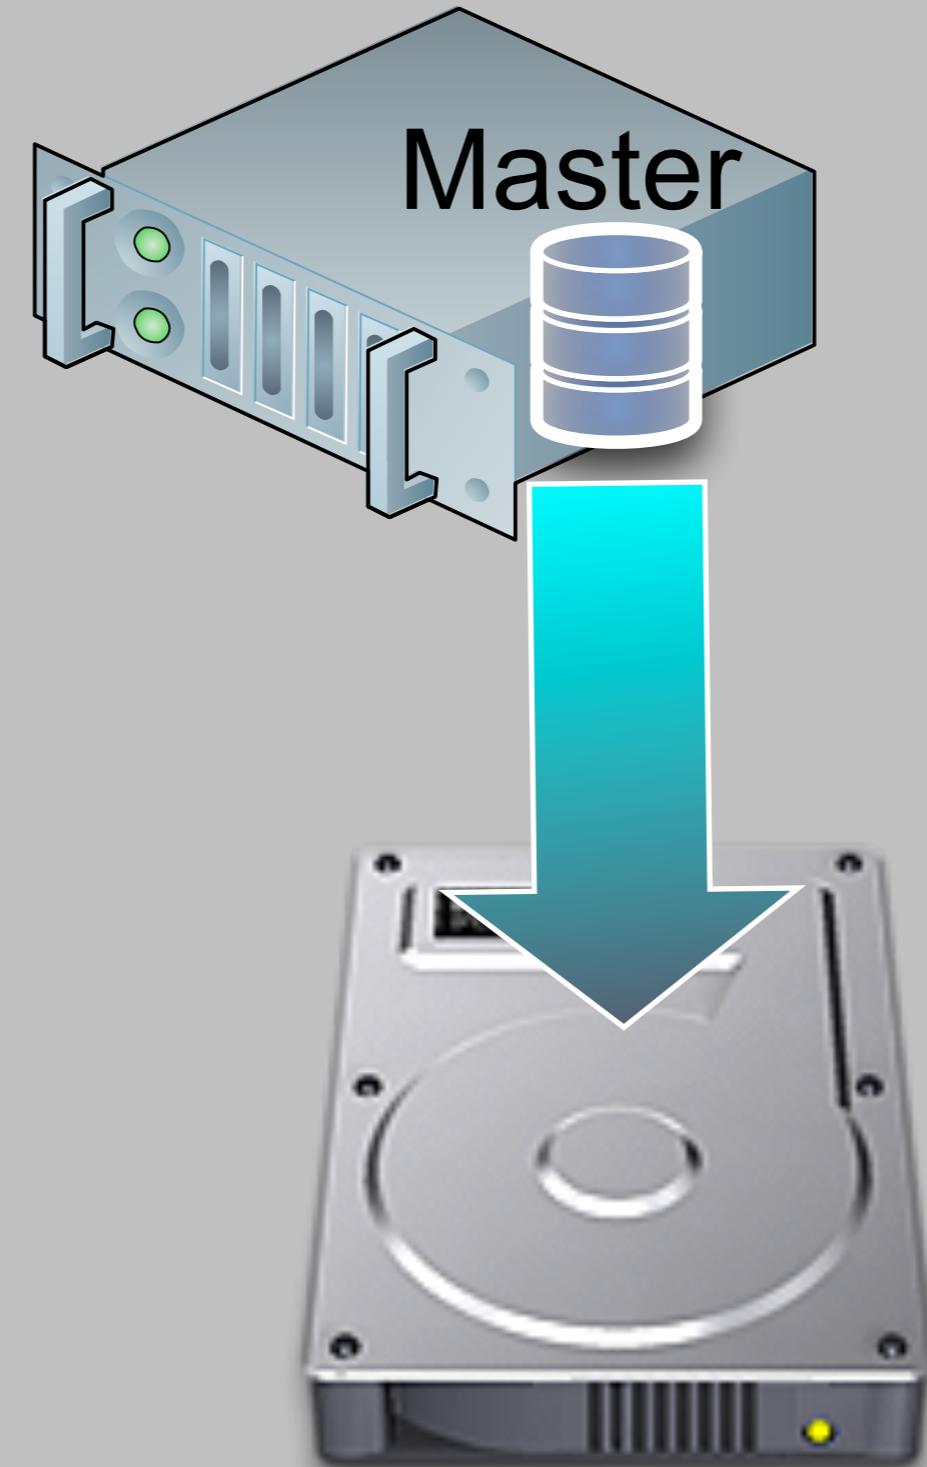
More info

1

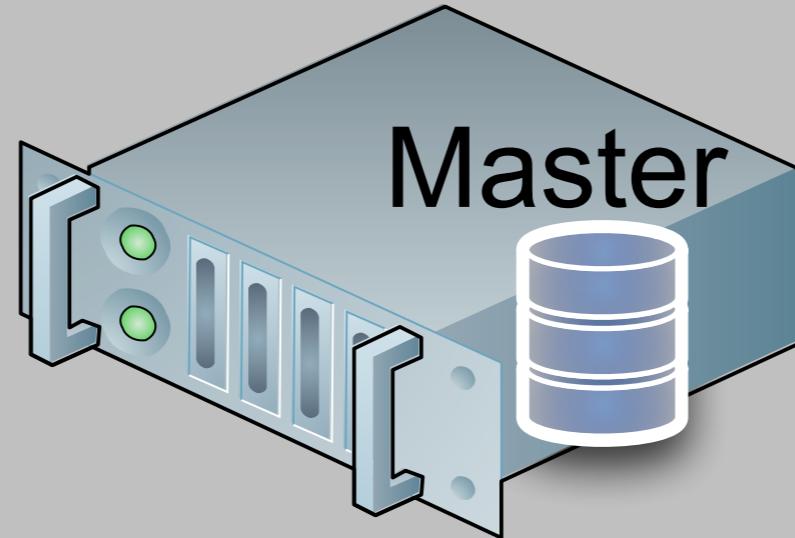
SHUT DOWN THE DATABASE SERVER



MAKE A BACKUP COPY



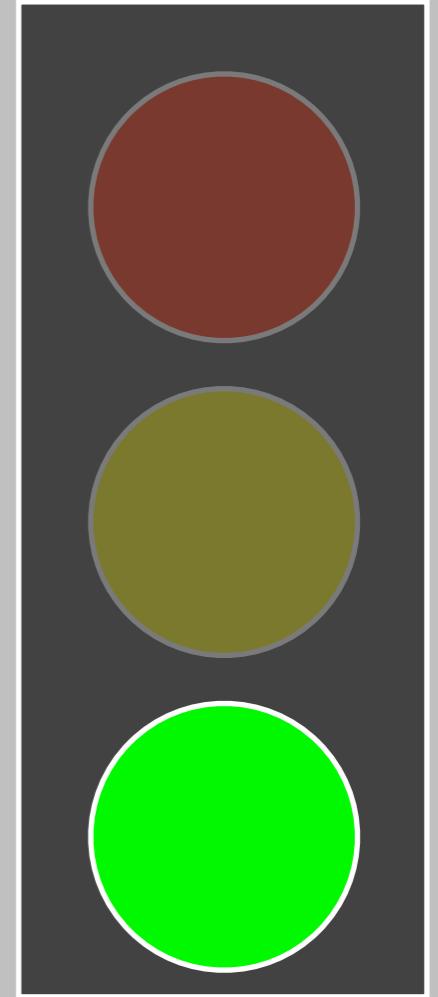
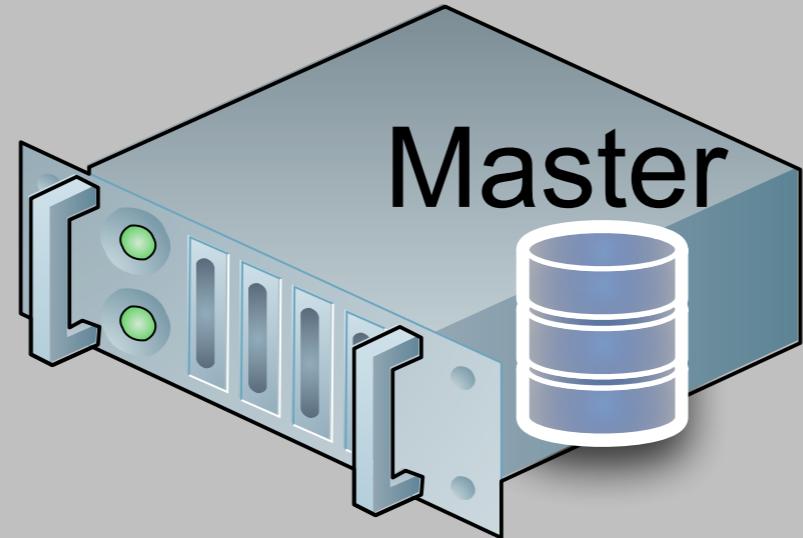
ENABLE THE MASTER



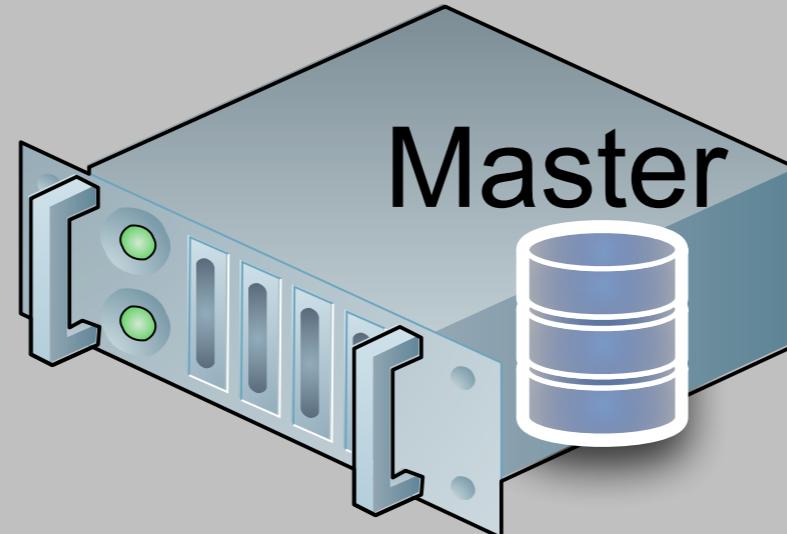
Configuration file

```
[mysqld]
log-bin=mysql-bin
server-id=1
```

RESTART THE MASTER



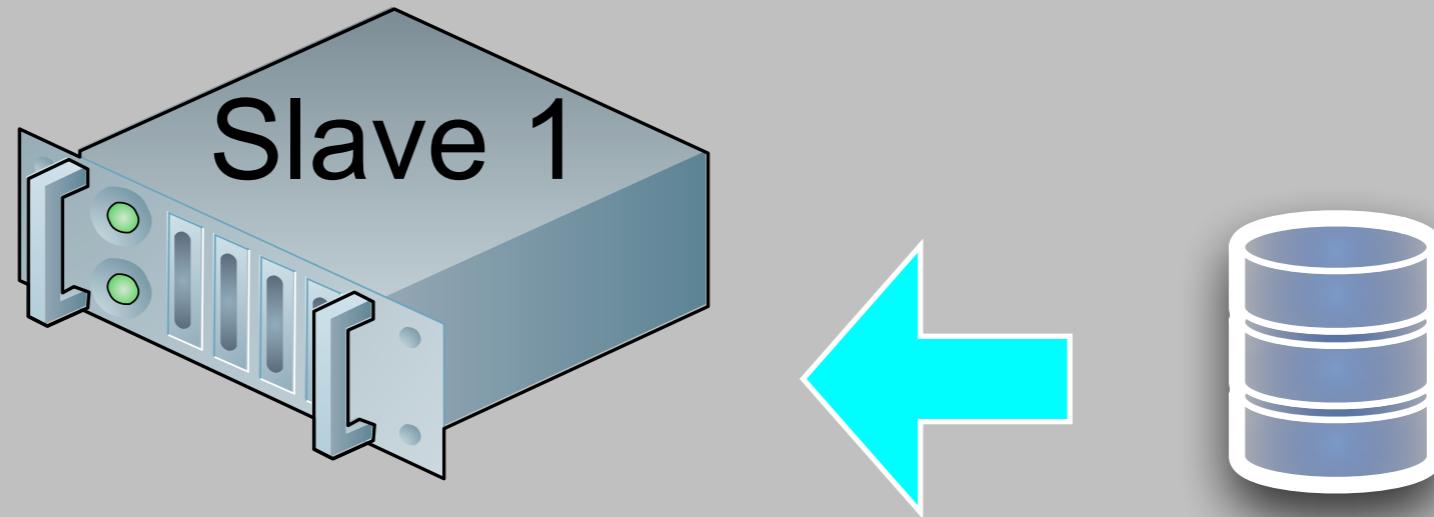
CREATE REPLICATION USER



SQL command

```
GRANT REPLICATION SLAVE ON *.*  
to 'slave_user@'10.10.100.%'  
IDENTIFIED BY 'slave_pass';
```

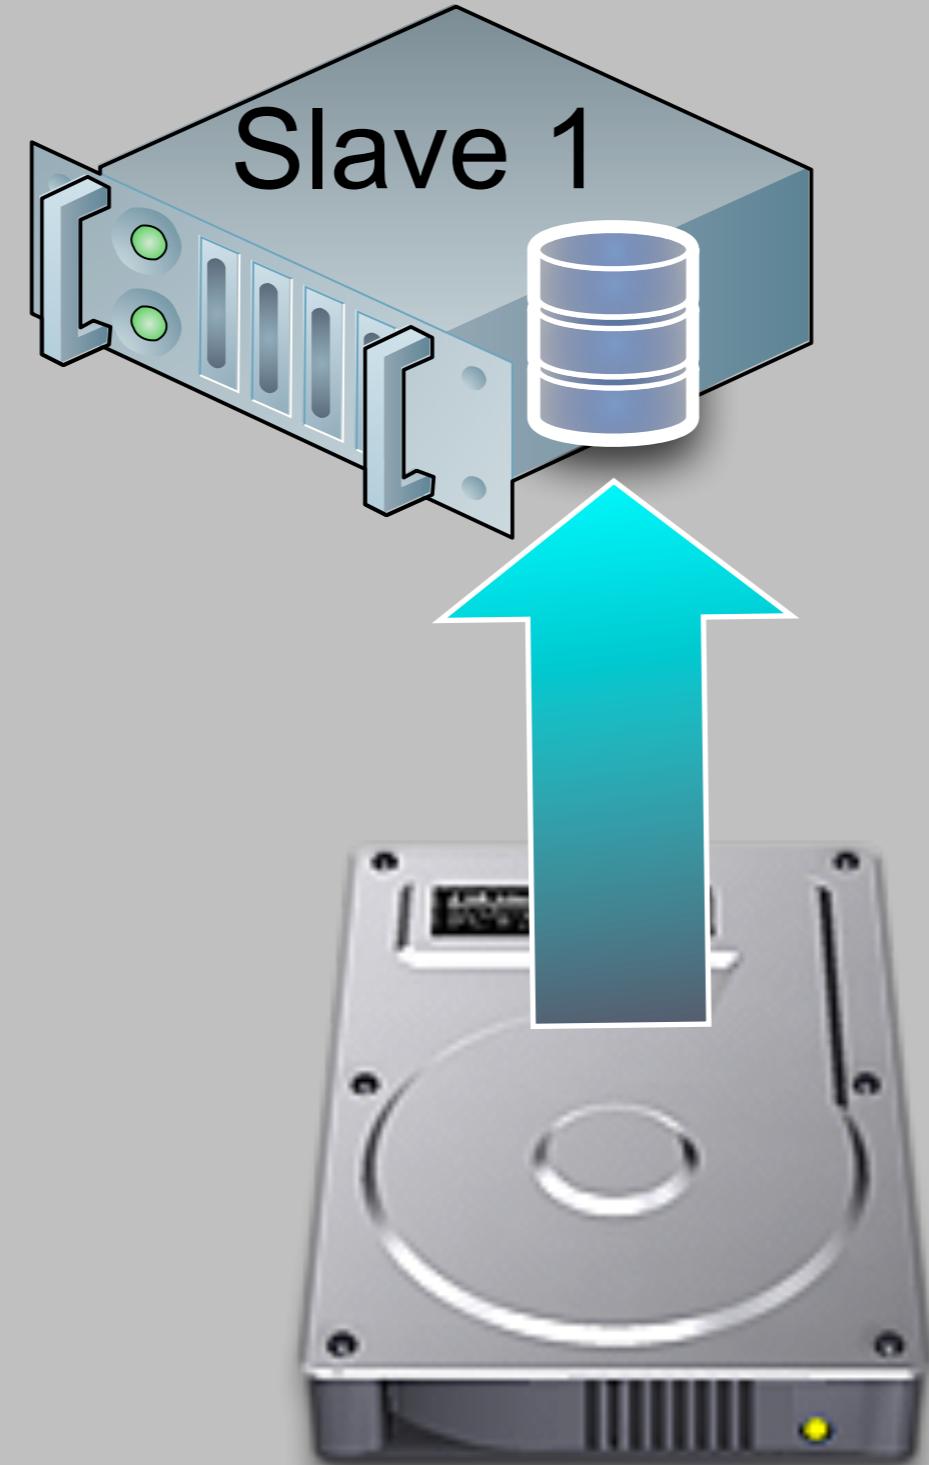
INSTALL MySQL on the slave



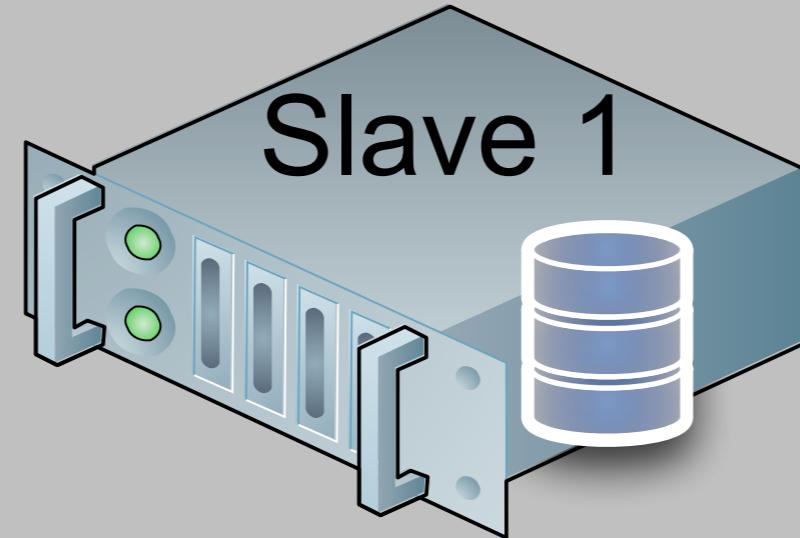
Make sure that:

- You're using the same version of MySQL
- You have the same directory structure
- The server is not started yet

COPY THE MASTER DATA to the slave



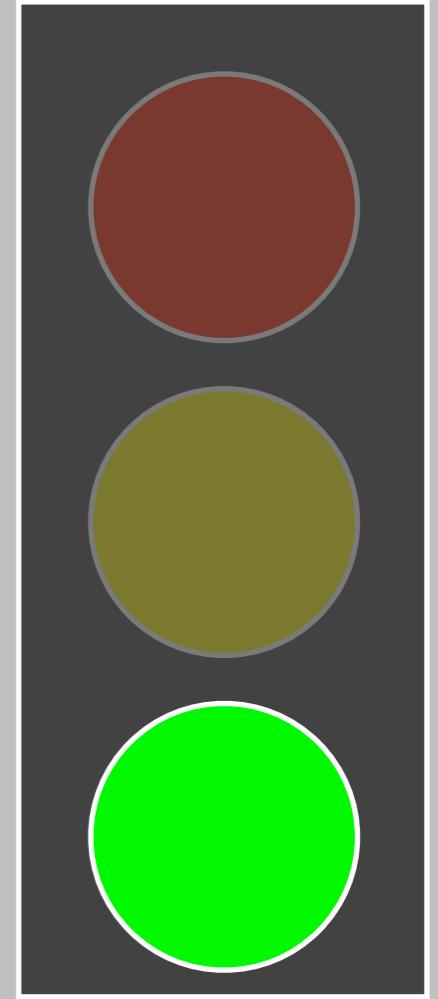
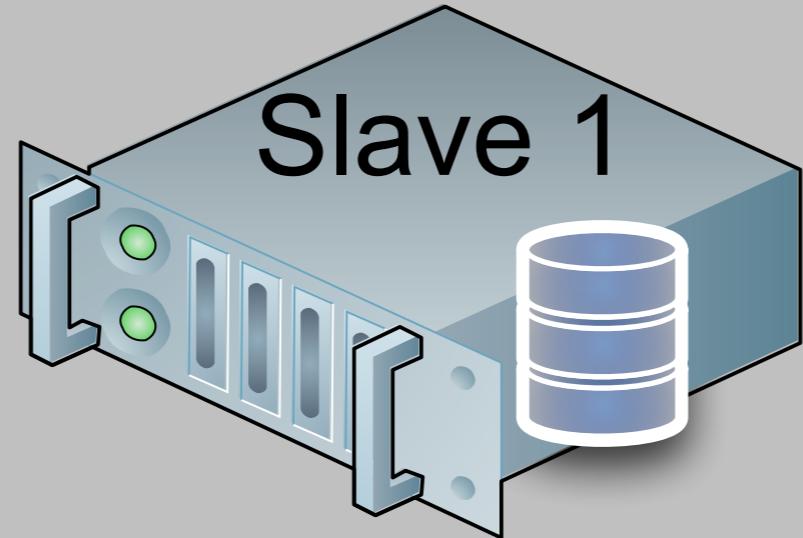
ENABLE THE SLAVE



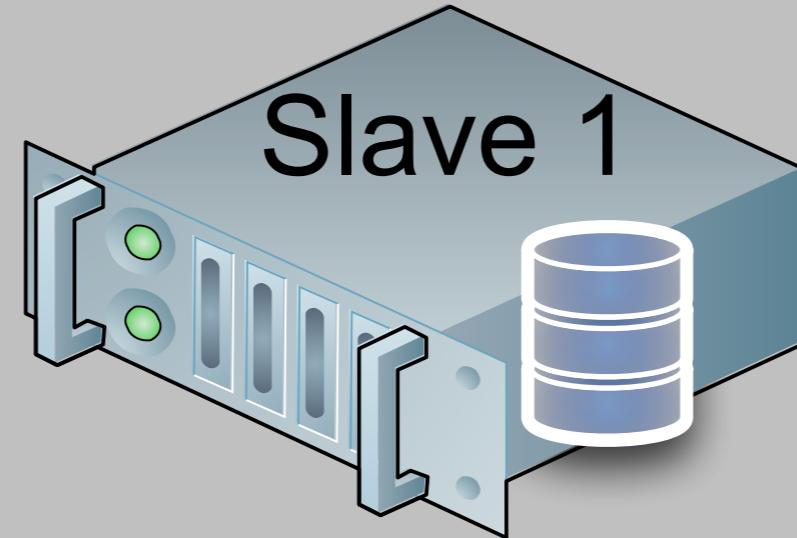
Configuration file

```
[mysqld]
server-id=2
relay-log=mysql-relay
read-only
# optional:
log-bin=mysql-bin
```

START THE SLAVE SERVER



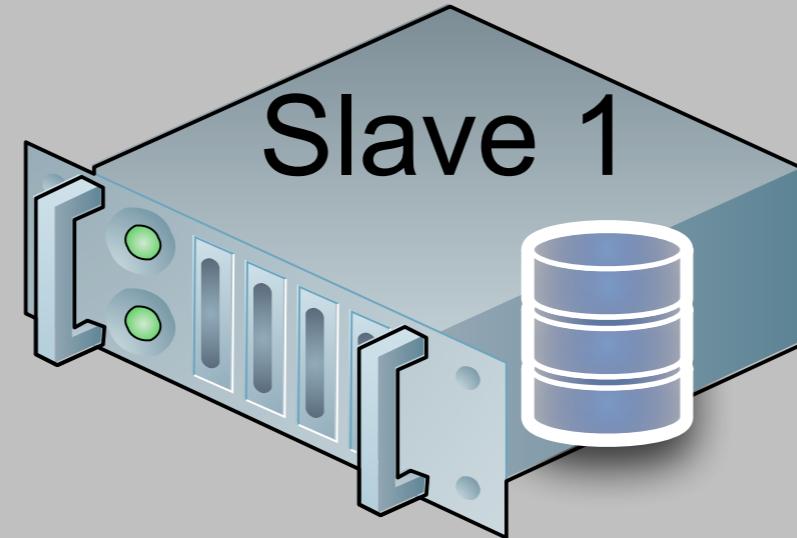
INITIALIZE THE SLAVE



SQL command

```
SET MASTER TO
MASTER_HOST=master_IP,
MASTER_PORT=3306,
MASTER_USER=slave_user,
MASTER_PASSWORD='slave_pwd';
```

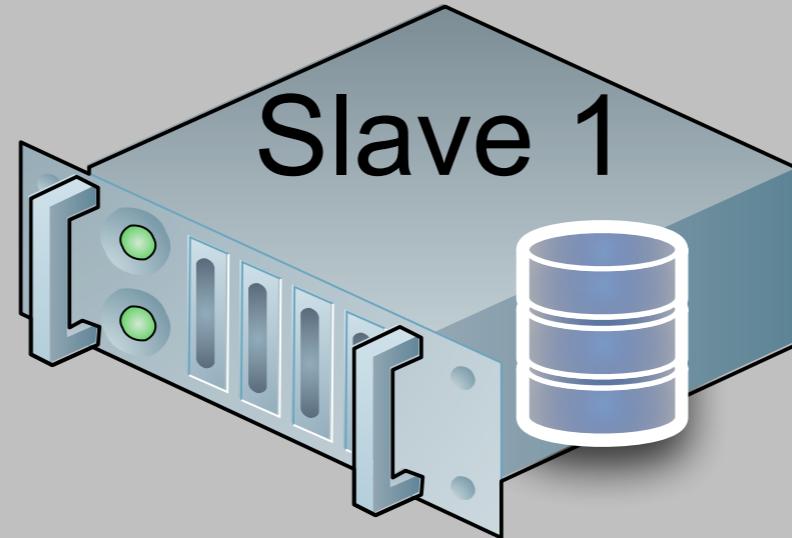
START THE SLAVE SERVICE



SQL command

```
START SLAVE ;
```

CHECK THE SLAVE



SQL command

```
SHOW SLAVE STATUS \G
```

...

```
Slave_IO_Running: Yes
Slave_SQL_Running: Yes
```

...

Troubleshooting

- **SHOW SLAVE STATUS says
SLAVE_IO_RUNNING=No**
 - Make sure that the slave host can connect to the master
 - Make sure that master and slave have different Server-id
 - Check the error log of both master and slave

Testing the slave

- Create a table in the master.
- Make sure that the slave has replicated the table.
- Insert data in the master
- read that data in the slave

What if the master was already logging?

- You have two options:
 - Physical copy
 - Logical copy

Physical copy

- stop master
- make copy
- remove binary log files and index
- start master
- (alternative: use xtrabackup)

Logical copy

- mysqldump --all-databases --master-data

Common replication commands

- CHANGE MASTER TO
- SHOW MASTER STATUS
- SHOW SLAVE STATUS
- START SLAVE
- STOP SLAVE
- RESET MASTER (caution!)
- RESET SLAVE (caution!)

CHANGE MASTER TO basic syntax

```
STOP SLAVE;
```

```
CHANGE MASTER TO
MASTER_HOST='hostname', # or IP
MASTER_PORT=3306,
MASTER_USER='slaveuser',
MASTER_PASSWORD='slavepassword',
MASTER_LOG_FILE='filename', # default: first binlog
MASTER_LOG_POS=123456; # default: from the beginning
```

```
START SLAVE;
```

AGENDA

Why replication

- The web economy
- Scaling out

How to set replication

- From single server to replication
- **Adding a slave**

Using Replication

- Binary log formats
- What gets replicated and how
- Replication awareness

Managing replication

- Monitoring
- Log management
- Replacing a slave
- Replacing a master

What Replication is for

- master replacement
- backup
- better reads

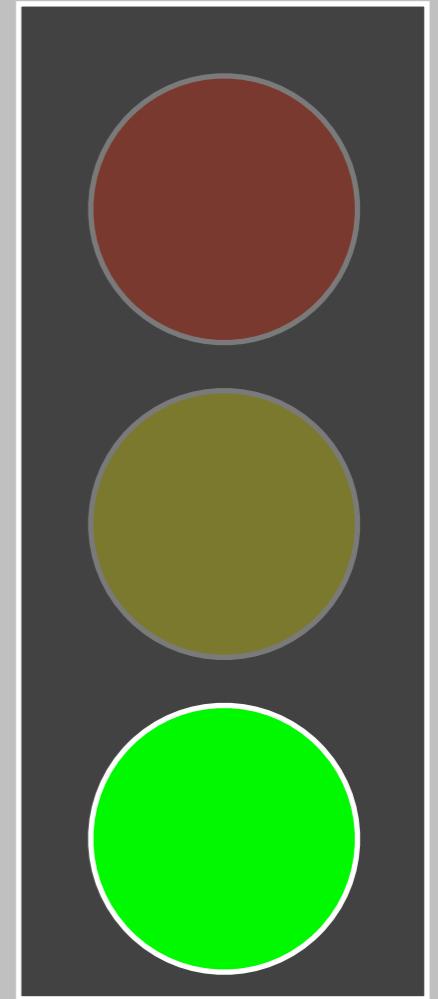
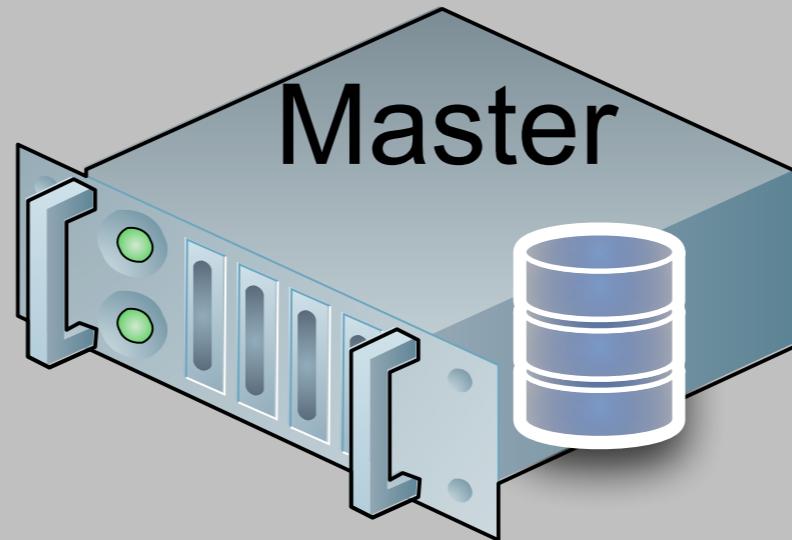
Gotchas, tips, and tricks

- What to read
- More replication sessions

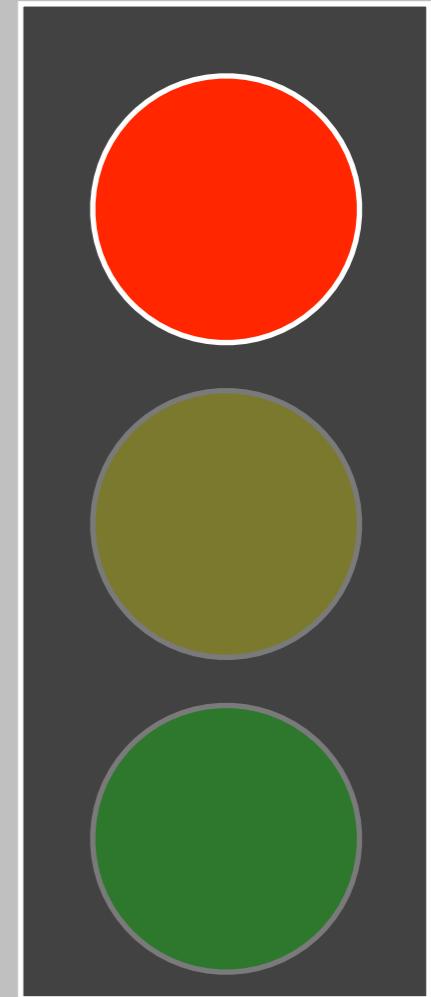
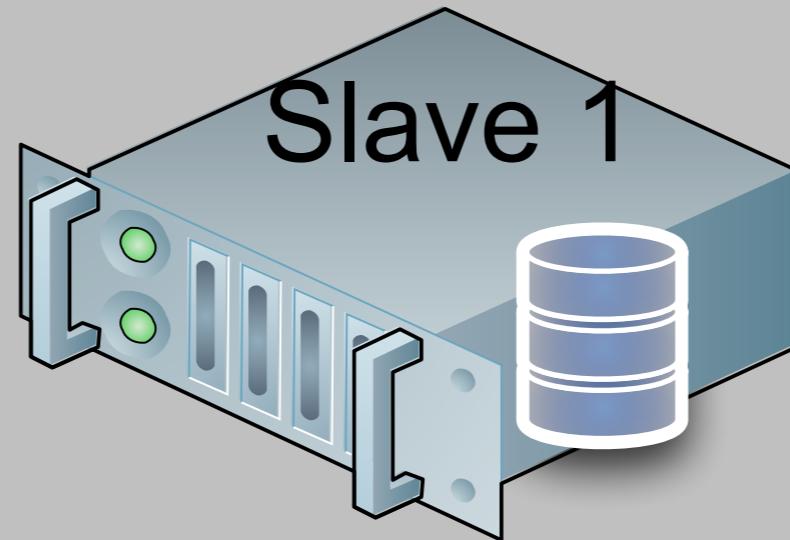
More info

1

NO NEED TO STOP THE MASTER!



STOP THE SLAVE

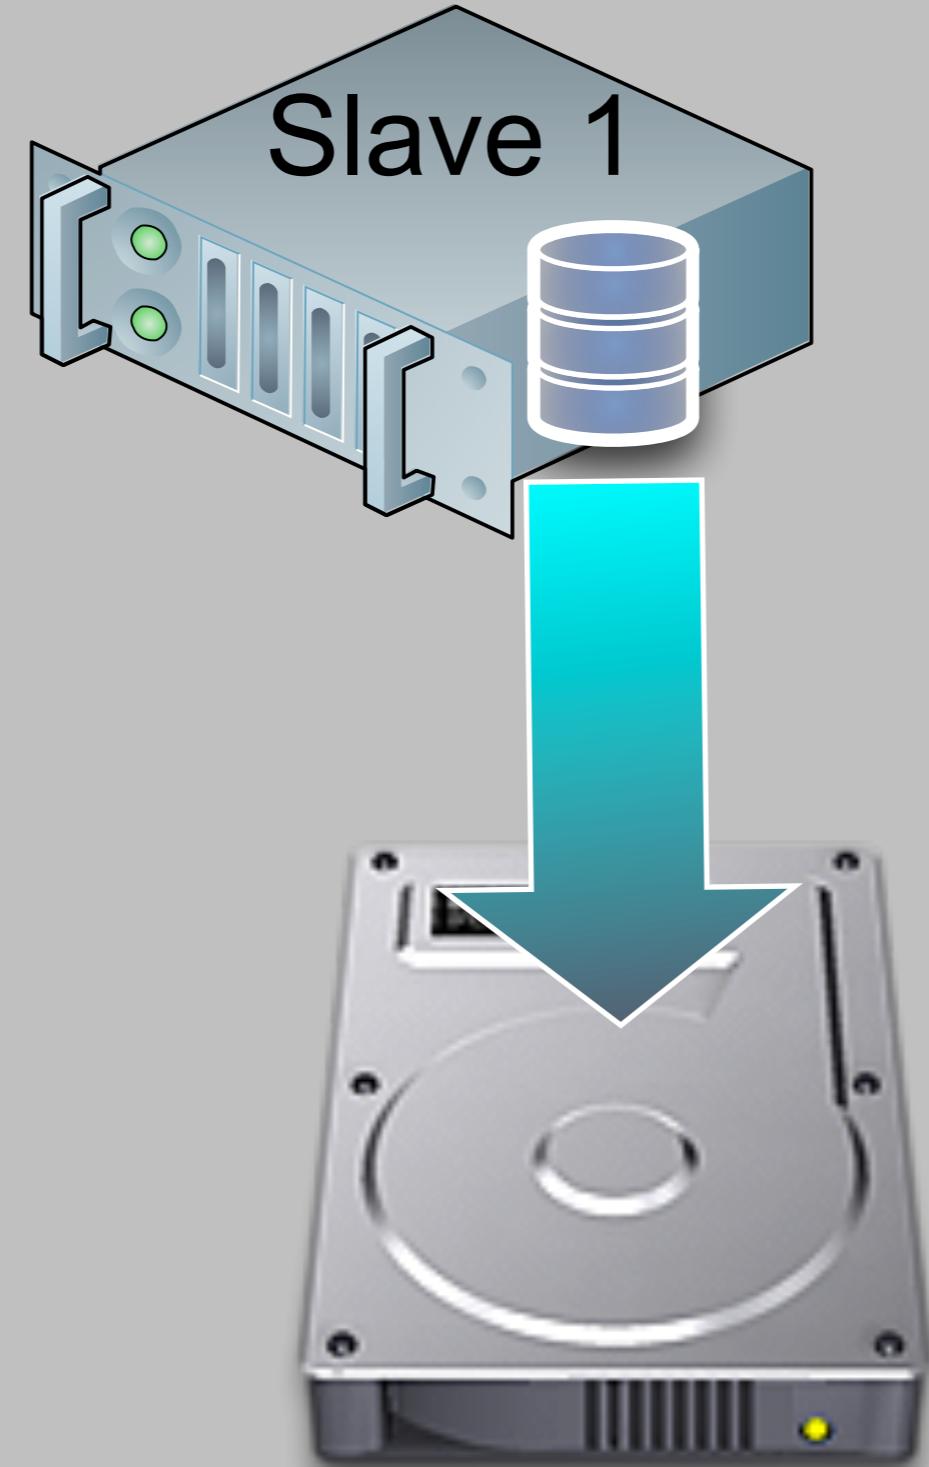


SQL command

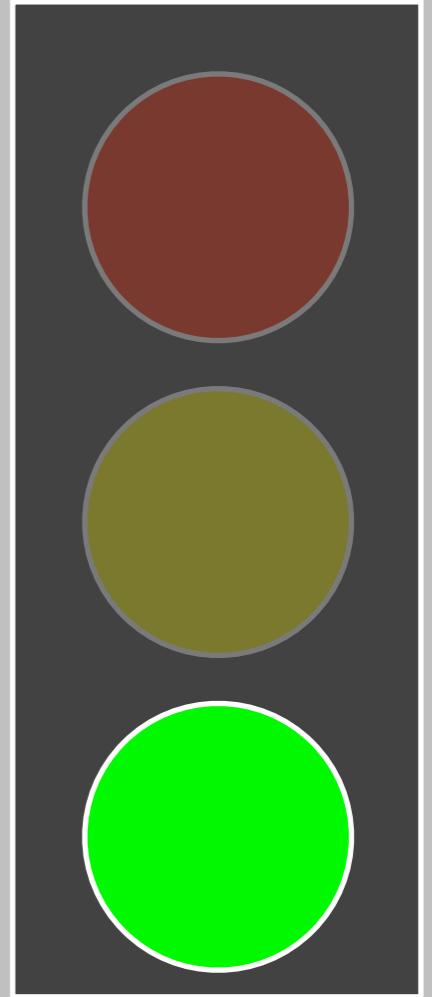
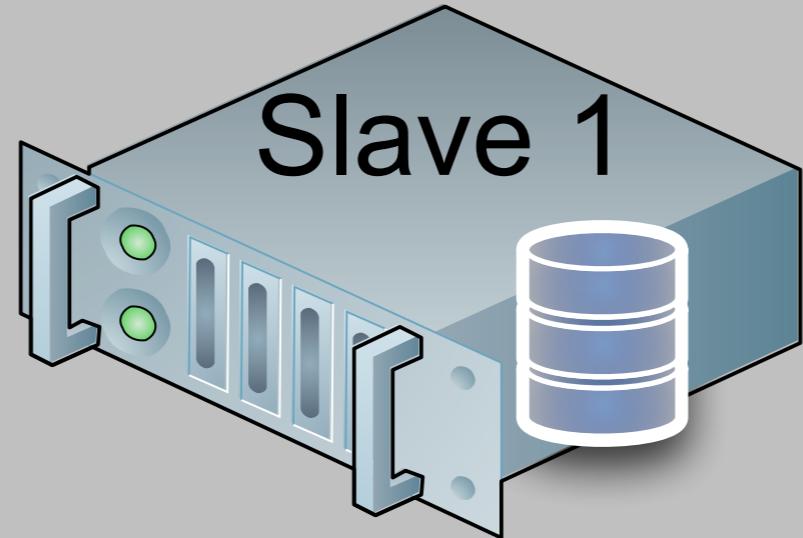
```
STOP SLAVE IO_THREAD;  
# wait until the SQL_THREAD  
# has done everything  
STOP SLAVE SQL_THREAD;  
# STOP THE SERVER
```

3

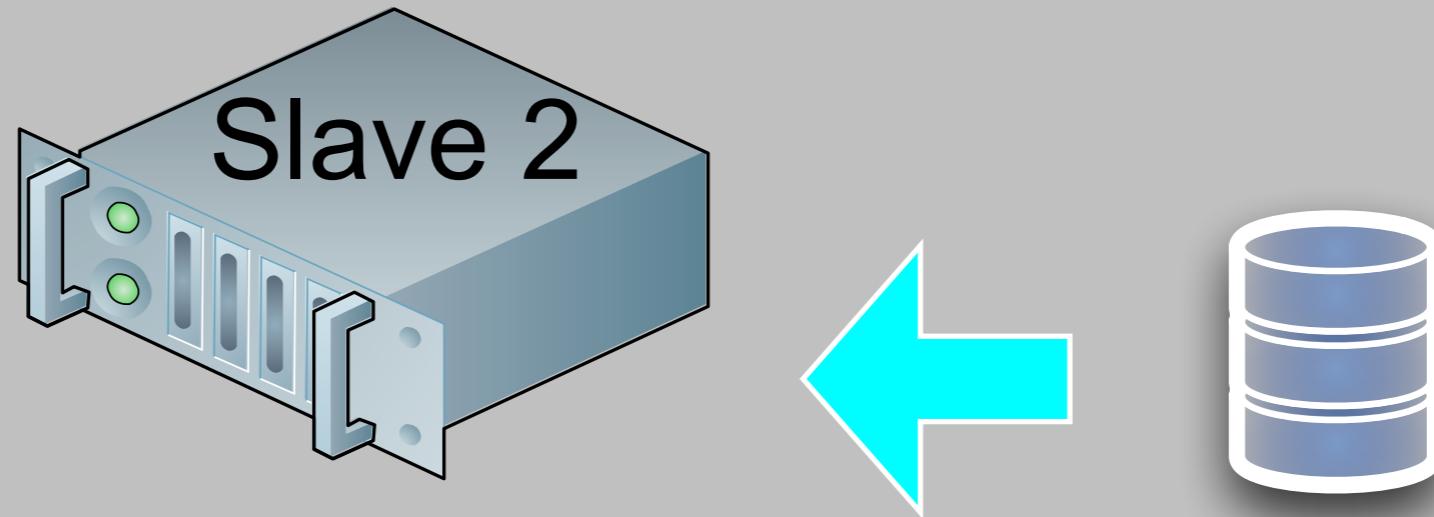
MAKE A COPY OF THE DATA DIRECTORY



RESTART THE SLAVE



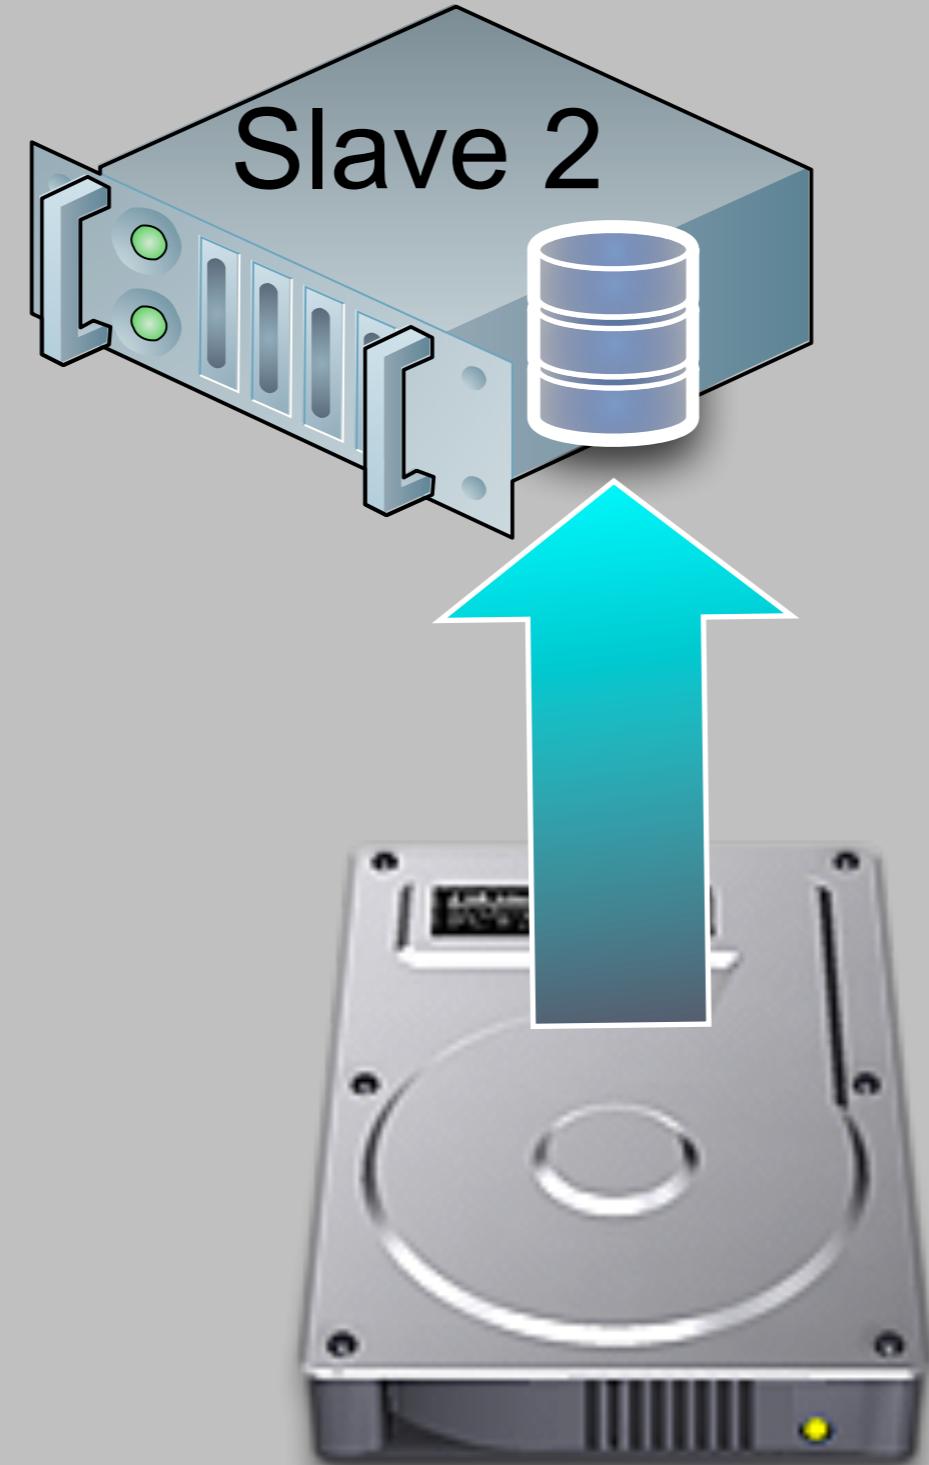
INSTALL MySQL on the new slave



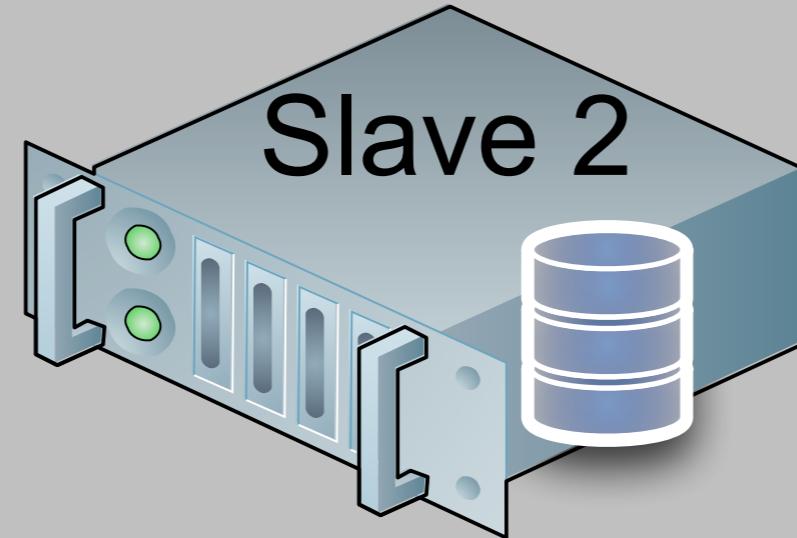
Make sure that:

- You're using the same version of MySQL
- You have the same directory structure
- The server is not started yet

COPY THE old slave DATA on the slave



ENABLE THE NEW SLAVE

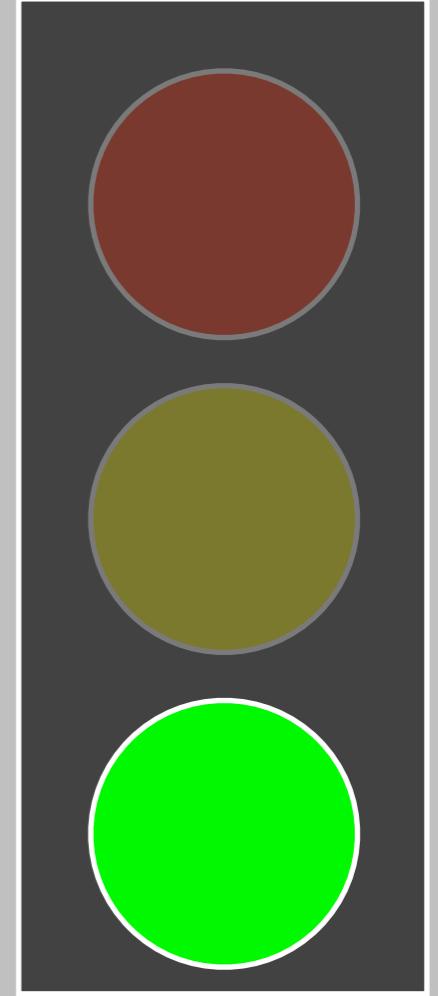
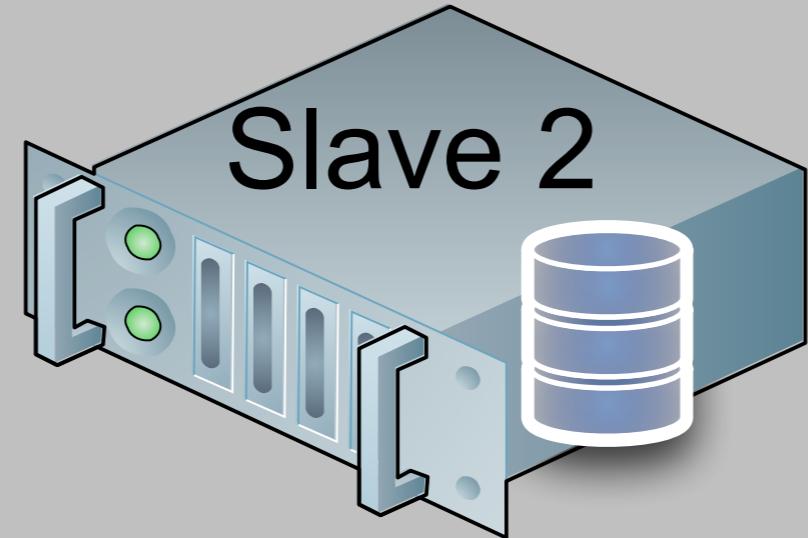


Configuration file

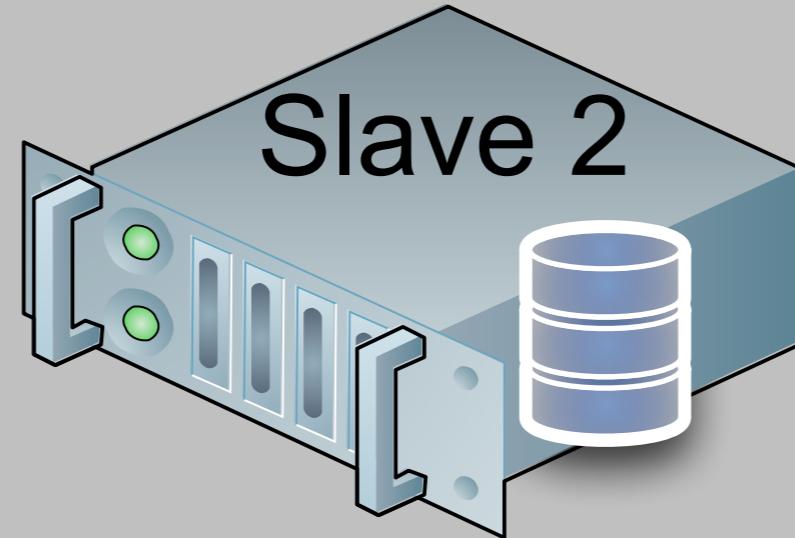
```
[mysqld]
server-id=3
relay-log=mysql-relay
read-only
# optional:
log-bin=mysql-bin
```

must be unique!

START THE NEW SLAVE



CHECK THE SLAVE



SQL command

```
SHOW SLAVE STATUS \G
```

...

```
Slave_IO_Running: Yes
Slave_SQL_Running: Yes
```

...

Why it works

- No need to issue a CHANGE MASTER TO command.
- Because we cloned the old slave
- The new slave gets its parameters from the .info files in the data directory

Starting and stopping replication

Starting and stopping replication

START SLAVE;

Starting and stopping replication

START SLAVE;

START SLAVE SQL_THREAD;

Starting and stopping replication

```
START SLAVE;
```

```
START SLAVE SQL_THREAD;
```

```
START SLAVE IO_THREAD;
```

Starting and stopping replication

```
START SLAVE;  
START SLAVE SQL_THREAD;  
START SLAVE IO_THREAD;  
START SLAVE UNTIL MASTER_LOG_FILE='filename',  
MASTER_LOG_POS=xxxx;
```

Starting and stopping replication

```
START SLAVE;  
START SLAVE SQL_THREAD;  
START SLAVE IO_THREAD;  
START SLAVE UNTIL MASTER_LOG_FILE='filename',  
MASTER_LOG_POS=xxxx;  
STOP SLAVE;
```

Starting and stopping replication

```
START SLAVE;  
START SLAVE SQL_THREAD;  
START SLAVE IO_THREAD;  
START SLAVE UNTIL MASTER_LOG_FILE='filename',  
MASTER_LOG_POS=xxxx;  
STOP SLAVE;  
STOP SLAVE SQL_THREAD;
```

Starting and stopping replication

```
START SLAVE;  
START SLAVE SQL_THREAD;  
START SLAVE IO_THREAD;  
START SLAVE UNTIL MASTER_LOG_FILE='filename',  
MASTER_LOG_POS=xxxx;  
STOP SLAVE;  
STOP SLAVE SQL_THREAD;  
STOP SLAVE IO_THREAD;
```

Hijacking the agenda

- MySQL Sandbox

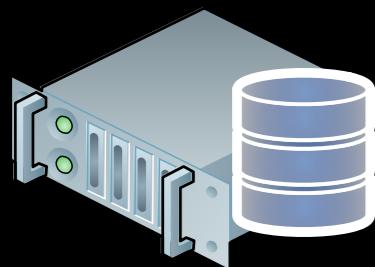
MySQL Sandbox

<http://mysqlsandbox.net>

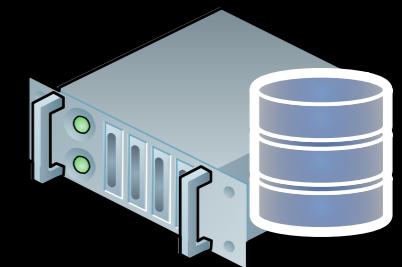
- Free software (Perl under GPL)
- One (unix) host
- Many database servers
- Single or multiple sandboxes
- Customized scripts to use the servers
- Standard or circular replication
- Installs **IN SECONDS**



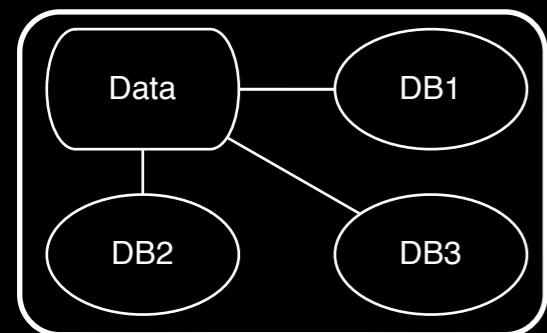
overview



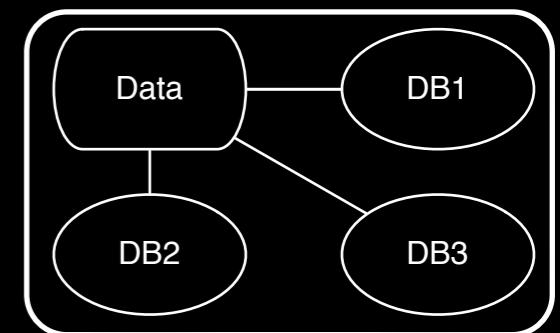
MySQL
server



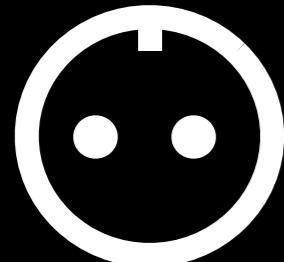
MySQL
server



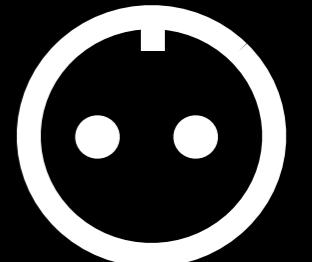
DATA DIRECTORY



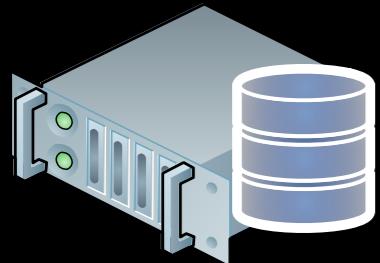
PORT



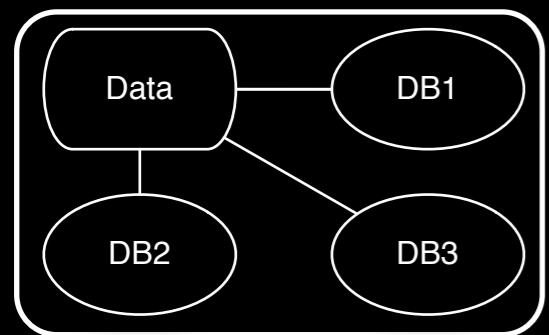
SOCKET



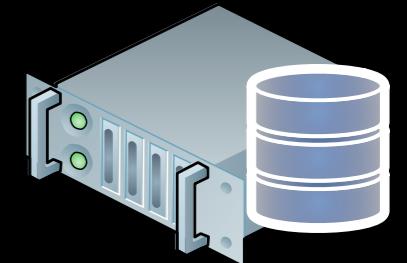
overview



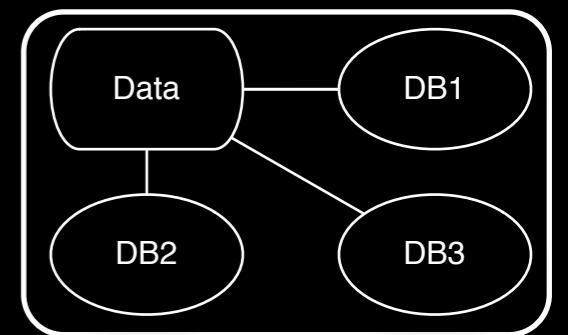
MySQL
server



`/var/lib/mysql`



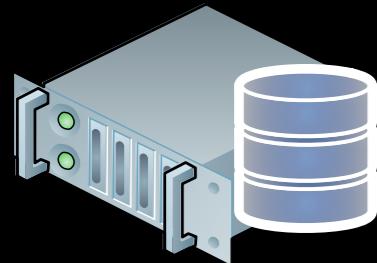
MySQL
server



`/var/lib/mysql`

DATA CORRUPTION

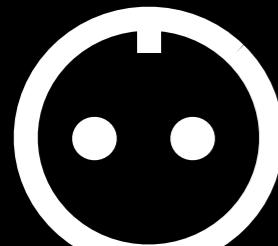
overview



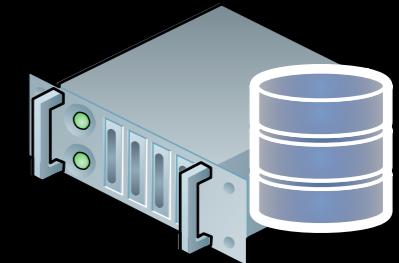
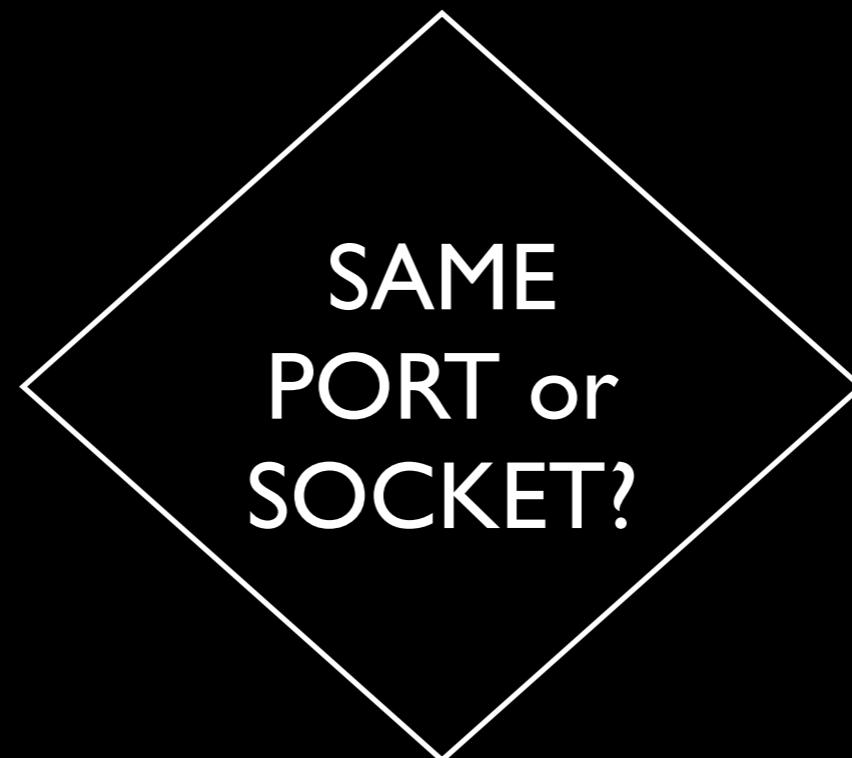
MySQL
server



3306



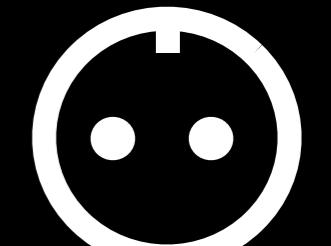
/tmp/mysql.sock



MySQL
server



3306



/tmp/mysql.sock

DOES NOT START

The hard way

The hard way

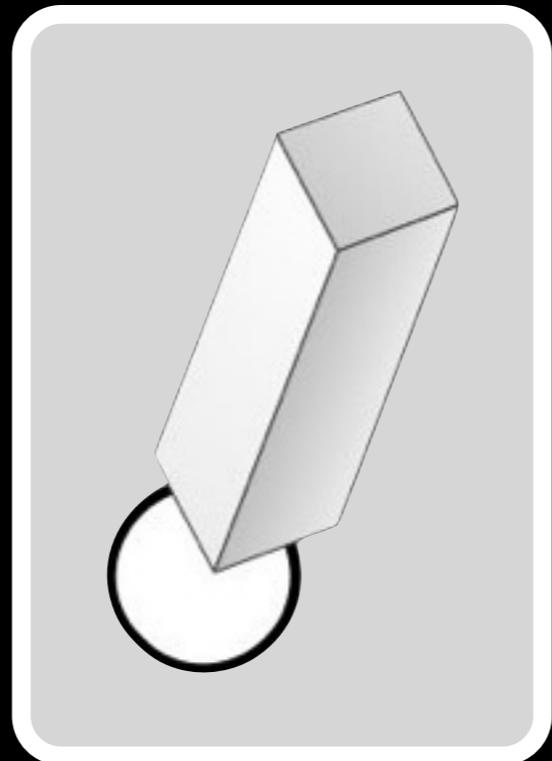
Read the manual



The hard way

Read the manual

try to figure out
what to change

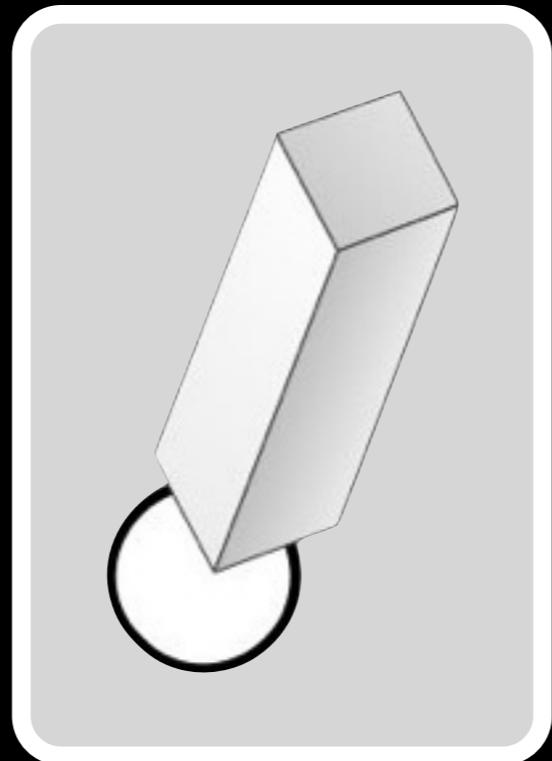
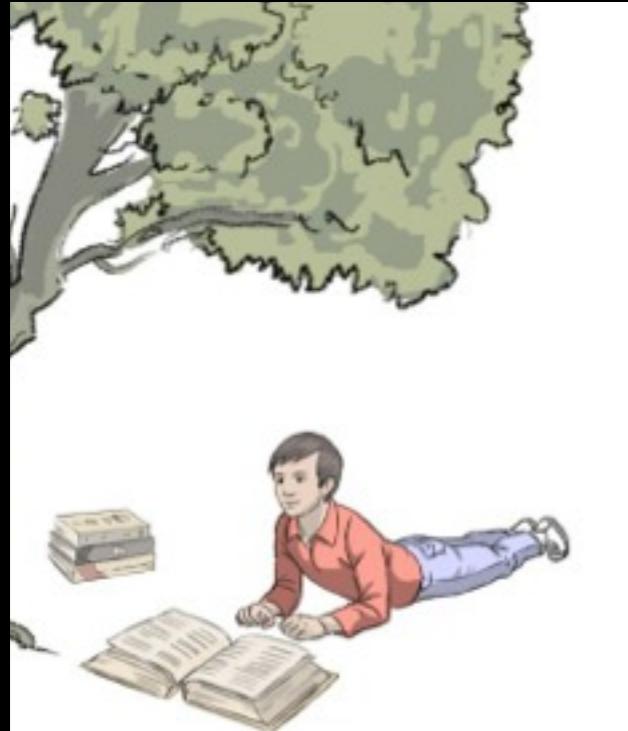


The hard way

Read the manual

try to figure out
what to change

Install



The easy way

MySQL Sandbox

```
$ make_sandbox \
  /path/to/mysql-5.1.54_linux.tar.gz

# it should work always
```

The easier way

Prepare once

Install many times

```
# some  
# preliminary  
# work
```



```
$ make_sandbox 5.1.54
```

The easiest way

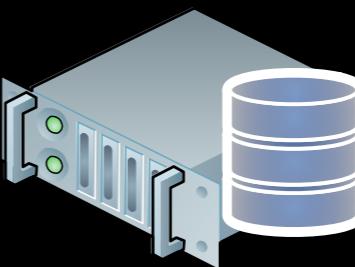
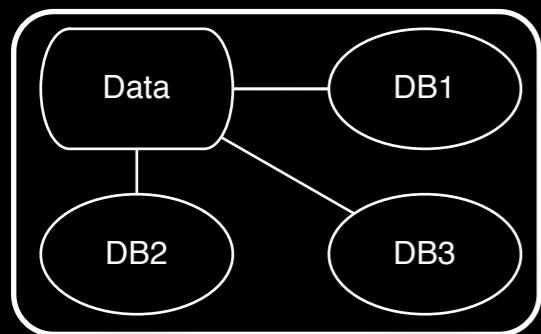
Prepare once

Install many times

```
# some  
# preliminary  
# work
```

```
$ sb 5.1.54
```

MySQL Sandbox



MySQL
server

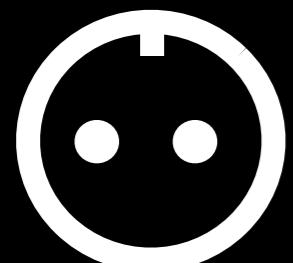
VERSION



`$SANDBOX_HOME/msb_VERSION/data`

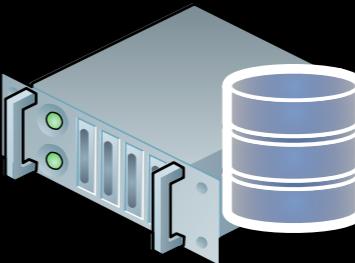
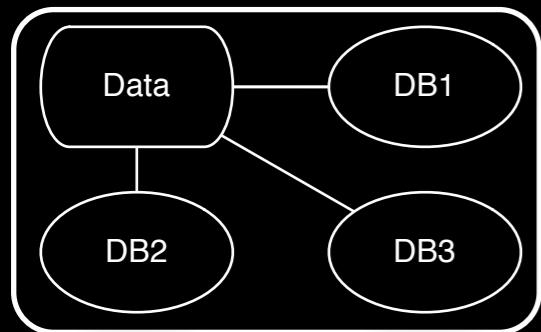


VERSION



`/tmp/mysql_VERSION.sock`

MySQL Sandbox



5.1.54

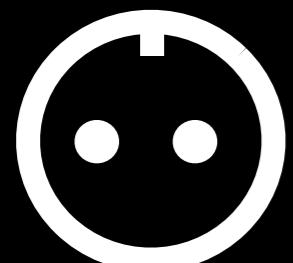
MySQL
server



`$SANDBOX_HOME/msb_5_1_54/data`

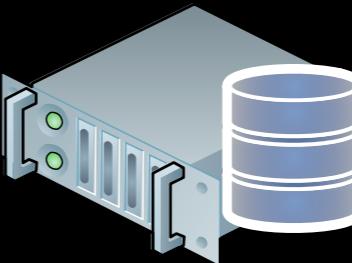
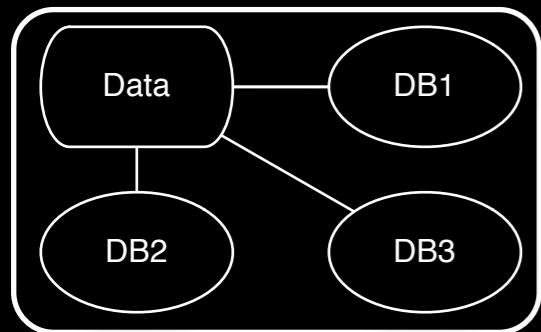


5154



`/tmp/mysql_5154.sock`

MySQL Sandbox



5.5.9

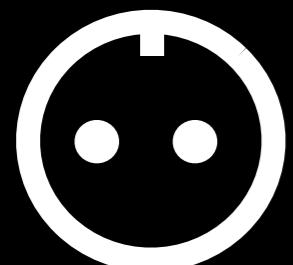
MySQL
server



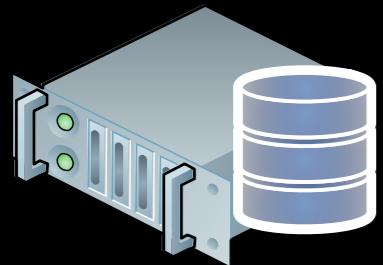
`$SANDBOX_HOME/msb_5_5_09/data`



5509



`/tmp/mysql_5509.sock`

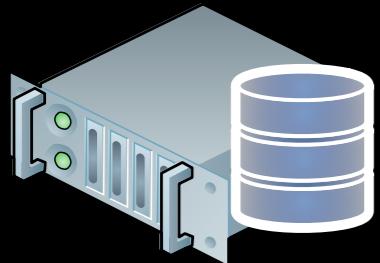


MySQL
server

Single Sandbox

customized scripts

```
start
stop
restart
status
clear
send_kill
use
```



MySQL
server

Multiple Sandbox

customized scripts

```
start_all
stop_all
restart_all
status_all
clear_all
send_kill_all
use_all
```

m	n1
s1	n2
s2	n3



Where do you get it

- from CPAN

```
sudo su -
```

```
cpan MySQL::Sandbox
```

- from launchpad

```
http://launchpad.net/mysql-sandbox
```

The easy replication way

MySQL Sandbox

```
$ make_replication_sandbox \
  /path/to/mysql-5.1.54_linux.tar.gz
```

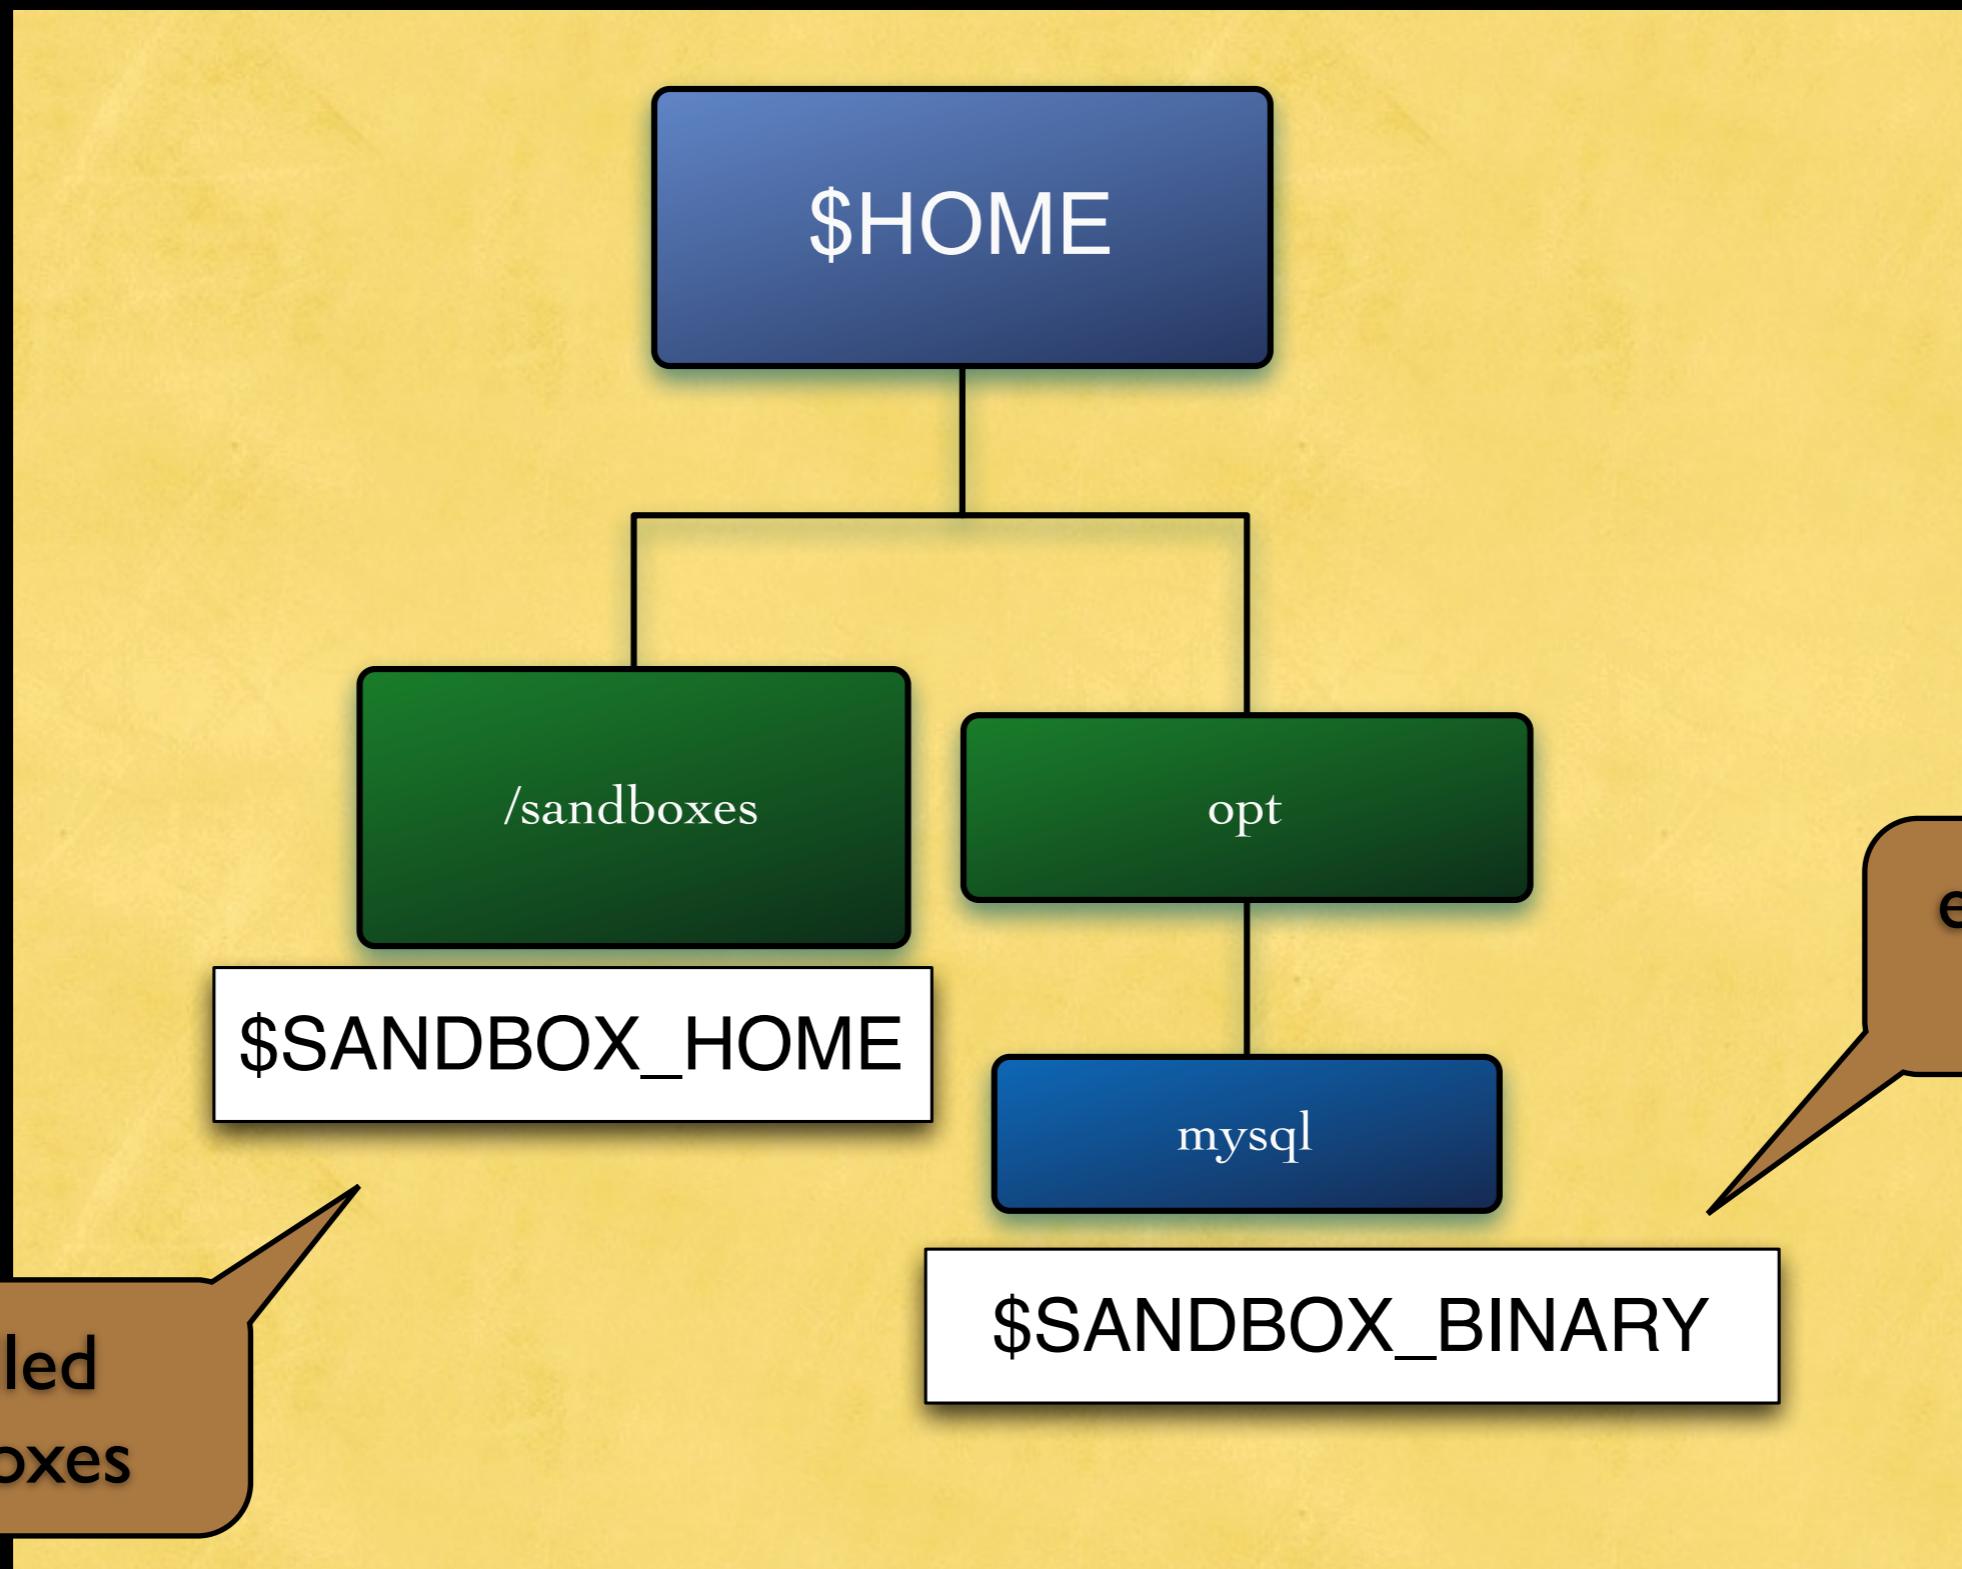
Prepare once

Install many times

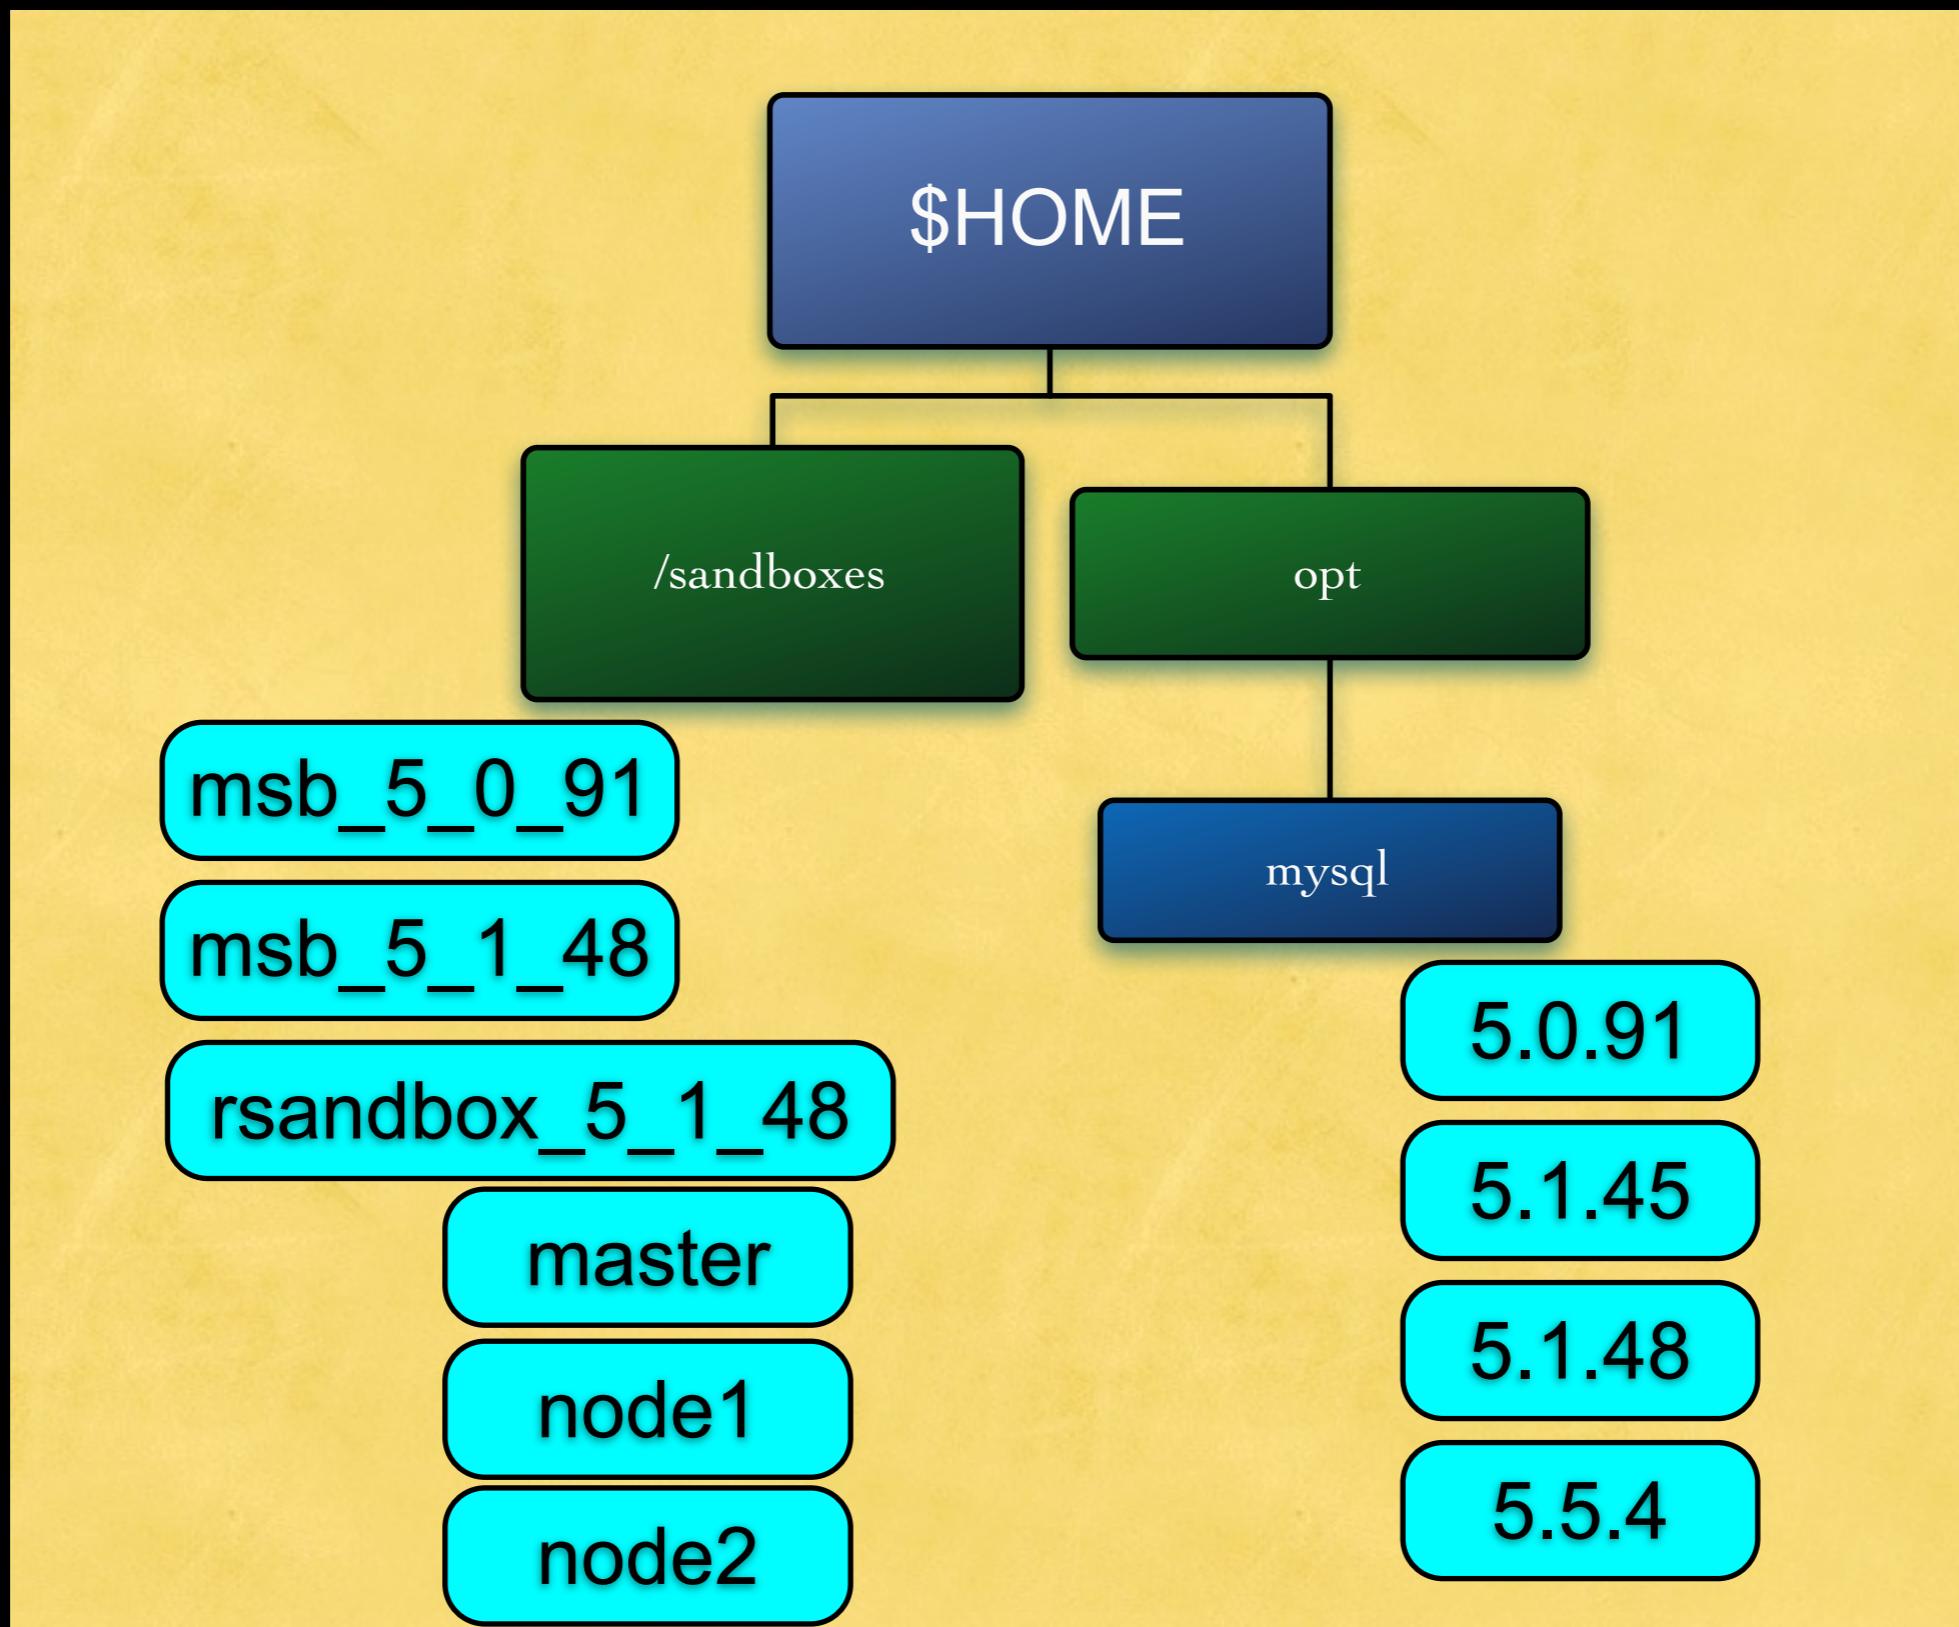
```
# some
# preparation
```

```
$ make_replication_sandbox
  5.1.54
```

default architecture



default architecture





creating a single sandbox

```
make_sandbox \  
/path/to/mysql-X.X.XX-OS.tar.gz
```

using a single sandbox

```
# after
# make_sandbox \
#   /path/to/mysql-X.X.XX-OS.tar.gz

$ cd $SANDBOX_HOME/msb_X_X_XX
$ ./use
```

creating a single sandbox with a specific options file

```
make_sandbox \
  /path/to/mysql-X.X.XX-OS.tar.gz \
  --my_file=/path/to/my.cnf
```

easily create a sandbox after the first one

The long way

```
$ cd $HOME/opt/mysql      # $SANDBOX_BINARY
$ gunzip -c \
/path/to/mysql-5.1.34-osx10.5-x86.tar.gz \
| tar -xf -
$ mv mysql-5.1.34-osx10.5-x86 5.1.34
$ make sandbox 5.1.34
```

easily create a sandbox after the first
one

The short way

```
$ make_sandbox \
path/to/mysql-5.1.34-osx10.5-x86.tar.gz \
--export_binaries
```

starting a single sandbox

```
$ cd $SANDBOX_HOME/msb_x_x_xx  
$ ./start
```

starting a single sandbox with temporary options

```
$ cd $SANDBOX_HOME/msb_x_x_xx
```

```
$ ./start --option=value
```

```
$ ./restart --option=value
```

```
$ ./start --key-buffer=20000000
```

creating a sandbox with custom port and directory

```
$ make_sandbox 5.1.34 \
  --sandbox_port=7800 \
  --sandbox_directory=mickeymouse
```

creating a sandbox with automatic port checking

```
$ make_sandbox 5.1.34 --check_port
```

```
# if 5.1.34 is free
#   port=5134
#   directory=msb_5_1_34
# else
#   port=5135 (or the first free)
#   directory=msb_5_1_34_a
```

create a replication sandbox

```
$ make_replication_sandbox \
  path/to/mysql-5.1.34-osx10.5-x86.tar.gz
```

create a circular replication sandbox

```
$ make_replication_sandbox \
--circular=4 \
path/to/mysql-5.1.34-osx10.5-x86.tar.gz
```

changing port to an existing sandbox

```
$ sbtool -o port \
  -s /path/to/source/sandbox \
  --new_port=XXXX
```

installing the innodb plugin

```
$ sbtool -o plugin \
--plugin=innodb \
-s /path/to/source/sandbox
```

creating a replication sandbox with new base port

```
$ make_replication_sandbox \
  --replication_directory=newwdir \
  --check_base_port 5.0.79

# Creates a replication directory under
# $SANDBOX_HOME/newdir
# The previous one is preserved.
# No conflicts happen
```

AGENDA

How to set up replication

DEMO

AGENDA

Why replication

- The web economy
- Scaling out

How to set replication

- From single server to replication
- Adding a slave

Using Replication

- **Binary log formats**
- What gets replicated and how
- Replication awareness

Managing replication

- Monitoring
- Log management
- Replacing a slave
- Replacing a master

What Replication is for

- master replacement
- backup
- better reads

Gotchas, tips, and tricks

More info

- What to read
- More replication sessions

Binary log formats

- **Statement** based replication
 - default
 - available since 3.23
- **Row** based replication
 - introduced in 5.1
 - ROW or MIXED

Viewing the binary logs

```
# statement based replication
```

```
$ mysqlbinlog binary-log-name
```

```
# or, as a SQL command
```

```
SHOW BINLOG EVENTS IN 'binary-log-name';
```

```
# row based replication
```

```
$ mysqlbinlog --verbose \  
--base64-output=decode-rows binary-log-name
```

binary log examples

```
# statement based replication
```

```
mysqlbinlog binary-log-name
```

```
# row based replication
```

```
mysqlbinlog --verbose \  
--base64-output=decode-rows binary-log-name
```

AGENDA

Why replication

- The web economy
- Scaling out

How to set replication

- From single server to replication
- Adding a slave

Using Replication

- Binary log formats
- **What gets replicated and how**
- Replication awareness

Managing replication

- Monitoring
- Log management
- Replacing a slave
- Replacing a master

What Replication is for

- master replacement
- backup
- better reads

Gotchas, tips, and tricks

More info

- What to read
- More replication sessions

what gets replicated

What

Statement-based

row-based

what gets replicated

What	Statement-based	row-based
insert/update/delete	statement	affected records

what gets replicated

What	Statement-based	row-based
insert/update/delete	statement	affected records
schema/table/view creation/drop	statement	statement

what gets replicated

What	Statement-based	row-based
insert/update/delete	statement	affected records
schema/table/view creation/drop	statement	statement
stored routine/trigger creation/drop	statement	statement

what gets replicated

What	Statement-based	row-based
insert/update/delete	statement	affected records
schema/table/view creation/drop	statement	statement
stored routine/trigger creation/drop	statement	statement
stored procedure call	statements executed inside SP	affected records

what gets replicated

What	Statement-based	row-based
insert/update/delete	statement	affected records
schema/table/view creation/drop	statement	statement
stored routine/trigger creation/drop	statement	statement
stored procedure call	statements executed inside SP	affected records
stored function call	function call	affected records

what gets replicated

What	Statement-based	row-based
insert/update/delete	statement	affected records
schema/table/view creation/drop	statement	statement
stored routine/trigger creation/drop	statement	statement
stored procedure call	statements executed inside SP	affected records
stored function call	function call	affected records
trigger execution	none: the slave runs the trigger	affected records moved. No trigger runs

what gets replicated

What	Statement-based	row-based
insert/update/delete	statement	affected records
schema/table/view creation/drop	statement	statement
stored routine/trigger creation/drop	statement	statement
stored procedure call	statements executed inside SP	affected records
stored function call	function call	affected records
trigger execution	none: the slave runs the trigger	affected records moved. No trigger runs
event creation	slaveside disabled event	slaveside disabled event

what gets replicated

What	Statement-based	row-based
insert/update/delete	statement	affected records
schema/table/view creation/drop	statement	statement
stored routine/trigger creation/drop	statement	statement
stored procedure call	statements executed inside SP	affected records
stored function call	function call	affected records
trigger execution	none: the slave runs the trigger	affected records moved. No trigger runs
event creation	slaveside disabled event	slaveside disabled event
event execution	statement	affected records

AGENDA

Binary logs in practice

DEMO

AGENDA

Why replication

- The web economy
- Scaling out

How to set replication

- From single server to replication
- Adding a slave

Using Replication

- Binary log formats
- What gets replicated and how
- **Replication awareness**

Managing replication

- Monitoring
- Log management
- Replacing a slave
- Replacing a master

What Replication is for

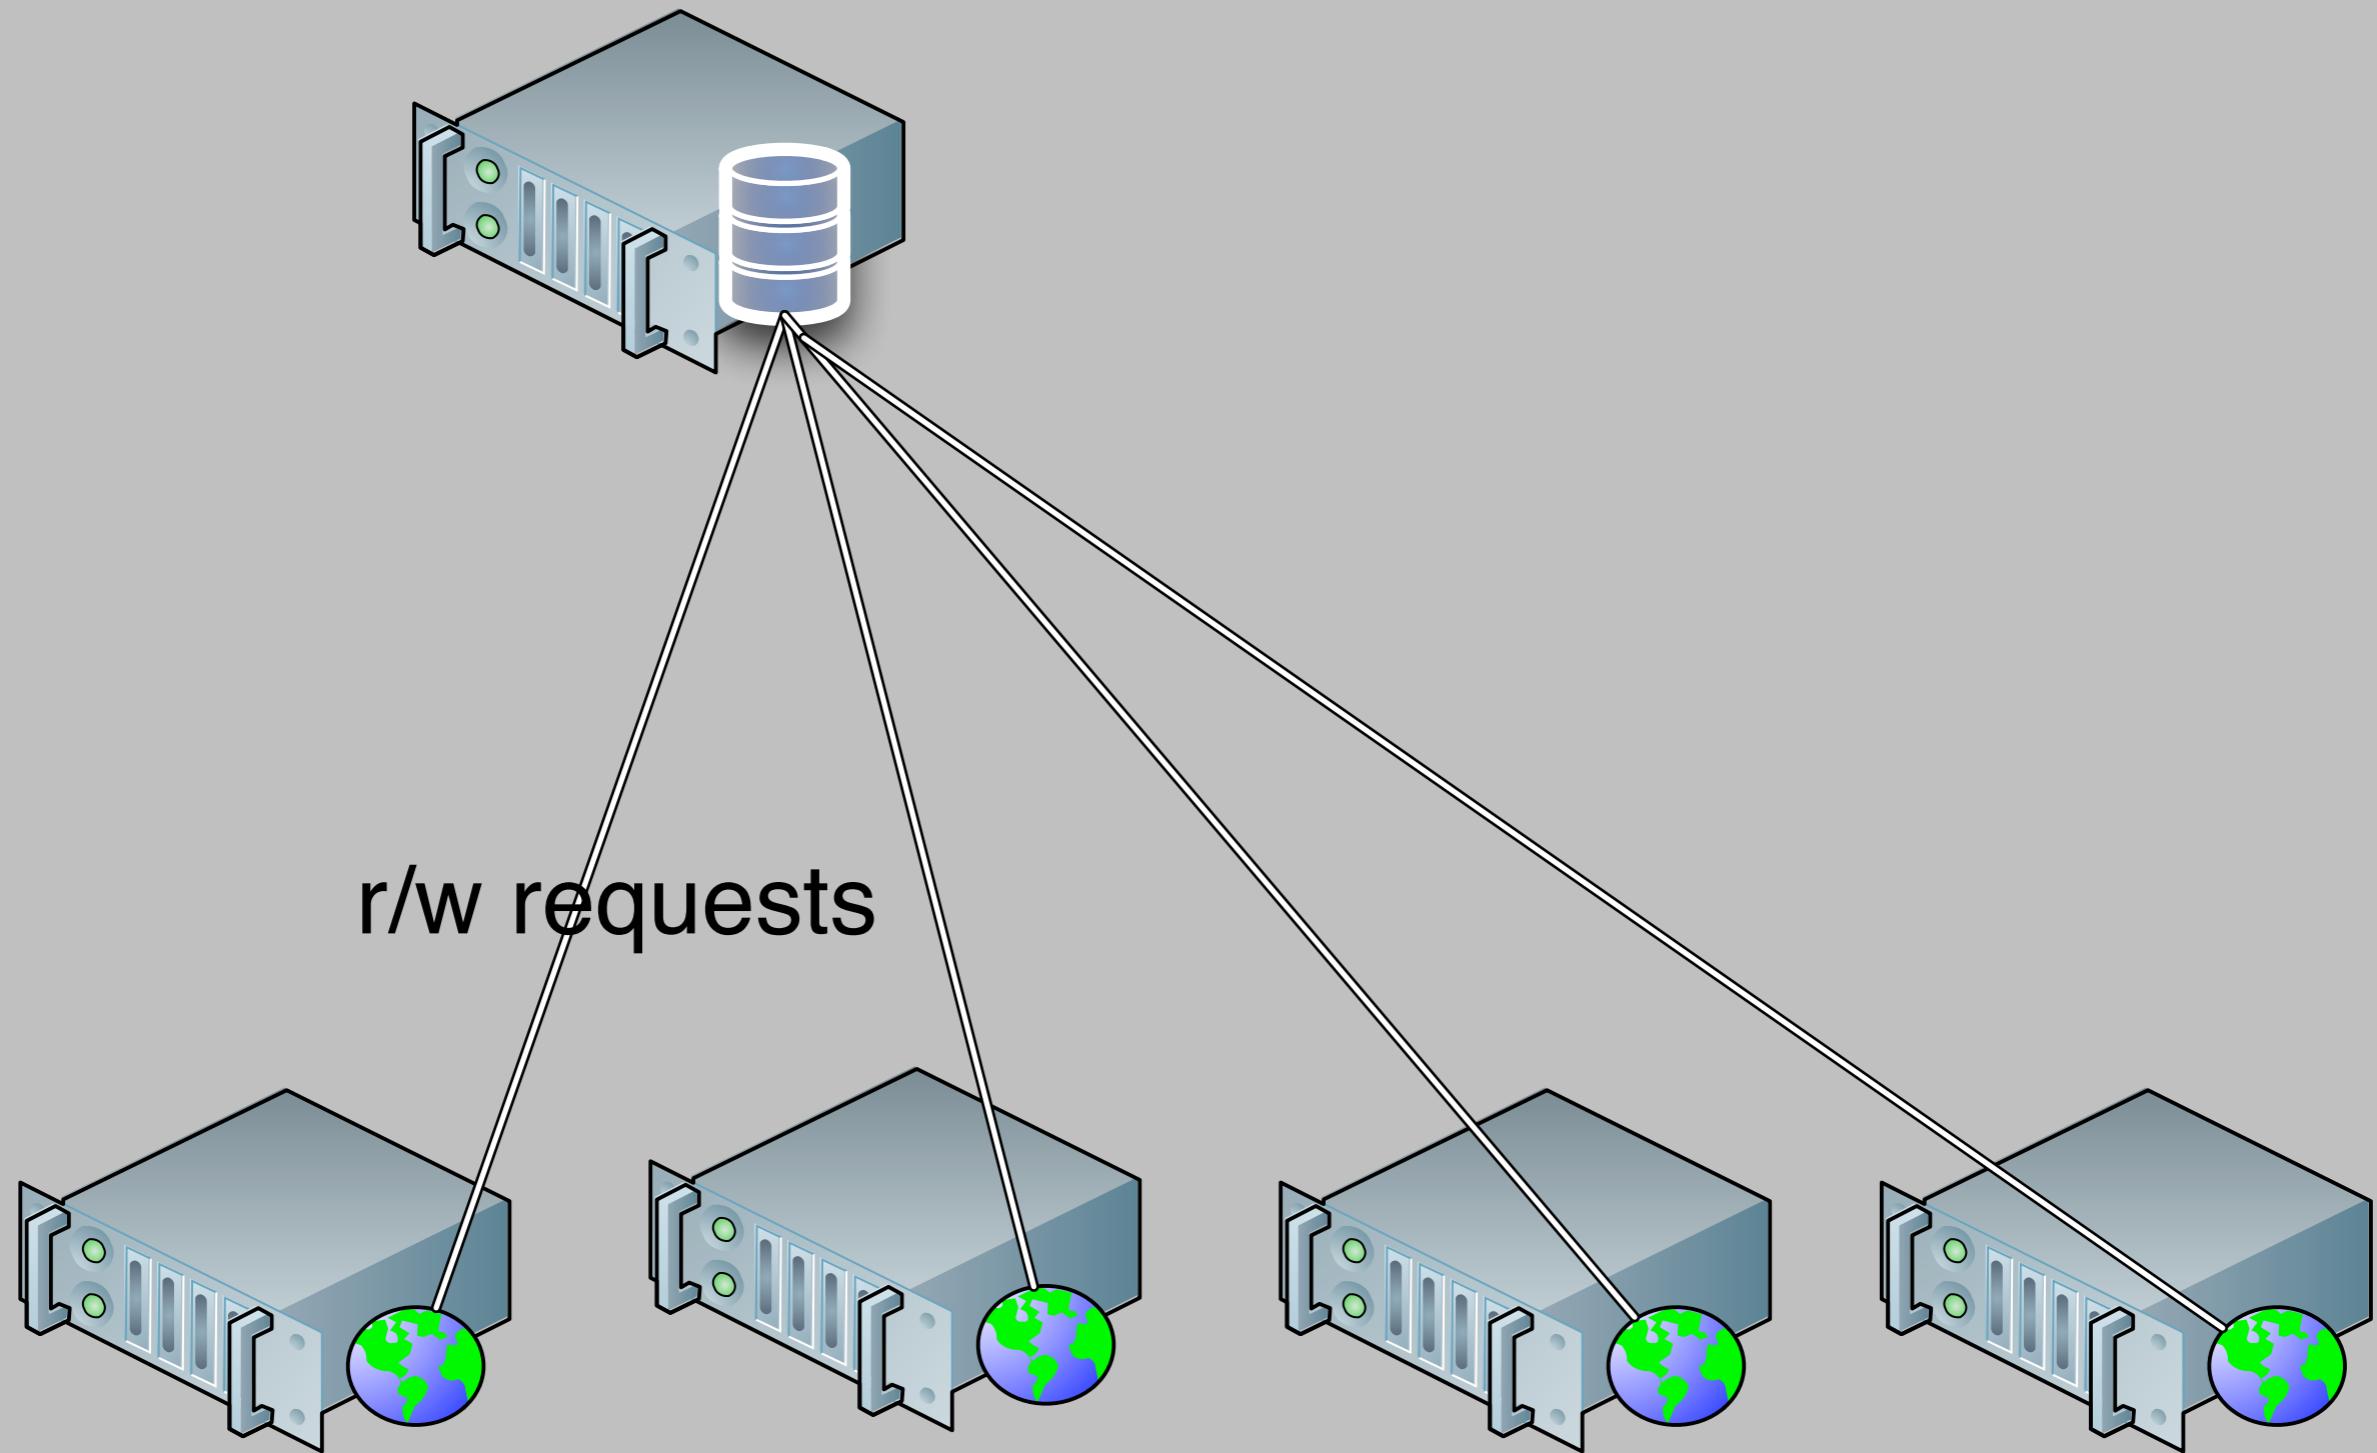
- master replacement
- backup
- better reads

Gotchas, tips, and tricks

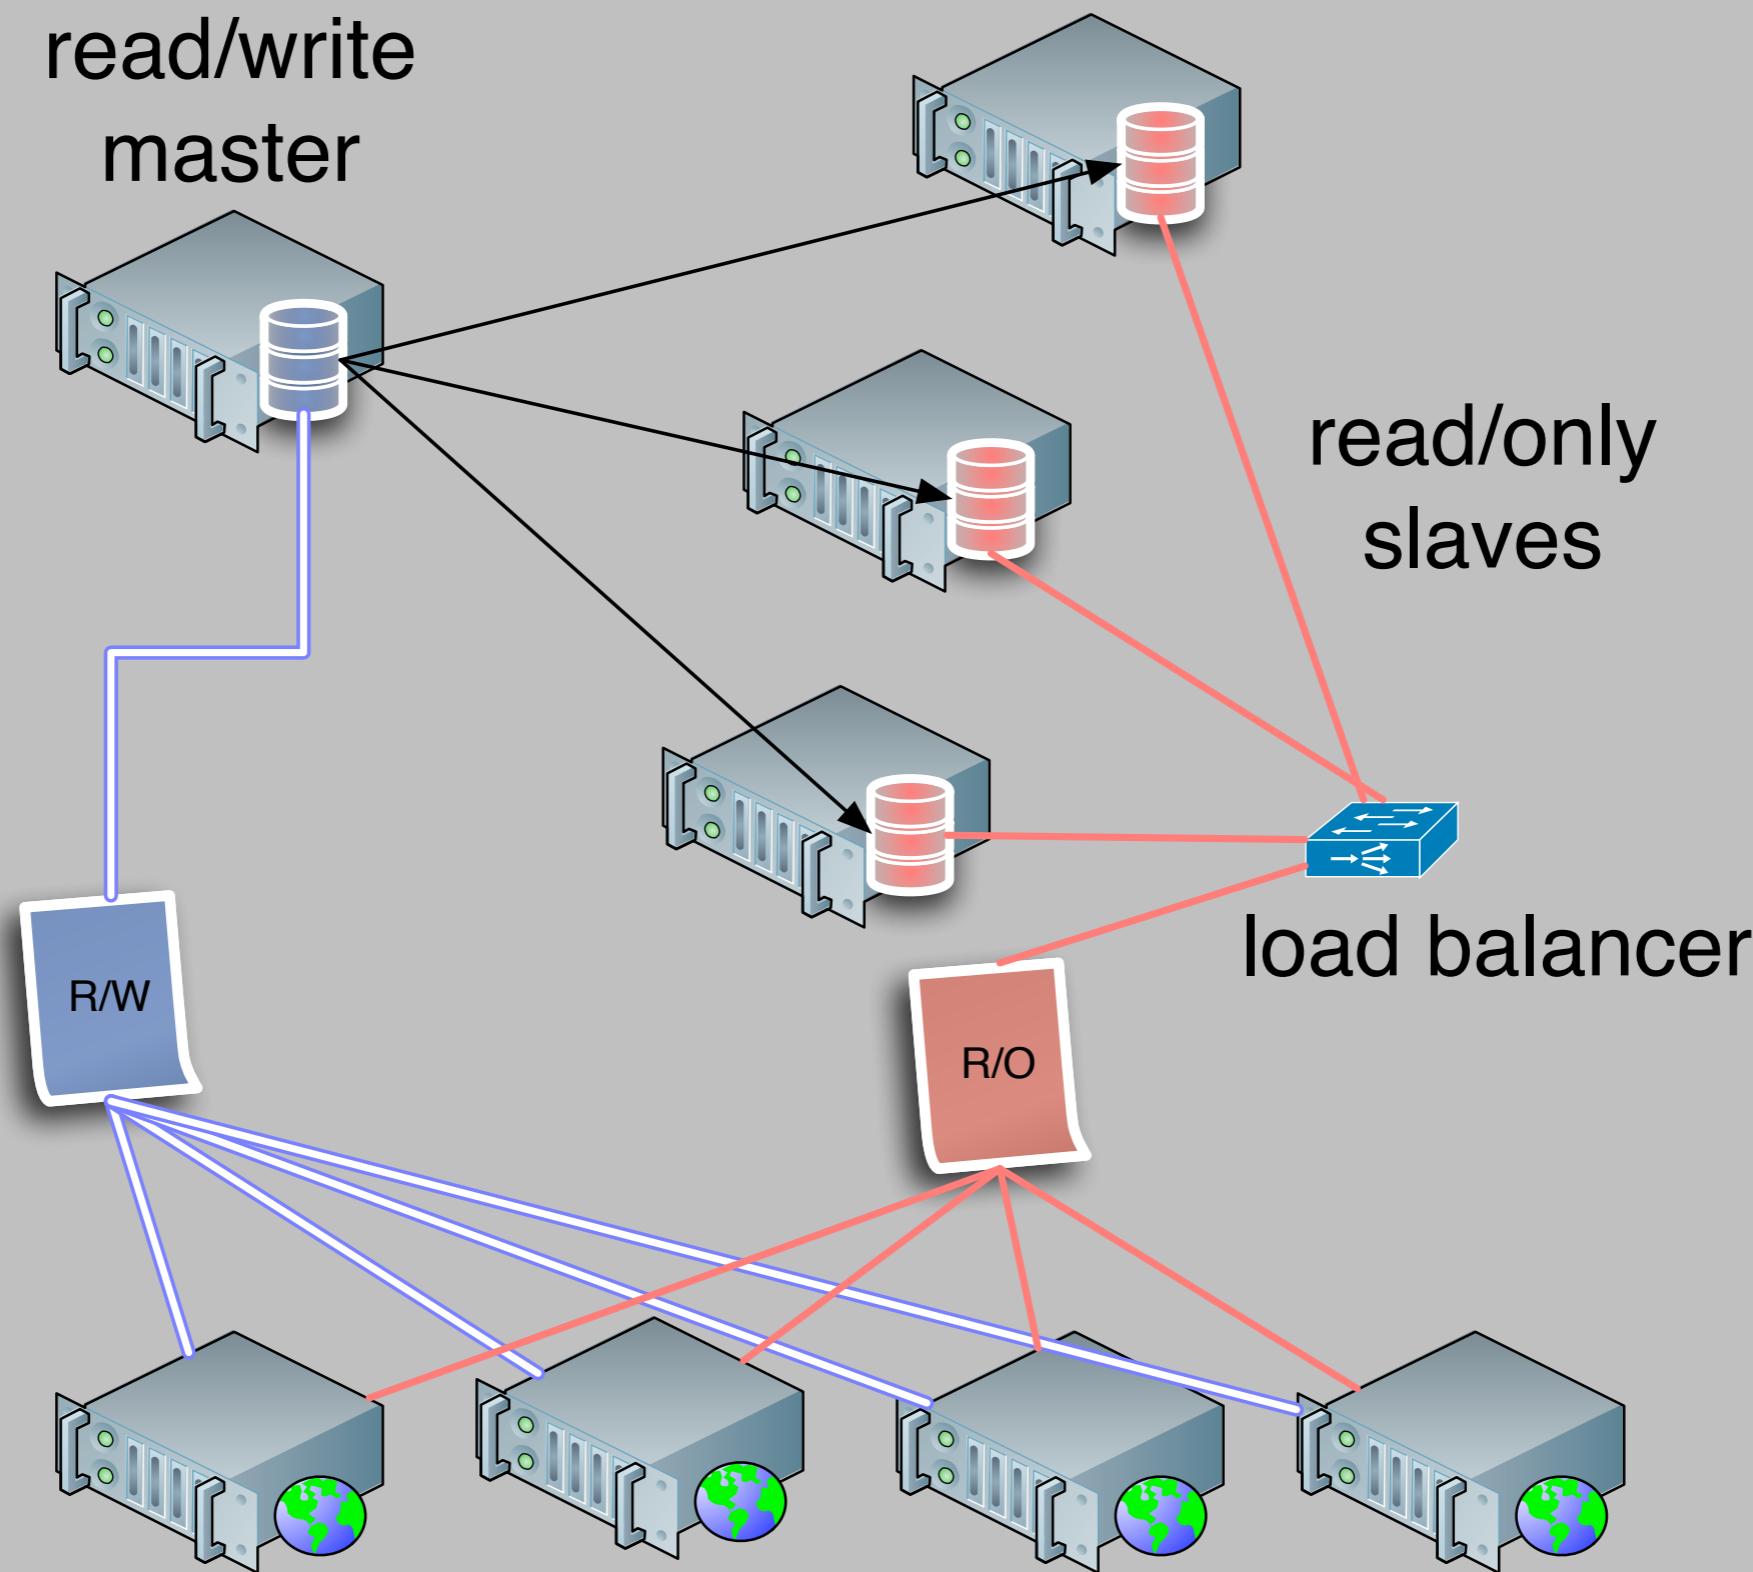
- What to read
- More replication sessions

More info

From single server application



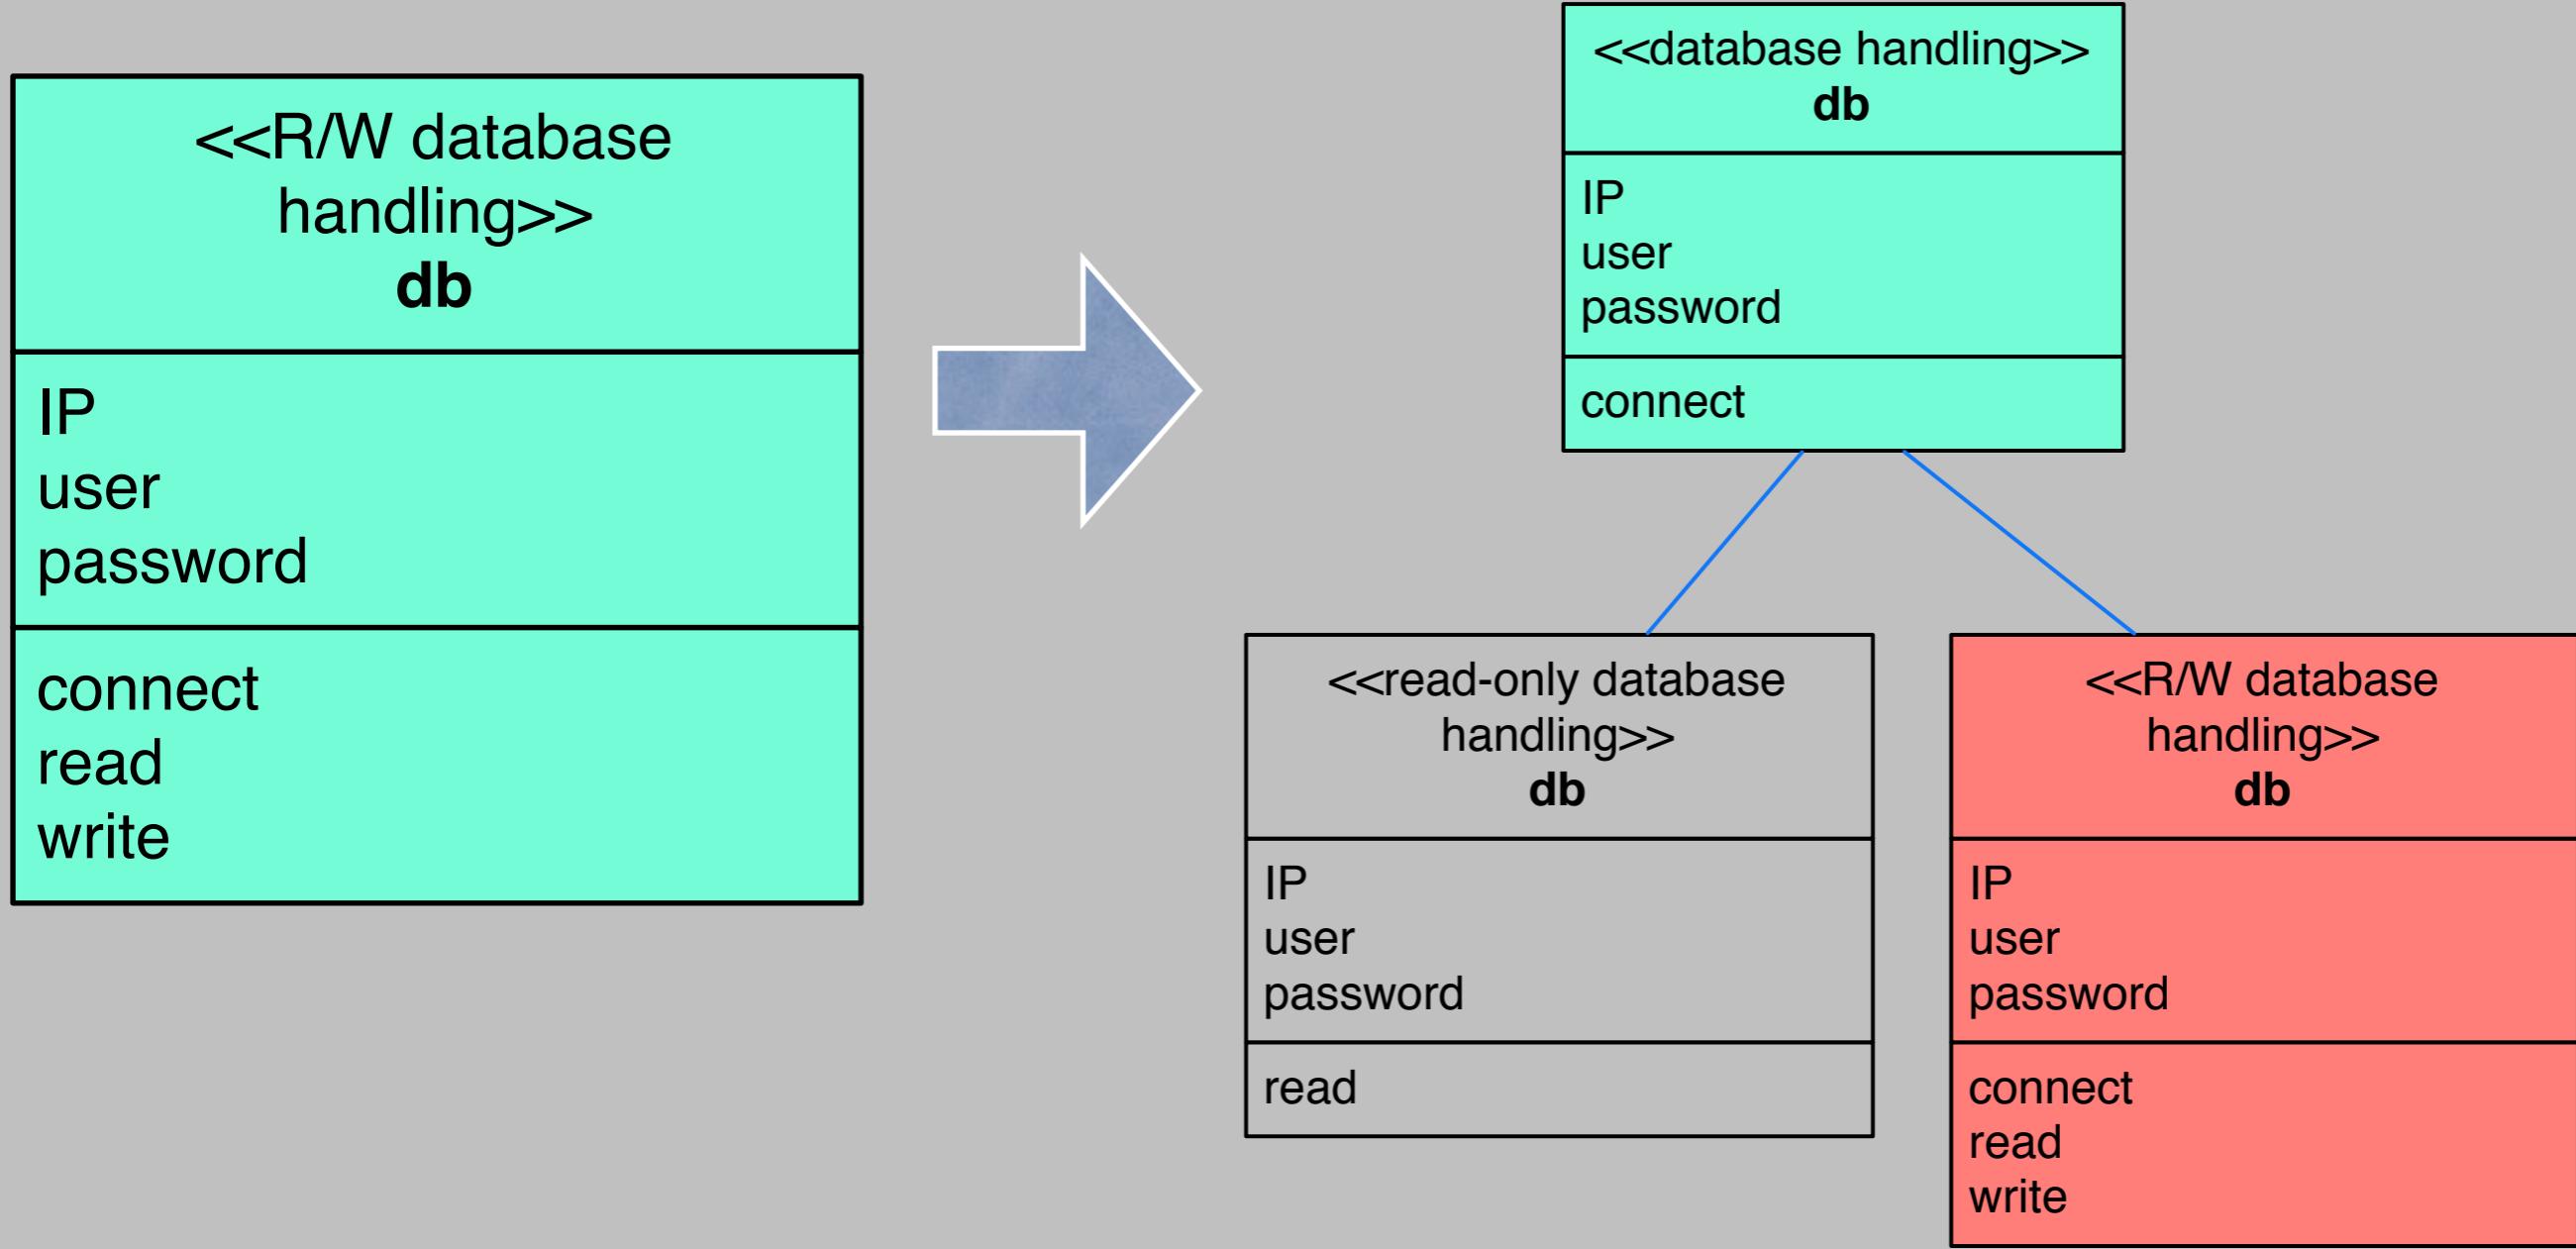
To replication-aware application



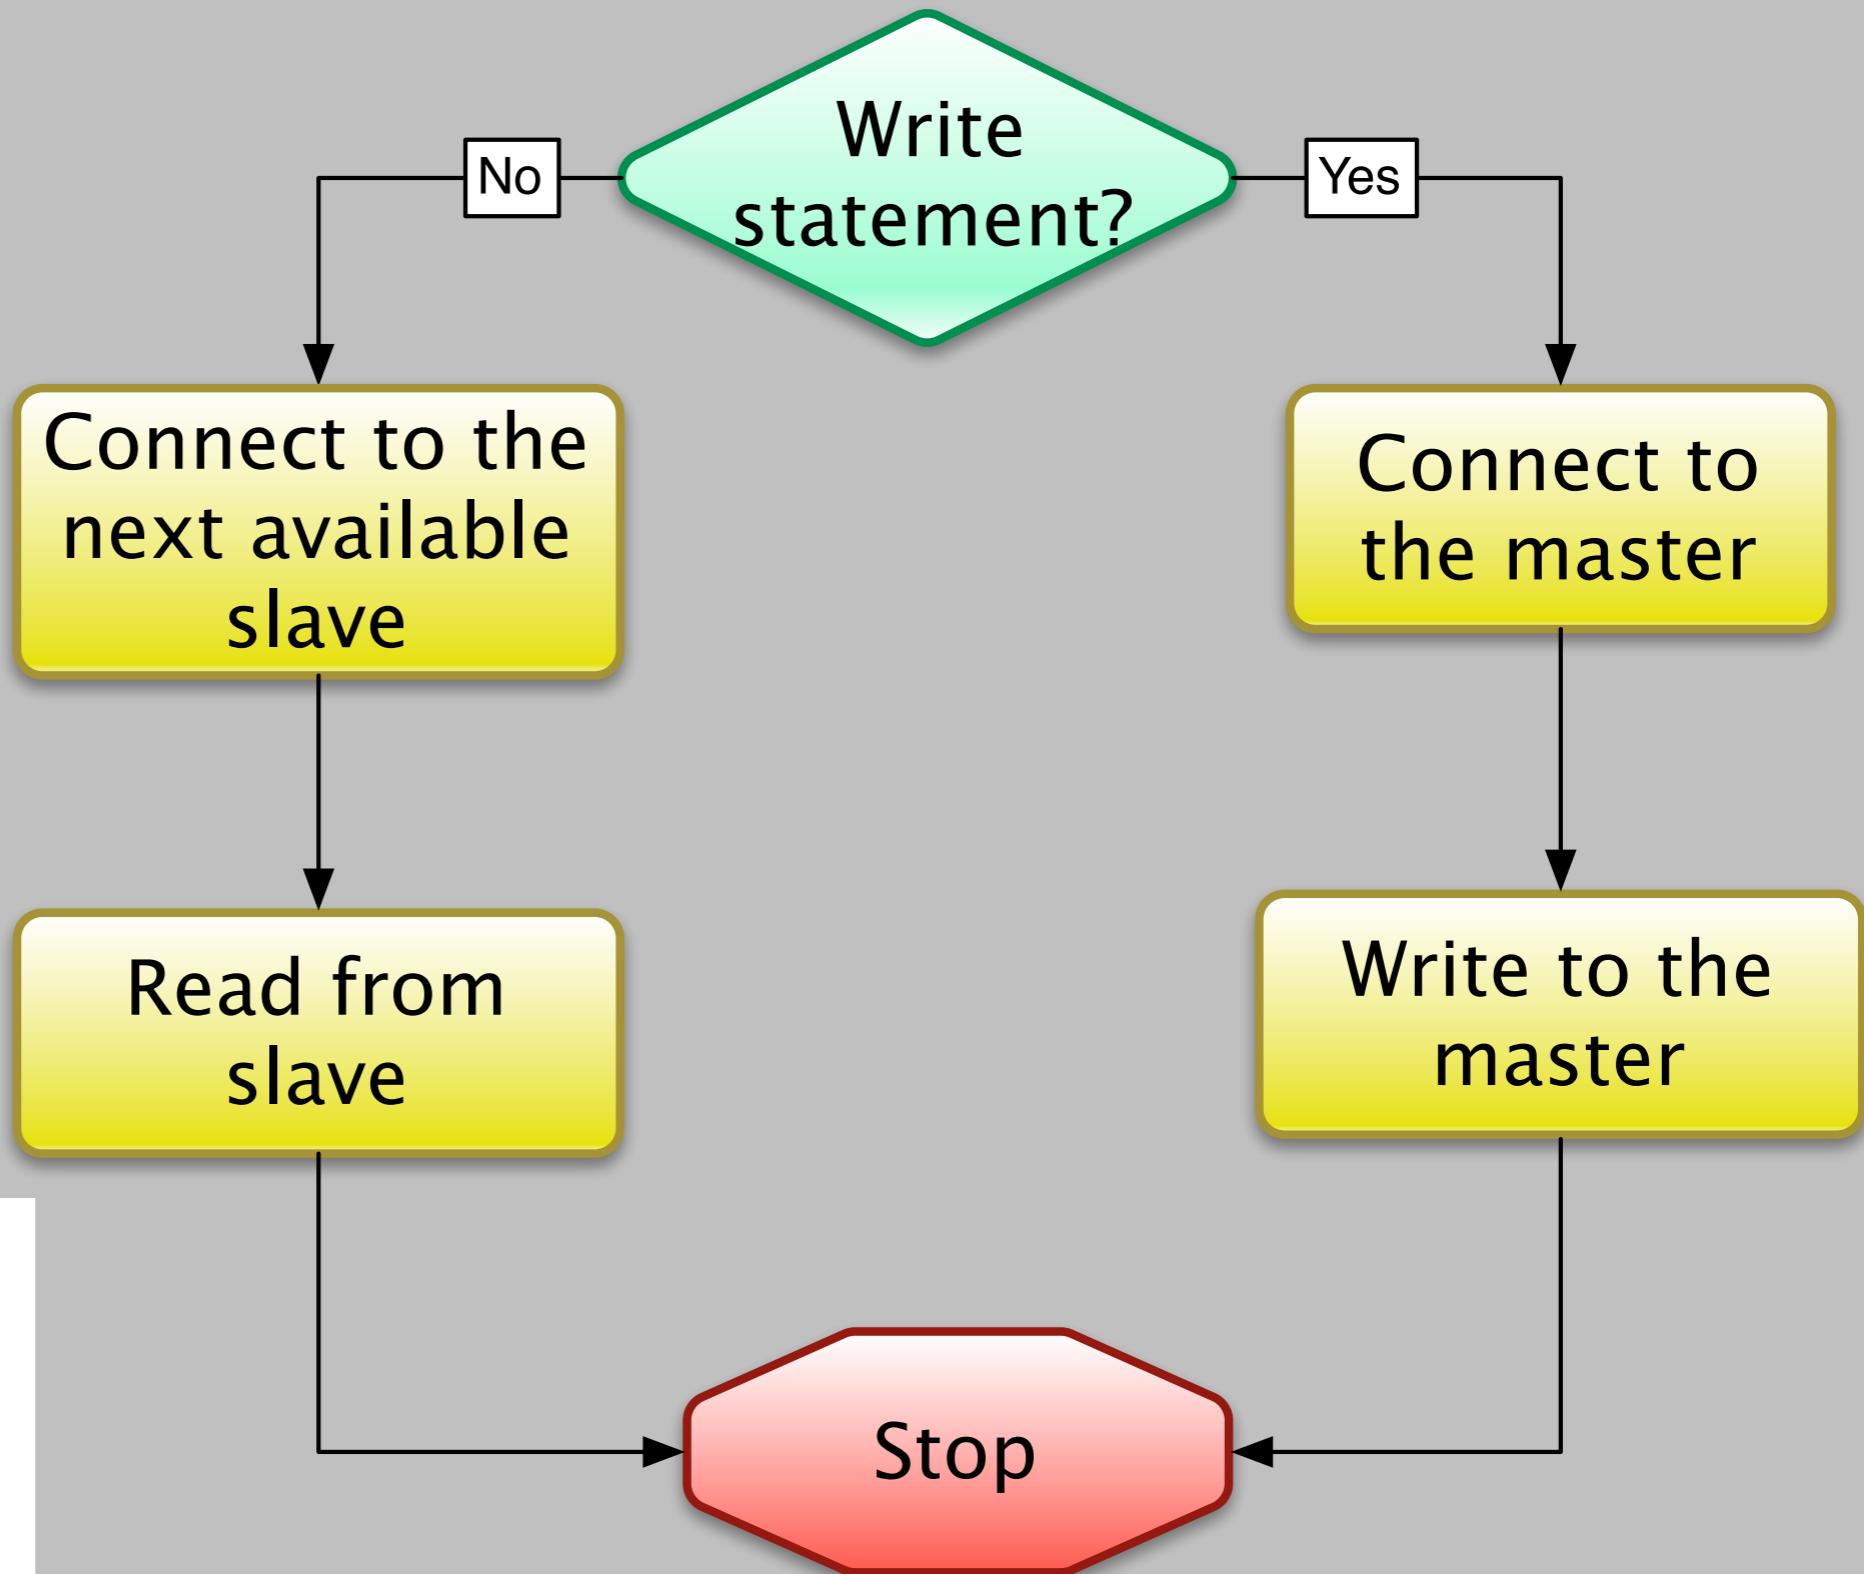
Single server application

```
$link = mysql_connect(  
    $server_IP,  
    'mysql_user',  
    'mysql_password'  
) ;  
$result = mysql_query(  
    'INSERT INTO table_name (x) VALUES (1)',  
    $link  
) ;  
$result = mysql_query(  
    'SELECT * FROM table_name WHERE x=1',  
    $link  
) ;
```

Making an application aware of replication



Using replication: the **WRONG** way

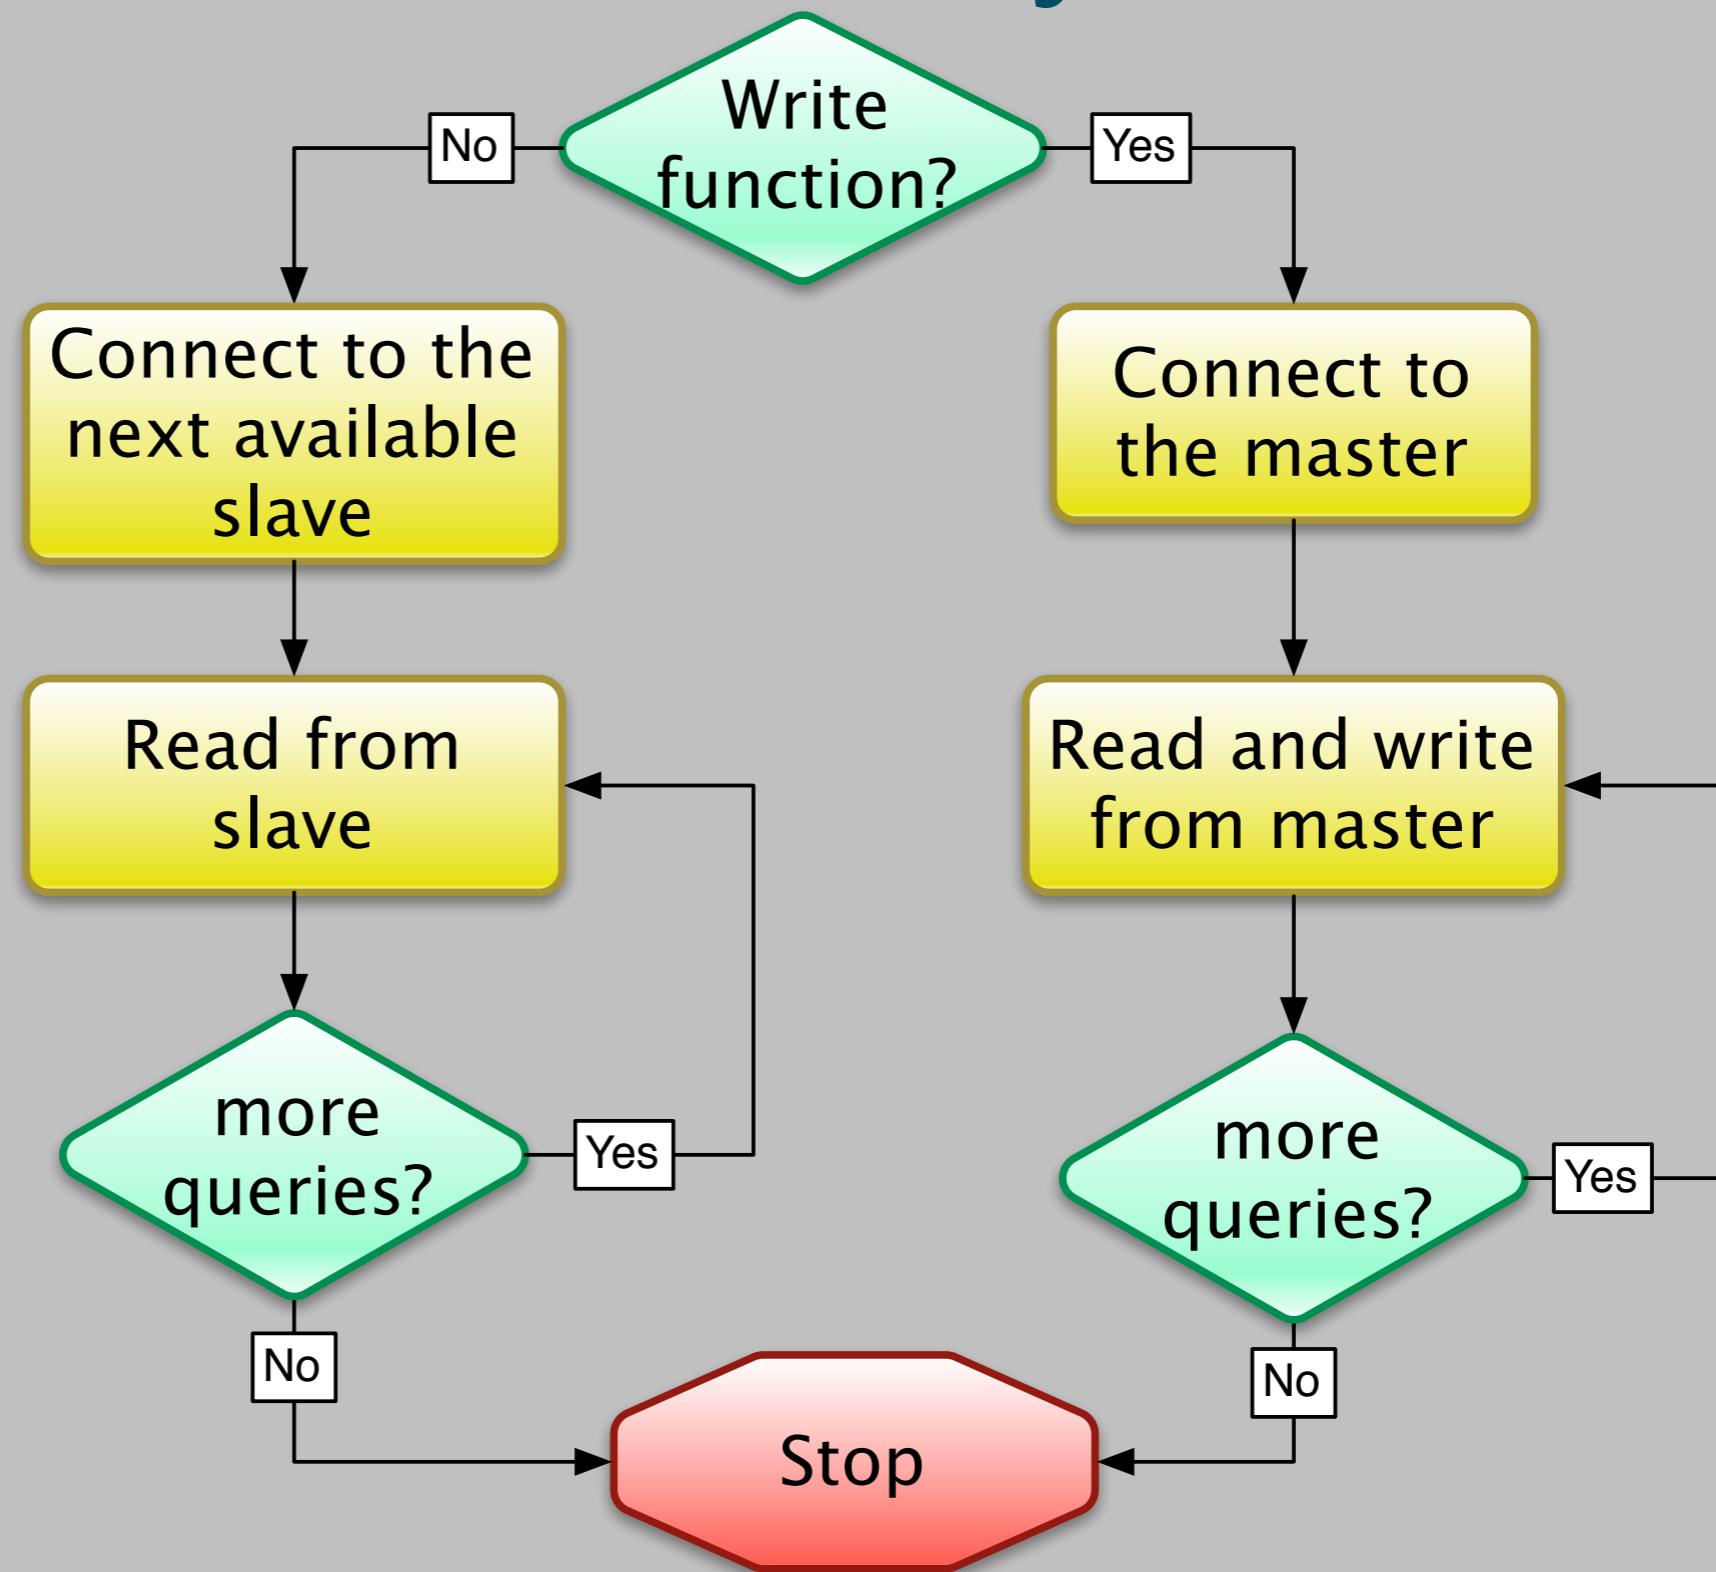


**R/W split
by
statement**

Why statement split is wrong

- Breaks transactions
- High risk of inconsistency
- Loses or corrupts data

Using replication: the RIGHT way

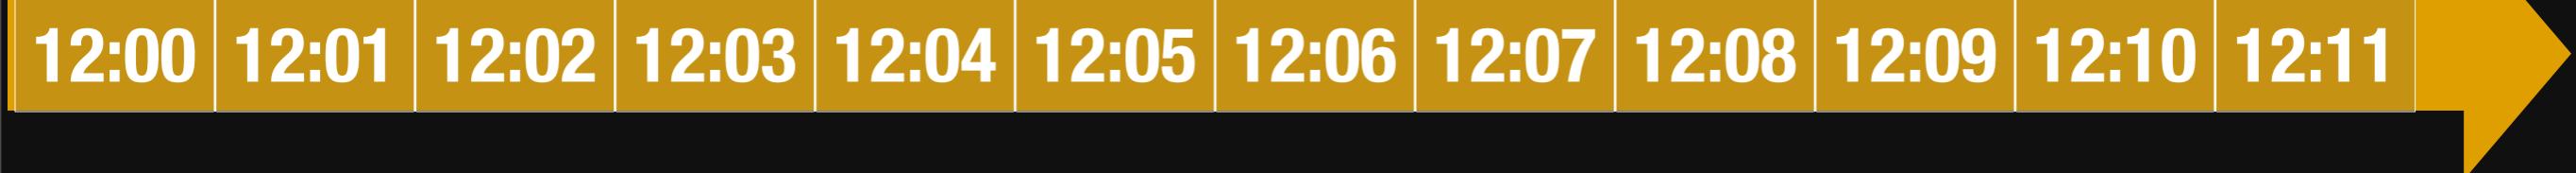


Replication is single threaded

- Let's assume you have two queries on the master
- one query takes 3 minutes to complete
- the other takes 5 minutes
- Both start at 12:00 noon

Single thread replication

MASTER



SLAVE

Single thread replication

MASTER

QUERY 1

QUERY 2

12:00 12:01 12:02 12:03 12:04 12:05 12:06 12:07 12:08 12:09 12:10 12:11

(does nothing)

SLAVE

Single thread replication

MASTER

QUERY 1

QUERY 2

12:00 12:01 12:02 12:03 12:04 12:05 12:06 12:07 12:08 12:09 12:10 12:11

(does nothing)

QUERY 1

106

Single thread replication

MASTER

QUERY 1

QUERY 2

12:00 12:01 12:02 12:03 12:04 12:05 12:06 12:07 12:08 12:09 12:10 12:11

(does nothing)

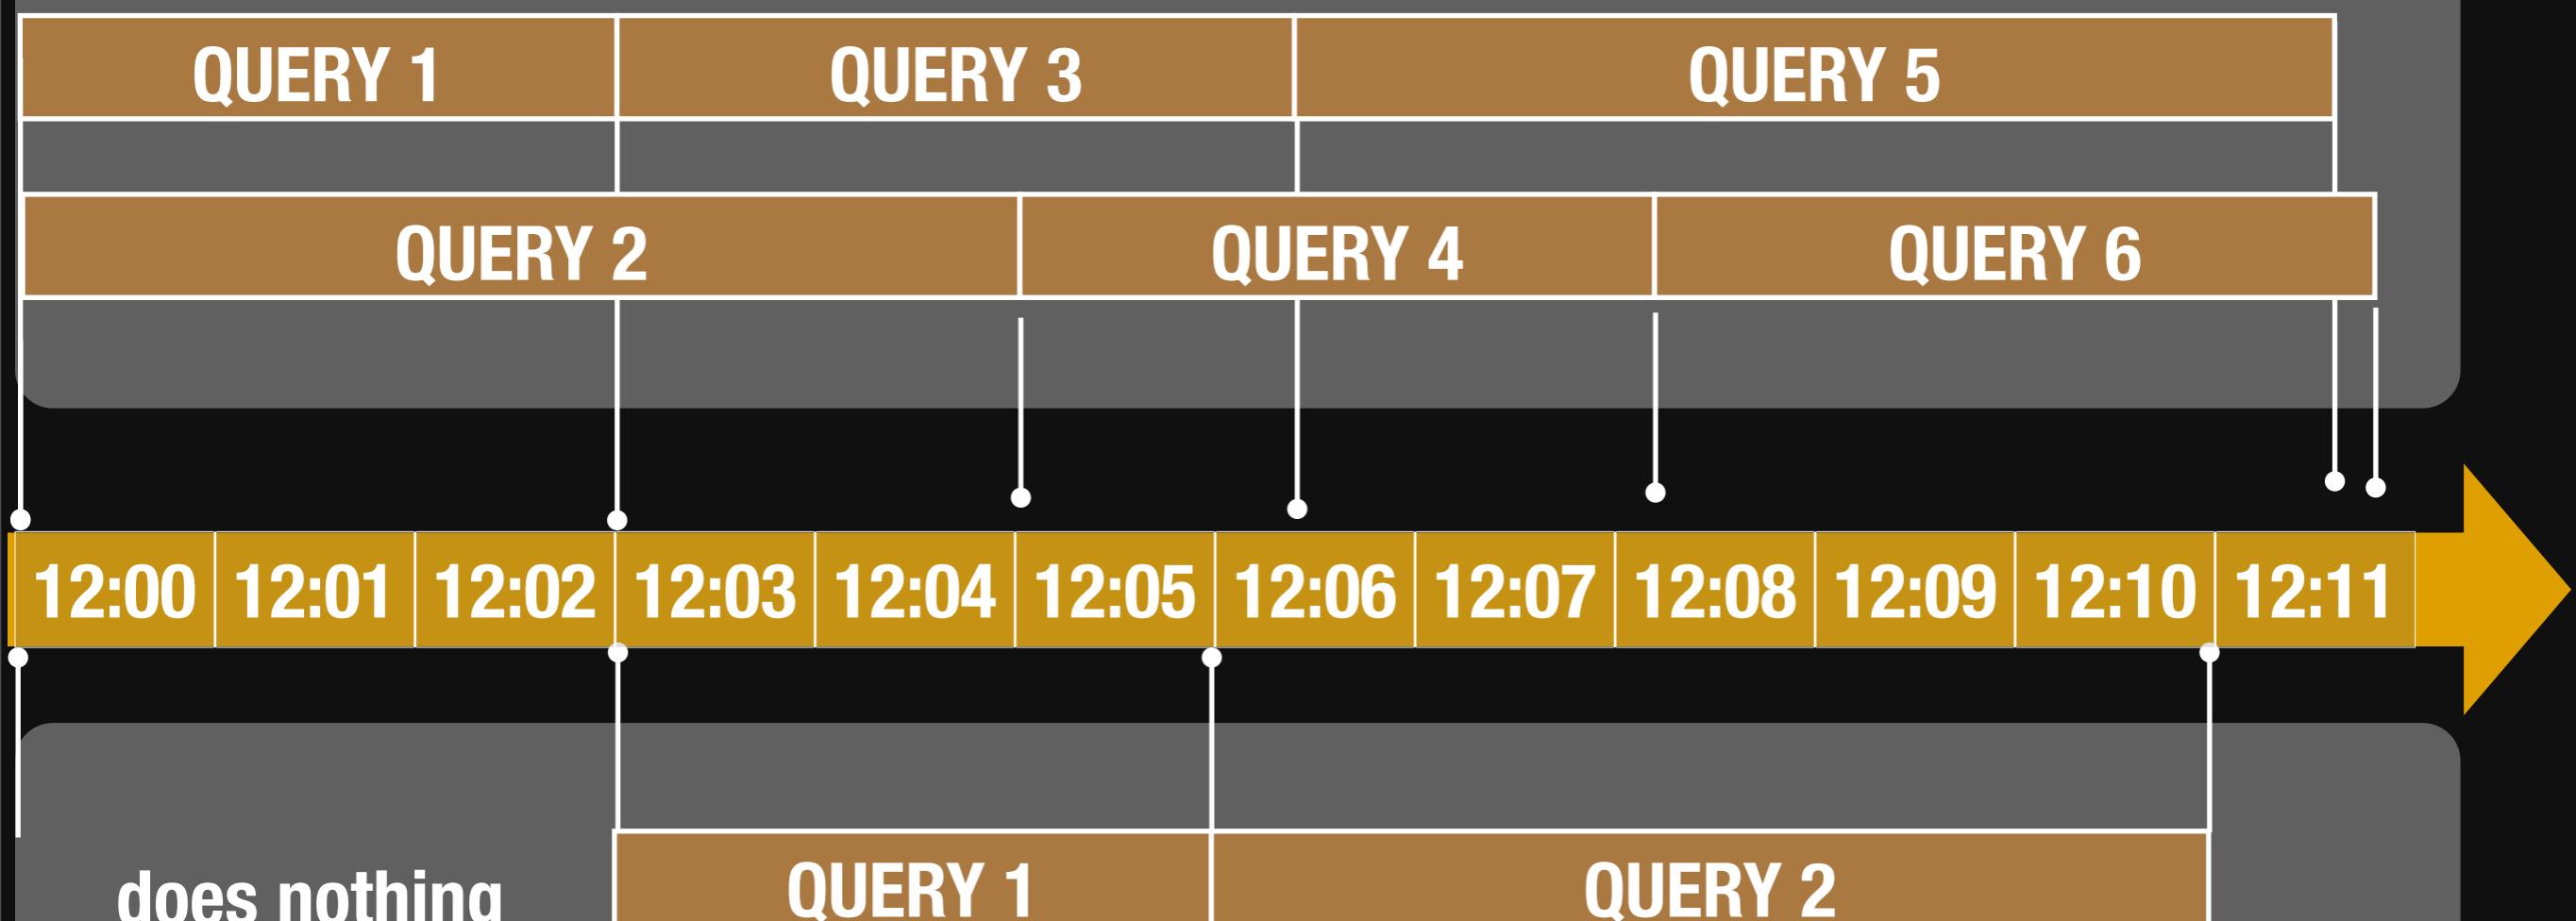
QUERY 1

QUERY 2

106

Single thread replication

MASTER



SLAVE

AGENDA

Why replication

- The web economy
- Scaling out

How to set replication

- From single server to replication
- Adding a slave

Using Replication

- Binary log formats
- What gets replicated and how
- Replication awareness

Managing replication

- **Monitoring**
- Log management
- Replacing a slave
- Replacing a master

What Replication is for

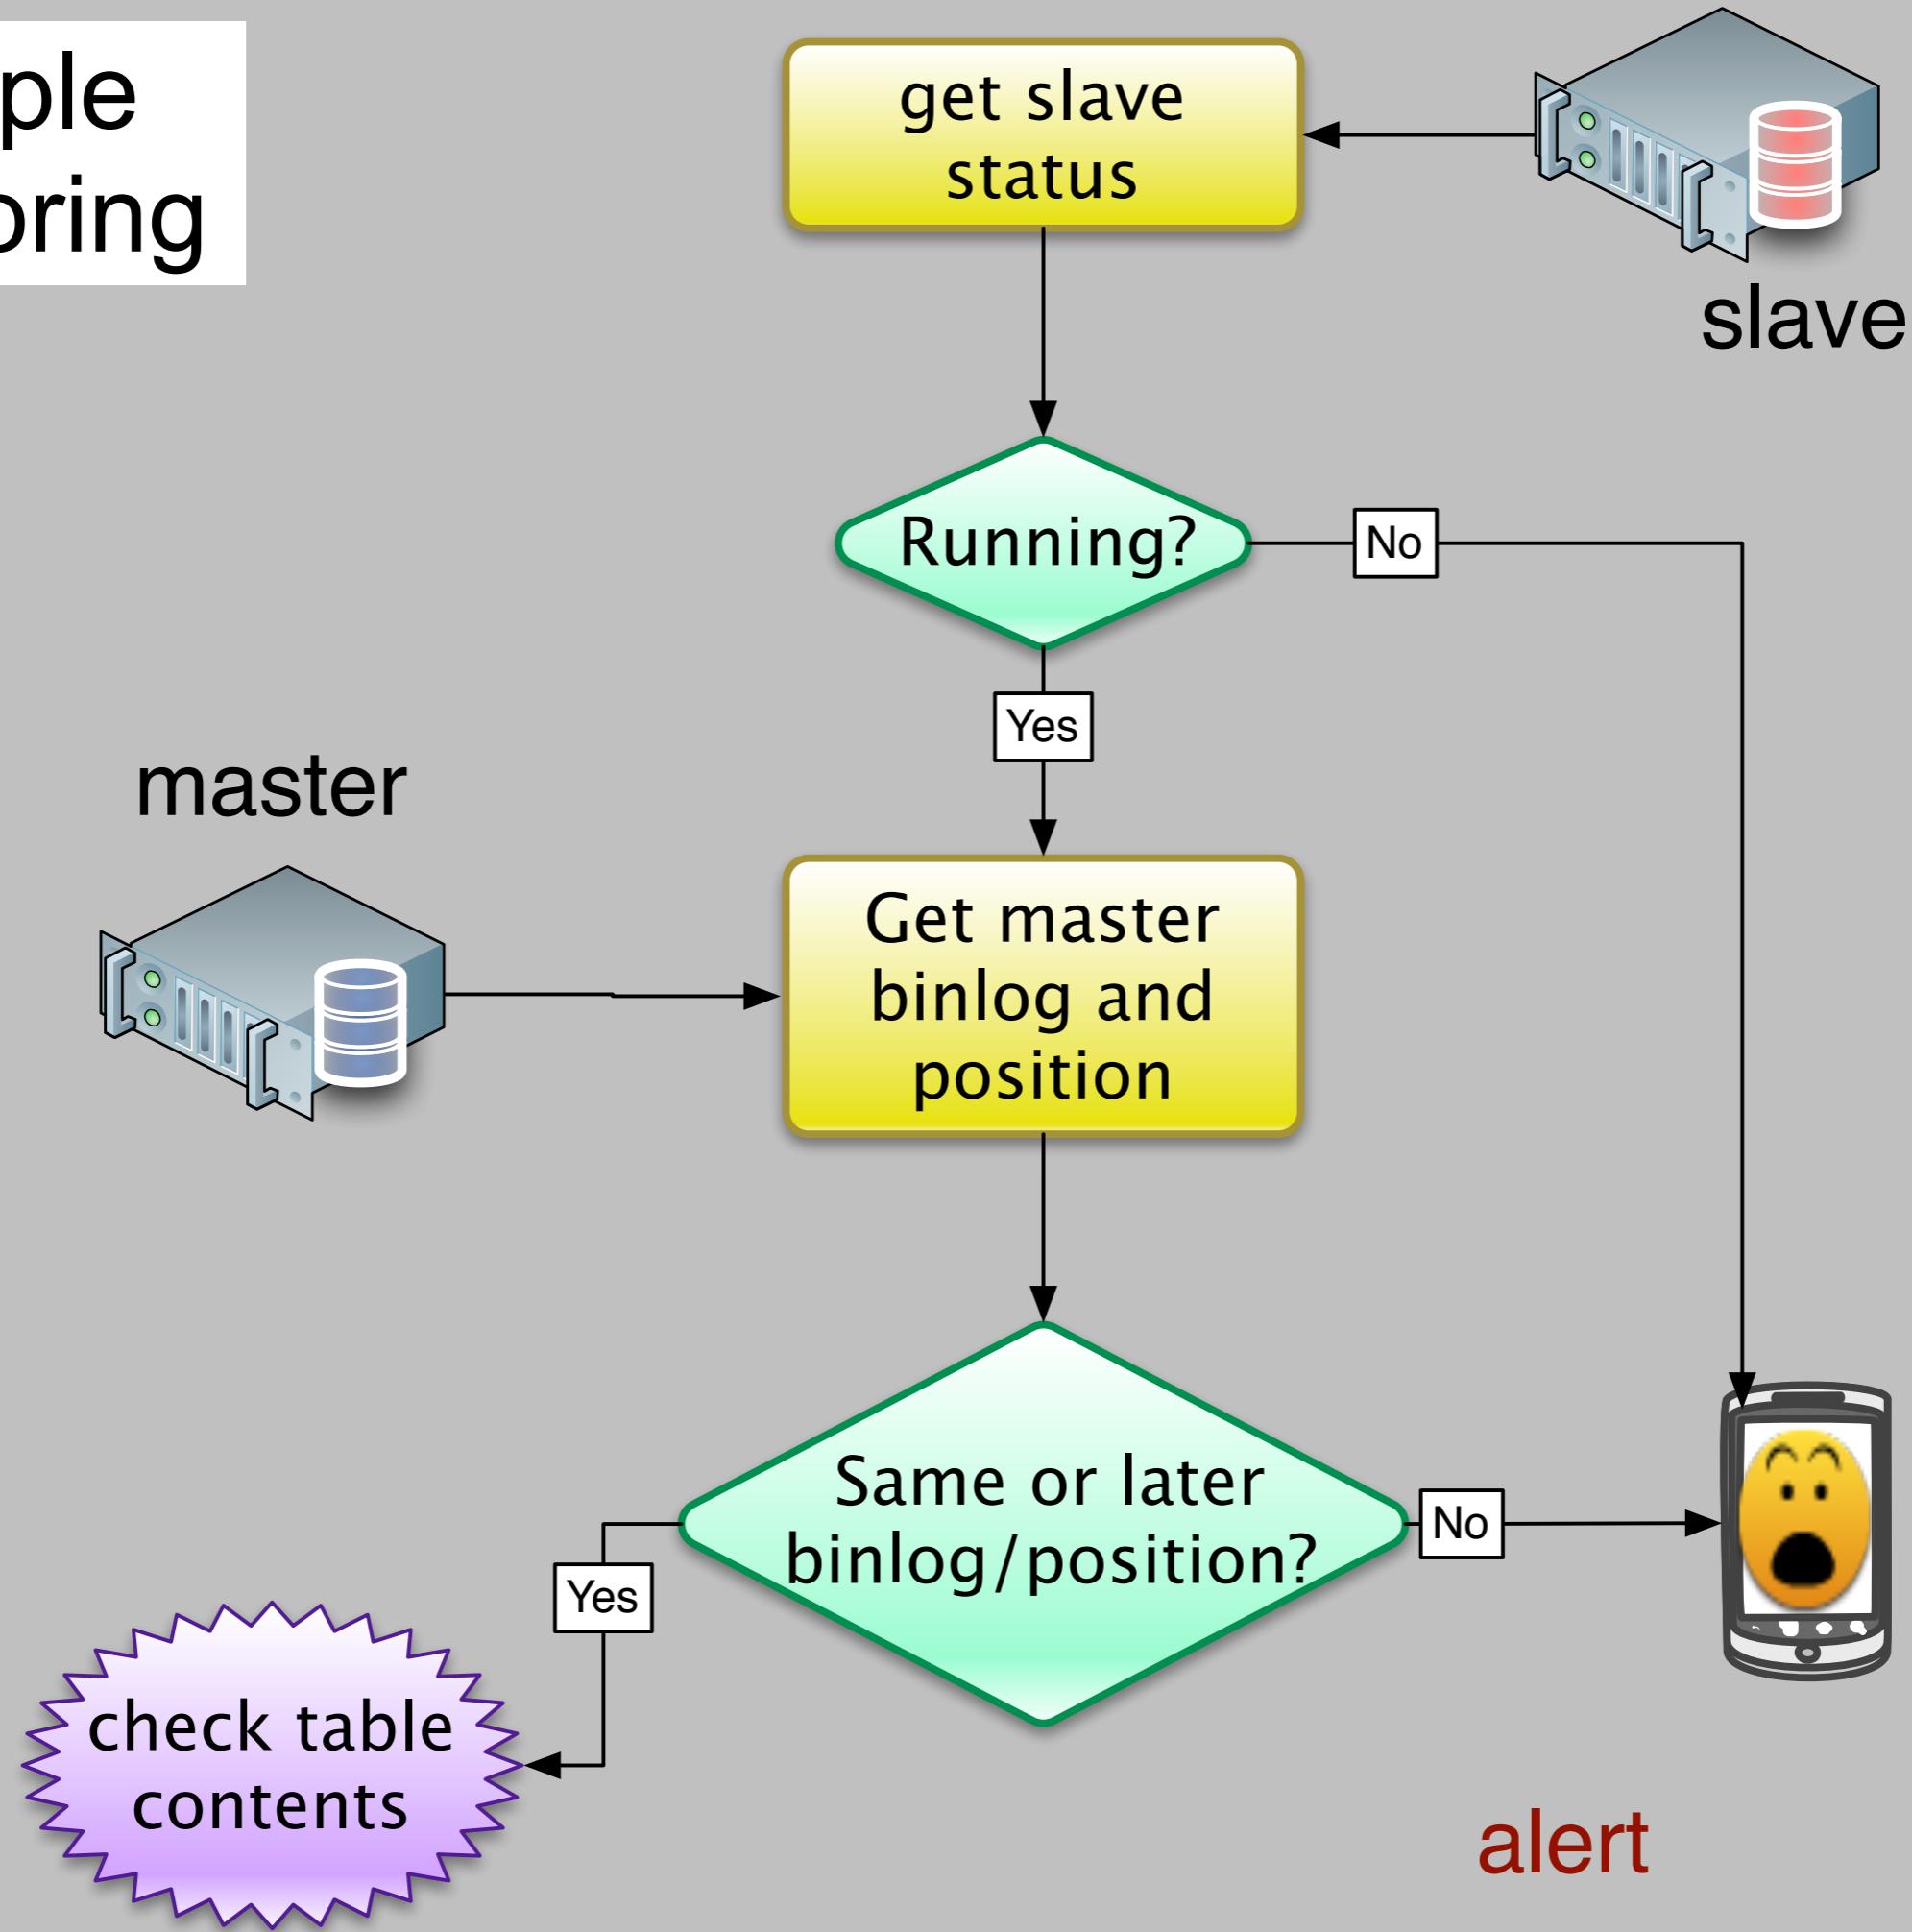
- master replacement
- backup
- better reads

Gotchas, tips, and tricks

More info

- What to read
- More replication sessions

Sample monitoring



commands for checking replication

master> SHOW MASTER STATUS

master> SHOW PROCESSLIST

slave> SHOW SLAVE STATUS

slave> SHOW PROCESSLIST

replication user seen in the master

```
master > show processlist\G
```

```
***** 1. row *****
```

```
Id: 2
```

```
User: rsandbox
```

```
Host: localhost:57011
```

```
db: NULL
```

```
Command: Binlog Dump
```

```
Time: 81625
```

```
State: Master has sent all binlog to slave; waiting  
for binlog to be updated
```

```
Info: NULL
```

```
[...]
```

IO-thread in the slave

```
slave > show processlist\G
```

```
***** 1. row *****
```

Id: 2

User: system user

Host:

db: NULL

Command: Connect

Time: 124801

State: Waiting for master to send event

Info: NULL

[...]

SQL thread in the slave

```
slave > show processlist\G
```

```
[...]
```

```
***** 2. row *****
```

```
Id: 3
```

```
User: system user
```

```
Host:
```

```
db: NULL
```

```
Command: Connect
```

```
Time: 124712
```

```
State: Slave has read all relay log; waiting for the  
slave I/O thread to update it
```

```
Info: NULL
```

```
[...]
```

monitoring replication

master> SHOW MASTER STATUS

slave> SHOW SLAVE STATUS

FULL SCRIPTS:

<http://datacharmer.blogspot.com/2011/04/refactored-again-poor-mans-mysql.html>

<http://forge.mysql.com/tools/tool.php?id=6>

show master status

File: mysql-bin.000002

Position: 78045744

Binlog_Do_DB:

Binlog_Ignore_DB:

show slave status

```
Slave_IO_State: Waiting for master to send event
  Master_Host: 127.0.0.1
  Master_User: rsandbox
  Master_Port: 27371
  Connect_Retry: 60
  Master_Log_File: mysql-bin.000002
  Read_Master_Log_Pos: 78045744
  Relay_Log_File: mysql_sandbox27372-relay-bin.000055
  Relay_Log_Pos: 78045889
  Relay_Master_Log_File: mysql-bin.000002
  Slave_IO_Running: Yes
  Slave_SQL_Running: Yes
```

show slave status

...

```
Replicate_Do_DB:  
Replicate_Ignore_DB:  
Replicate_Do_Table:  
Replicate_Ignore_Table:  
Replicate_Wild_Do_Table:  
Replicate_Wild_Ignore_Table:
```

...

show slave status

...

Last_Error: 0

Last_Error:

Skip_Counter: 0

Exec_Master_Log_Pos: 78045744

Relay_Log_Space: 78046100

...

Seconds_Behind_Master: 0

...

Last_IO_Error: 0

Last_IO_Error:

Last_SQL_Error: 0

Last_SQL_Error:

AGENDA

Breaking (and fixing) replication

DEMO

AGENDA

Why replication

- The web economy
- Scaling out

How to set replication

- From single server to replication
- Adding a slave

Using Replication

- Binary log formats
- What gets replicated and how
- Replication awareness

Managing replication

- Monitoring
- **Log management**
- Replacing a slave
- Replacing a master

What Replication is for

- master replacement
- backup
- better reads

Gotchas, tips, and tricks

More info

- What to read
- More replication sessions

Logs rotation

```
# server variables
```

```
max-binlog-size
```

```
expire-log-days
```

```
# logs commands
```

```
SHOW MASTER LOGS;
```

```
PURGE MASTER LOGS TO 'filename';
```

```
FLUSH [BINARY] LOGS;
```

AGENDA

Why replication

- The web economy
- Scaling out

How to set replication

- From single server to replication
- Adding a slave

Using Replication

- Binary log formats
- What gets replicated and how
- Replication awareness

Managing replication

- Monitoring
- Log management
- **Replacing a slave**
- Replacing a master

What Replication is for

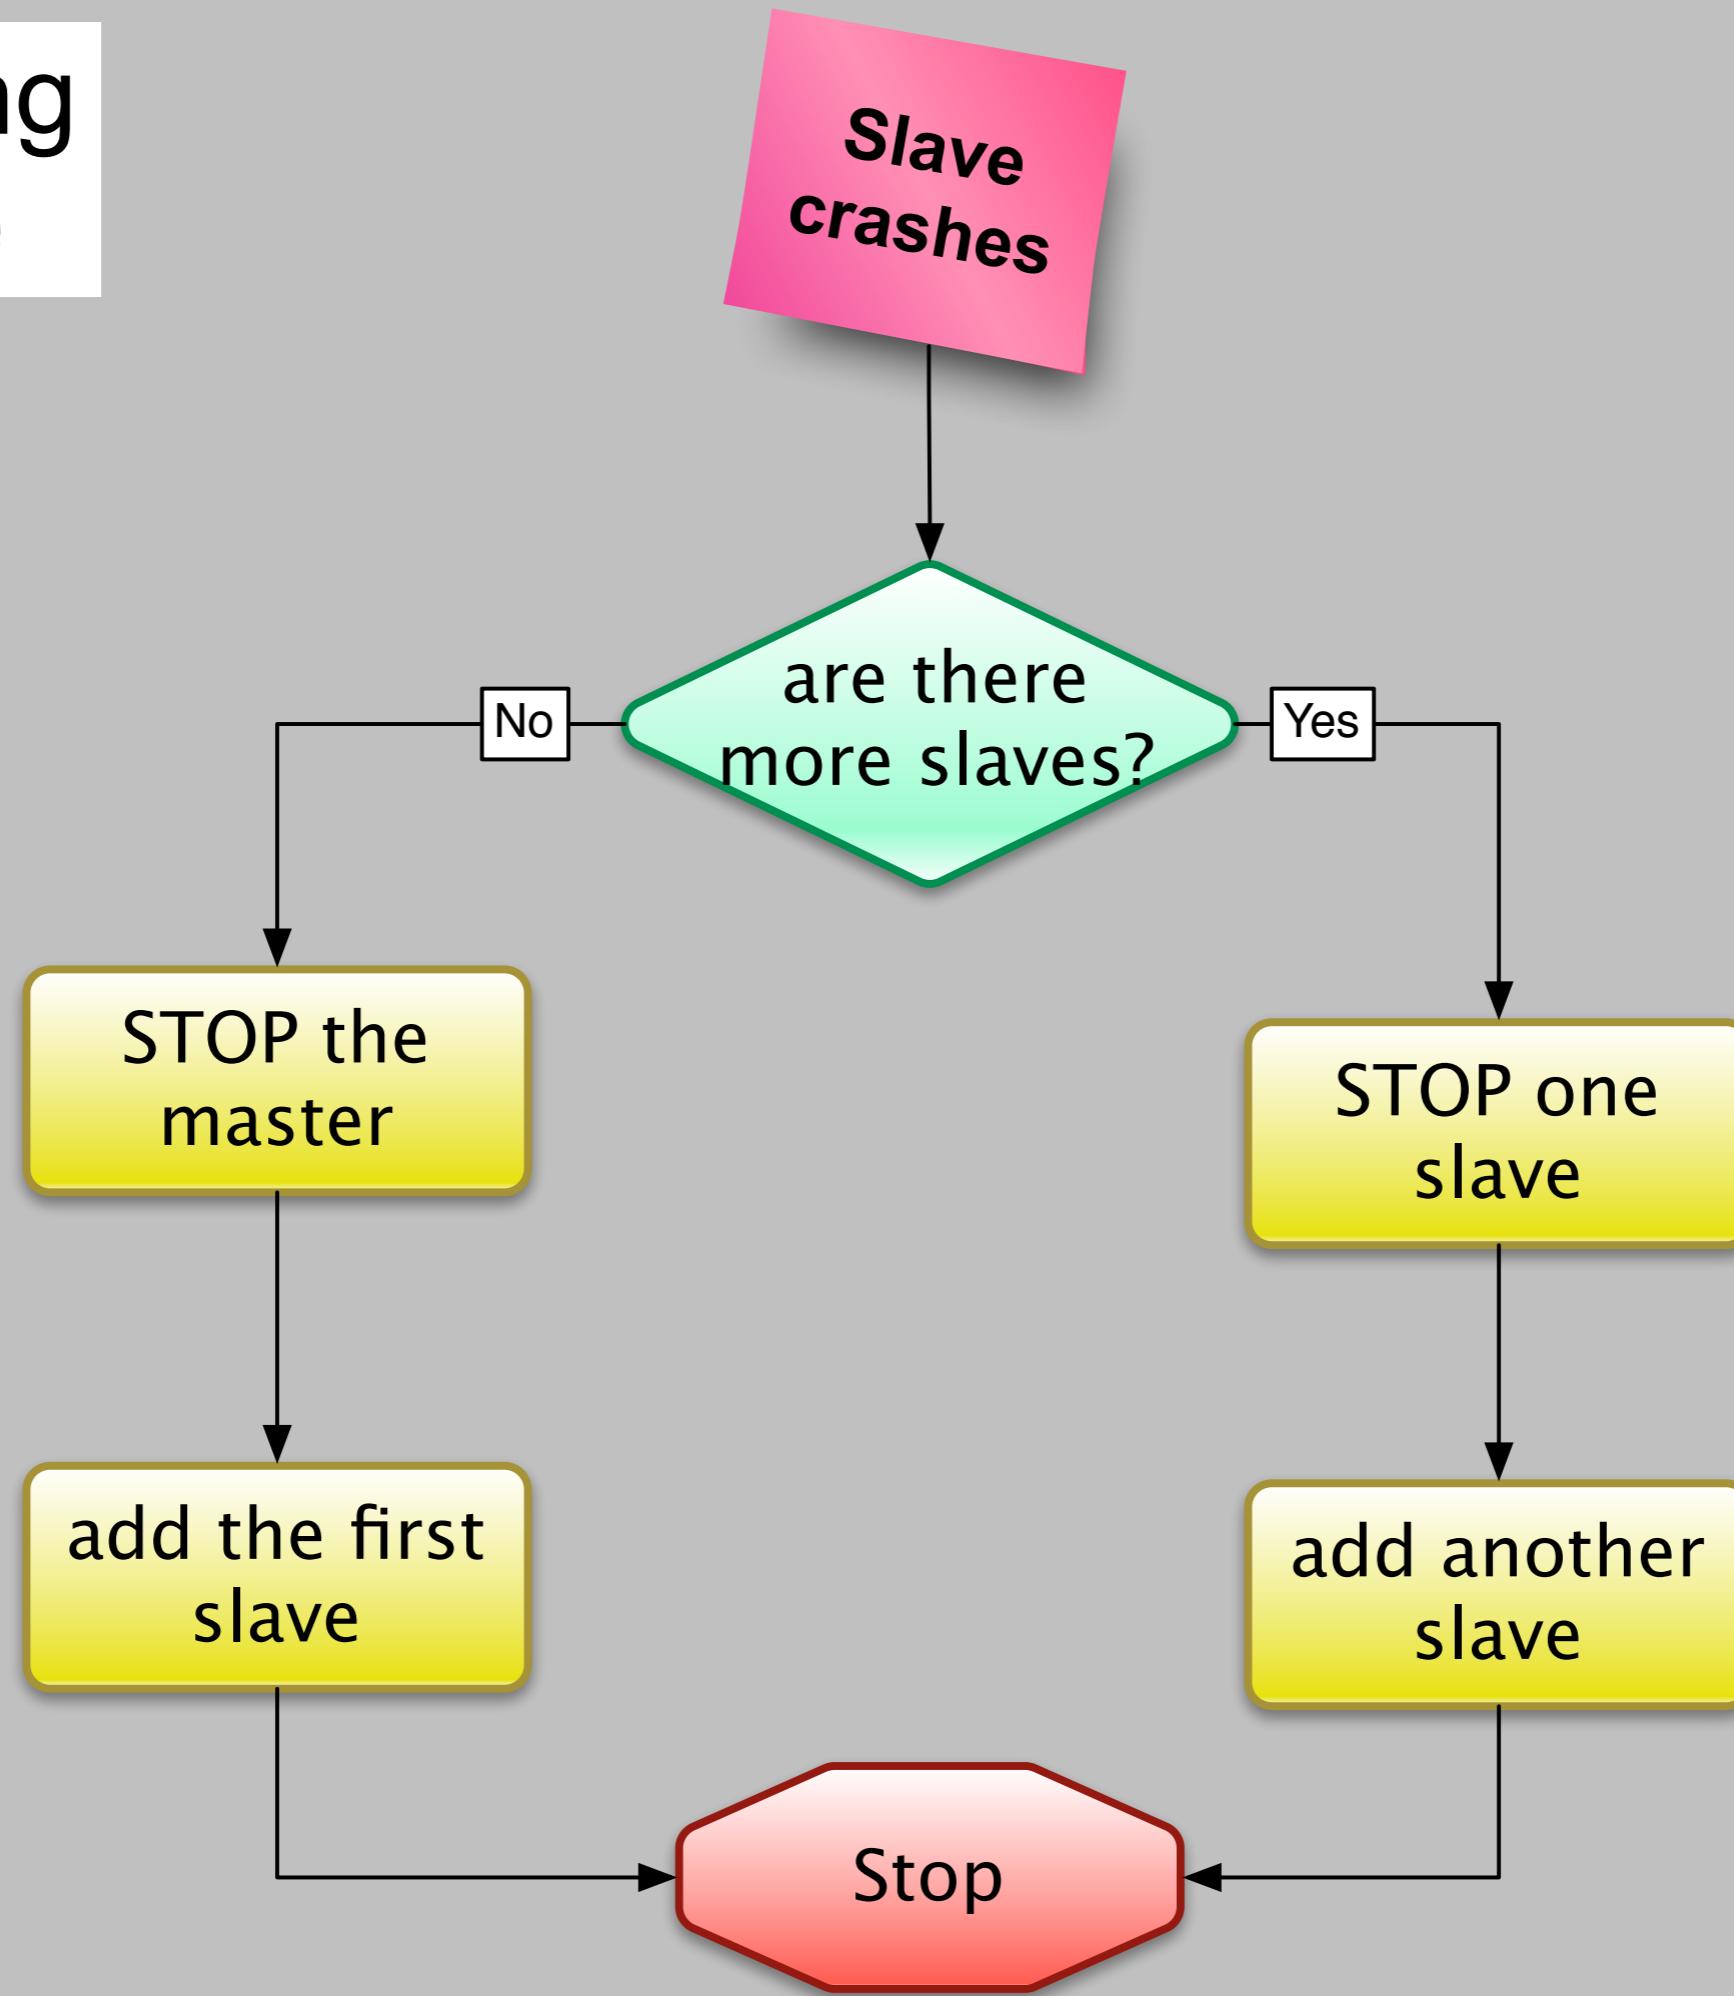
- master replacement
- backup
- better reads

Gotchas, tips, and tricks

More info

- What to read
- More replication sessions

Replacing a slave



AGENDA

Why replication

- The web economy
- Scaling out

How to set replication

- From single server to replication
- Adding a slave

Using Replication

- Binary log formats
- What gets replicated and how
- Replication awareness

Managing replication

- Monitoring
- Log management
- Replacing a slave
- **Replacing a master**

What Replication is for

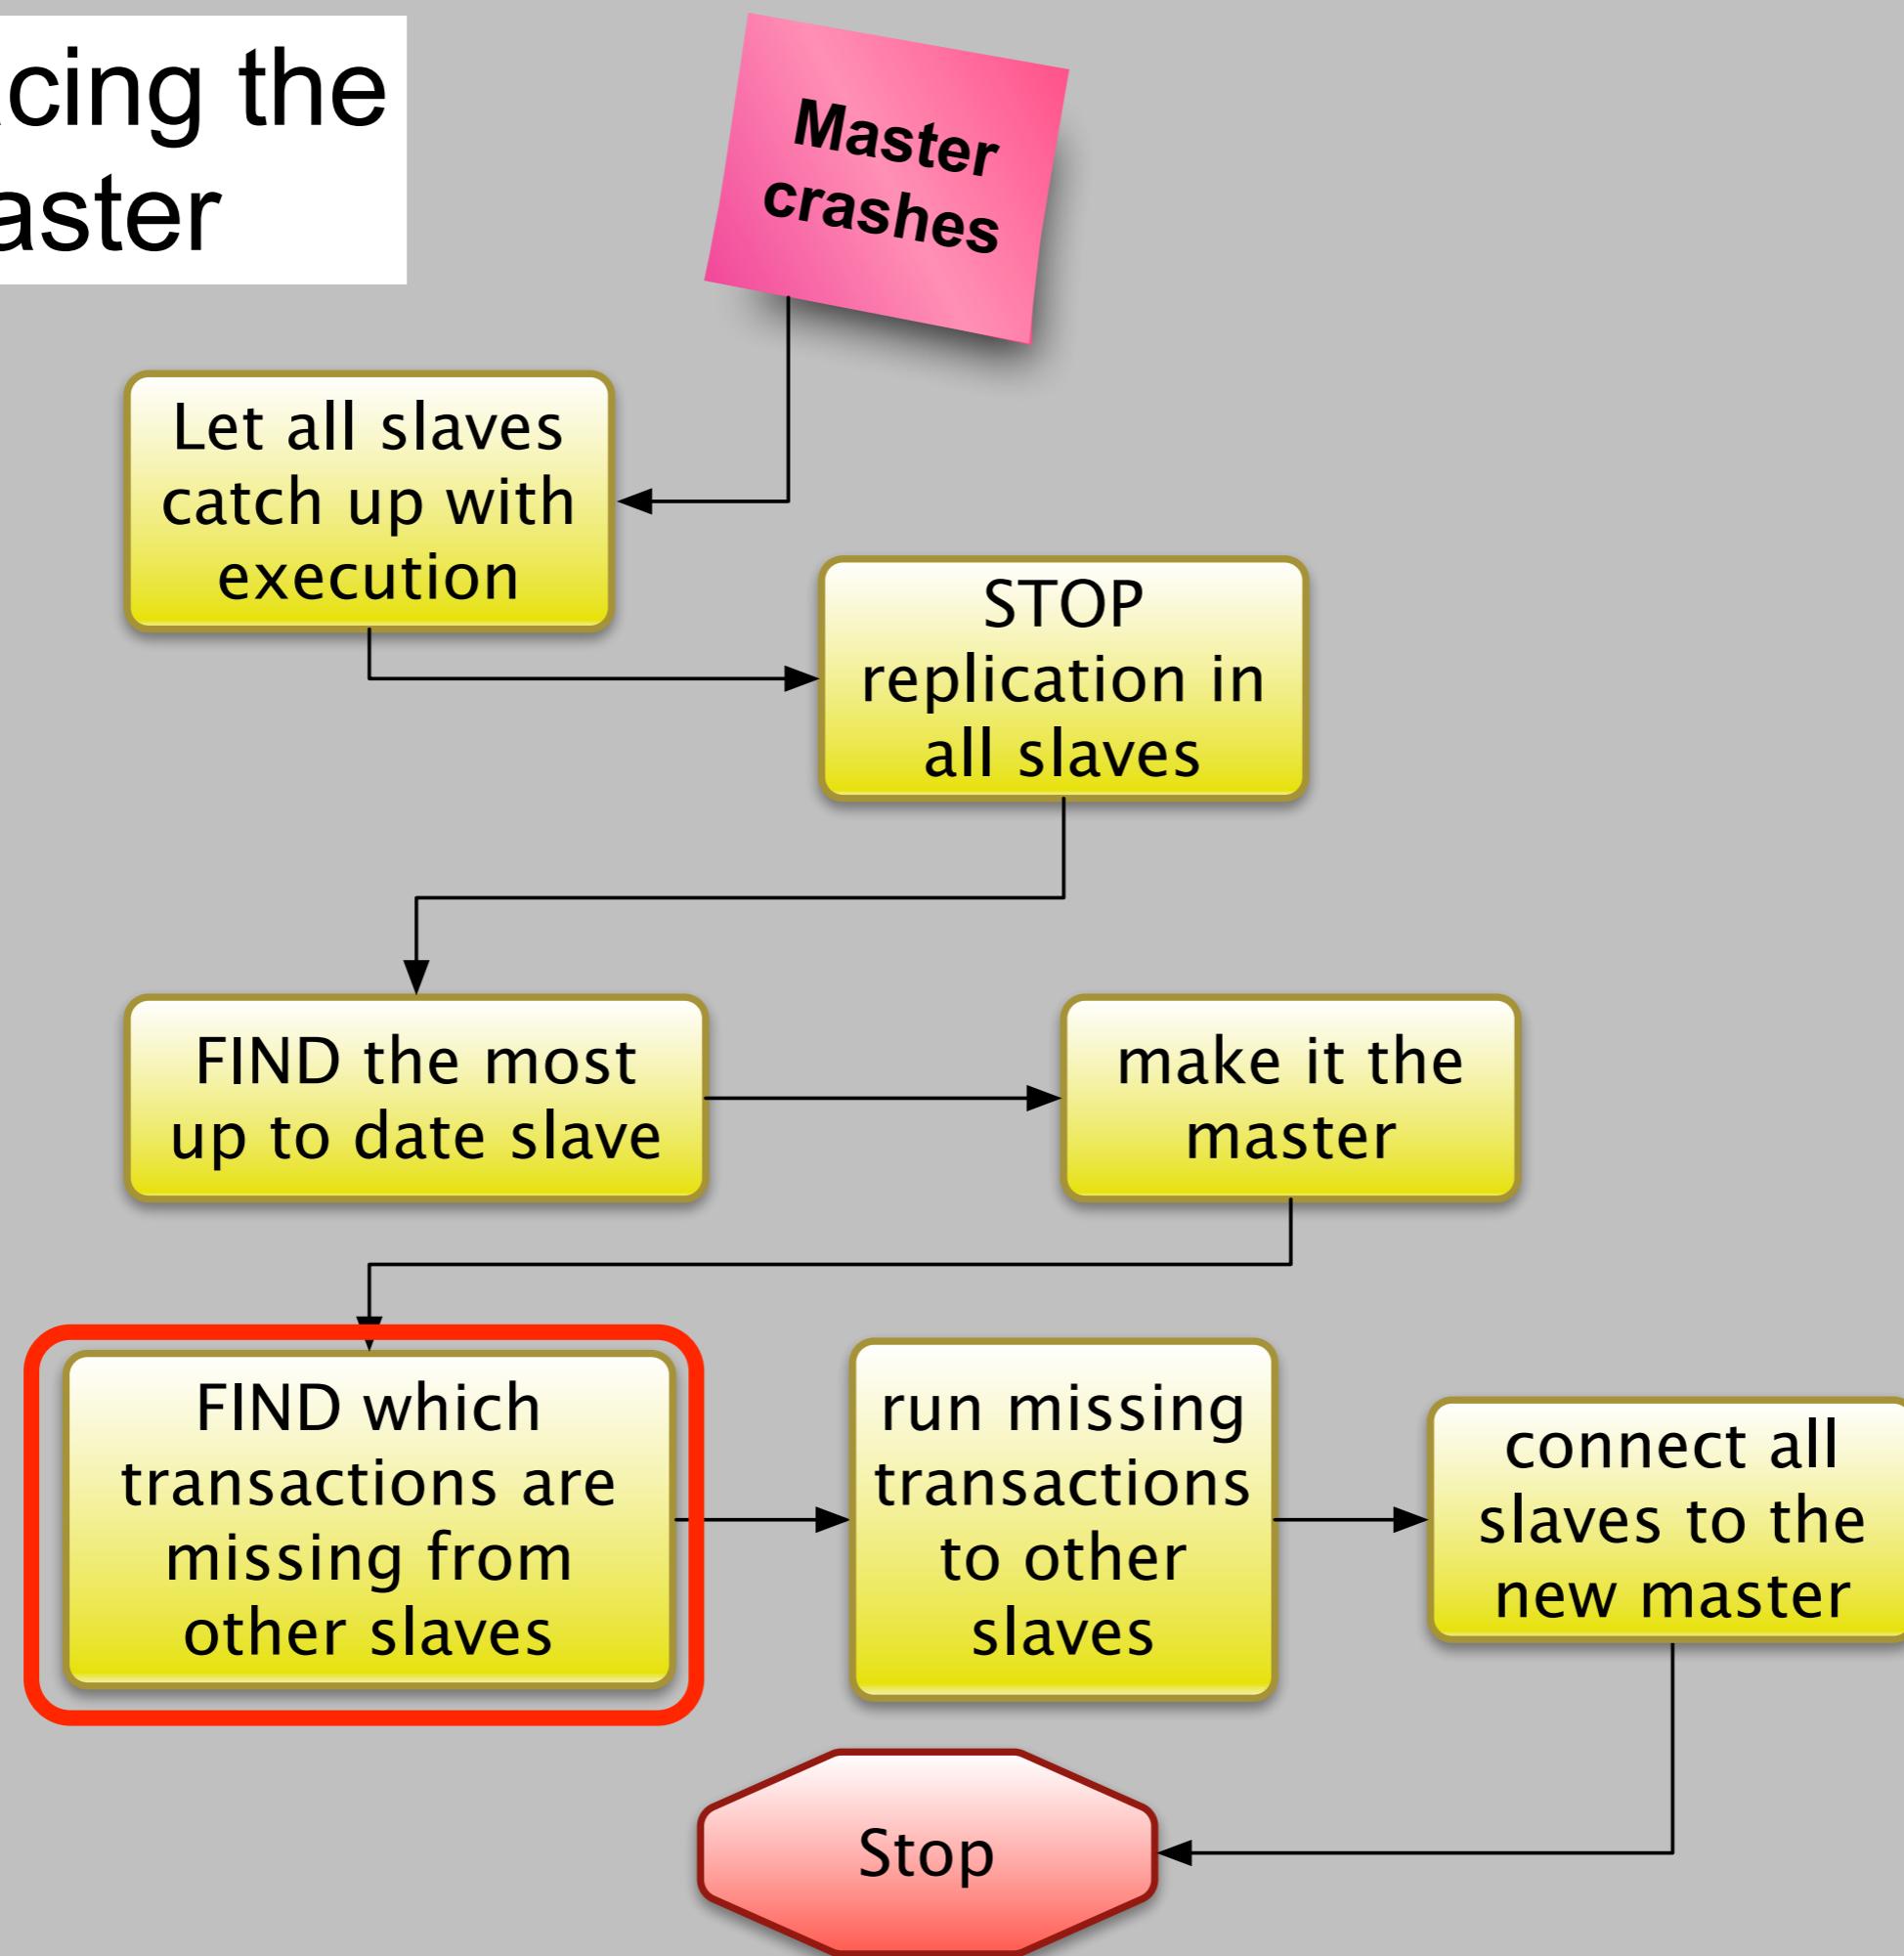
- master replacement
- backup
- better reads

Gotchas, tips, and tricks

More info

- What to read
- More replication sessions

Replacing the master



Planned master switch

- Stop accepting writes
- Wait until all slaves have caught up
- Stop replication in all slaves
- Promote a slave to master
- Point all slaves to the new master

Changing a failed master

- Pre-requisite:
- `log_bin` and `log_slave_updates` must be enabled in all the slaves
- If not, there is more manual labor

Changing a failed master (1)

- Wait until all slaves have caught up
- Identify the most advanced slave
- Make that slave the new master
- ... so far, so good

Changing a failed master (2)

- For each remaining slave:
- Find the missing statements
- Find the **LAST** statement replicated by the slave
- Find the same statement in the new master binlog (*)
- get the position of the **NEXT** statement

(*) if `log_slave_updates` was not enabled, you need to convert the relay log statements to SQL and do the next step manually

Changing a failed master (3)

- For each remaining slave:
- Apply the missing statements
 - `CHANGE MASTER TO`
`master_host="new_master_hostname",`
`master_port=new_master_port,`
`master_log_file="mysql-bin.xxxxxxx",`
`master_log_pos=YYYYY`

Reasons for complexity

- No global transaction ID

What is a global transaction ID

- A unique identifier of a transaction
- Unique for the whole cluster, not for each node
- Generated by the ultimate source (the master)
- Does not change when the transaction goes through an intermediate master

Why should you care

- Failure recovery
- MySQL DOES NOT have it
- What can you do? Either wait for MySQL 5.6 or use Tungsten Replicator.

AGENDA

Why replication

- The web economy
- Scaling out

How to set replication

- From single server to replication
- Adding a slave

Using Replication

- Binary log formats
- What gets replicated and how
- Replication awareness

Managing replication

- Monitoring
- Log management
- Replacing a slave
- Replacing a master

What Replication is for

- **master replacement**
- backup
- better reads

Gotchas, tips, and tricks

More info

- What to read
- More replication sessions

Slave = master replacement

- **if:**
 - there are no filters;
 - you are monitoring replication
 - you make sure data is consistent

Slave != backup

- There is no replacement for a good backup.
- Data loss due to a mistake in the master will propagate to the slaves.
- Disasters are always smarter than replication.
- Do your backups!
- Session on Wednesday at 12pm: "**Be a Data Management Hero with Good Backups!**"

AGENDA

Why replication

- The web economy
- Scaling out

How to set replication

- From single server to replication
- Adding a slave

Using Replication

- Binary log formats
- What gets replicated and how
- Replication awareness

Managing replication

- Monitoring
- Log management
- Replacing a slave
- Replacing a master

What Replication is for

- master replacement
- **backup**
- better reads

Gotchas, tips, and tricks

More info

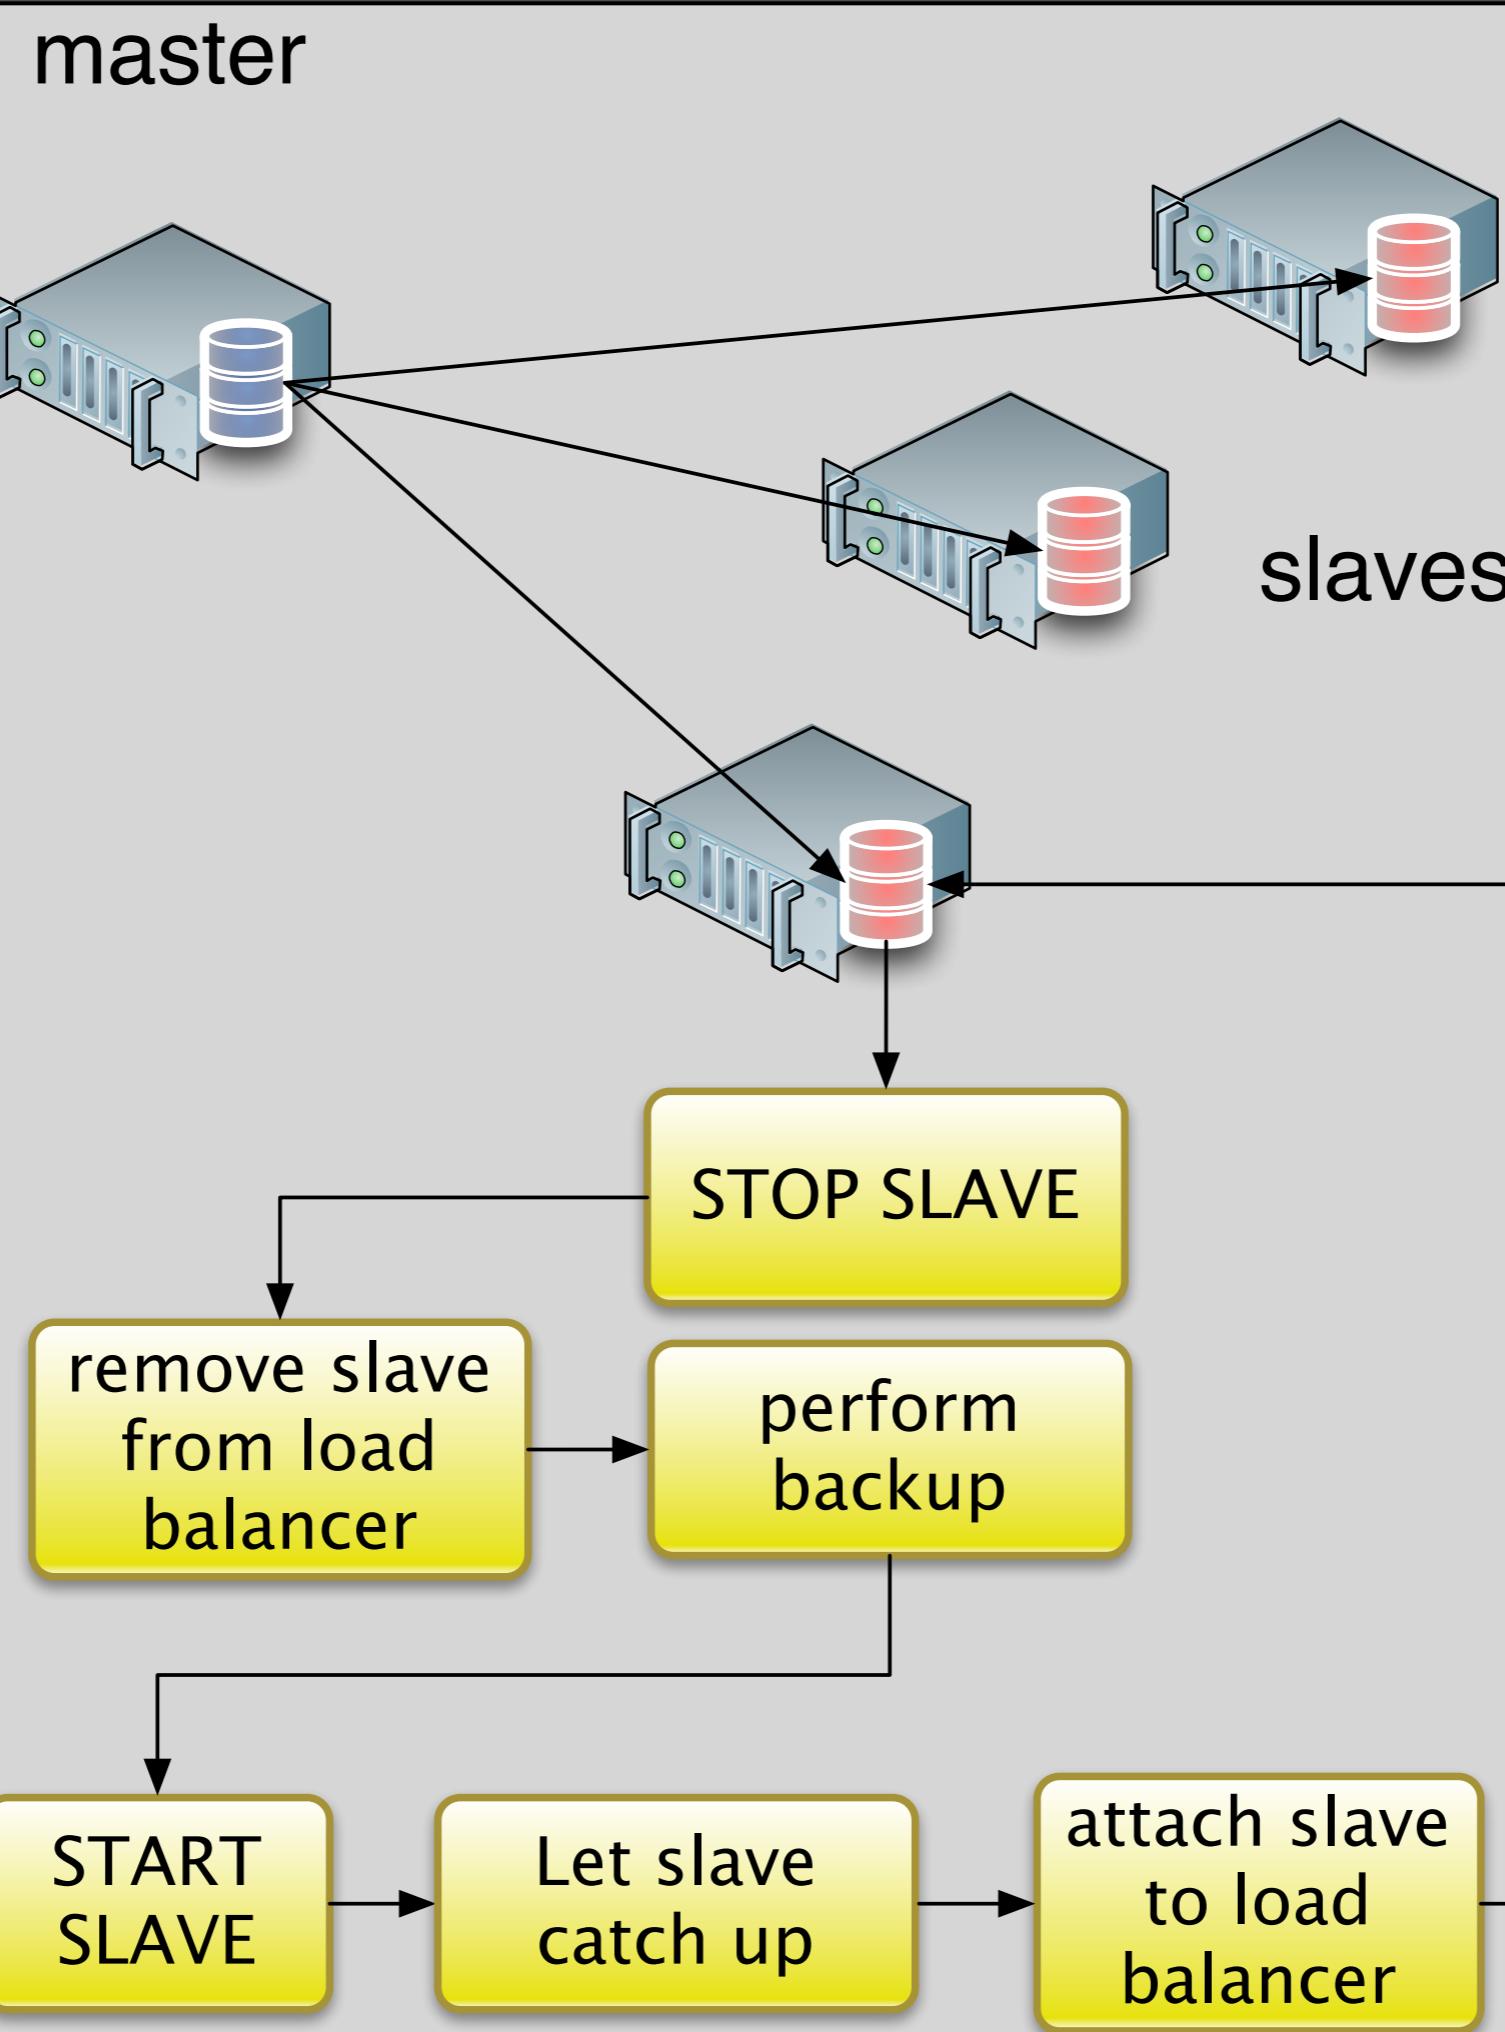
- What to read
- More replication sessions

Replication

=

let the slave do the dirty work

backup



AGENDA

Why replication

- The web economy
- Scaling out

How to set replication

- From single server to replication
- Adding a slave

Using Replication

- Binary log formats
- What gets replicated and how
- Replication awareness

Managing replication

- Monitoring
- Log management
- Replacing a slave
- Replacing a master

What Replication is for

- master replacement
- backup
- **better reads**

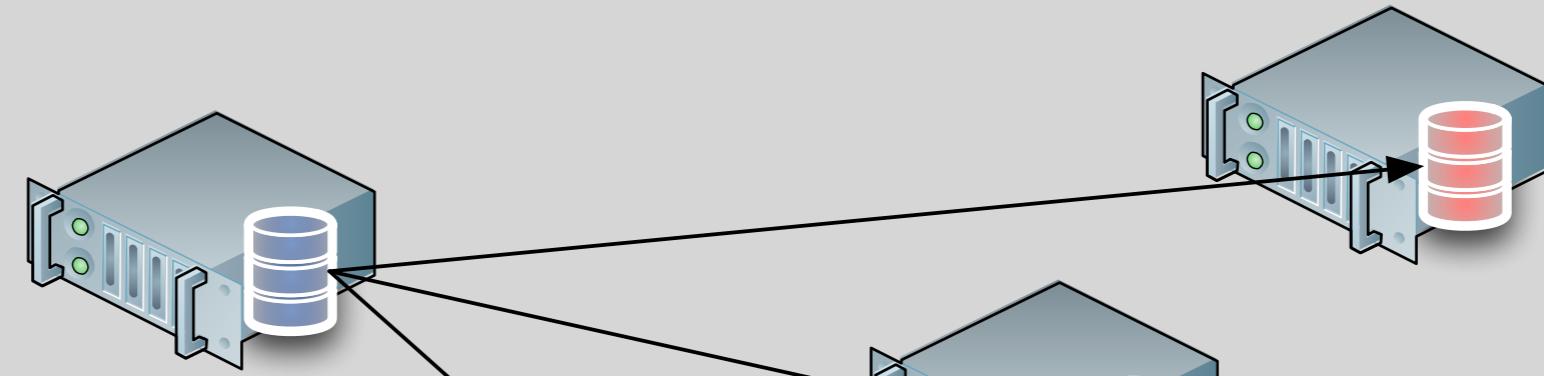
Gotchas, tips, and tricks

- What to read
- More replication sessions

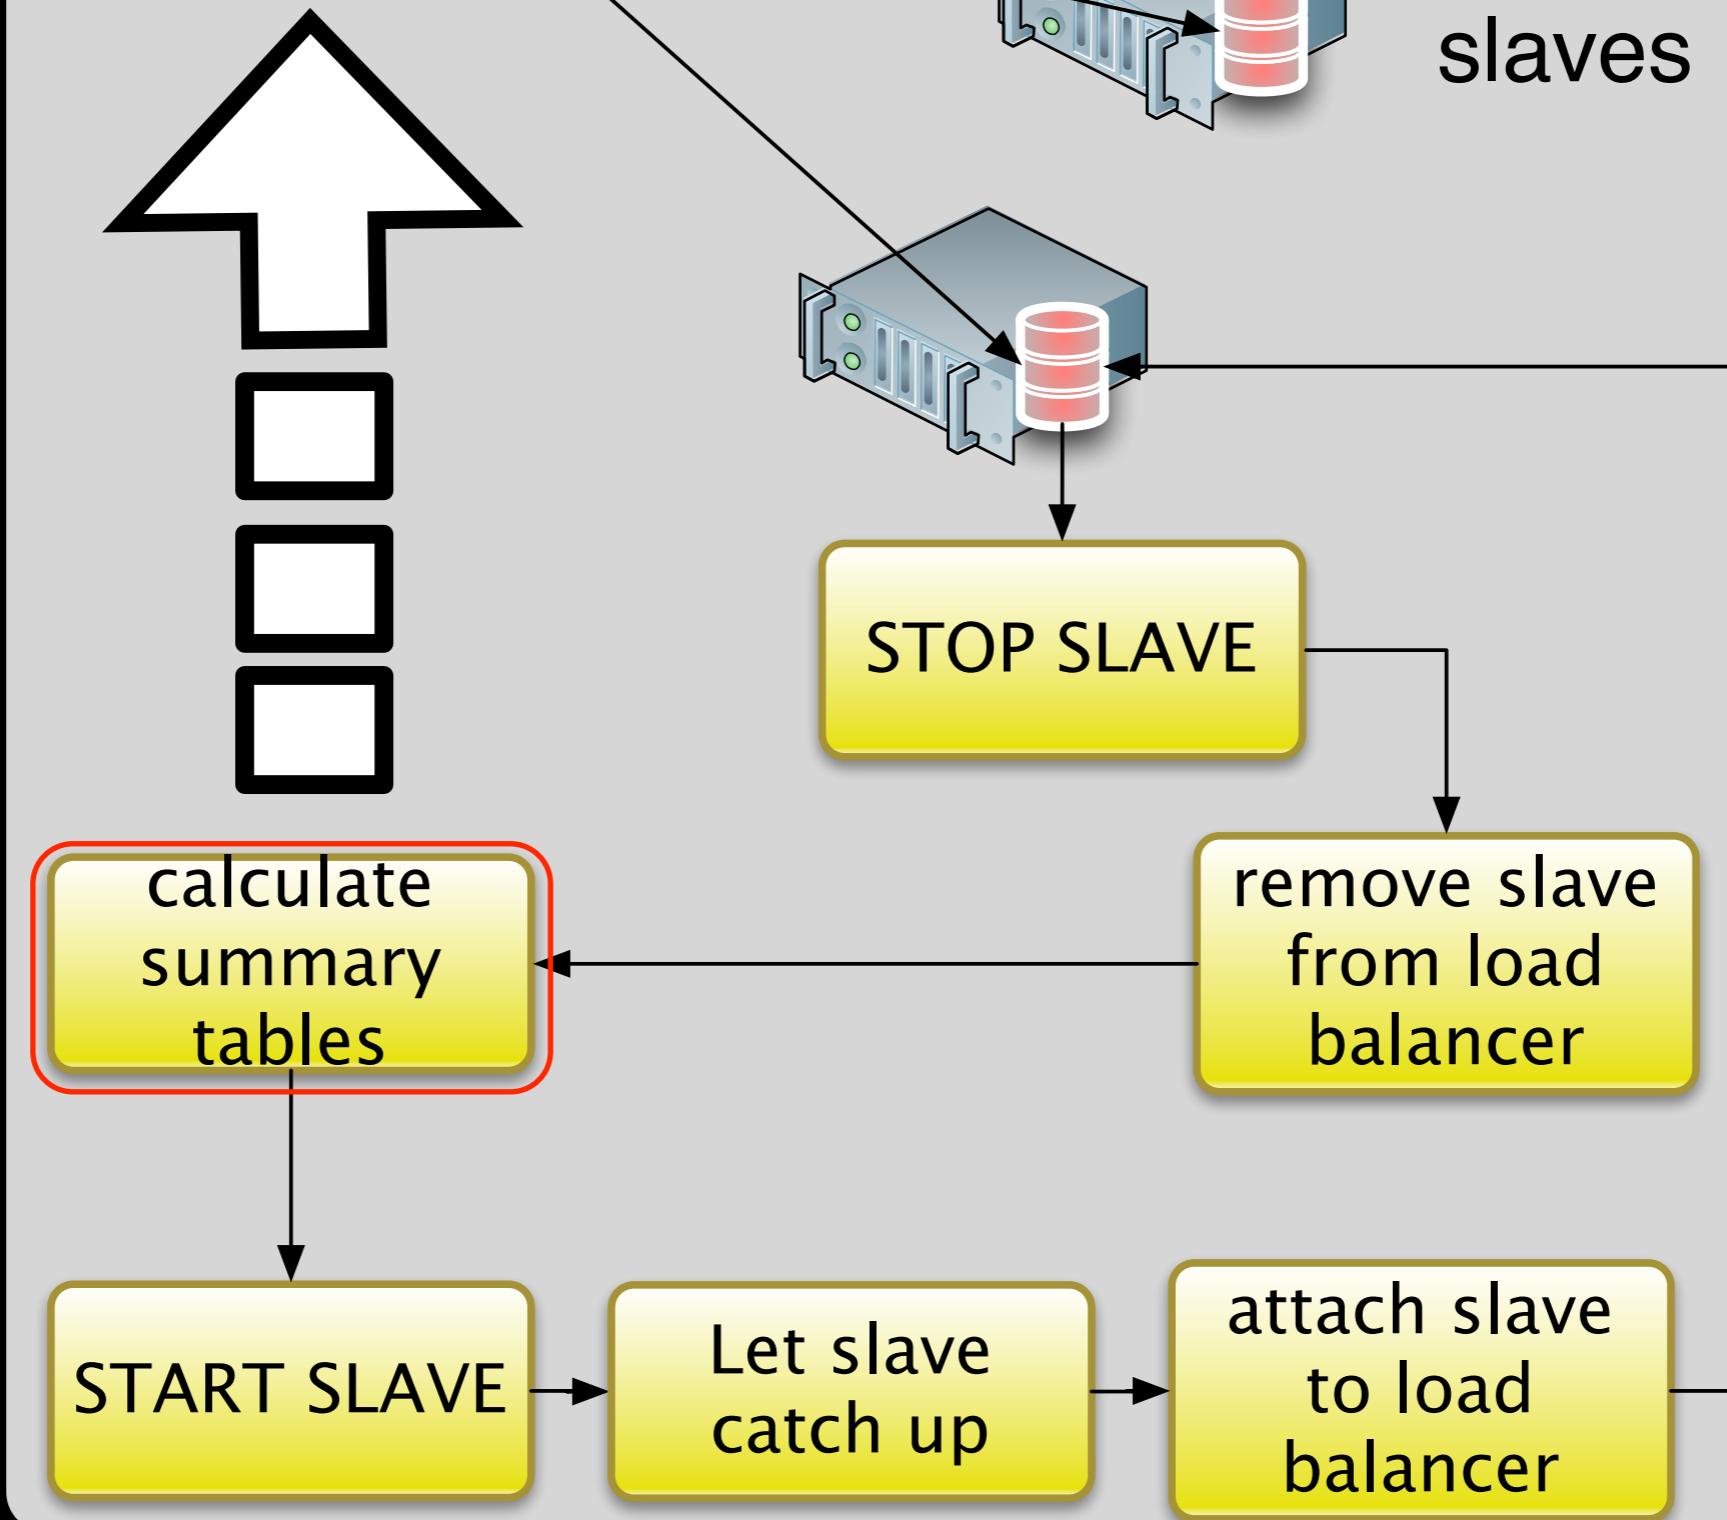
More info

make summary tables

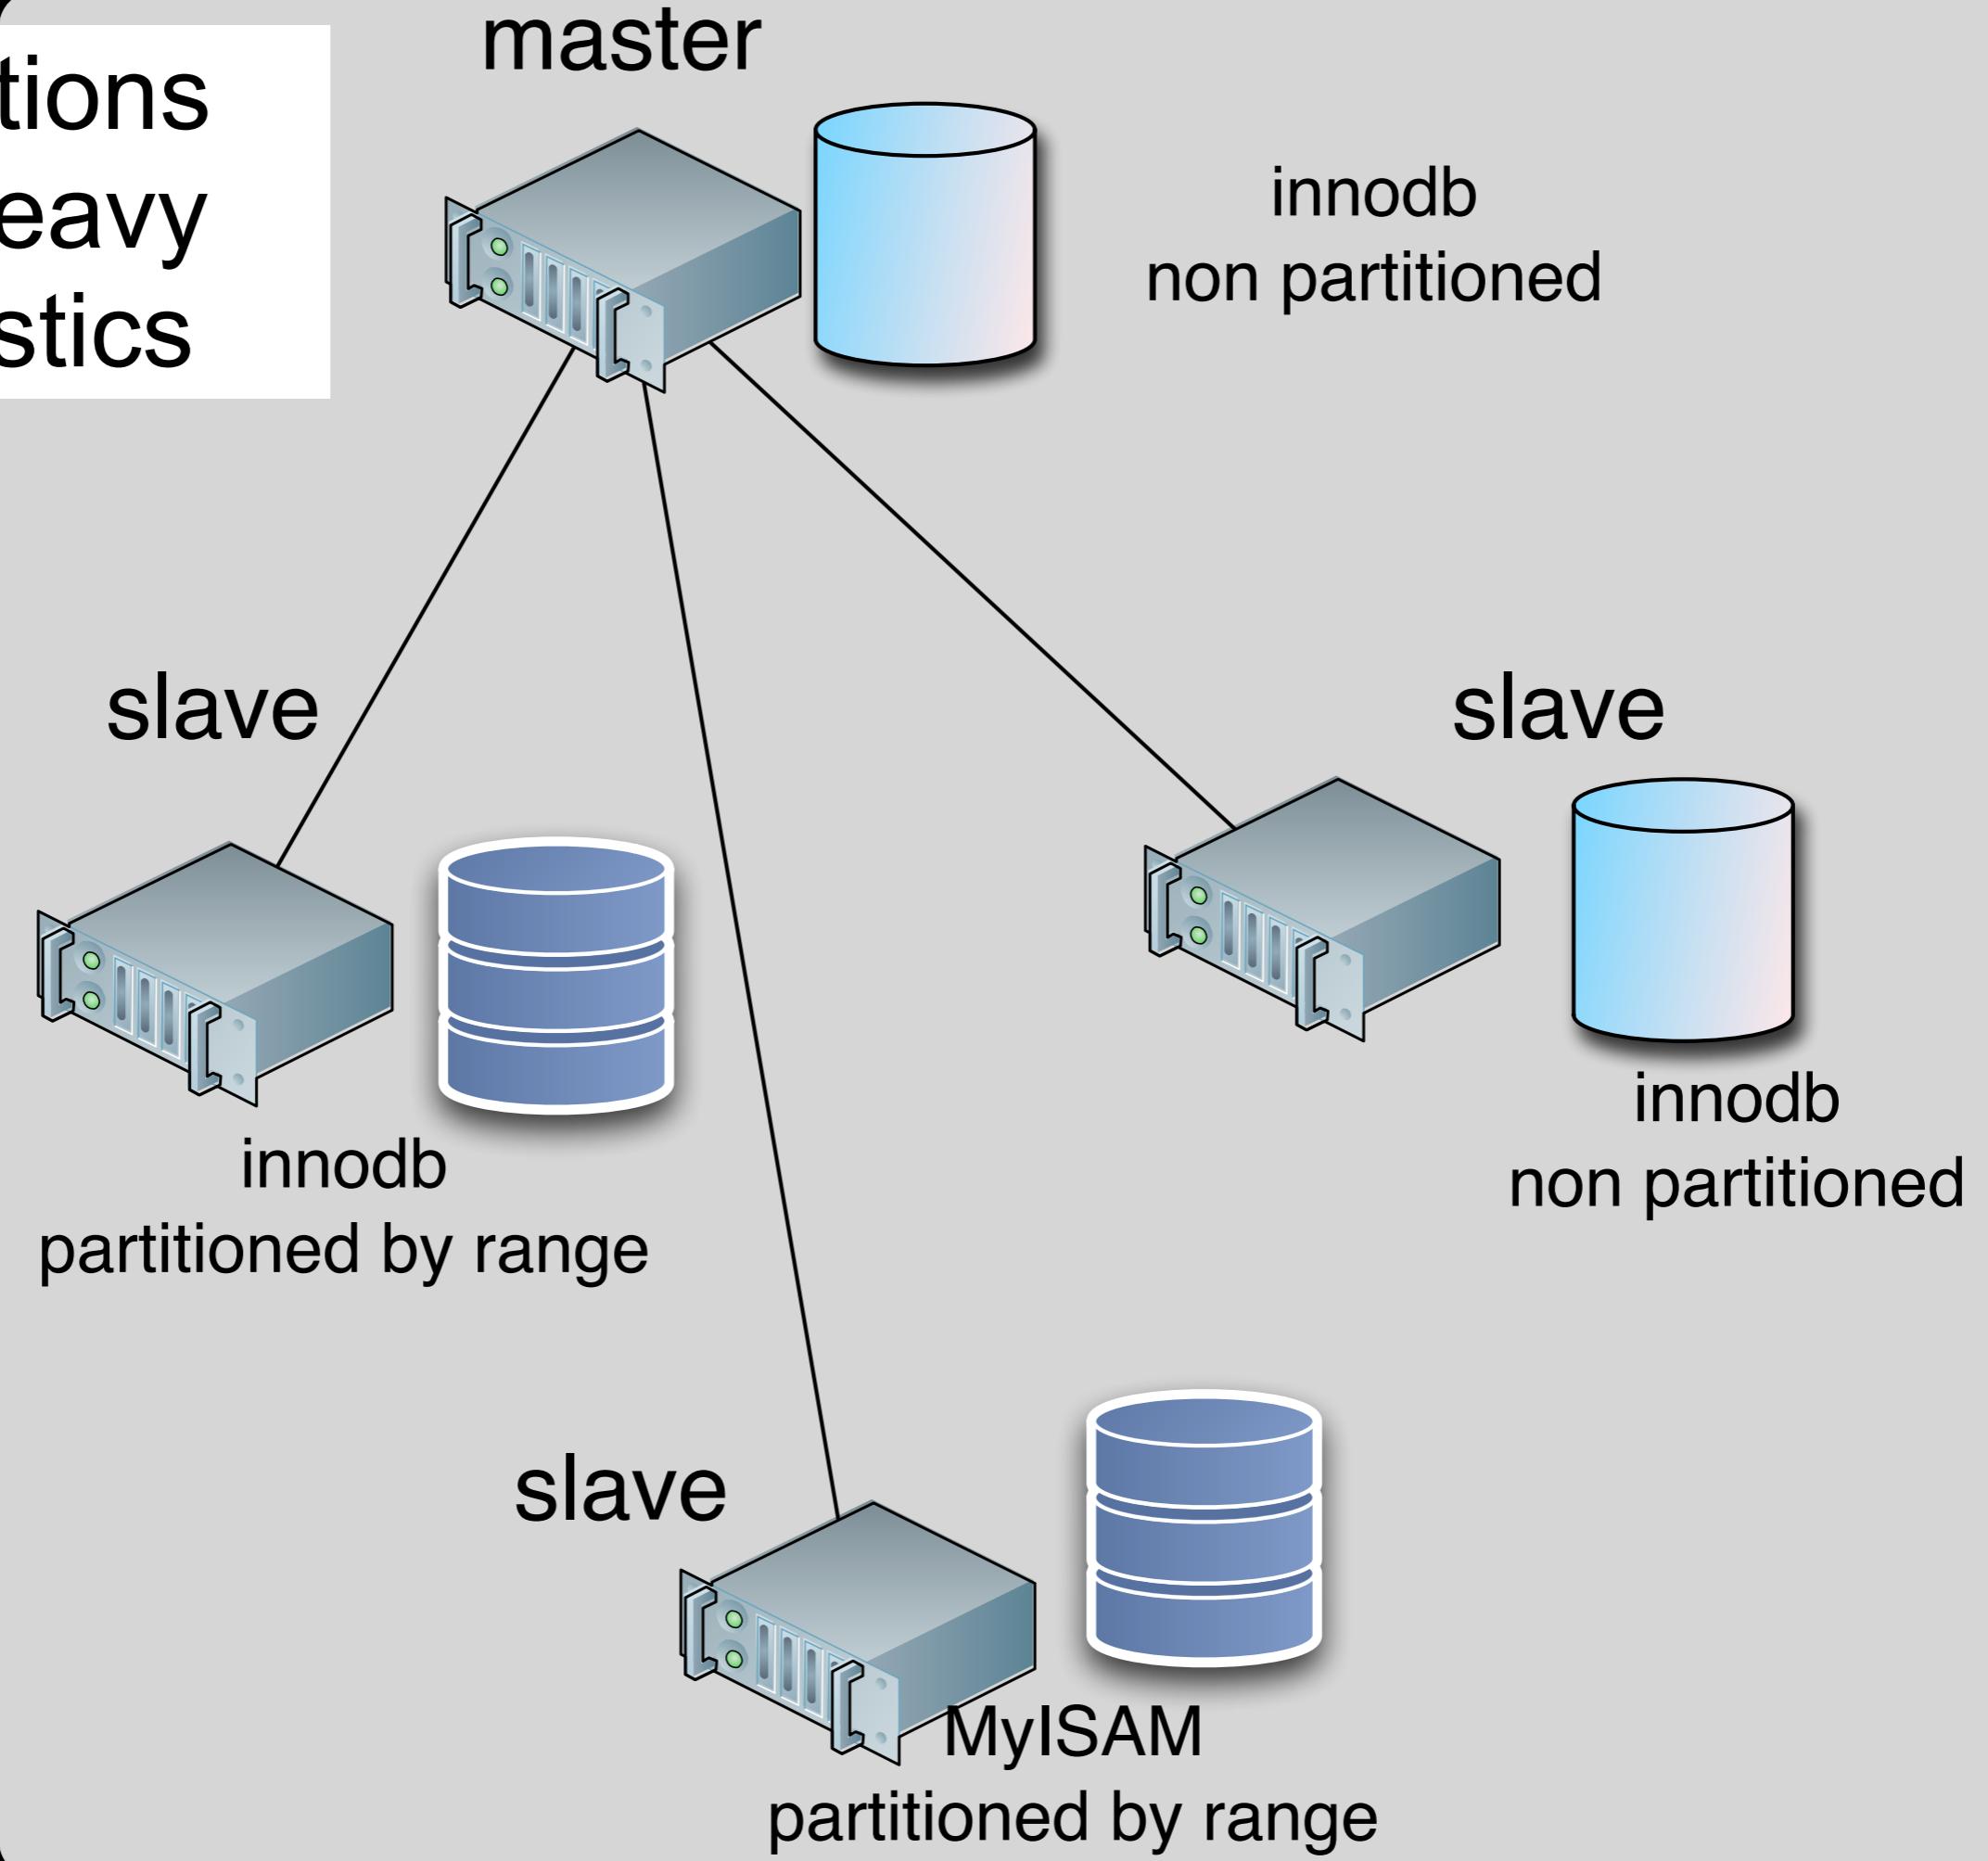
master



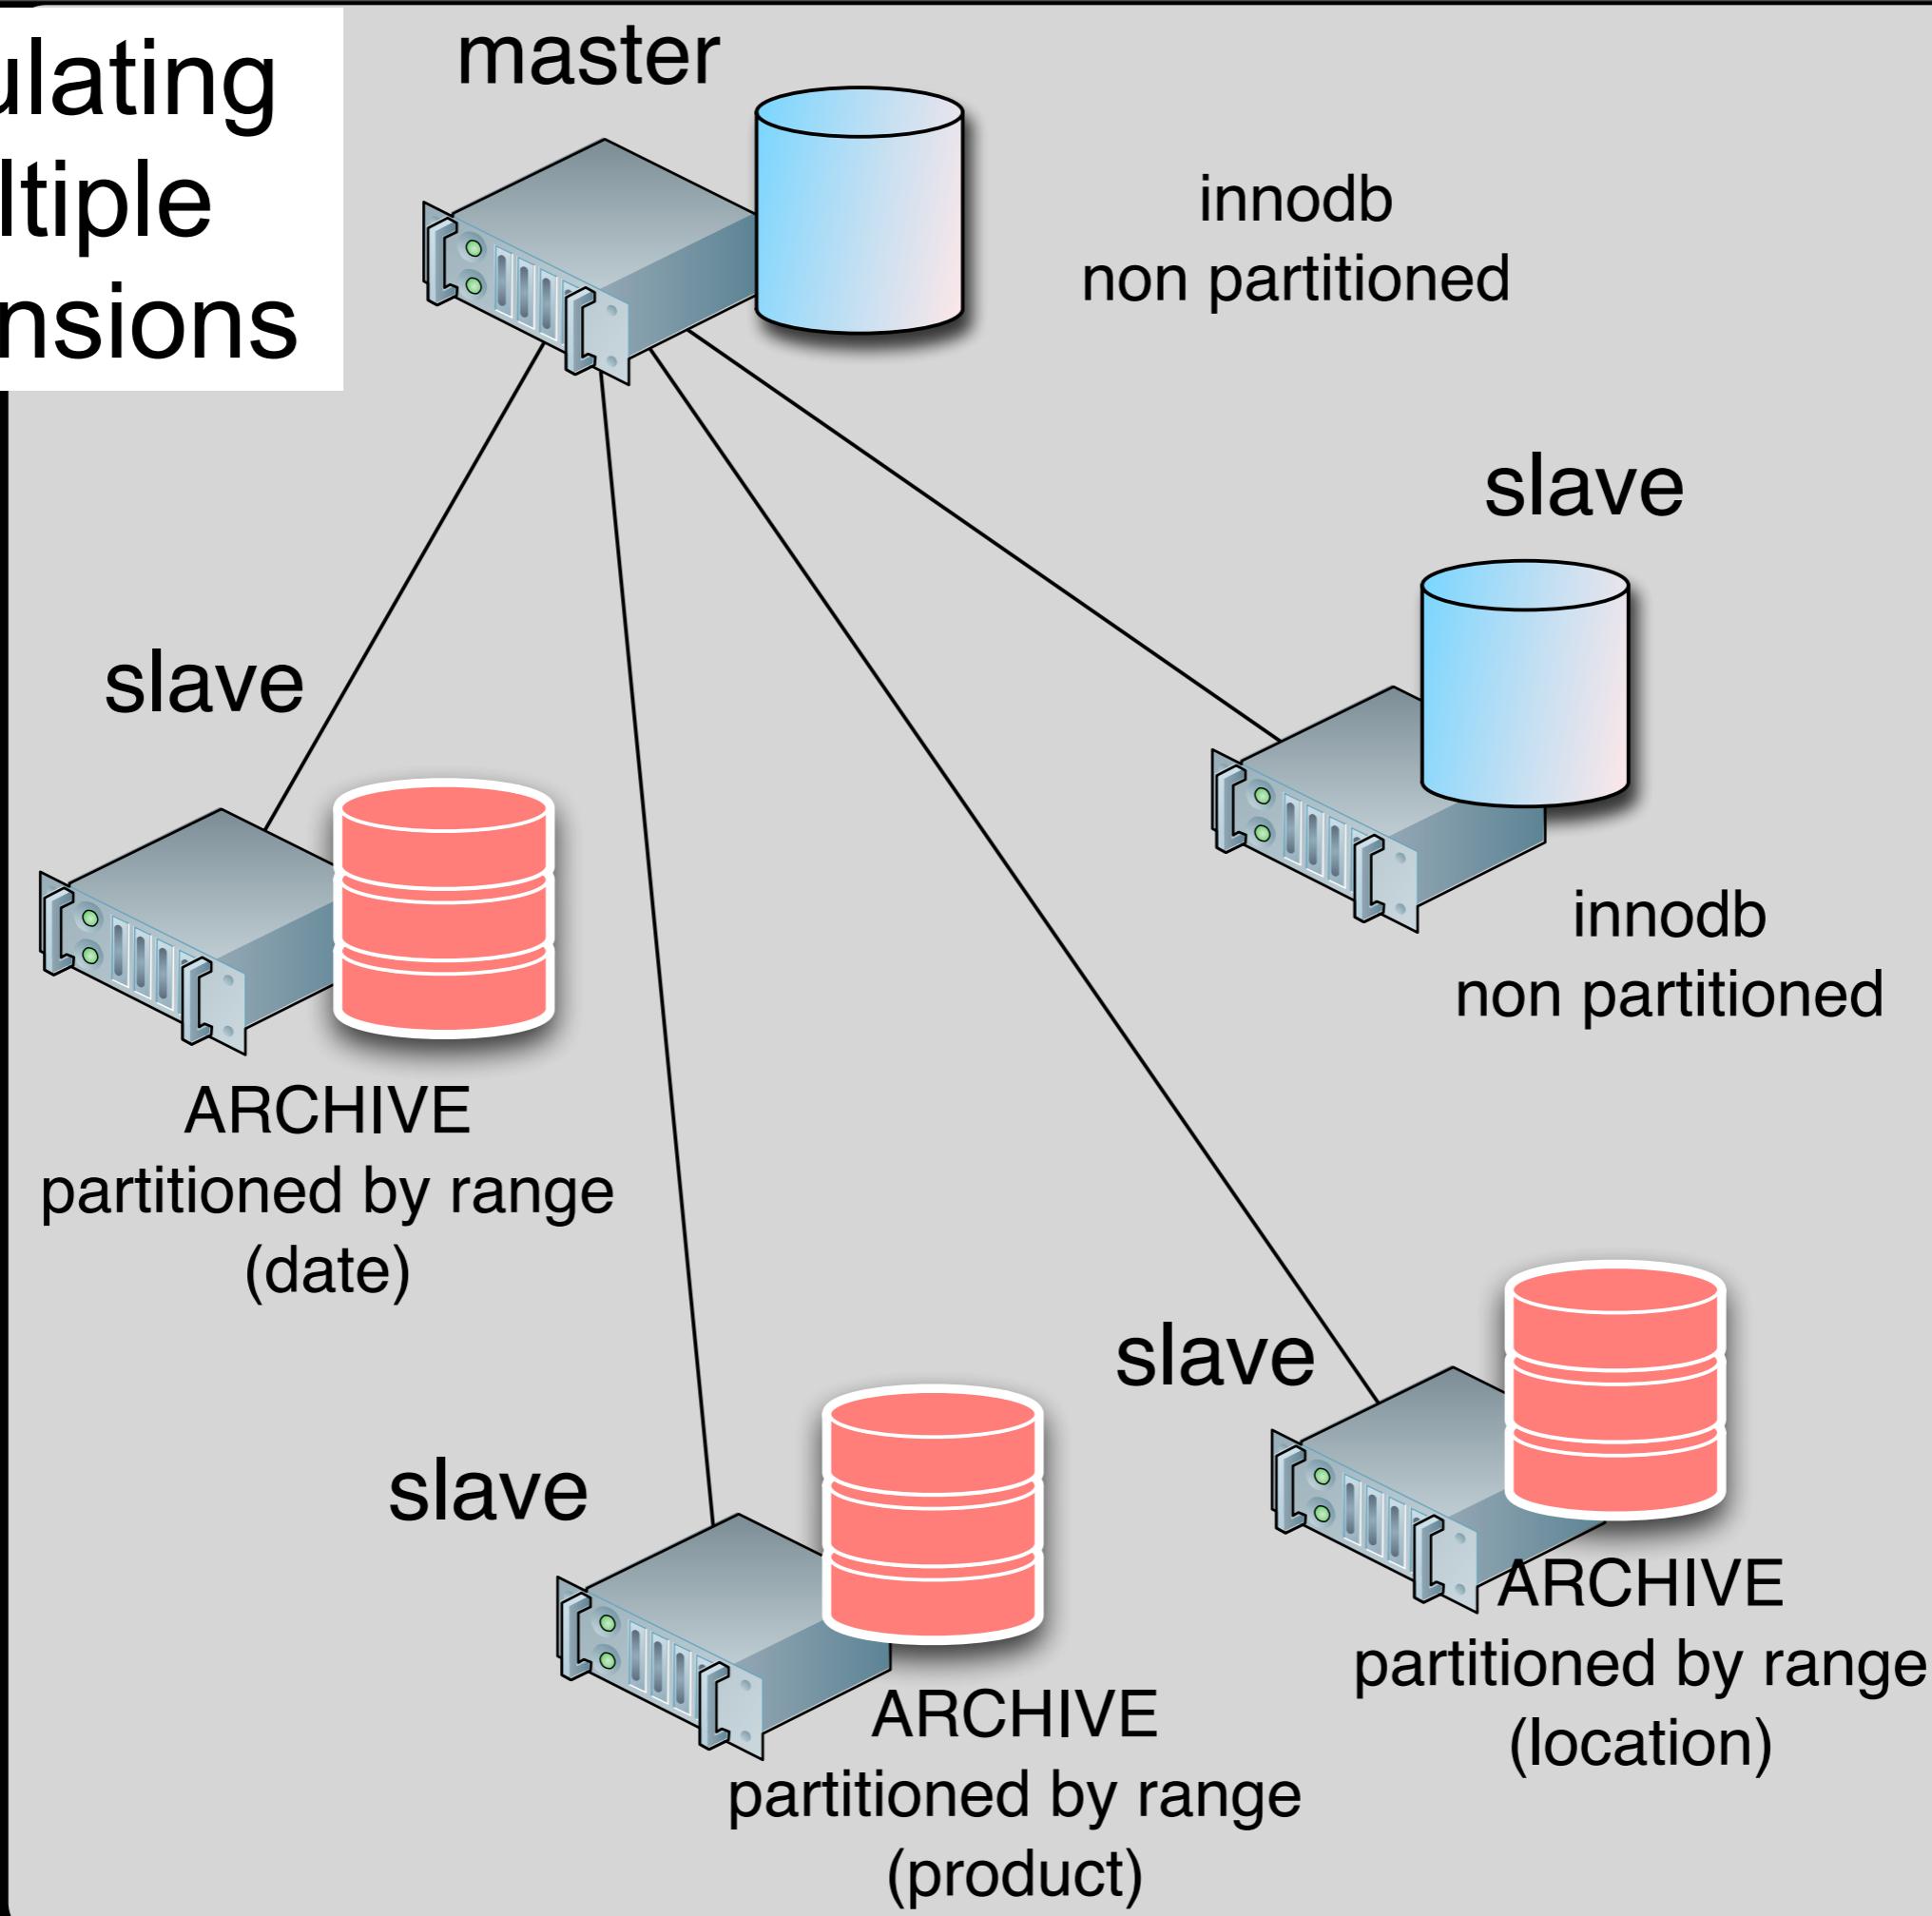
slaves



Partitions for heavy statistics



Simulating multiple dimensions



AGENDA

Why replication

- The web economy
- Scaling out

How to set replication

- From single server to replication
- Adding a slave

Using Replication

- Binary log formats
- What gets replicated and how
- Replication awareness

Managing replication

- Monitoring
- Log management
- Replacing a slave
- Replacing a master

What Replication is for

- master replacement
- backup
- better reads

Gotchas, tips, and tricks

- What to read
- More replication sessions

More info

READ-ONLY slave

- Good practice: make a slave read-only
- Caveat 1: if promoting a slave, you need to remove the read-only option;
- Caveat 1: a user with SUPER privilege can write to a read-only slave

Default engine

- Caveat: default storage engine is not replicated.
- Example: master with default engine=Innodb and slave with default engine=MyISAM
- a CREATE TABLE without the ENGINE clause will use different engines on master and slave

AGENDA

master-to-master and circular replication

DEMO

Common filters

- on the master
 - binlog-do-db
 - binlog ignore-db
- on the slave
 - replicate-do-db
 - replicate-ignore-db

Common filters

- on the master
 - binlog-do-db
 - binlog ignore-db
- on the slave
 - replicate-do-db
 - replicate-ignore-db

**DO NOT
USE!**

do Filters

What happens when everything goes right

transactions from client

all data comes to the master

binlog-do-*

The master logs only some of the transactions

**replicate-
do-***

The slave only gets some of the above

ignore Filters

transactions from client

all data comes to the master

binlog-ignore-*

The master logs all the transactions, except the ones that should be ignored

**replicate-
ignore-***

The slave further filters some of the above

filters

What happens when something goes wrong

Master my.cnf:
binlog-do-db=foo



`use foo;`
`insert into bar.t1 values (1)`
it is replicated (breaks replication)



`use bar;`
`insert into foo.t1 values (1)`
is NOT replicated

filters

What happens when something goes wrong

slave my.cnf:
replicate-do-db=foo



`use foo;`
`insert into bar.t1 values (1)`
it is replicated (breaks replication)



`use bar;`
`insert into foo.t1 values (1)`
is NOT replicated

"Safer" filters

- on the slave:
 - replicate-wild-do-table=db_name.%
 - replicate-wild-do-table=foo%.bar%
 - replicate-wild-ignore-table=db_name.%
 - replicate-wild-ignore-table=foo%.bar%

"Safer" filters

- on the slave:
 - replicate-wild-do-table=db_name.%
 - replicate-wild-do-table=foo%.bar%
 - replicate-wild-ignore-table=db_name.%
 - replicate-wild-ignore-table=foo%.bar%

**DON'T MIX "do" and
"ignore" filters**

General rules of replication filters

- DON'T USE filters on the master.
- If you apply slave filters, the slave is not suitable for replacing a master or taking backups;
- Filters can break replication;
- Slave filters don't save bandwidth.

AGENDA

Why replication

- The web economy
- Scaling out

How to set replication

- From single server to replication
- Adding a slave

Using Replication

- Binary log formats
- What gets replicated and how
- Replication awareness

Managing replication

- Monitoring
- Log management
- Replacing a slave
- Replacing a master

What Replication is for

- master replacement
- backup
- better reads

Gotchas, tips, and tricks

More info

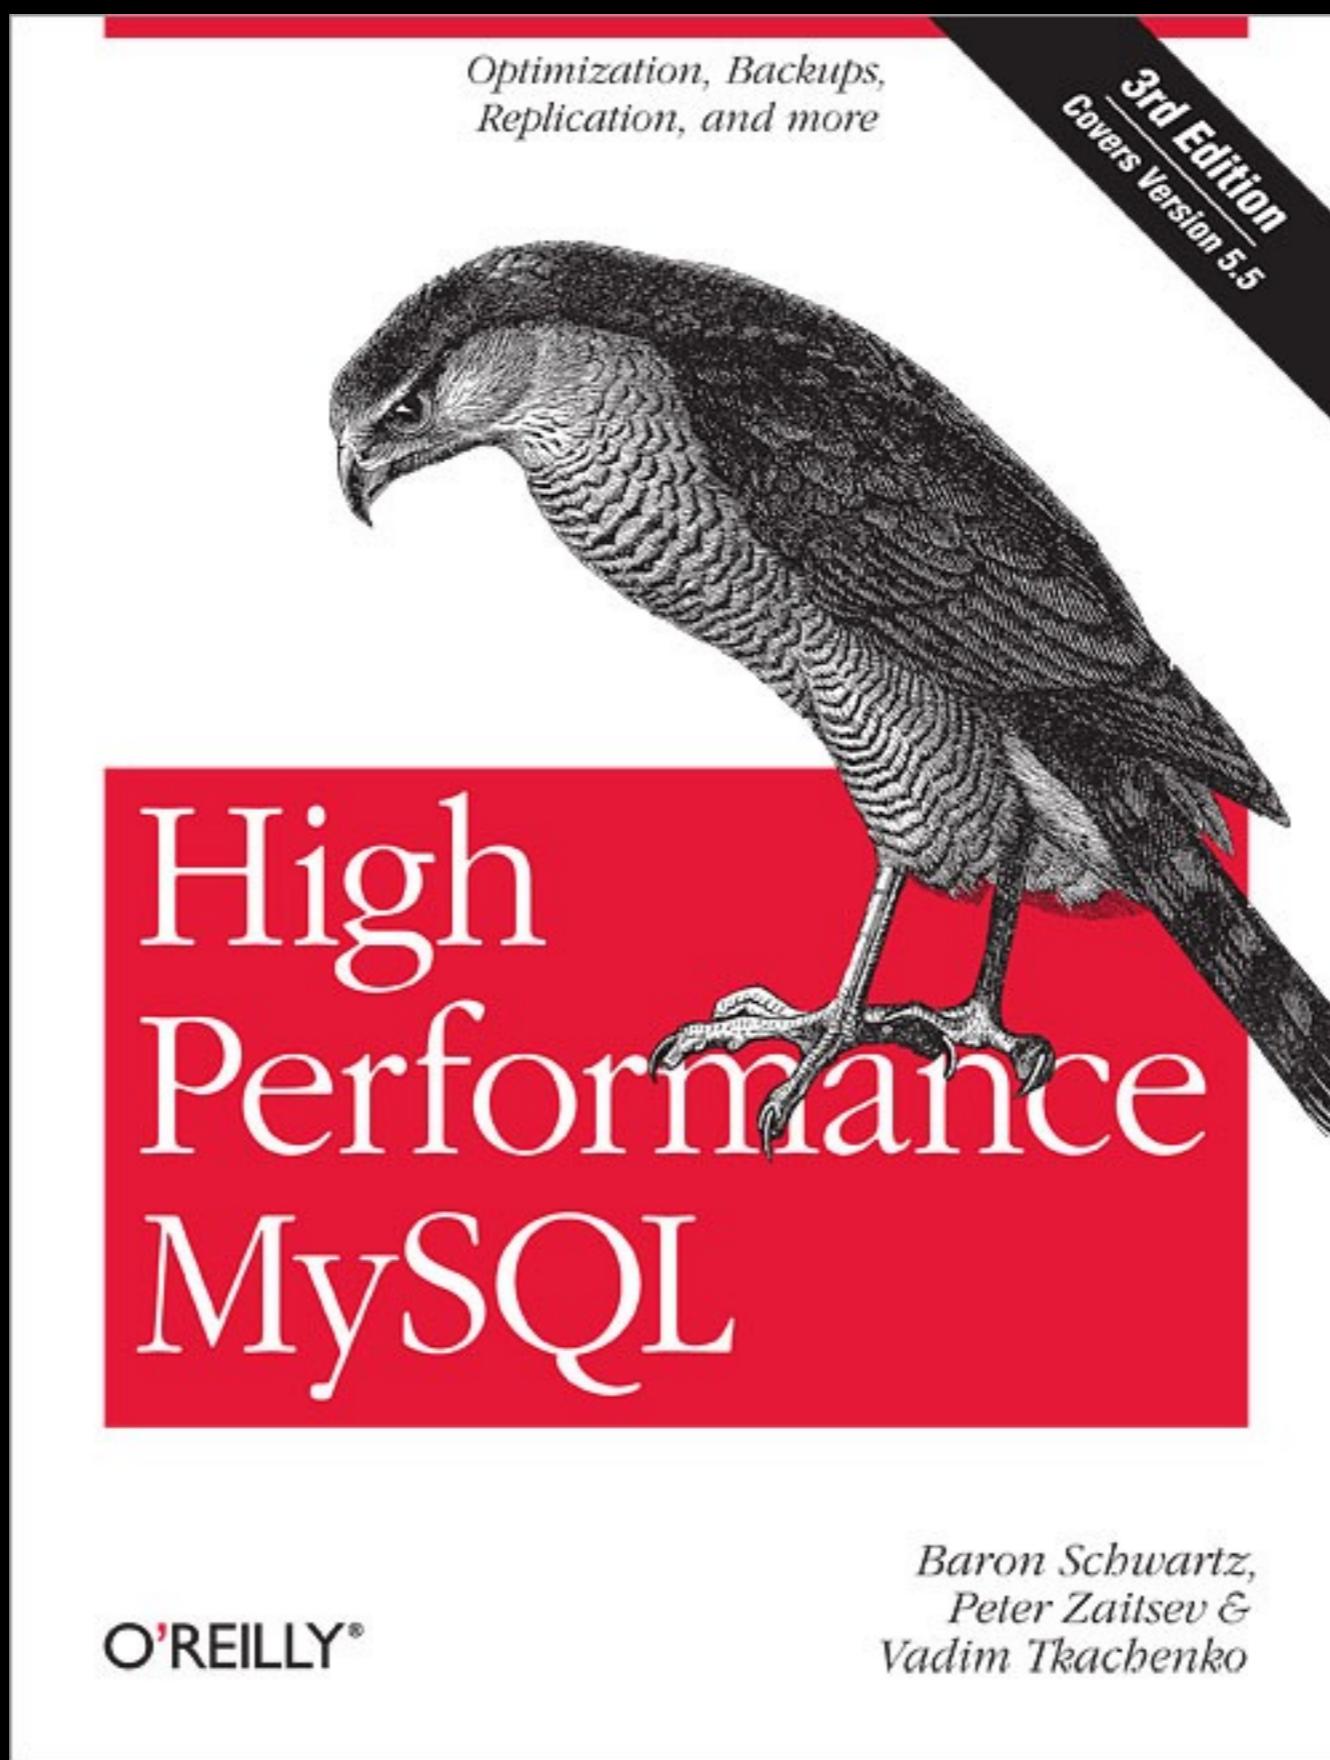
- **What to read**
- More replication sessions

Read more

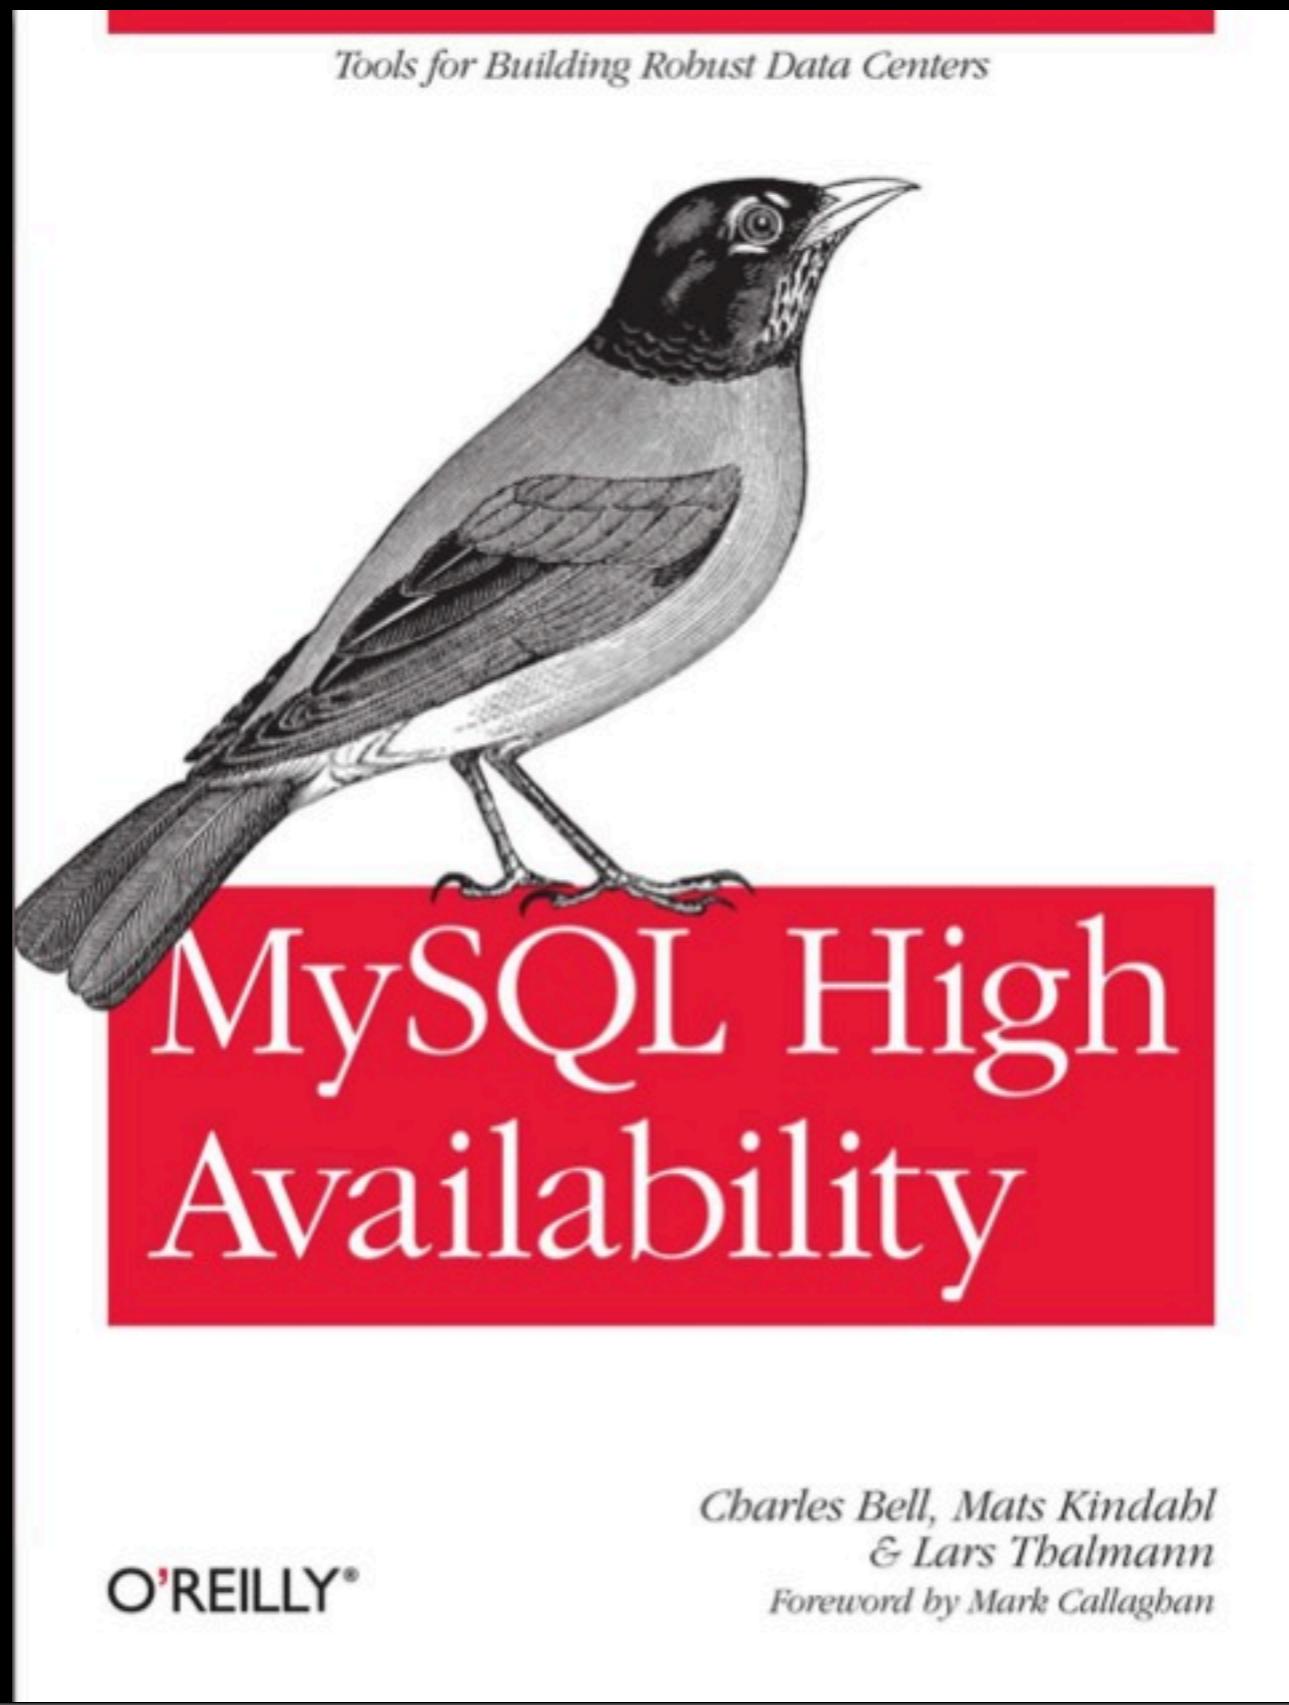
- The MySQL online manual

<http://dev.mysql.com/doc>

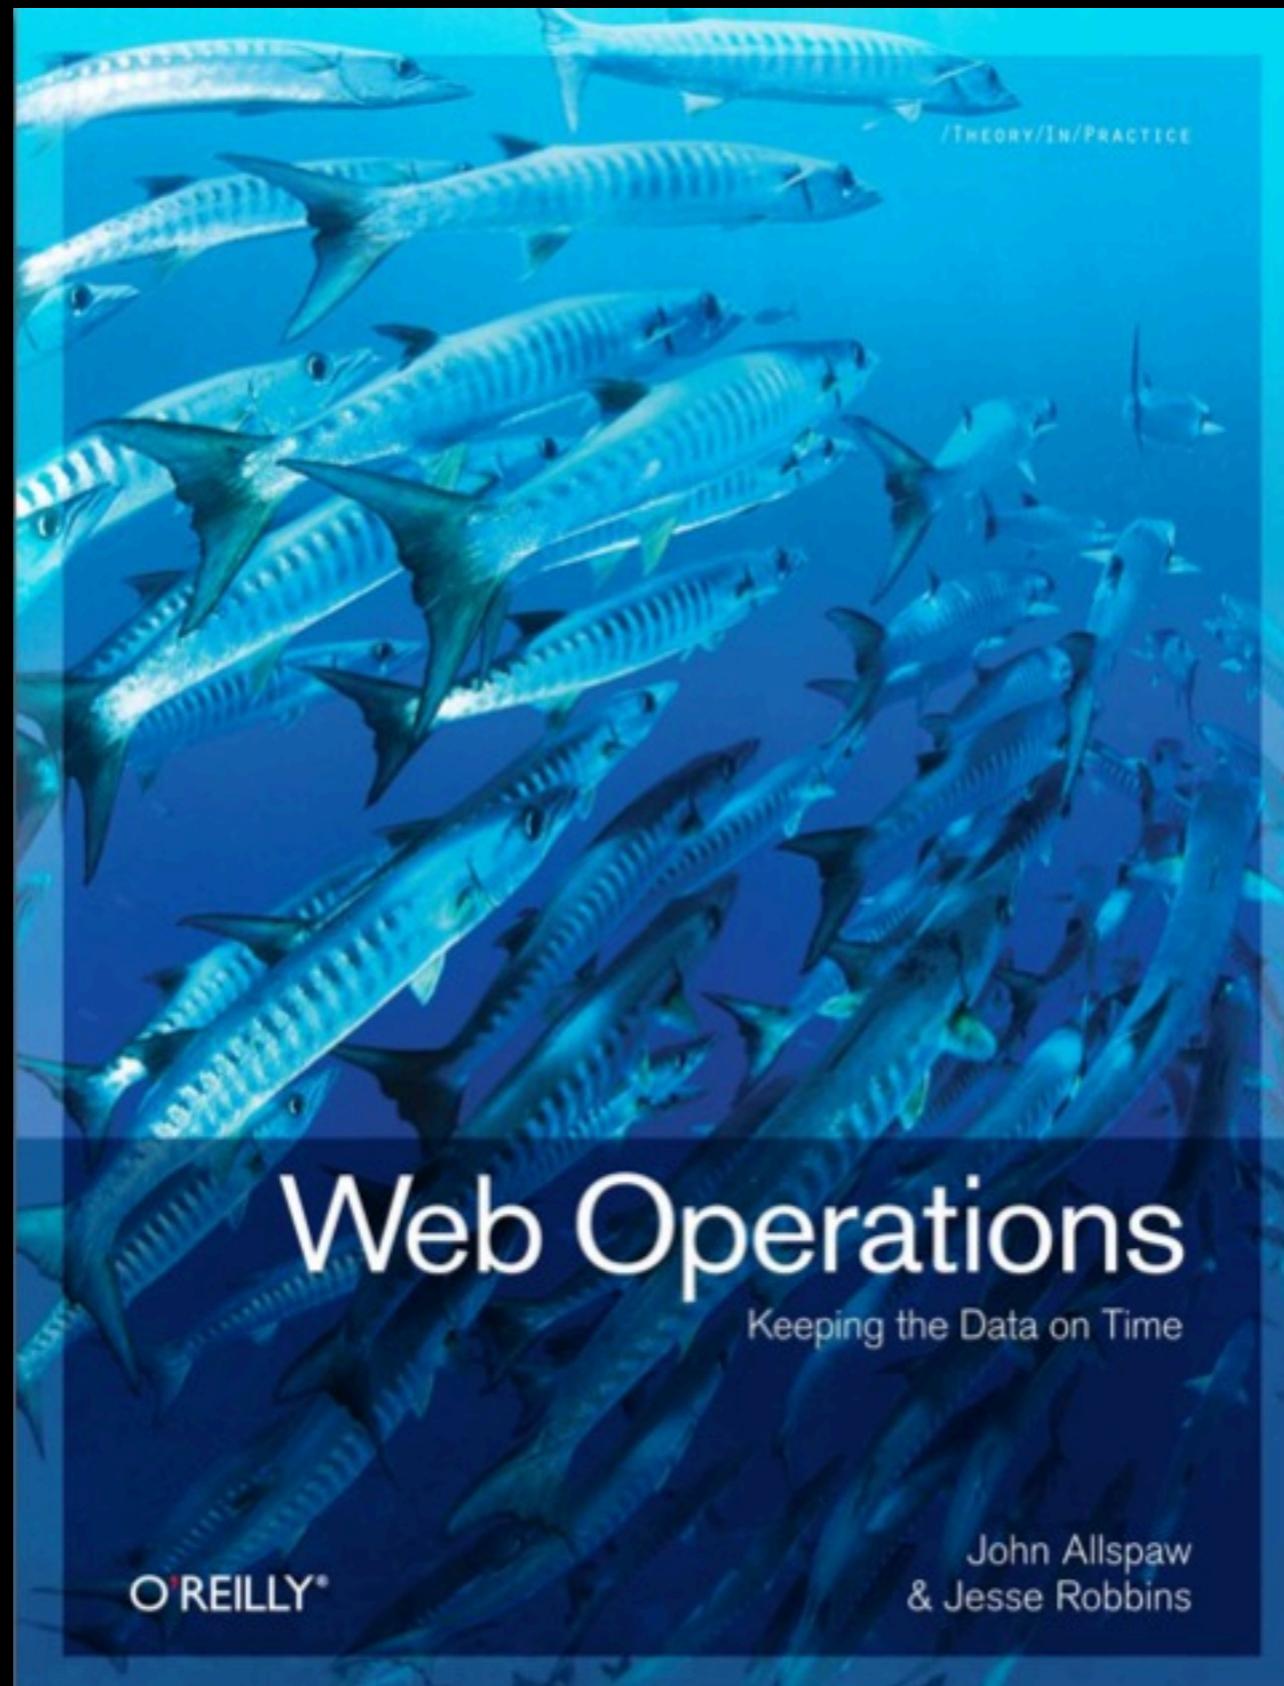
High Performance MySQL



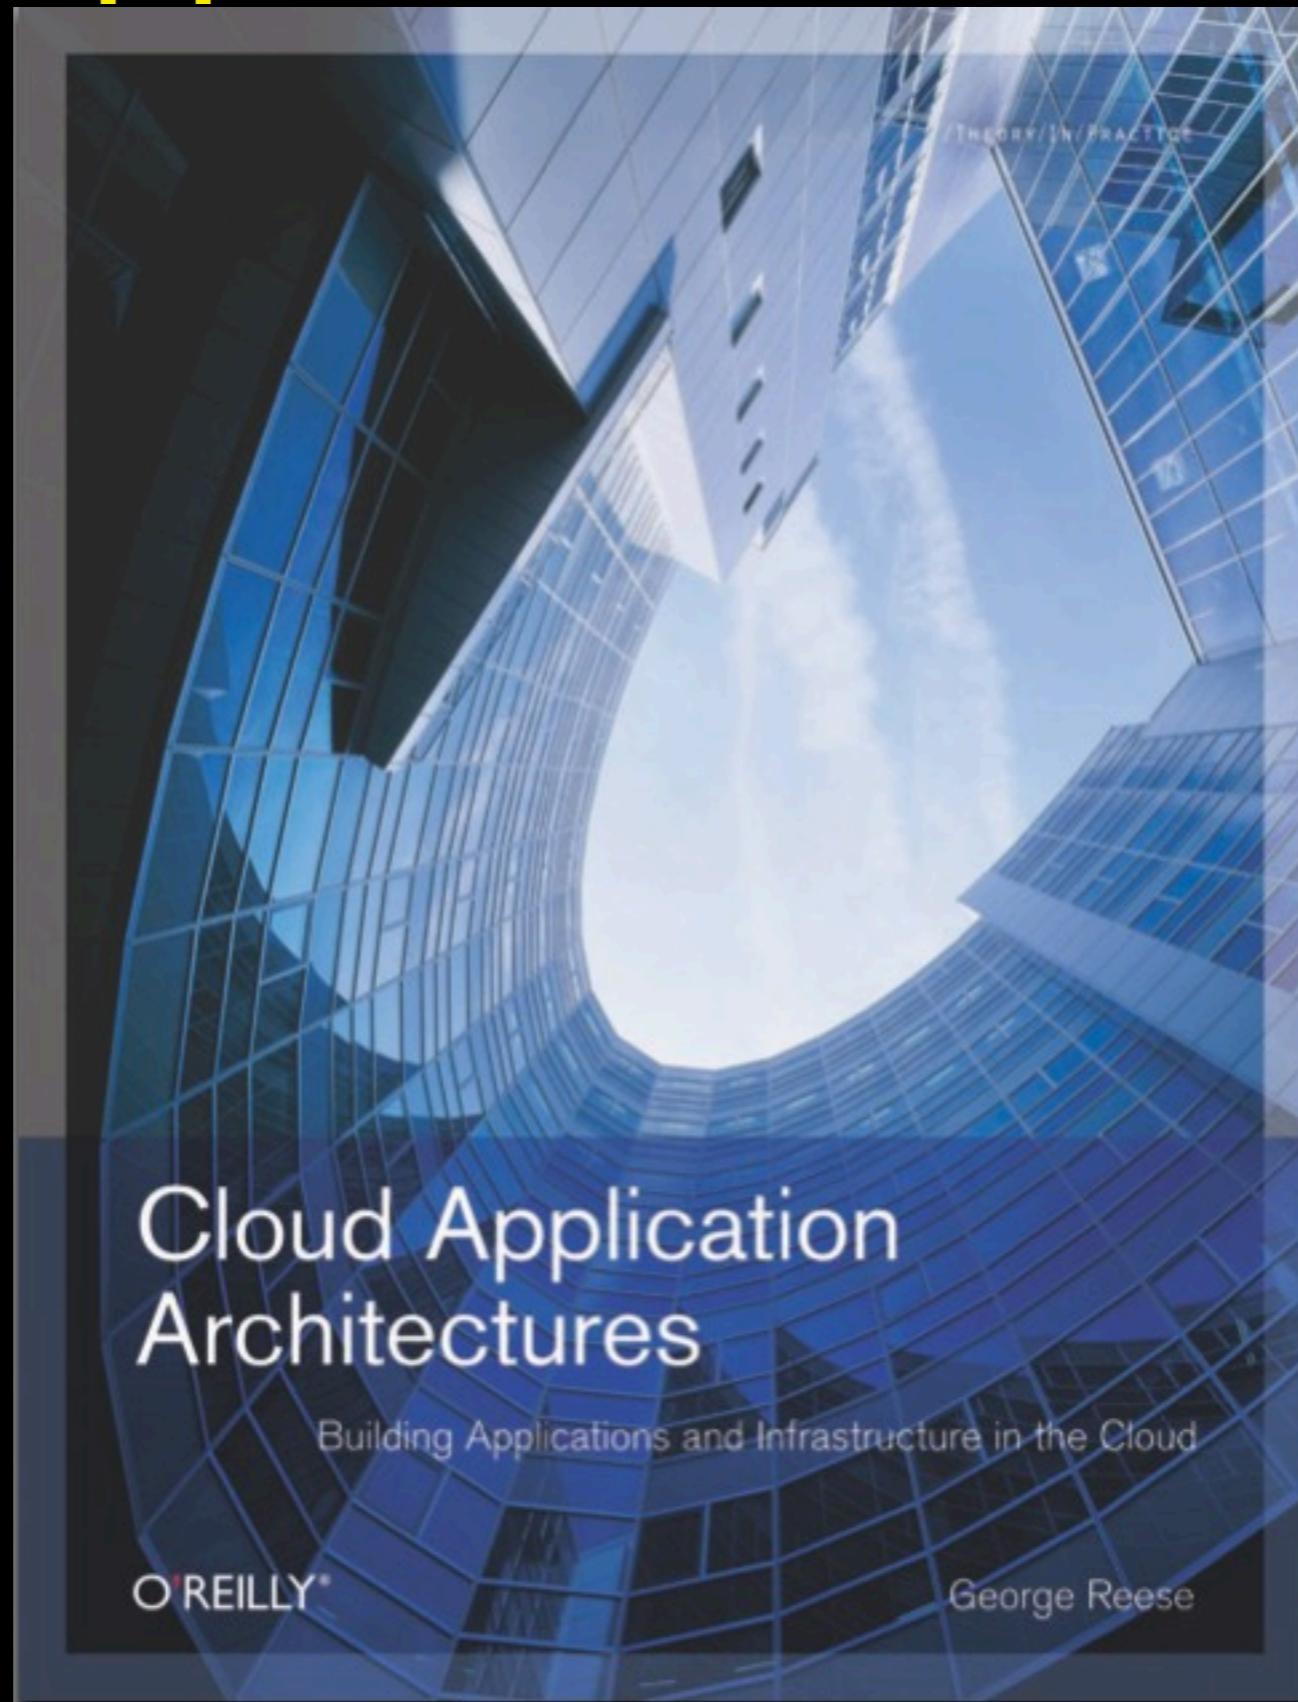
MySQL High Availability



Web Operations



Cloud Application Architectures



AGENDA

Why replication

- The web economy
- Scaling out

How to set replication

- From single server to replication
- Adding a slave

Using Replication

- Binary log formats
- What gets replicated and how
- Replication awareness

Managing replication

- Monitoring
- Log management
- Replacing a slave
- Replacing a master

What Replication is for

- master replacement
- backup
- better reads

Gotchas, tips, and tricks

More info

- What to read
- **More replication sessions**

More replication sessions - Wednesday

- 11am Building a multi-master, multi-region database infrastructure in Amazon EC2
- 11am Performance practices for minimizing replication delay
- 1pm Diagnosing & Fixing MySQL Replication
- Replaying database load with Percona Playback
- 2pm Build simple and complex replication clusters with Tungsten Replicator
- 3:30pm What's new in MySQL 5.5 and 5.6 Replication

More replication sessions - Thursday

- 11am MySQL Replication: Pros and Cons
- 2pm Boost Your Replication Throughput with Parallel Apply, Prefetch, and Batching
- 3pm Verifying MySQL Replication Safely With pt-table-checksum 2.0



continuent

Database replication and clustering

WE ARE HIRING!

- Implementation/support engineer
- QA engineer
- Documentation writer

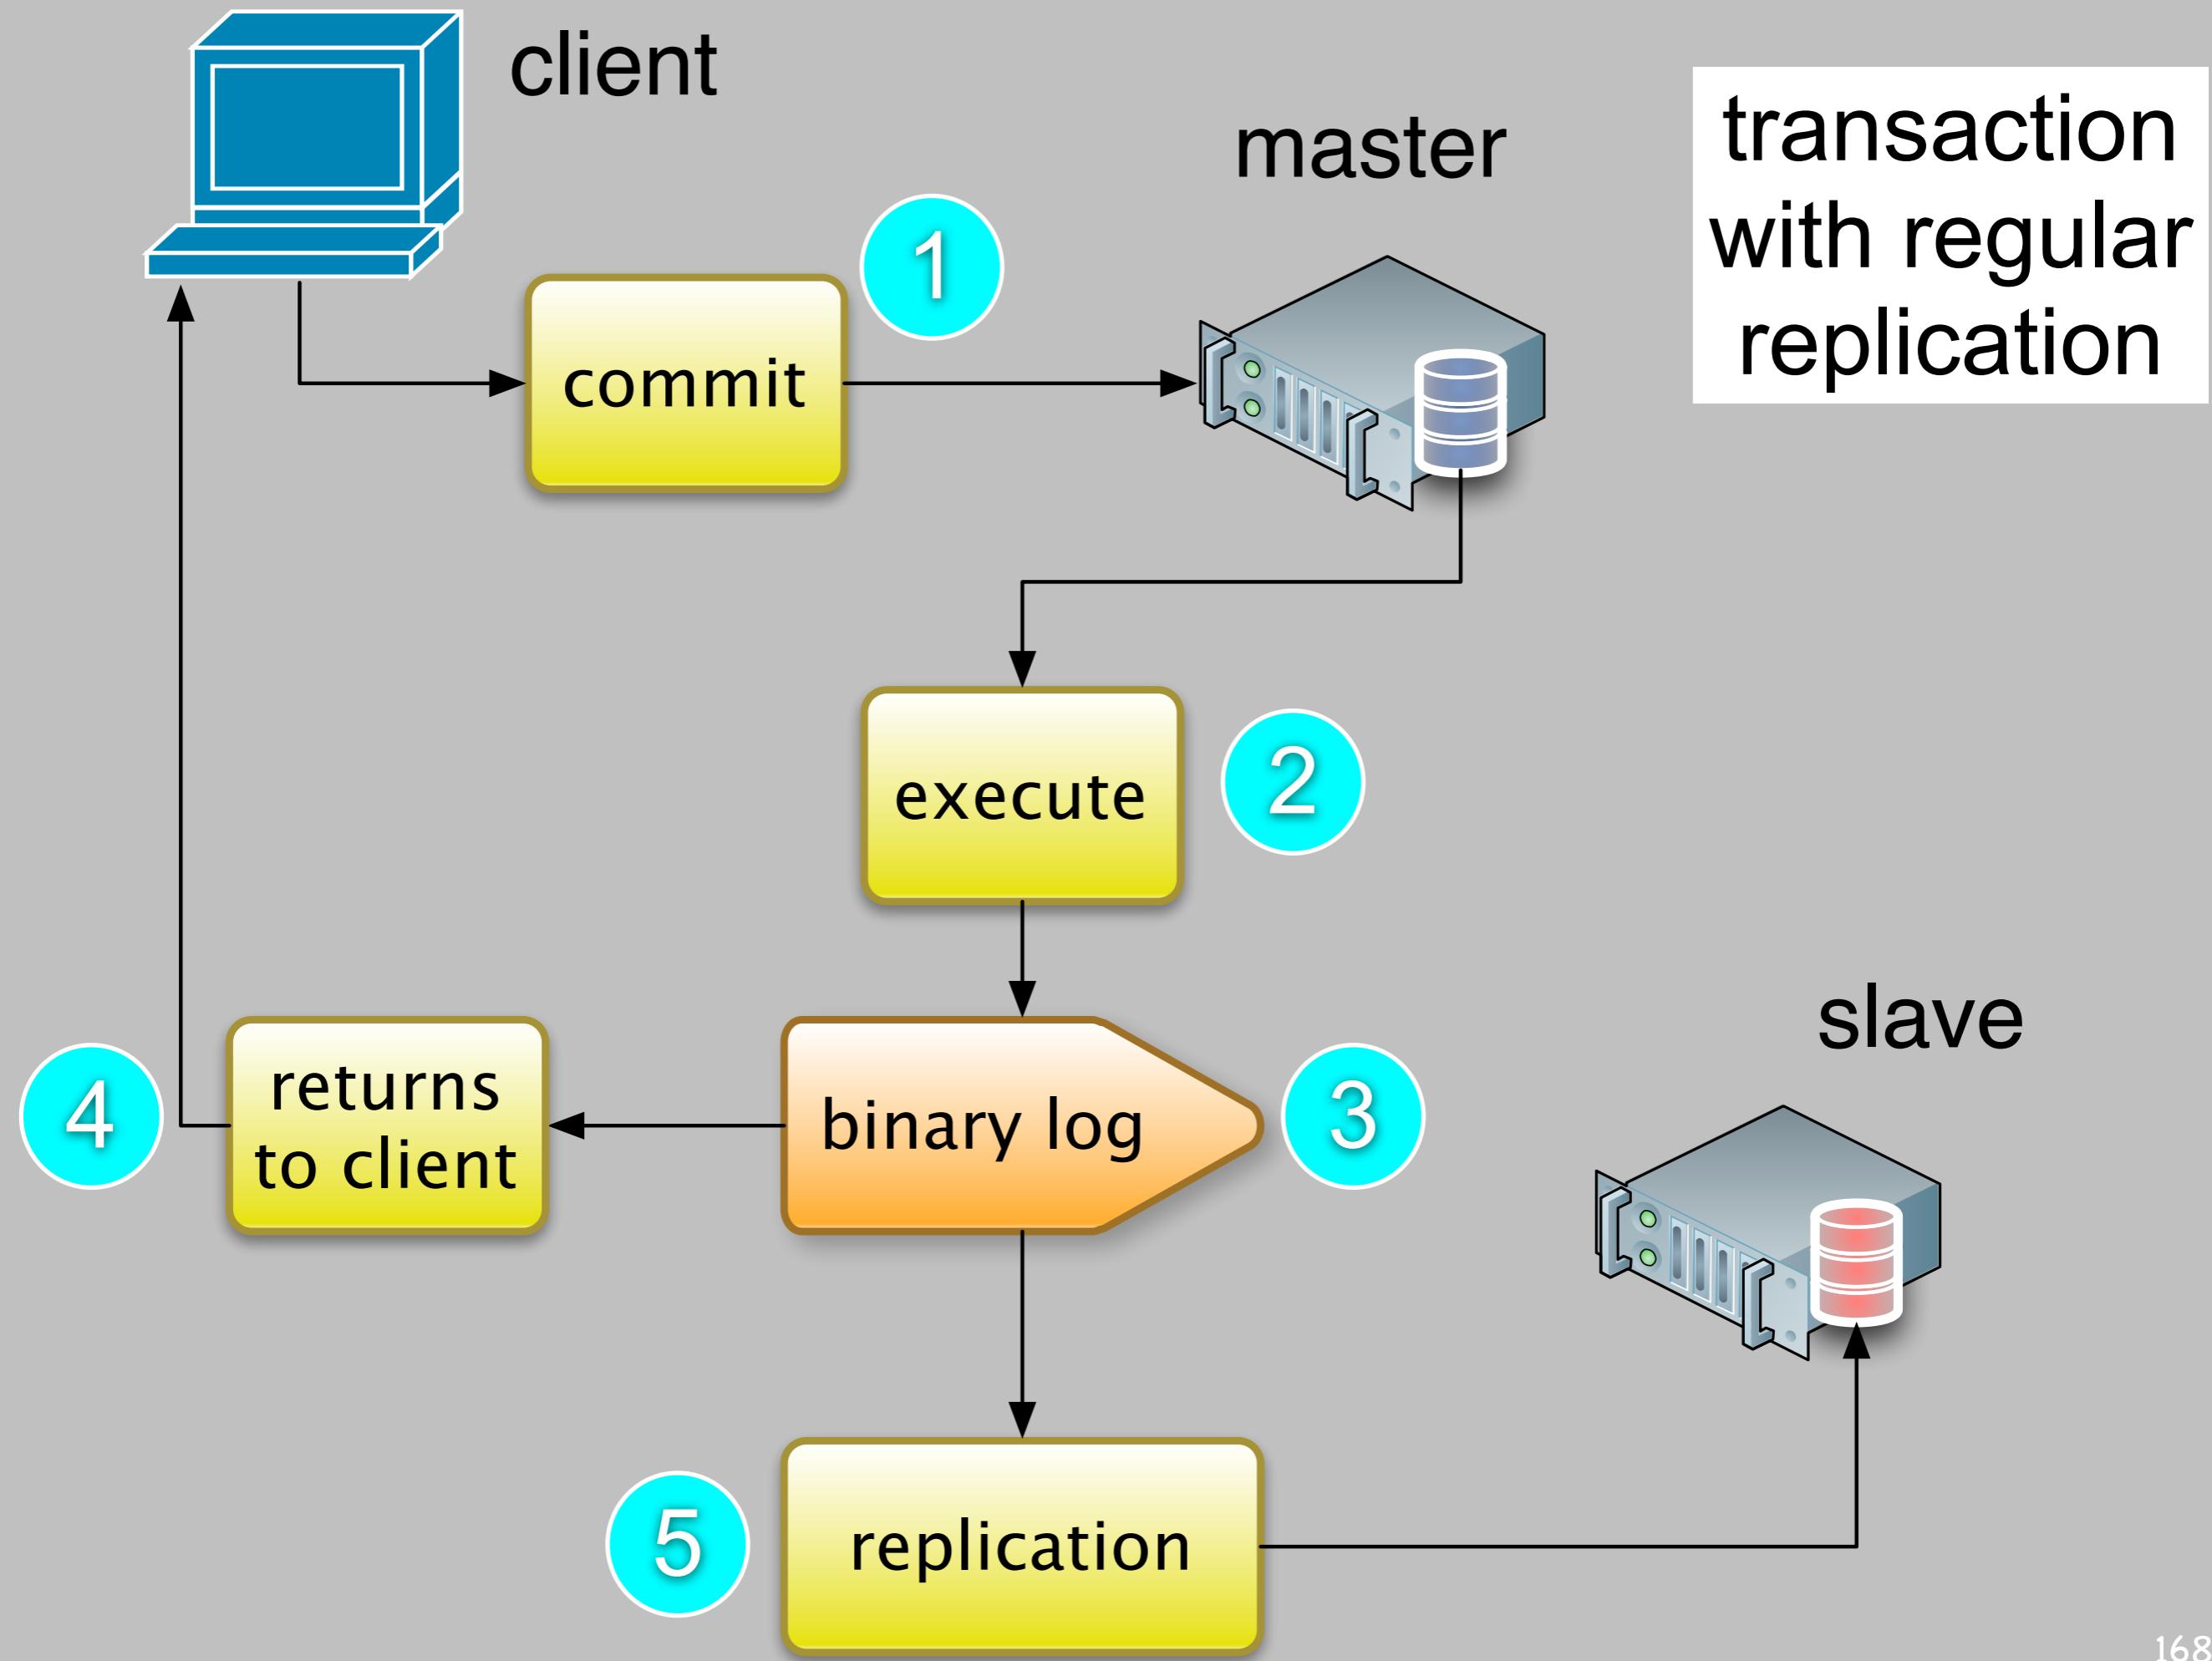
- <http://www.slideshare.net/datacharmer>

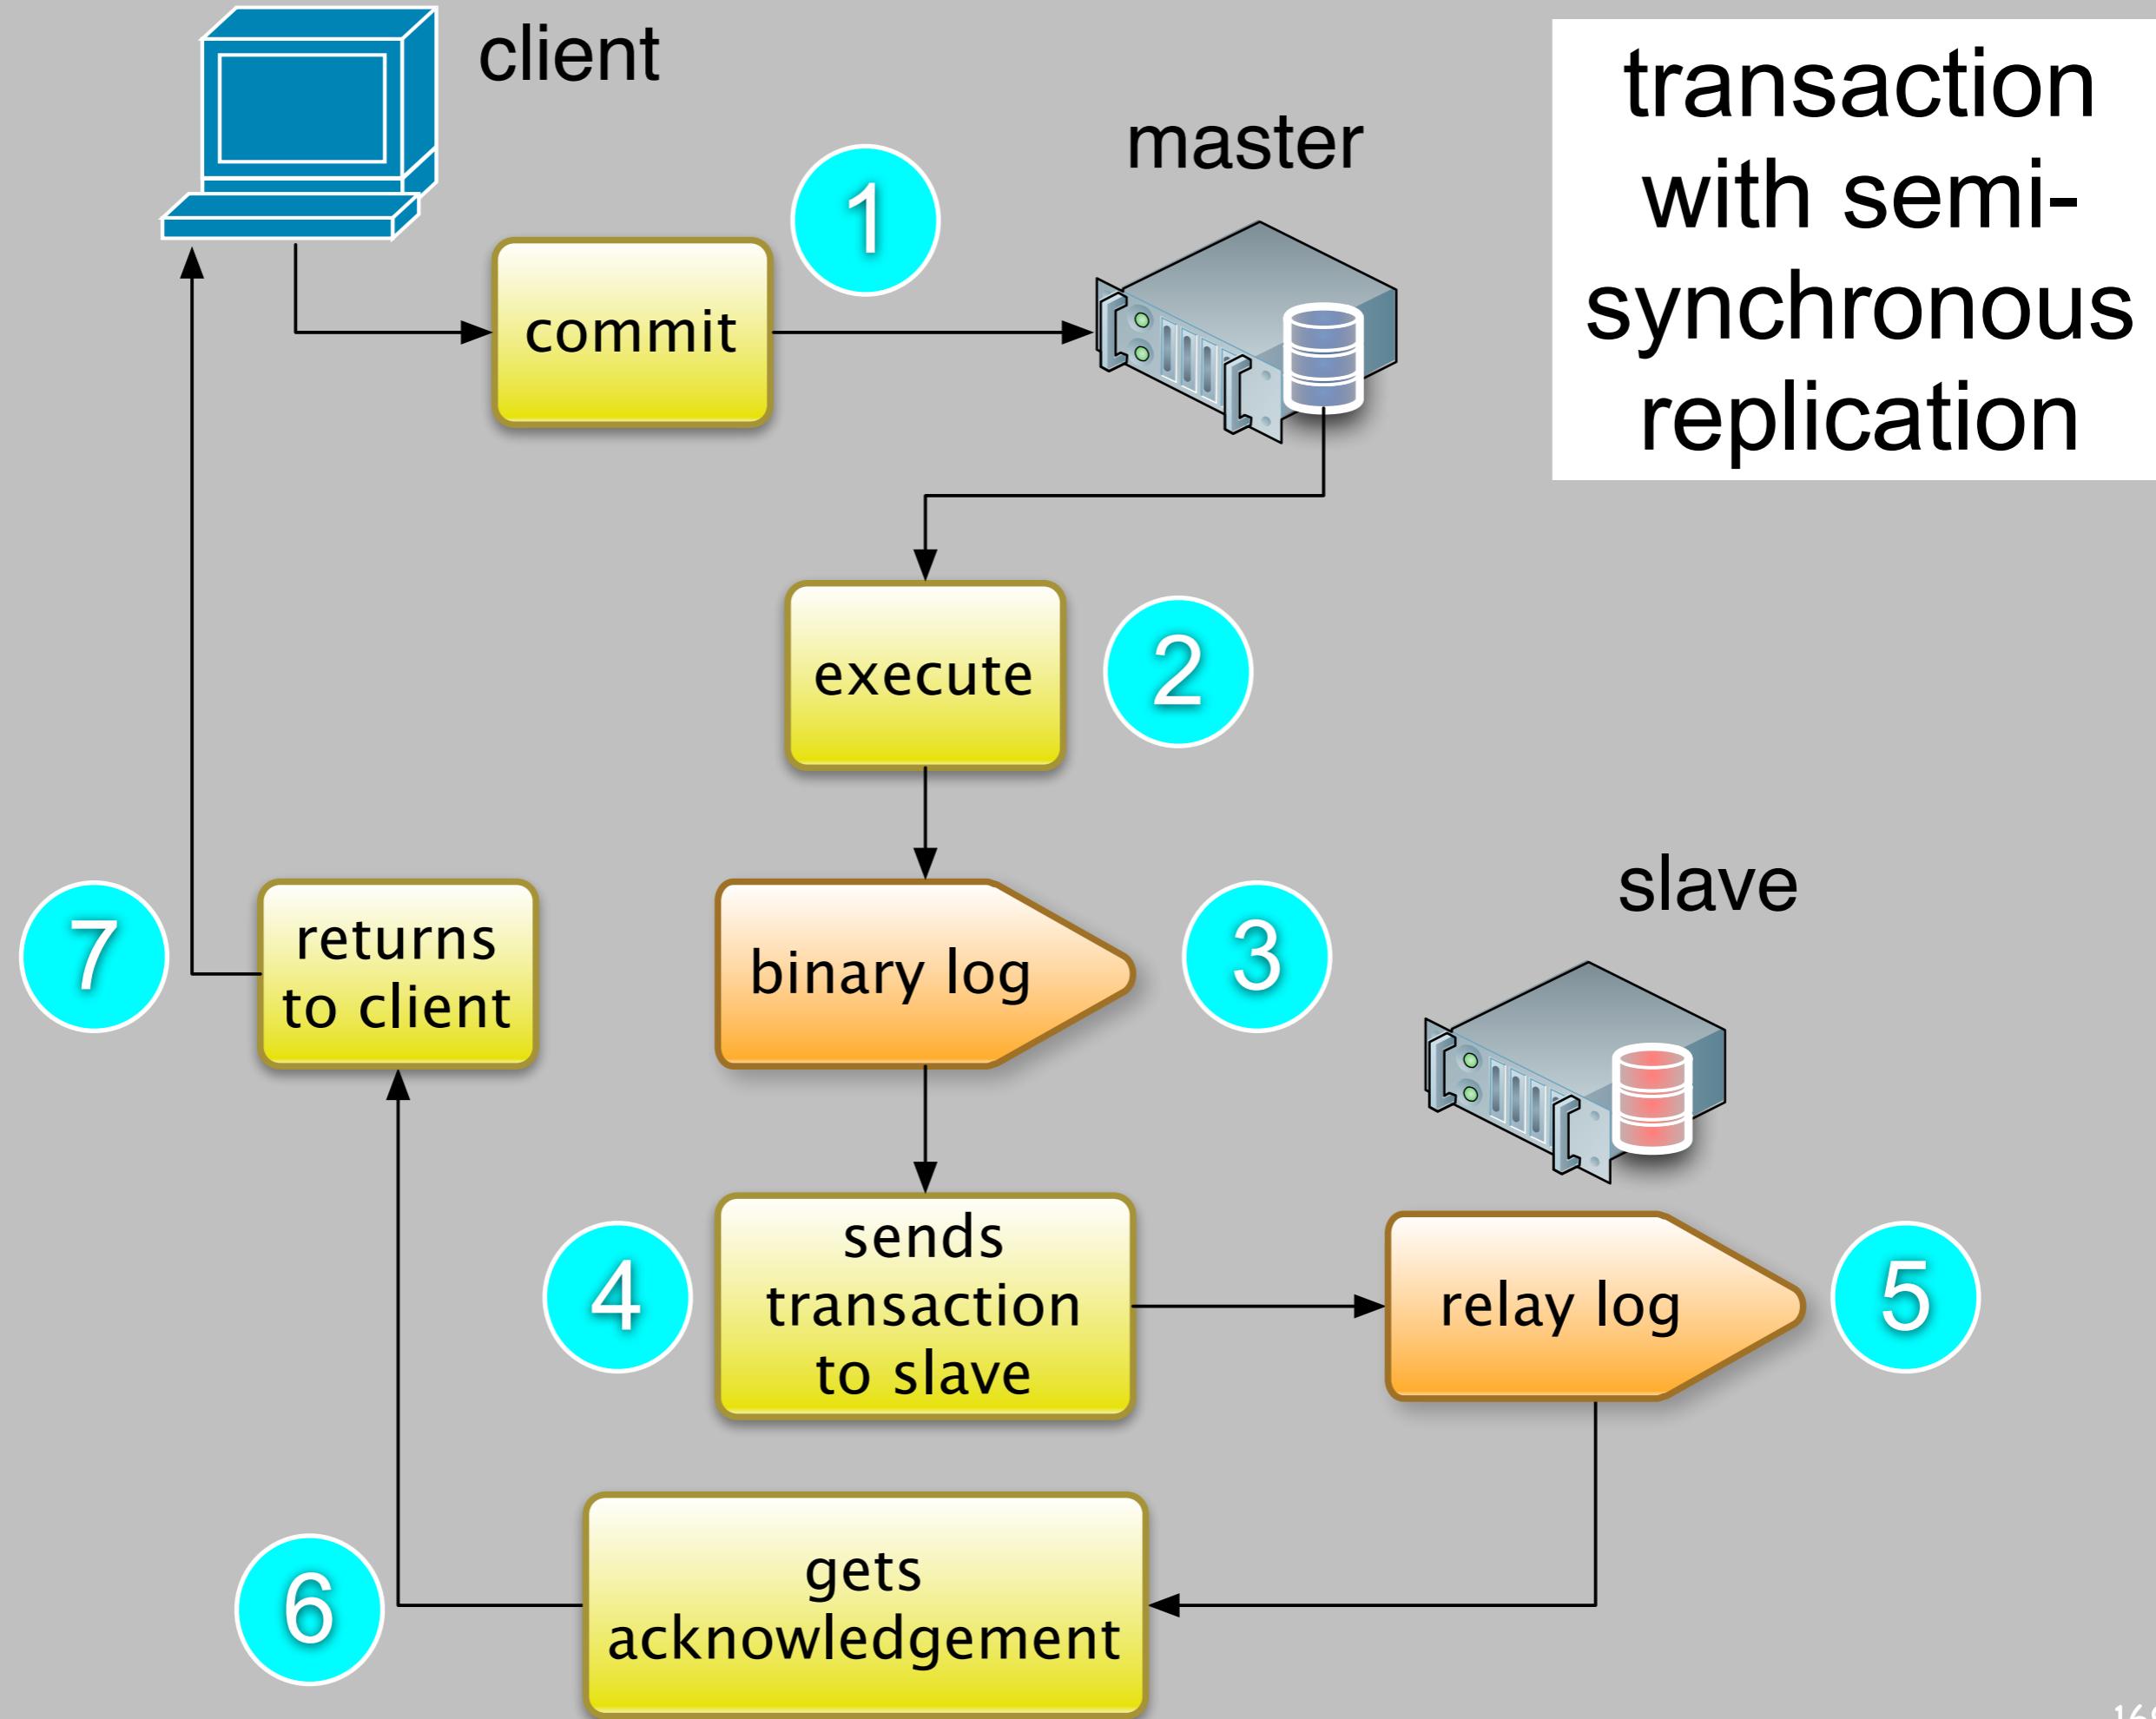
Bonus slides

- Semi-synchronous replication

semi-synchronous replication

- Available in 5.5 and higher
- Makes sure that at least one slave has copied the data.
- Increases reliability





semi-synchronous replication in practice

- installation:
 - it's a plugin.
 - Actually, two plugins

semi-synch replication install

```
# in the master
plugin-load=rpl_semi_sync_master=semisync_master.so
rpl_semi_sync_master_enabled=1

# in each slave
plugin-load=rpl_semi_sync_slave=semisync_slave.so
rpl_semi_sync_slave_enabled=1

# restart all servers
```

semi-synch replication check

```
# in the master
show variables like 'rpl_semi%';
+-----+-----+
| Variable_name | Value |
+-----+-----+
| rpl_semi_sync_master_enabled | ON   |
| rpl_semi_sync_master_timeout | 10000 |
| rpl_semi_sync_master_trace_level | 32   |
| rpl_semi_sync_master_wait_no_slave | ON   |
+-----+-----+
```

semi-synch replication check

```
show status like "rpl_semi_%tx";
```

variable_name	value
RPL_SEMI_SYNC_MASTER_NO_TX	0
RPL_SEMI_SYNC_MASTER_YES_TX	0

semi-synch replication test

```
master> create table t1 ( i int);  
Query OK, 0 rows affected (0.01 sec)
```

```
master> show status like "rpl_semi_%tx";  
+-----+-----+  
| Variable_name | Value |  
+-----+-----+  
| Rpl_semi_sync_master_no_tx | 0 |  
| Rpl_semi_sync_master_yes_tx | 1 |  
+-----+-----+
```

disabling semi-synch

```
# for each slave

set global rpl_semi_sync_slave_enabled=0;
stop slave io_thread;
start slave io_thread;
```

disabled semi-synch replication test

```
master> insert into t1 values (1);  
Query OK, 1 row affected (10.00 sec)
```

```
master> show status like "rpl_semi_%tx";  
+-----+-----+  
| Variable_name | Value |  
+-----+-----+  
| Rpl_semi_sync_master_no_tx | 1 |  
| Rpl_semi_sync_master_yes_tx | 1 |  
+-----+-----+  
2 rows in set (0.00 sec)
```

disabled semi-synch replication test

```
master> insert into t1 values (2);  
Query OK, 1 row affected (0.01 sec)
```

```
master> show status like "rpl_semi_%tx";  
+-----+-----+  
| Variable_name | Value |  
+-----+-----+  
| Rpl_semi_sync_master_no_tx | 2 |  
| Rpl_semi_sync_master_yes_tx | 1 |  
+-----+-----+  
2 rows in set (0.00 sec)
```

re-enabling semi-synch

```
# in one slave
```

```
set global rpl_semi_sync_slave_enabled=1;  
stop slave io_thread;  
start slave io_thread;
```

reenabled semi-synch replication test

```
master> insert into t1 values (3);  
Query OK, 1 row affected (0.01 sec)
```

```
master> show status like "rpl_semi_%tx";  
+-----+-----+  
| Variable_name | Value |  
+-----+-----+  
| Rpl_semi_sync_master_no_tx | 2 |  
| Rpl_semi_sync_master_yes_tx | 2 |  
+-----+-----+  
2 rows in set (0.00 sec)
```