

Replicating MySQL Data to TiDB

For Near Real-Time Analytics

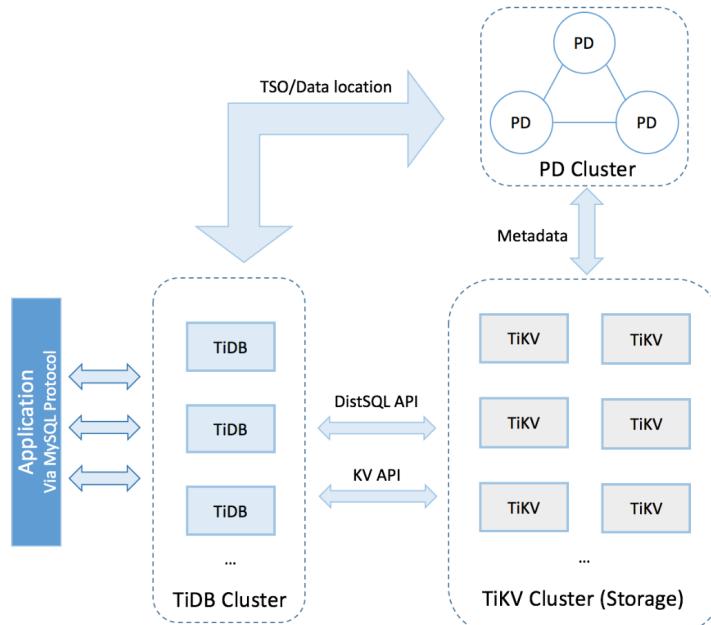


28-30 MAY 2019
AUSTIN, TEXAS

Who Are We

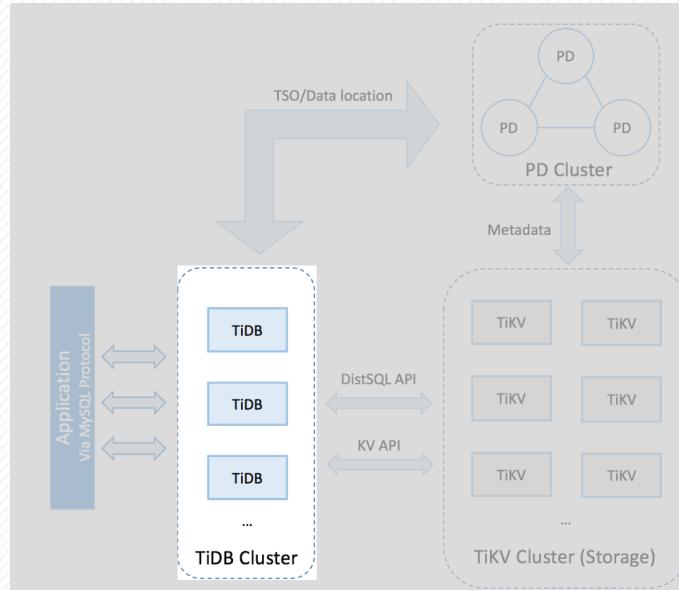
- Jervin Real, Architect
- Francisco Bordenave, Architect

TiDB at a Glance



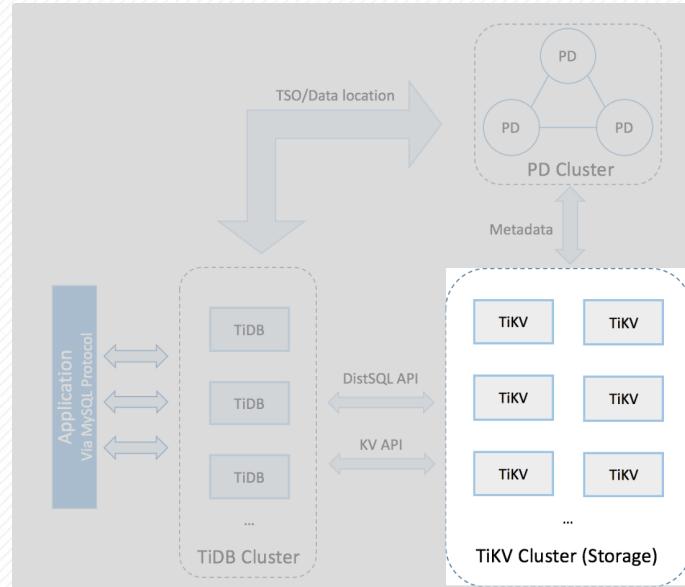
TiDB at a Glance

- TiDB cluster server process SQL and interacts with clients
 - MySQL Compatibility



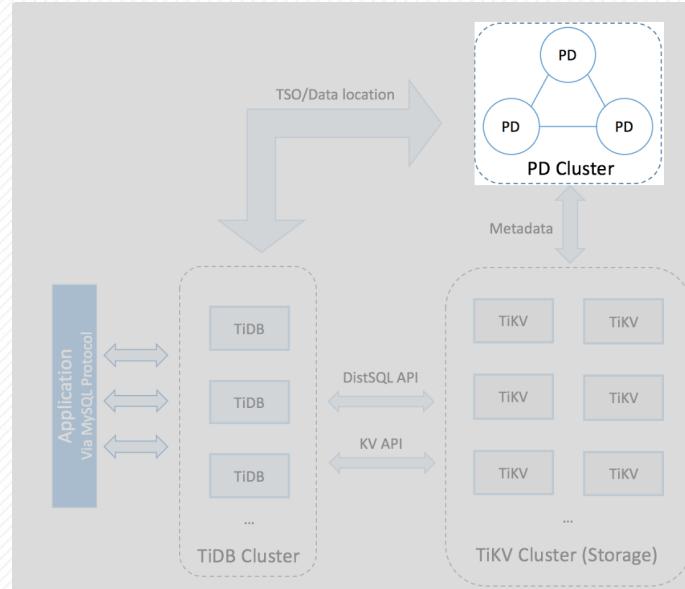
TiDB at a Glance

- TiKV cluster distributed storage, RocksDB



TiDB at a Glance

- PD cluster stores metadata, load balancing and monitor storage



MySQL Compatibility

- MySQL clients and drivers
- DML, DDL compatibility - some limitations
- Easy to start querying DB

Replication Goals

- Consistency
 - Resume from previous checkpoint
 - Verifying consistency
- Performance
 - Can we avoid replication lag?
- Observability
 - Is the process broken?
 - Are we querying stale data?





Toolbox

Ready To Use Tools



Toolbox

- Data Migration (full suite)
 - dm-master
 - dm-worker
 - mydumper
 - loader
 - syncer
- tidb-lightning
- DIY



mydumper

- <https://github.com/maxbube/mydumper> fork
 - Additional customizations
- Uncompressed dump only for loader

```
mydumper -B percona -o . -C -t 3 -h 10.3.4.4 \
-u tidb -p tidb -v 3
```



loader (1)

- Loads mydumper data in parallel, ~1MB chunk per load, resumable
- Uncompressed dump only
- `myloader` is possible too
 - Be aware of trx size limits (i.e. use `-q 1`)*
 - Not resumable
- Database/table filters/routing.

```
loader -d ./ -h 127.0.0.1 -P 4000 -u root -t 3 \
-rm-checkpoint -L info
```

```
myloader -d ./ -B percona -h 127.0.0.1 -u root \
-P 4000 -t 3 -q 1
```

*<https://github.com/tikv/tikv/issues/2986>

loader (2)

- Database/table filters/routing.

```
{  
  "log-level": "info",  
  "log-file": "",  
  "status-addr": ":8272",  
  "pool-size": 3,  
  "dir": "./",  
  "db": {  
    "host": "127.0.0.1",  
    "user": "root",  
    "port": 4000,  
    "sql-mode": "@DownstreamDefault"  
  },  
  "checkpoint-schema": "tidb_loader",  
  "config-file": "",  
  "route-rules": null,  
  "do-table": null,  
  "do-db": null,  
  "ignore-table": null,  
  "ignore-db": null,  
  "rm-checkpoint": false  
}
```



loader Status

- Available from `tidb_loader.checkpoint` by default

```
mysql> select * from tidb_loader.checkpoint \G
***** 1. row *****
    id: 13917f
  filename: percona.history.sql
 cp_schema: percona
  cp_table: history
    offset: 55998837
   end_pos: 2685020785
create_time: 2019-05-29 02:56:14
update_time: 2019-05-29 03:08:53
***** 2. row *****
    id: 13917f
  filename: percona.search.sql
 cp_schema: percona
  cp_table: search
    offset: 59999013
   end_pos: 812604151
create_time: 2019-05-29 02:56:14
update_time: 2019-05-29 03:08:52
...
```



syncer (1)

- Resumes from loader (or any replication coordinates)
- Runs per table, per event (not at transaction level)
- Allows routing/filtering per table
- Allows DML filtering
- Multiple worker threads, but still in commit order.
 - Eventually consistent

```
# cat config.toml
...
#[[skip-dmls]]
# db-name = "foo"
# tbl-name = "bar"
# type = "delete"
...

syncer -config config.toml
```



syncer (2)

```
{  
  "log-level": "info",  
  "log-file": "syncer.log",  
  "log-rotate": "day",  
  "status-addr": ":8271",  
  "server-id": 1234567890,  
  "meta": "./syncer.meta",  
  "persistent-dir": "",  
  "flavor": "mysql",  
  "worker-count": 16,  
  "batch": 100,  
  "max-retry": 100,  
  "replicate-do-table": null,  
  "replicate-do-db": null,  
  "replicate-ignore-table": null,  
  "replicate-ignore-db": null,  
  "skip-ddls": [],  
  "skip-dmls": null,  
  "route-rules": null,  
  ...  
}  
  
...  
  "from": {  
    "host": "10.3.4.4",  
    "user": "tidb",  
    "port": 3306  
  },  
  "to": {  
    "host": "127.0.0.1",  
    "user": "root",  
    "port": 4000  
  },  
  "enable-gtid": false,  
  "auto-fix-gtid": false,  
  "disable-detect": false,  
  "safe-mode": false,  
  "config-file": "config.toml",  
  "stop-on-ddl": false,  
  "execute-ddl-timeout": "3h",  
  "execute-dml-timeout": "1m",  
  "execute-queue-length": 5000,  
  "enable-ansi-quotes": false,  
  "timezone": ""  
}
```



syncer Status

- HTTP Interface

```
me@tidb:~/git/dm/syncer# curl localhost:8271/metrics \
    | egrep '[syncer_binlog_file|syncer_binlog_pos]\{node\}'

syncer_binlog_file{node="master"} 3
syncer_binlog_file{node="syncer"} 3
syncer_binlog_pos{node="master"} 7.31884075e+08
syncer_binlog_pos{node="syncer"} 7.32175424e+08
```

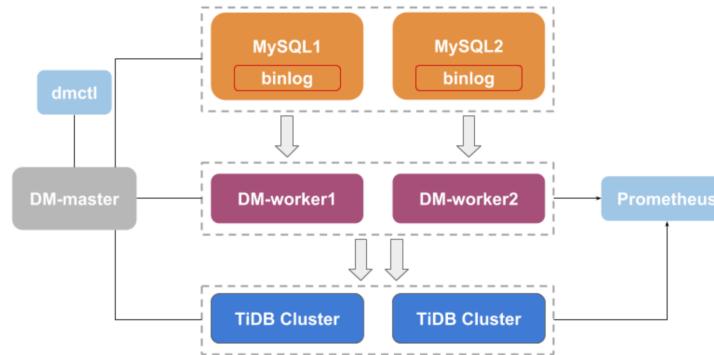
- Log file

```
2019/05/24 16:09:35 syncer.go:951: [info] [syncer]total events = 407220,
      tps = 719.666667, master-binlog = (mysql-bin.000009, 220613462),
      master-binlog-gtid=, syncer-binlog = (mysql-bin.000006, 306307142),
      syncer-binlog-gtid =
```

- TiDB checkpoint table when used with dm-woker

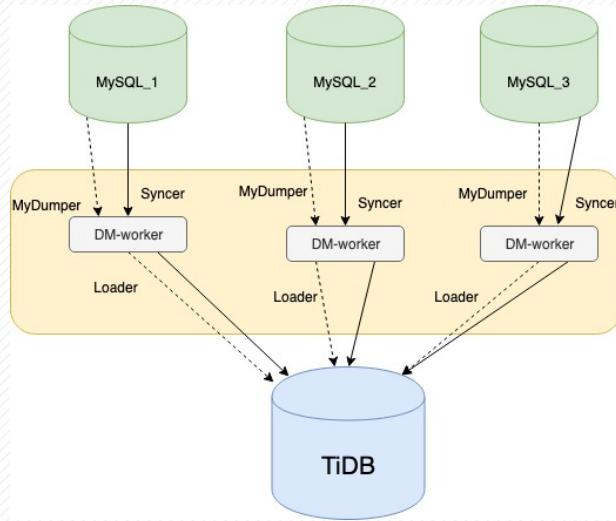
dm - Data Migration Tool

- dm-master, dm-worker
 - One dm-worker per replication source
 - Sharded table to single table



dm - Data Migration Tool

- dm-worker = (mydumper + loader + syncer)
 - binlog-[ignore|replicate|do]-[db|table]
 - skip-dmls (db|table|type)



lighting - super fast loader tool

- tikv-importer + tidb-lighting
- Parse mydumper or CSV files and convert into kv files
- Loads data directly to tikv server
- Disable compaction and set cluster in 'import' mode (makes writes way faster!)
 - In case of import failure cluster needs to be set back to 'normal' mode
- During load cluster service is limited



How Others Do It

Replicating to Another Platform

To Cassandra

- Apache Spark
- Striim <https://www.striim.com/integrations/mysql-cassandra/>
- Tungsten replicator

All solutions requires some type of ETL process. None of them are easy to implement as requires model de-normalization

To Clickhouse

- Altinity <https://github.com/altinity/clickhouse-mysql-data-reader>
- Tungsten replicator (of course)
- CSV to clickhouse-client (batch)
- DIY for tricky situations
 - See how at 5:40PM at Bid Bend CD
 - Jervin's next talk!

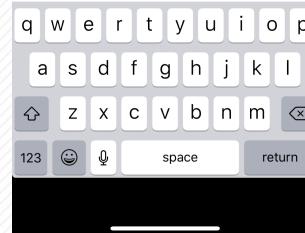
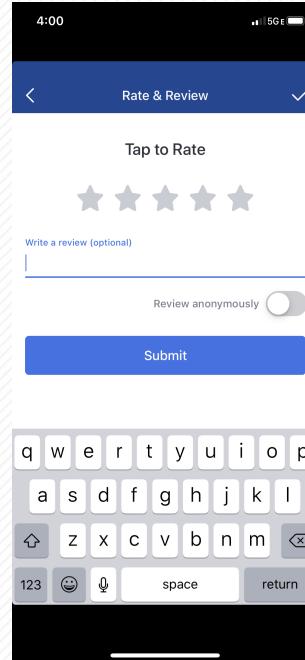
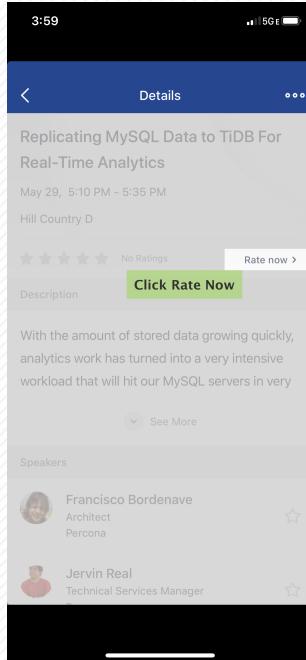
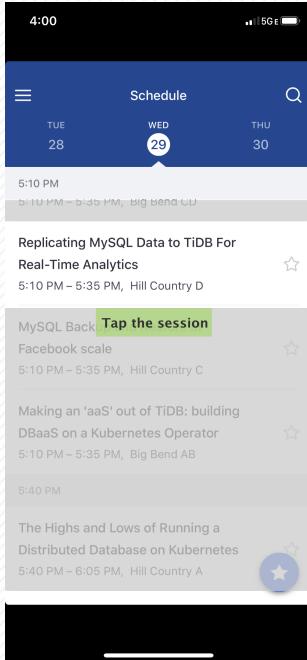
MariaDB ColumnStore

- Straight forward is distributed with MariaDB out of the box (not fully true)

Some nice benchmarks:

<https://www.percona.com/blog/2017/03/17/column-store-database-benchmarks-mariadb-columnstore-vs-clickhouse-vs-apache-spark/>

Rate Our Talk



Thank you for joining us this year!

... and thanks to our sponsors!



ProxySQL



aiven

Pythian
love your data®



Bloomberg
Engineering



PERCONA
LIVE

26 / 27



Questions?

